

# Efficient Computing for AI and Robotics: From Hardware Accelerators to Algorithm Design

Vivienne Sze ( @eems\_mit)

Massachusetts Institute of Technology

*In collaboration with Luca Carlone, Yu-Hsin Chen, Joel Emer, Sertac Karaman, Tushar Krishna, Peter Li, Yi-Lun Liao, Fangchang Ma, Angshu Parashar, Amr Suleiman, Po-An Tsai, Diana Wofk, Nellie Wu, Tien-Ju Yang, Zhengdong Zhang*

Slides available at

<https://tinyurl.com/ISCAS2021Sze>

# Processing at “Edge” instead of the “Cloud”



Communication



Privacy



Latency

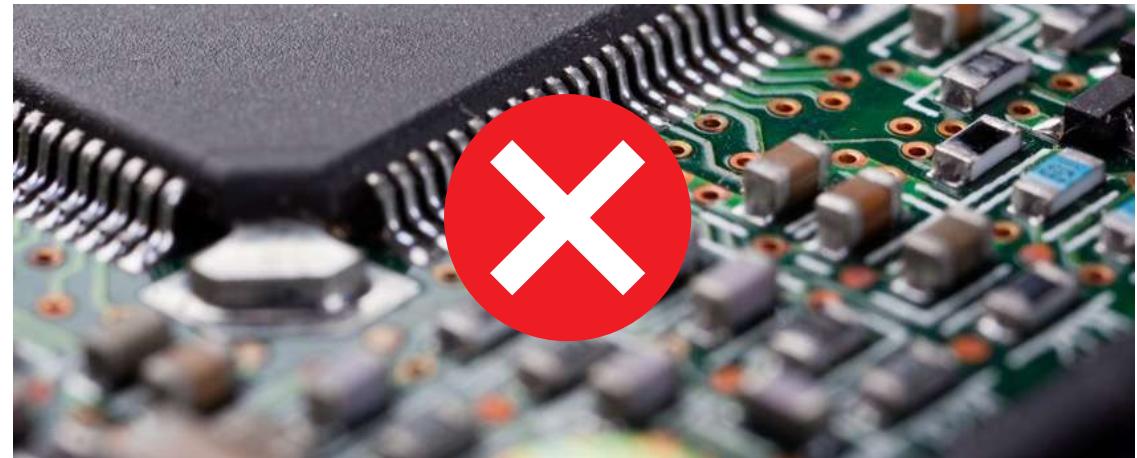
# Existing Processors Consume Too Much Power



< 1 Watt

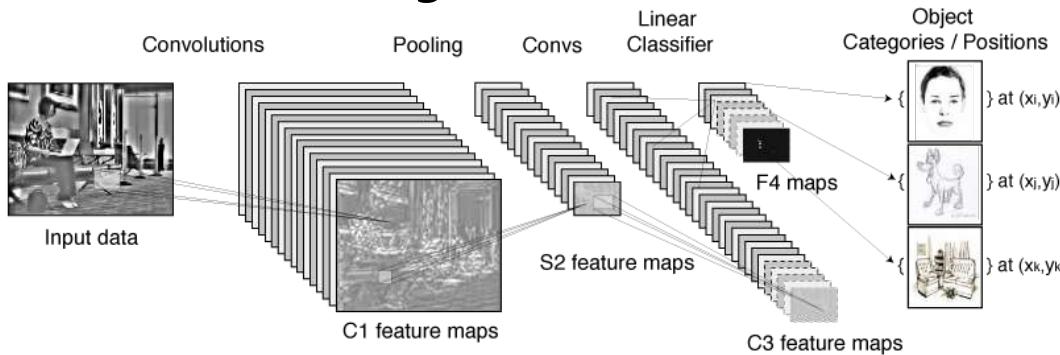


> 10 Watts



# Efficient Computing with Cross-Layer Design

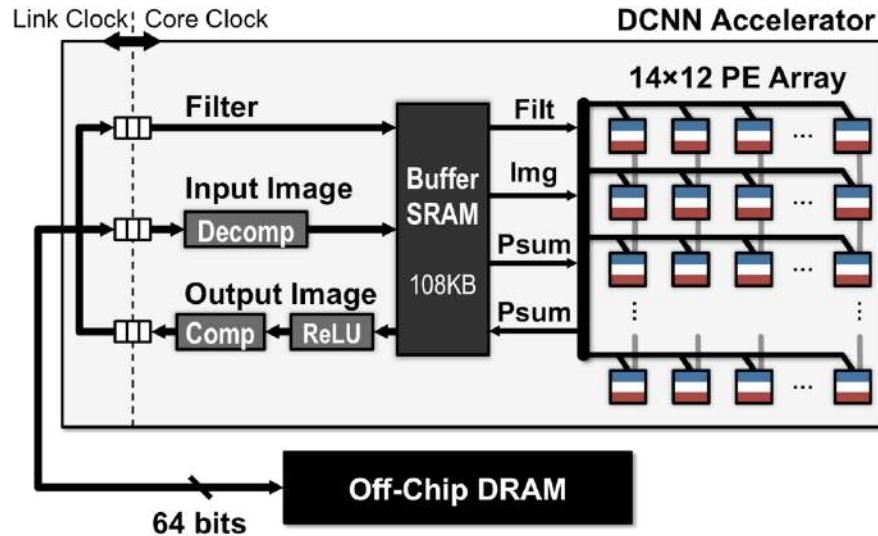
## Algorithms



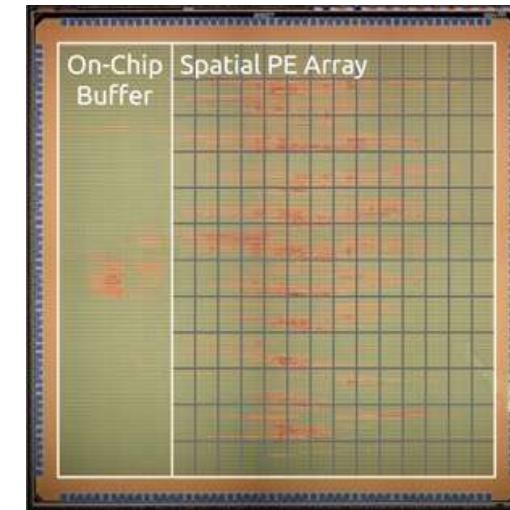
## Systems



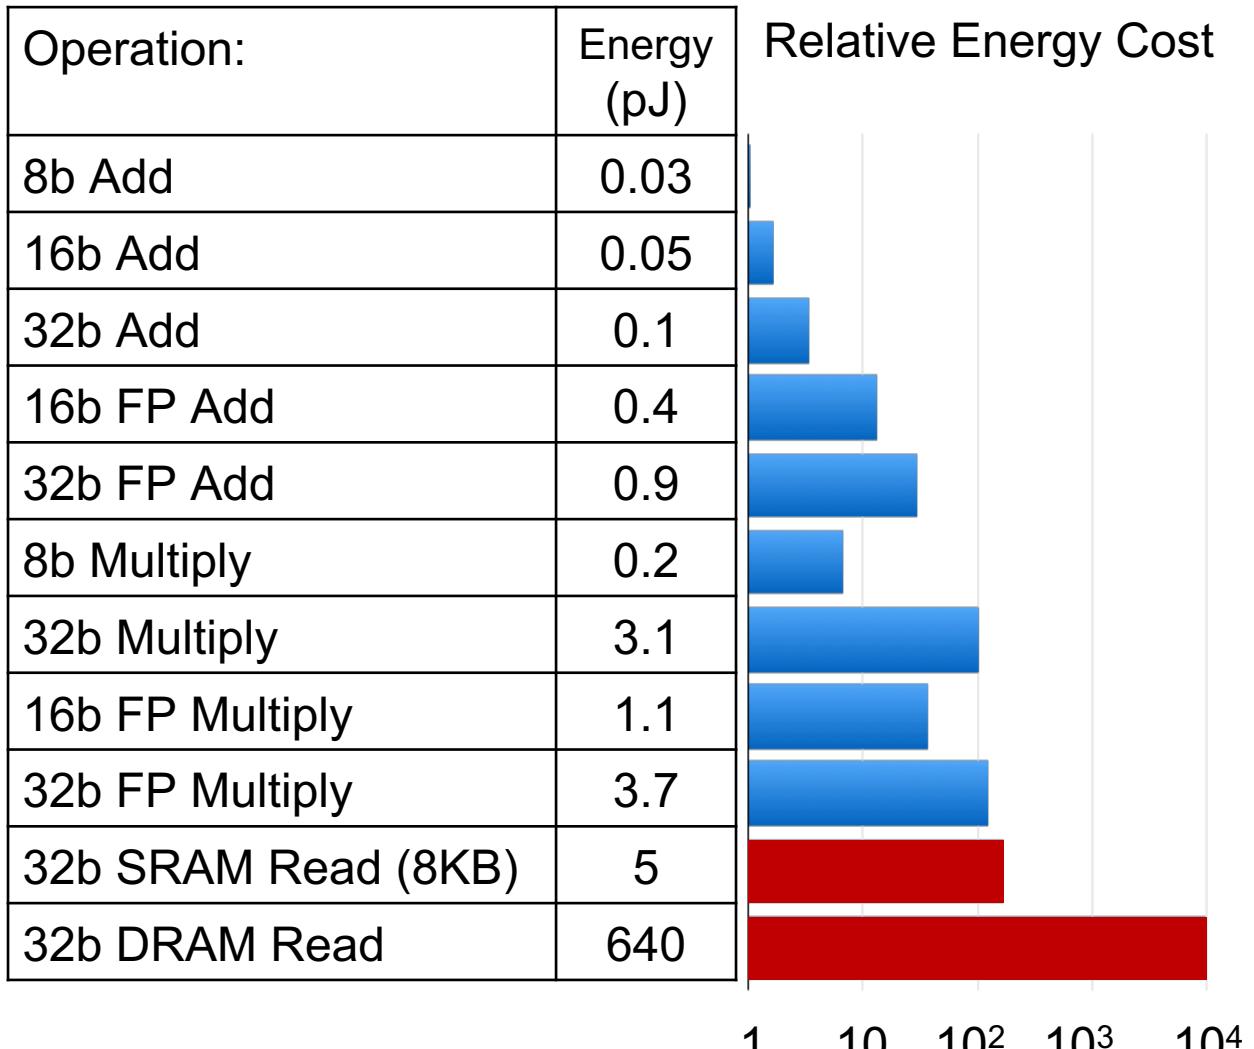
## Architectures



## Circuits



# Energy Dominated by Data Movement



Memory access is **orders of magnitude** higher energy than compute

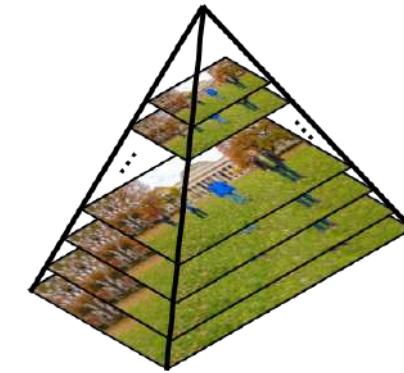
# Autonomous Navigation Uses a Lot of Data

## Semantic Understanding

- High frame rate
- Large resolutions
- Data expansion



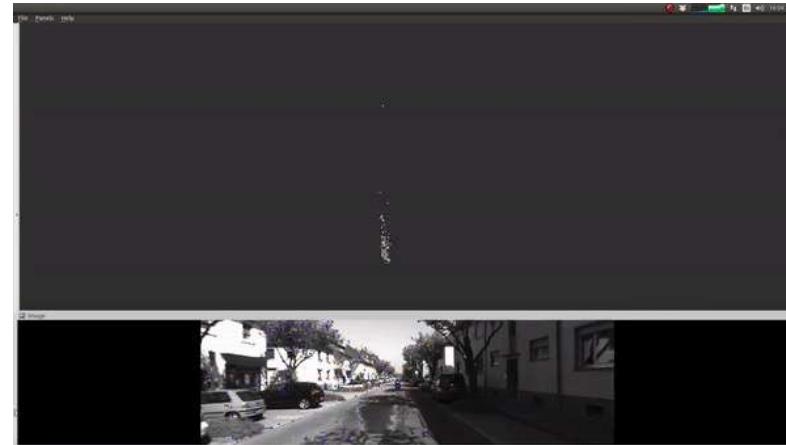
2 million pixels



10x-100x more pixels

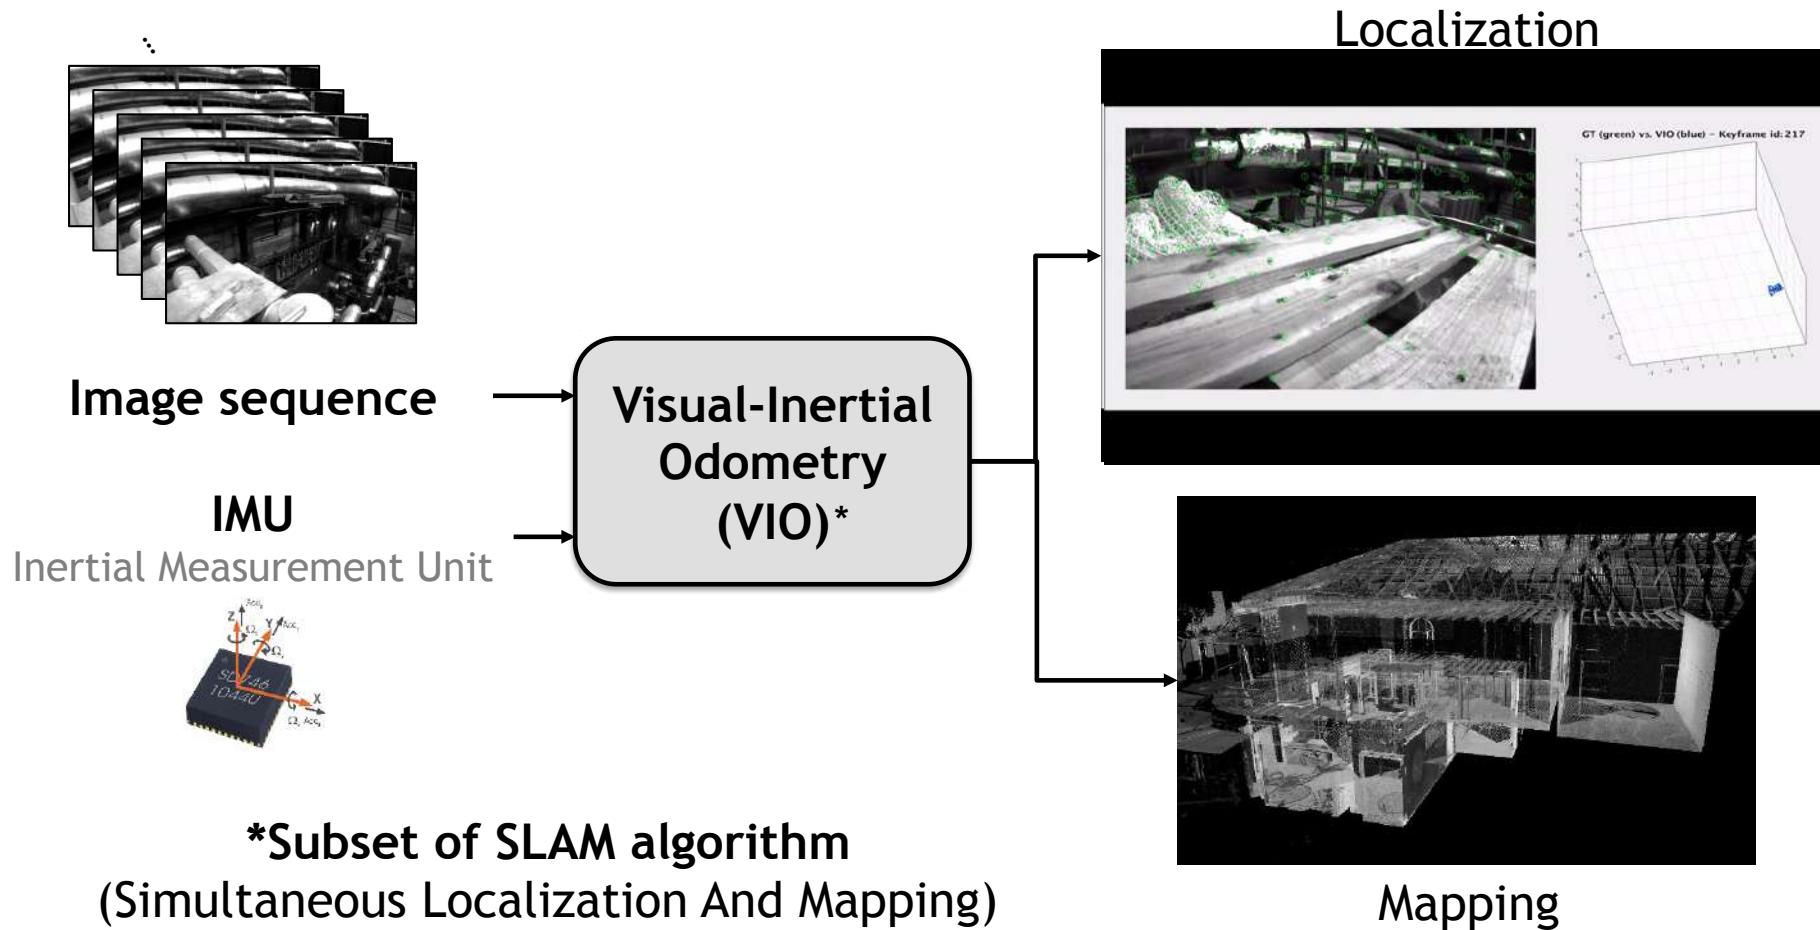
## Geometric Understanding

- Growing map size



# Visual-Inertial Localization

Determines location/orientation of robot from images and IMU  
(also used by headset in Augmented Reality and Virtual Reality)



# Localization at Under 25 mW

*First chip* that performs  
**complete** Visual-Inertial Odometry

Front-End for camera  
(*Feature detection, tracking, and outlier elimination*)

Front-End for IMU  
(*pre-integration of accelerometer and gyroscope data*)

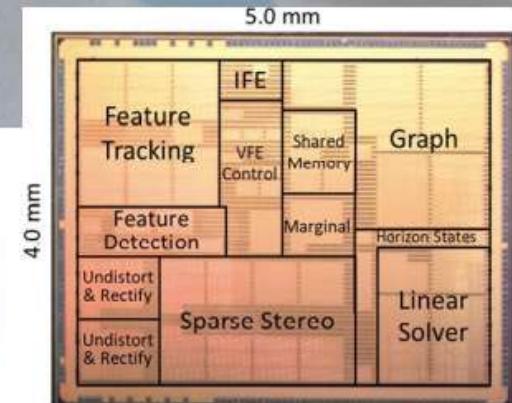
Back-End Optimization of Pose Graph

Consumes **684×** and **1582×** less energy than mobile and desktop CPUs, respectively



Navion

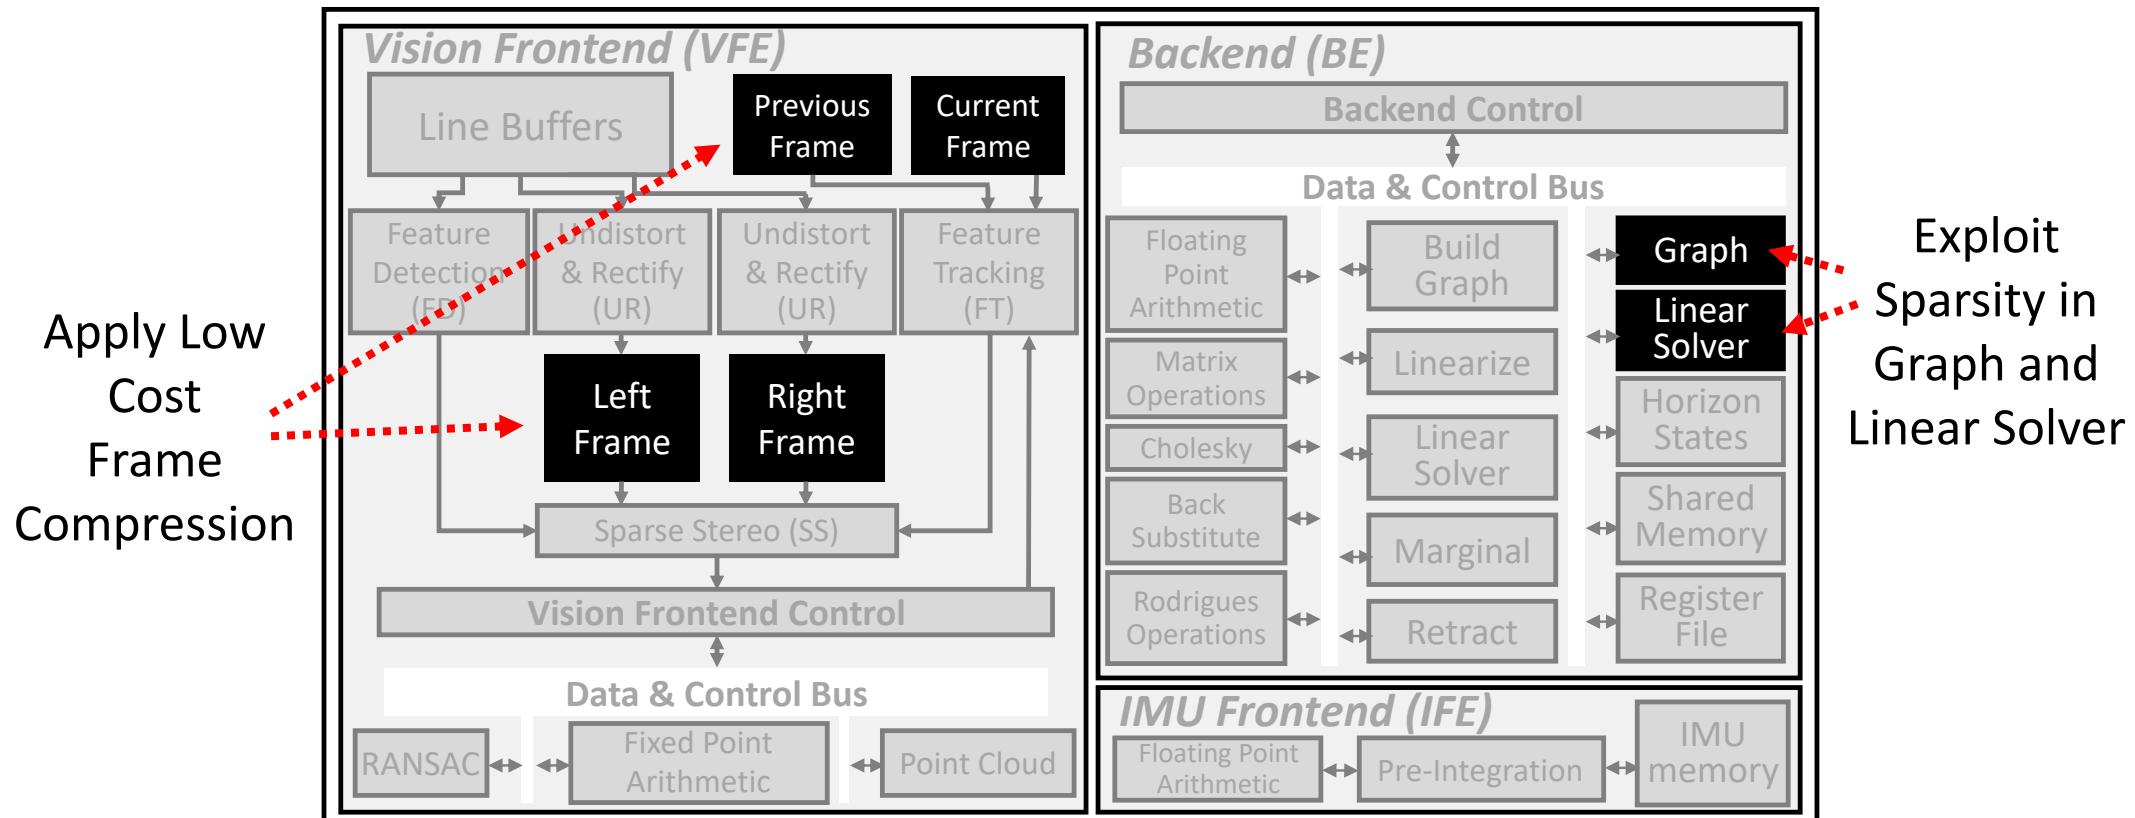
Technology	65nm CMOS	Supply	1 V
Chip area (mm <sup>2</sup> )	4.0 x 5.0	Resolution	752x480
Core area (mm <sup>2</sup> )	3.54 x 4.54	Camera rate	28 - 171 fps
Logic gates	2,043 kgates	Keyframe rate	16 - 90 fps
SRAM	854KB	Average Power	24 mW
VFE Frequency	62.5 MHz	GOPS	10.5 - 59.1
BE Frequency	83.3 MHz	GFLOPS	1 - 5.7



[Joint work with Sertac Karaman (AeroAstro)]

# Key Methods to Reduce Data Size

**Navion:** Fully integrated system – no off-chip processing or storage

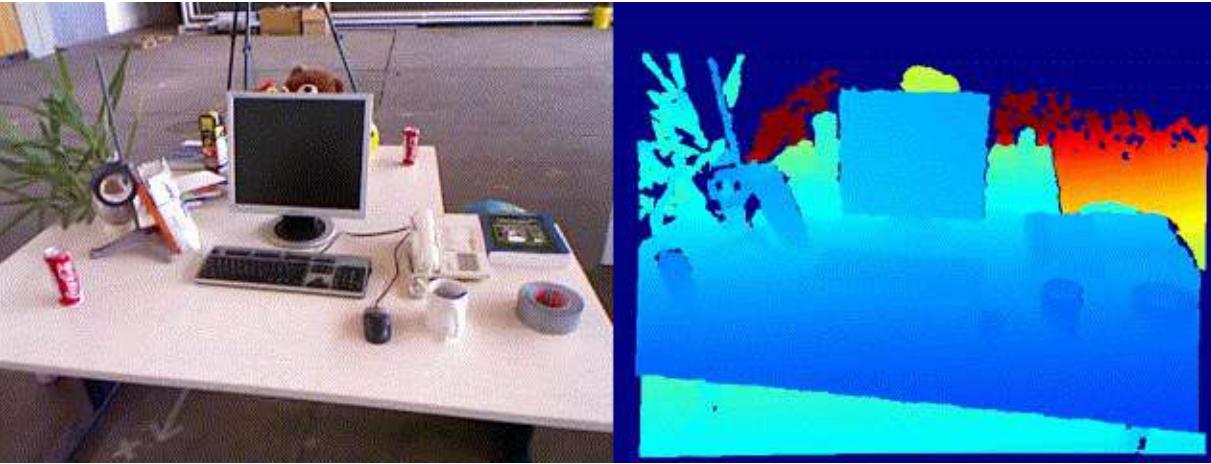


Use **compression** and **exploit sparsity** to reduce memory down to 854kB

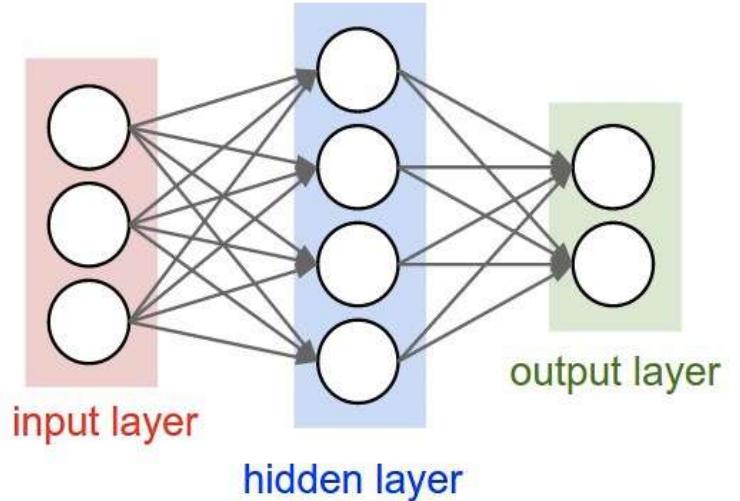
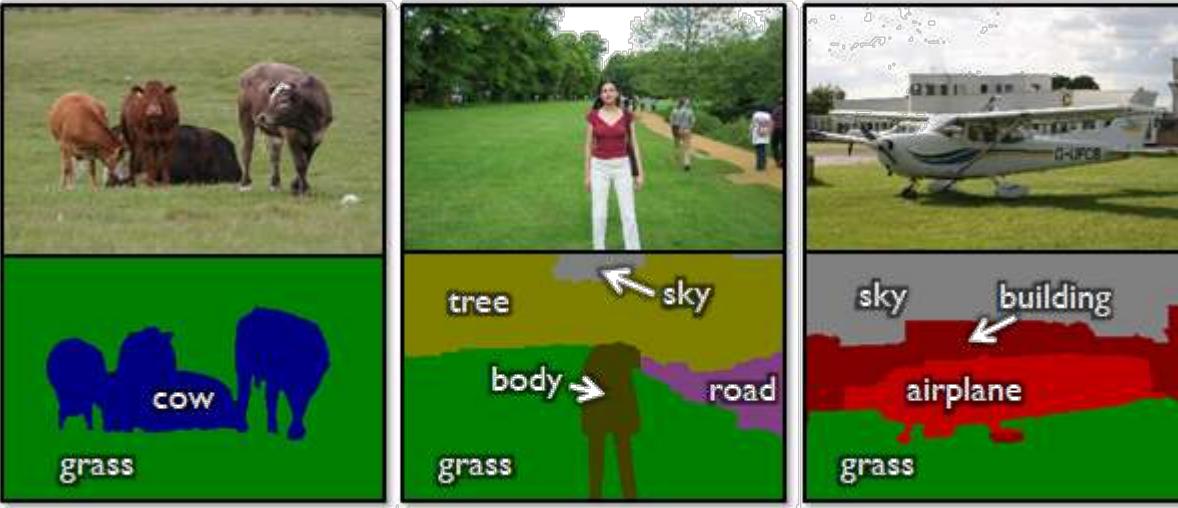
Navion Project Website: <http://navion.mit.edu>

# Understanding the Environment

## Depth Estimation



## Semantic Segmentation

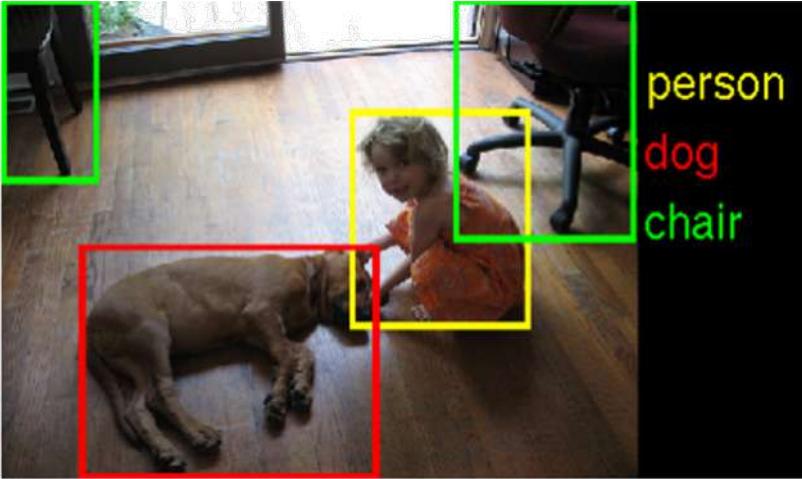


State-of-the-art approaches use **Deep Neural Networks**, which require up to several hundred millions of operations and weights to compute!  
*>100x more complex than video compression*

# Deep Neural Networks

*Deep Neural Networks (DNNs) have become a cornerstone of AI*

**Computer Vision**



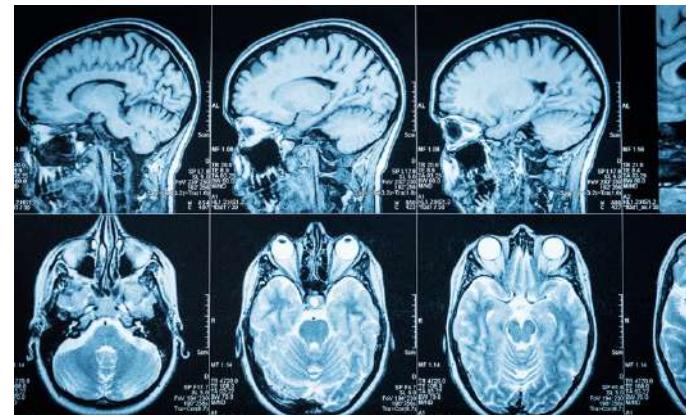
**Speech Recognition**



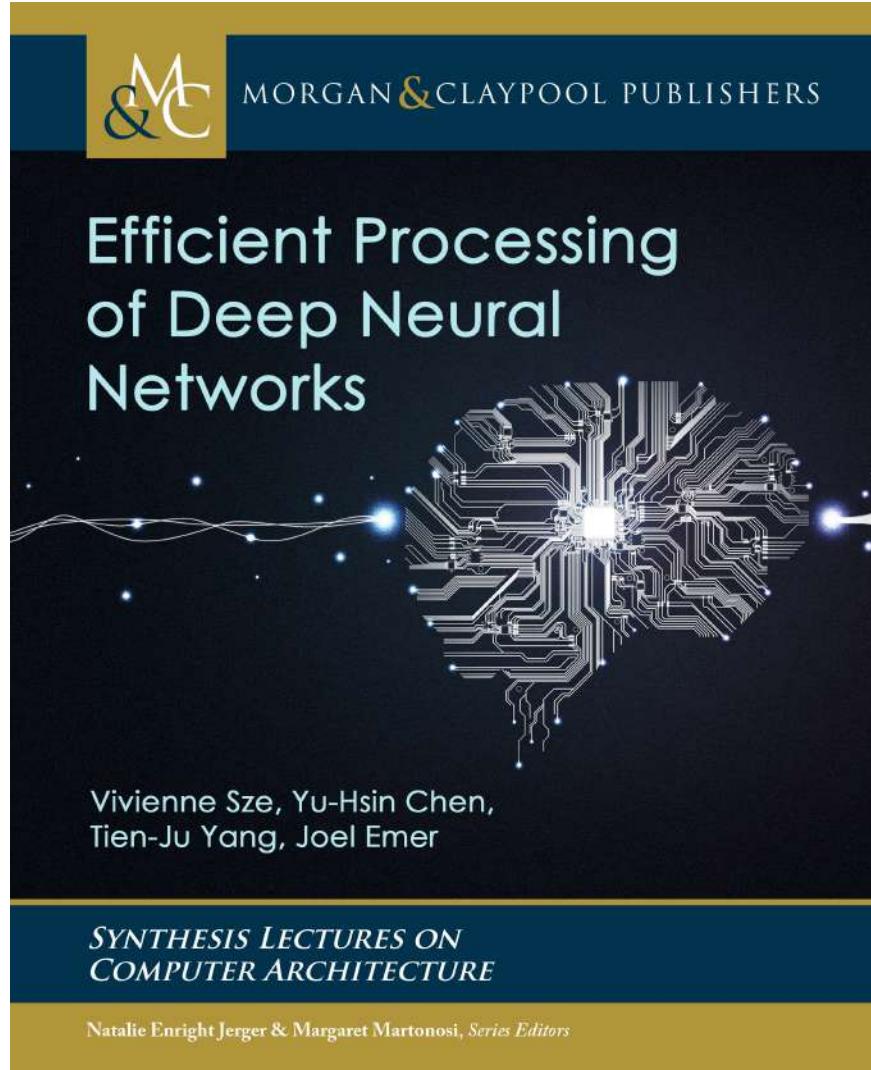
**Game Play**



**Medical**



# Book on Efficient Processing of DNNs



## ***Part I Understanding Deep Neural Networks***

*Introduction*

*Overview of Deep Neural Networks*

## ***Part II Design of Hardware for Processing DNNs***

*Key Metrics and Design Objectives*

*Kernel Computation*

*Designing DNN Accelerators*

*Operation Mapping on Specialized Hardware*

## ***Part III Co-Design of DNN Hardware and Algorithms***

*Reducing Precision*

*Exploiting Sparsity*

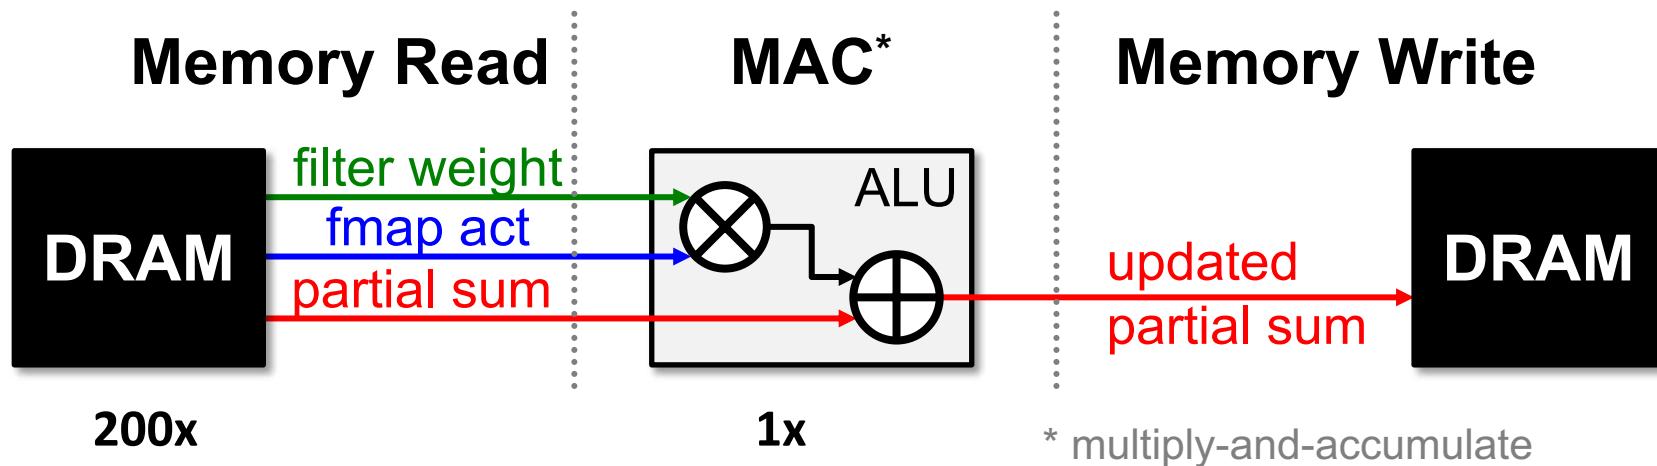
*Designing Efficient DNN Models*

*Advanced Technologies*

<https://tinyurl.com/EfficientDNNBook>

# Properties We Can Leverage

- Operations exhibit **high parallelism**  
→ **high throughput** possible
- Memory Access is the Bottleneck

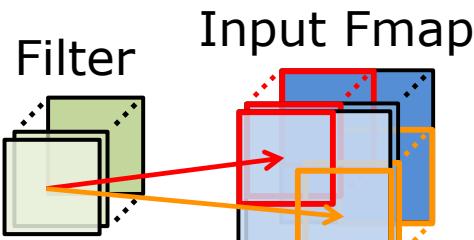


Worst Case: all memory R/W are **DRAM** accesses

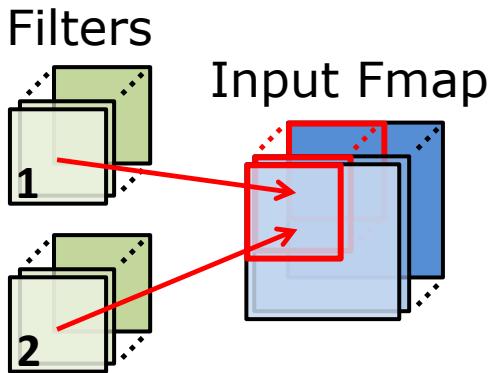
- Example: AlexNet has **724M** MACs  
→ **2896M** DRAM accesses required

# Properties We Can Leverage

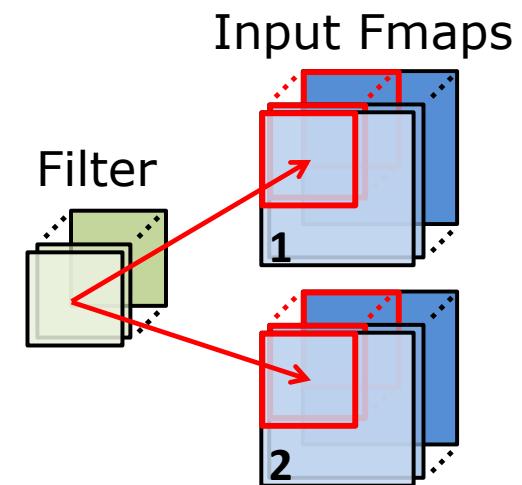
- Operations exhibit **high parallelism**  
→ high throughput possible
- Input data reuse** opportunities (**up to 500x**)



**Convolutional Reuse**  
(Activations, Weights)  
CONV layers only  
(sliding window)

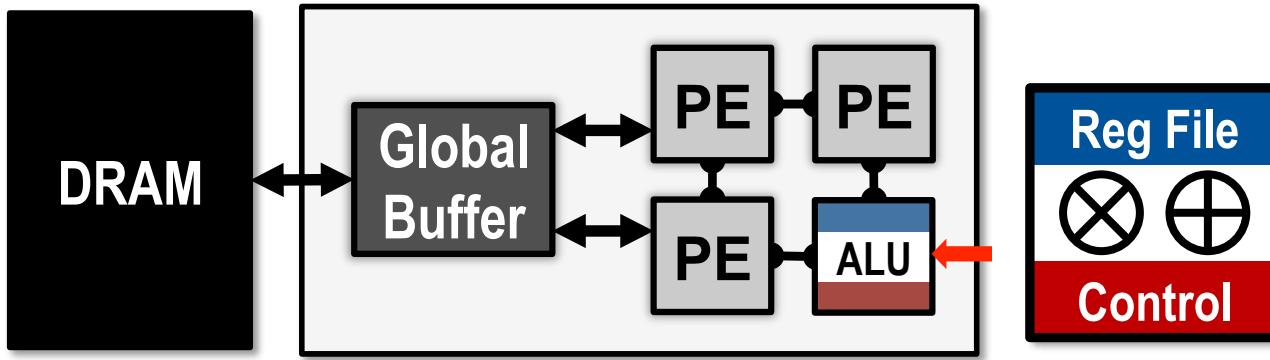


**Fmap Reuse**  
(Activations)  
CONV and FC layers

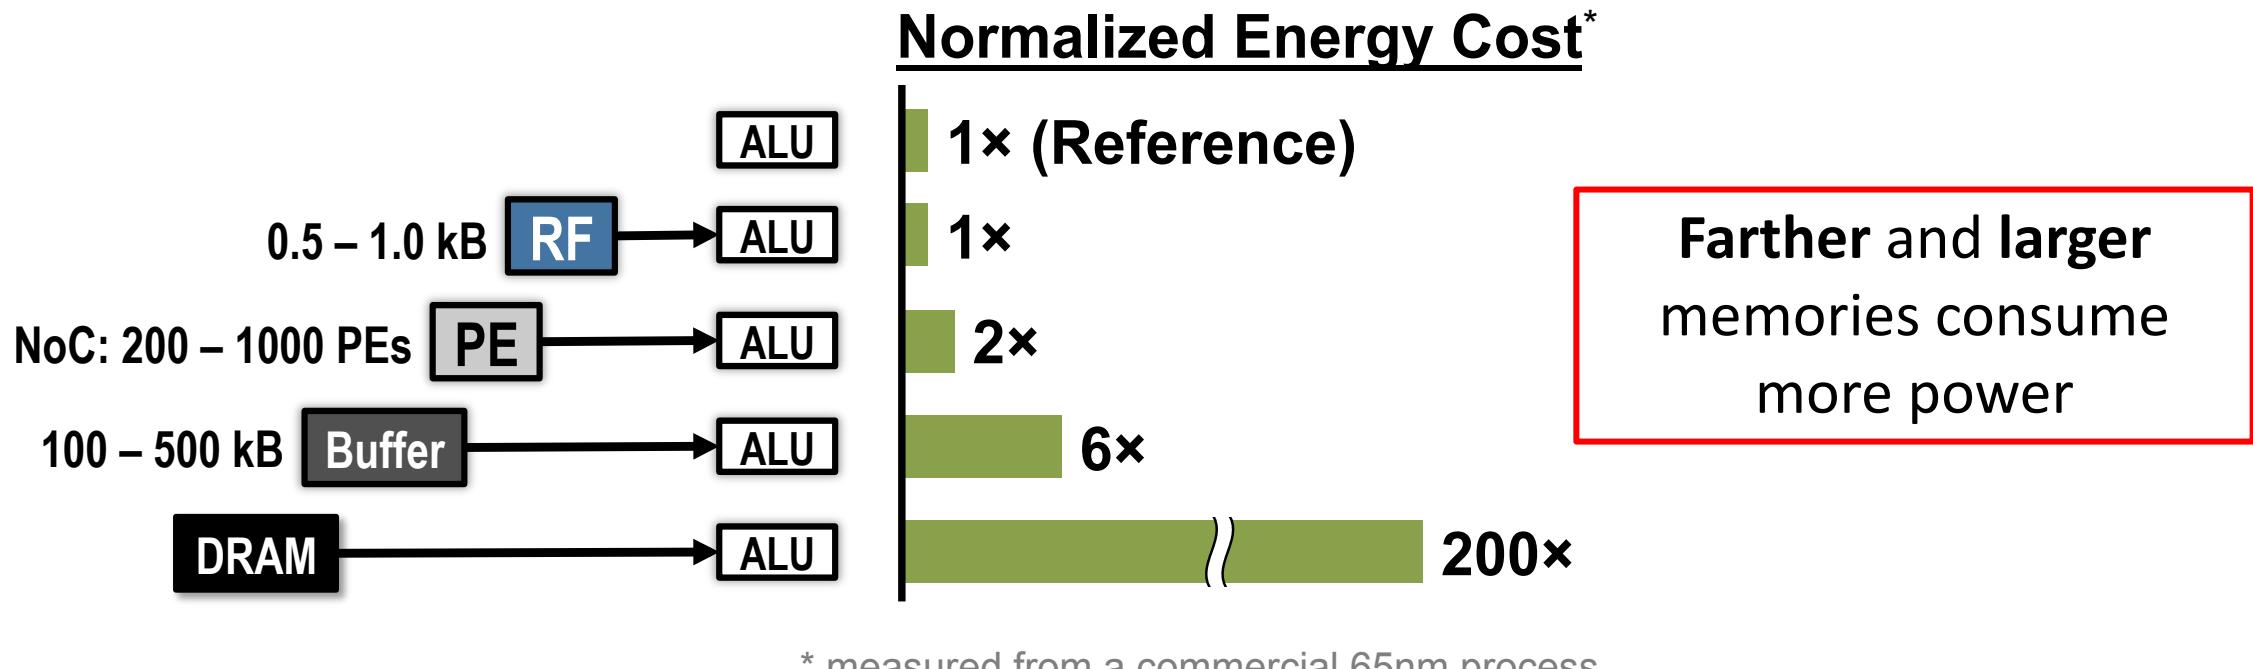


**Filter Reuse**  
(Weights)  
CONV and FC layers  
(batch size > 1)

# Exploit Data Reuse at Low-Cost Memories

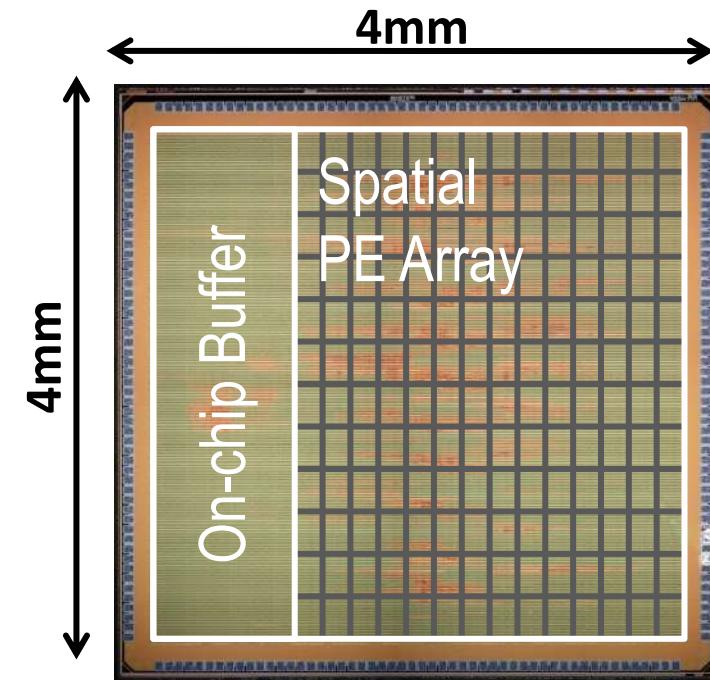
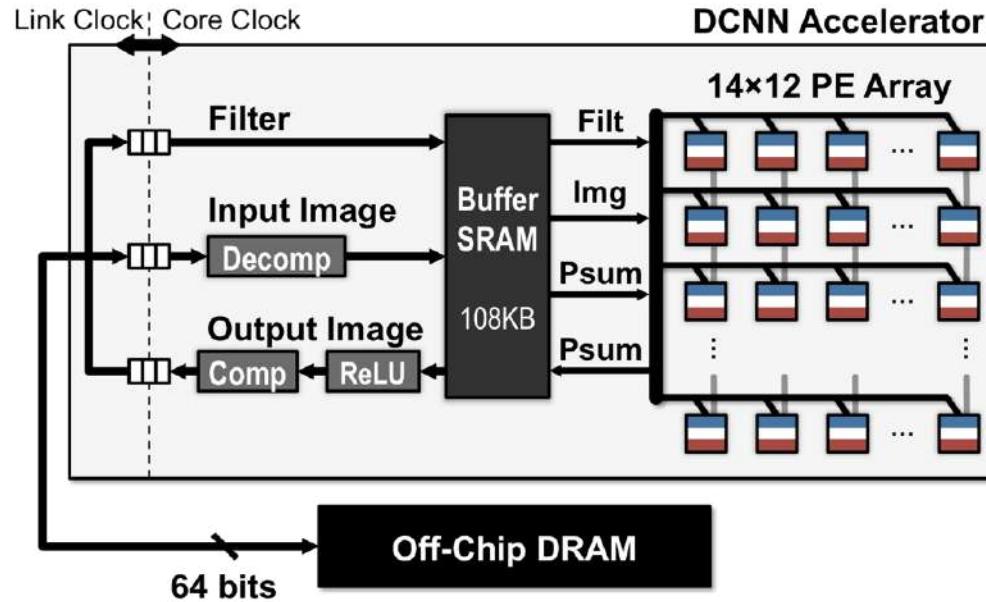


Specialized hardware with small (< 1kB) low cost memory near compute



# Energy-Efficient Dataflow for Deep Neural Networks

## Eyeriss: Row-Stationary Dataflow



[Chen, ISSCC 2016]

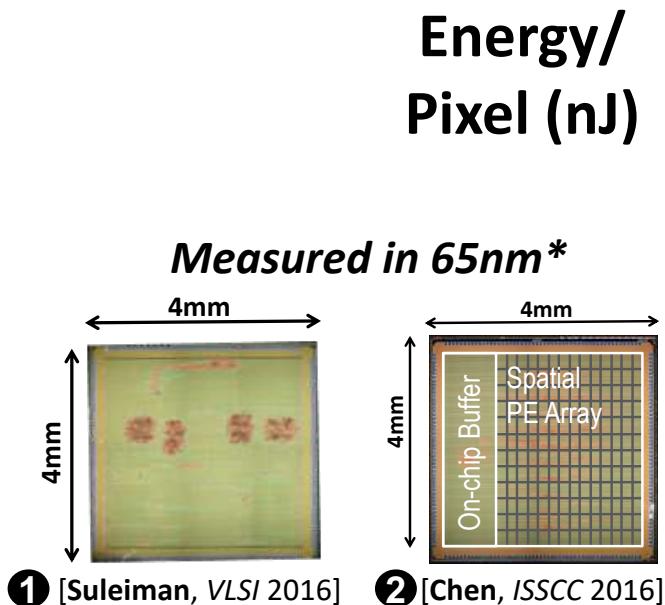
*Exploits data reuse for 100x reduction in memory accesses from global buffer and 1400x reduction in memory accesses from off-chip DRAM*

Overall >10x energy reduction compared to a mobile GPU (Nvidia TK1)

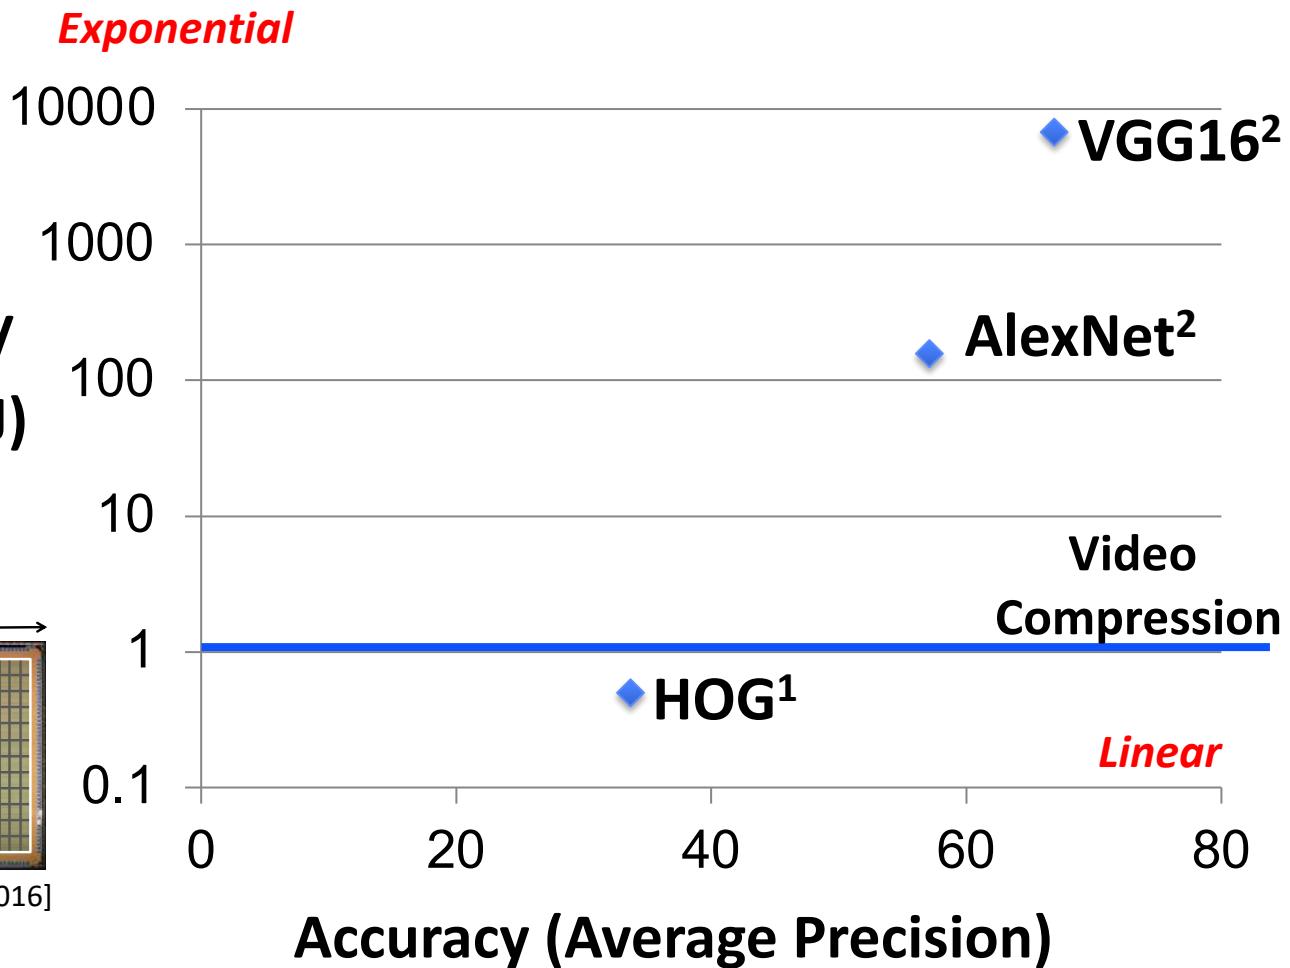
Eyeriss Project Website: <http://eyeriss.mit.edu>

Results for AlexNet

# Features: Energy vs. Accuracy



\* Only feature extraction. Does not include data, classification energy, augmentation and ensemble, etc.

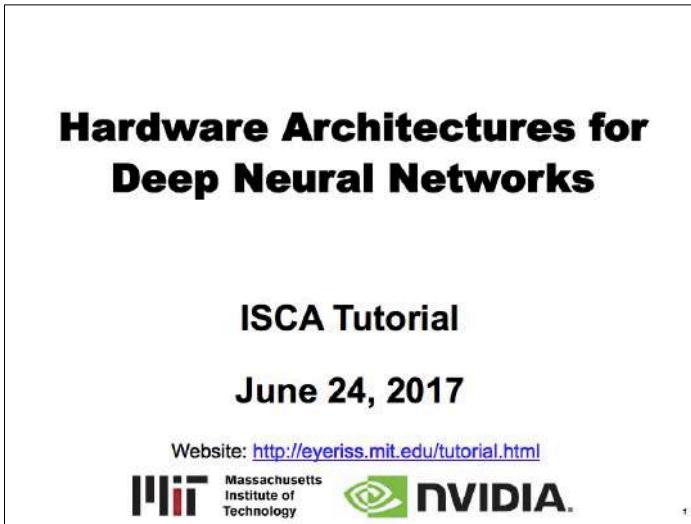


*Measured in on VOC 2007 Dataset*

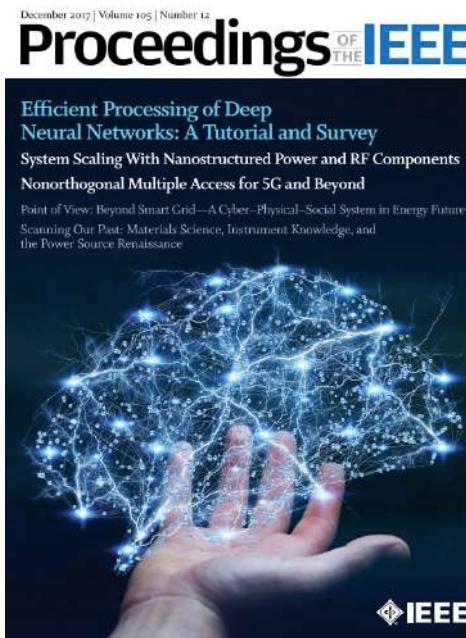
1. DPM v5 [Girshick, 2012]
2. Fast R-CNN [Girshick, CVPR 2015]

# Energy-Efficient Processing of DNNs

A significant amount of algorithm and hardware research on energy-efficient processing of DNNs



<http://eyeriss.mit.edu/tutorial.html>



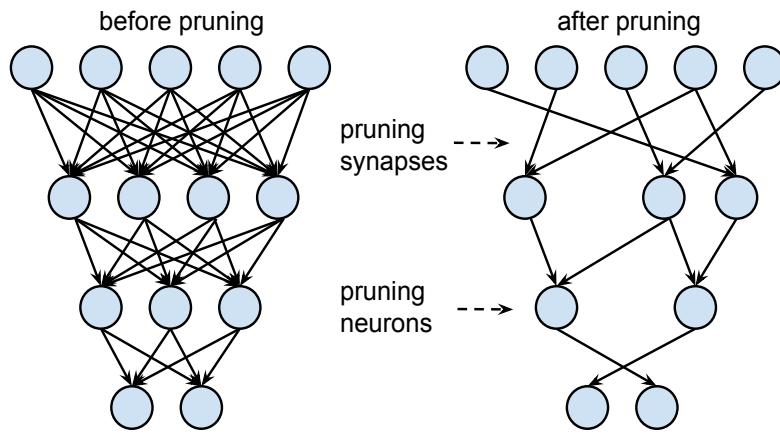
V. Sze, Y.-H. Chen,  
T.-J. Yang, J. Emer,  
*"Efficient Processing of  
Deep Neural Networks:  
A Tutorial and Survey,"*  
Proceedings of the IEEE,  
Dec. 2017

We identified various limitations to existing approaches

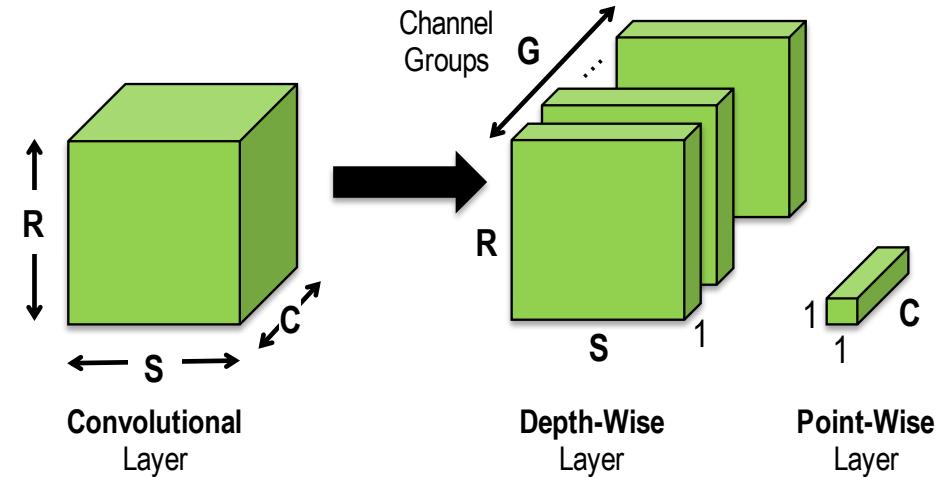
# Design of Efficient DNN Algorithms

Popular efficient DNN algorithm approaches

## Network Pruning



## Efficient Network Architectures



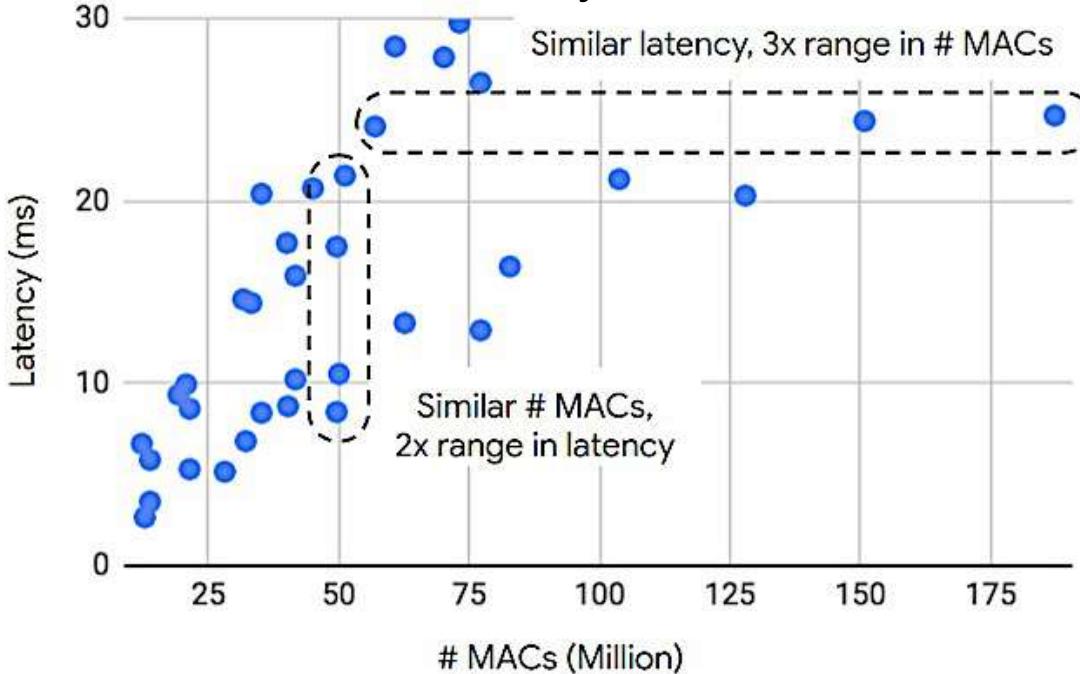
Examples: SqueezeNet, MobileNet

*... also reduced precision*

- Focus on reducing **number of MACs and weights**
- **Does it translate to energy savings and reduced latency?**

# Number of MACs and Weights are Not Good Proxies

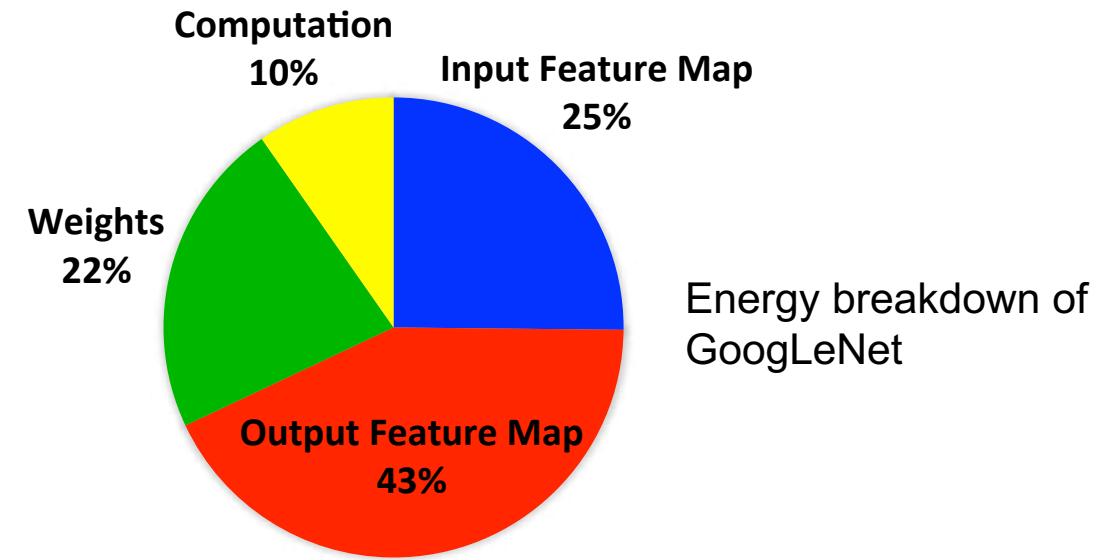
# of operations (MACs) does not approximate latency well



Source: Google

(<https://ai.googleblog.com/2018/04/introducing-cvpr-2018-on-device-visual.html>)

# of weights **alone** is not a good metric for energy  
(**All data types** should be considered)



<https://energyestimation.mit.edu/>

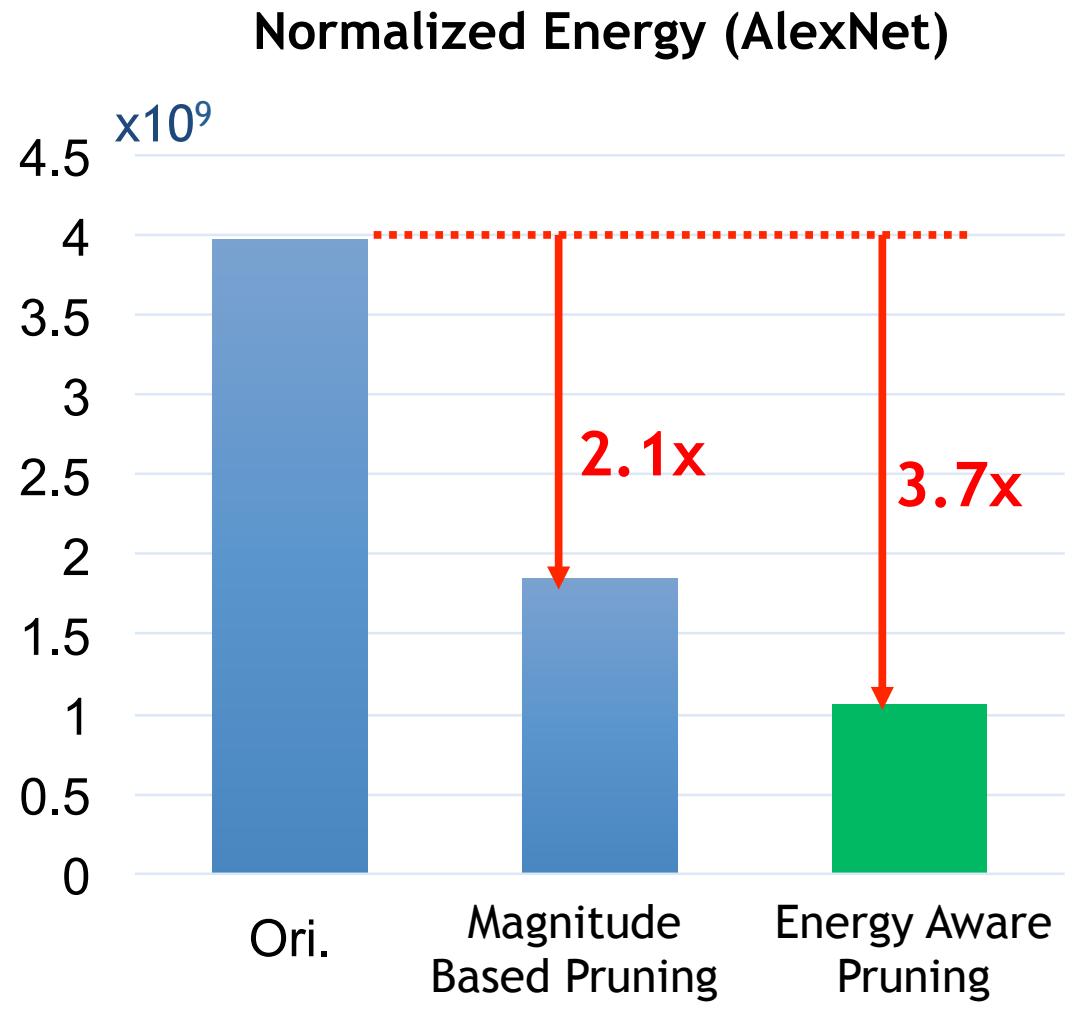
[Yang, CVPR 2017]

# Energy-Aware Pruning

**Directly target energy**  
and incorporate it into the  
optimization of DNNs to provide  
greater energy savings

- Sort layers based on energy and prune layers that consume the most energy first
- **Energy-aware pruning** reduces AlexNet energy by **3.7x** w/ similar accuracy
- Outperforms magnitude-based pruning by **1.7x**

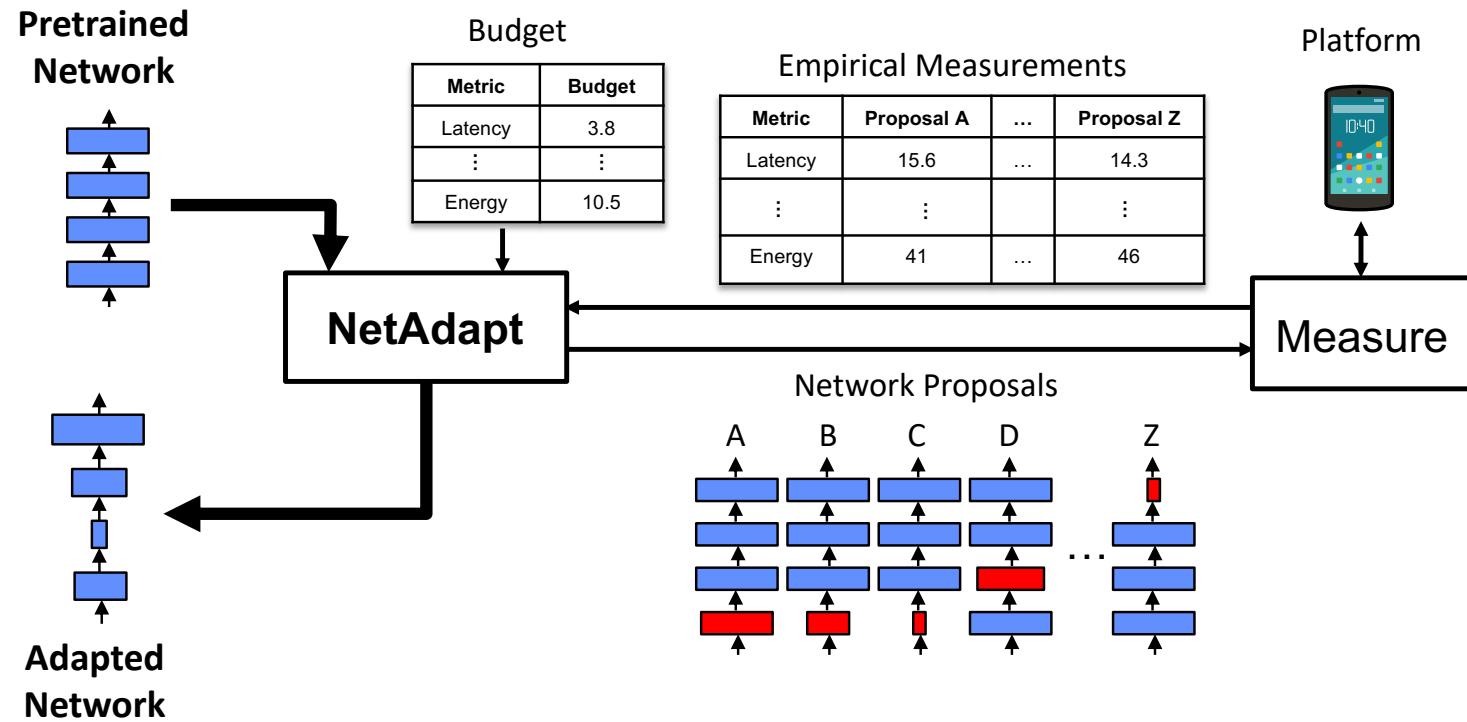
[Yang, CVPR 2017]



Pruned models available at  
<http://eyeriss.mit.edu/energy.html>

# NetAdapt: Platform-Aware DNN Adaptation

- **Automatically adapt DNN** to a mobile platform to reach a target latency or energy budget
- Use **empirical measurements** to guide optimization (avoid modeling of tool chain or platform architecture)
- **Few hyperparameters** to reduce tuning effort
- **>1.7x speed up** on MobileNet w/ similar accuracy

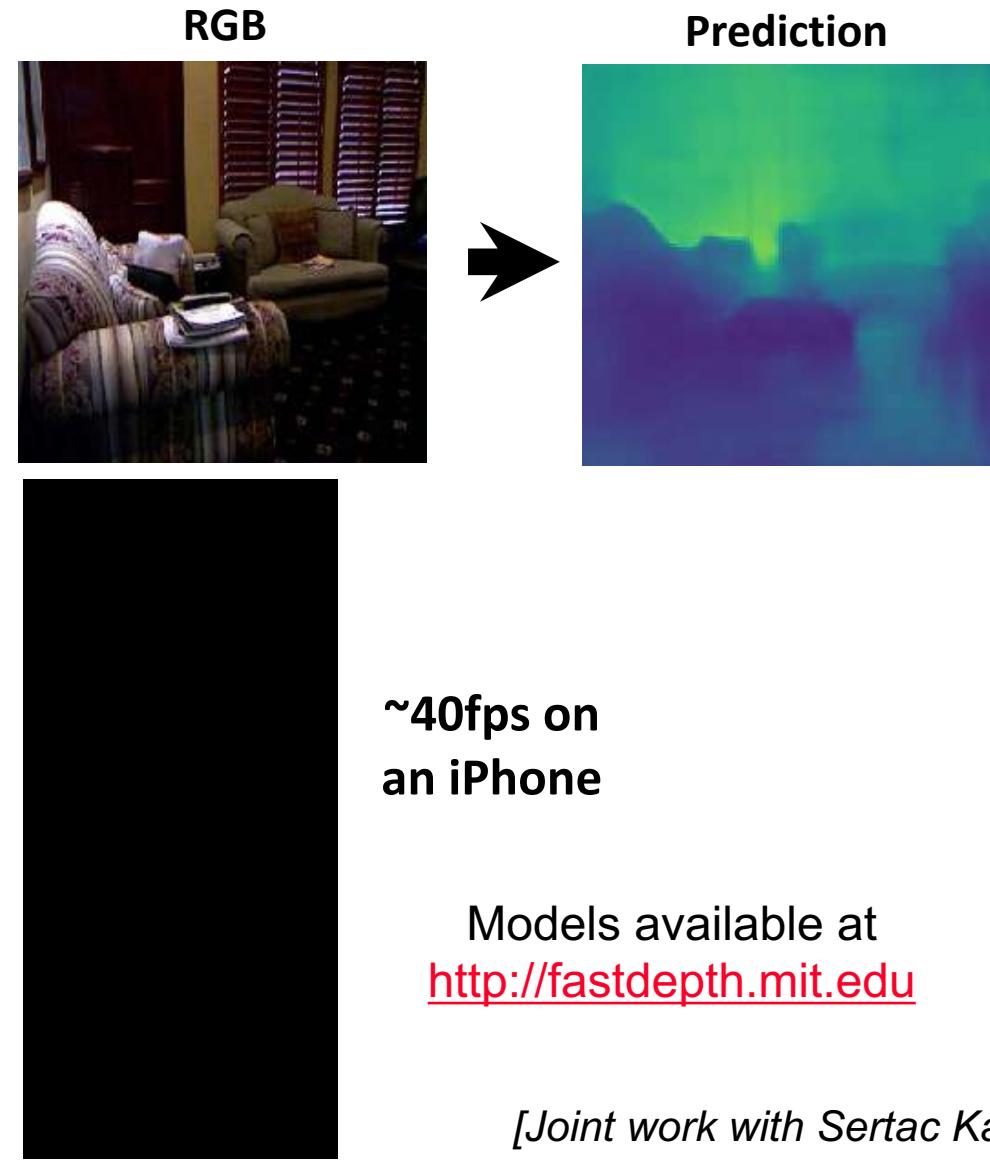
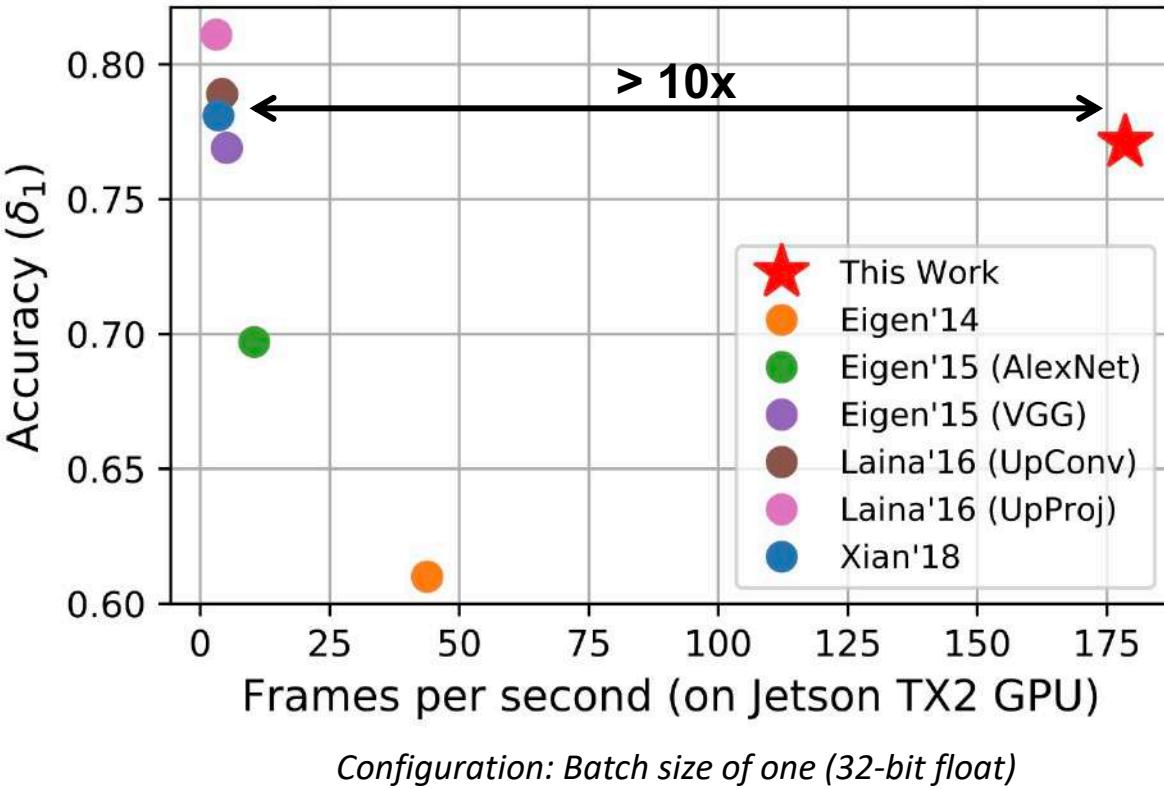


Code available at  
<http://netadapt.mit.edu>

*[In collaboration with Google's Mobile Vision Team]*

# FastDepth: Fast Monocular Depth Estimation

Depth estimation from a single RGB image desirable, due to the relatively low cost and size of monocular cameras.



# NetAdapt v2: Reduce Adaption Time

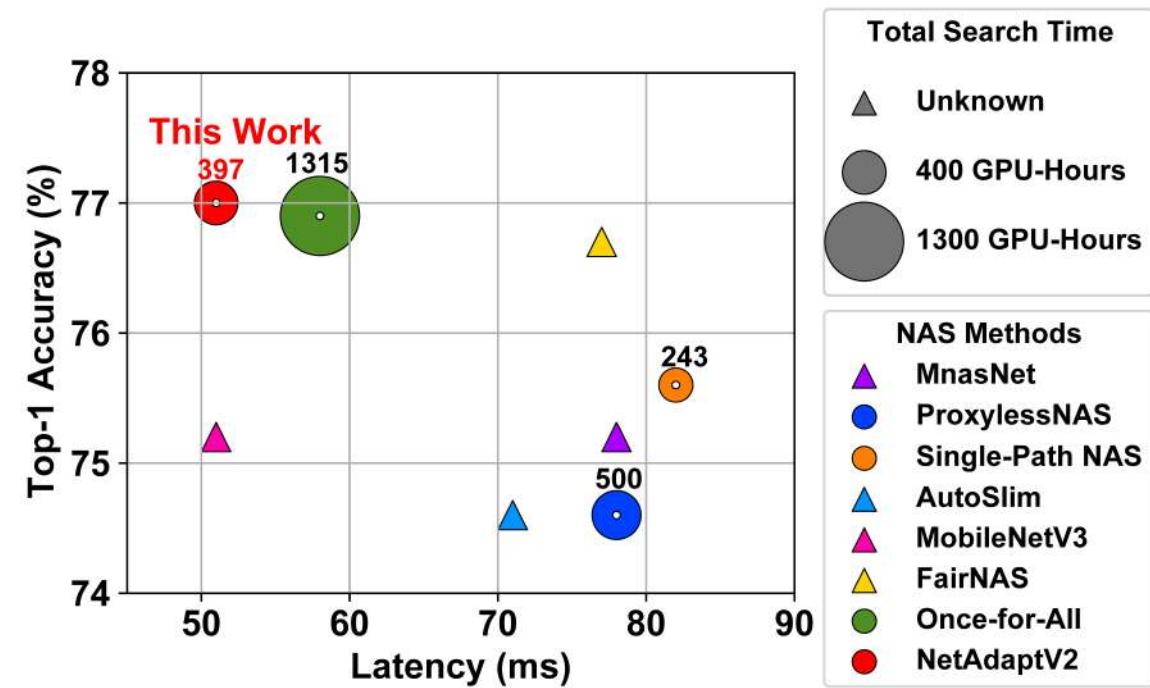
Reduce time to find efficient DNN that adapts to hardware by up to 5.8x

## Typical Steps in Neural Architecture Search (NAS):

- 1) Train super-network (search space of DNNs)
- 2) Sample and evaluate different DNNs
- 3) Fine tune the final DNN

## Contributions

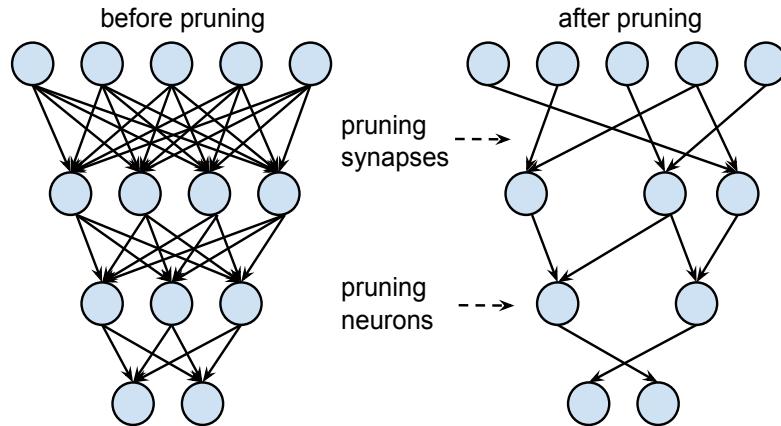
- ***Ordered dropout***: train multiple DNNs in *single* forward pass (reduce step 1)
- ***Channel-level bypass***: merge layer depth and channel width into a *single* search dimension (reduce step 2)
- ***Multi-layer coordinate descent optimizer***: consider joint effect of multiple layers (reduce step 2 & support non-differentiable metrics, e.g., latency)



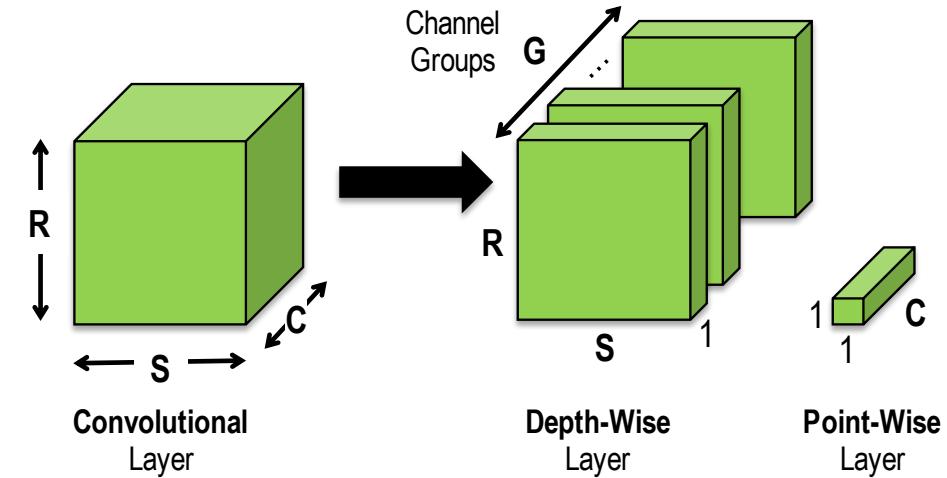
More info at <http://netadapt.mit.edu>

# Many Efficient DNN Design Approaches

## Network Pruning



## Compact Network Architectures



## Reduce Precision

**32-bit float** 1|0|1|0|0|1|0|1|0|0|0|0|0|0|0|0|0|0|0|0|0|0|0|0|0|0|0|1|0|0

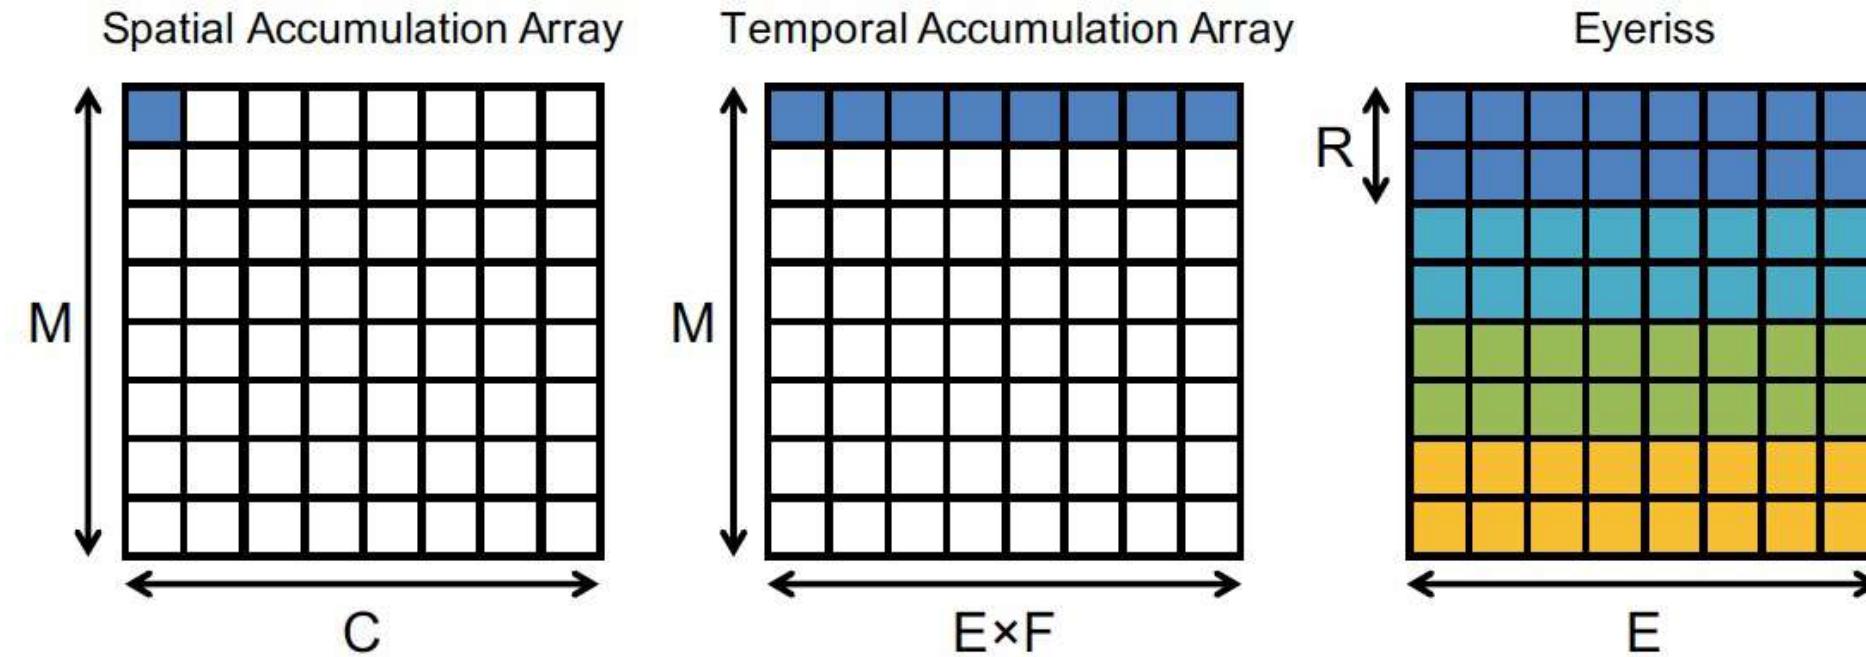
**8-bit fixed** 0|1|1|0|0|1|1|0

**Binary** 0

No guarantee that DNN algorithm designer will use a given approach.  
**Need flexible hardware!**

# Need Flexible Dataflow & Mapping

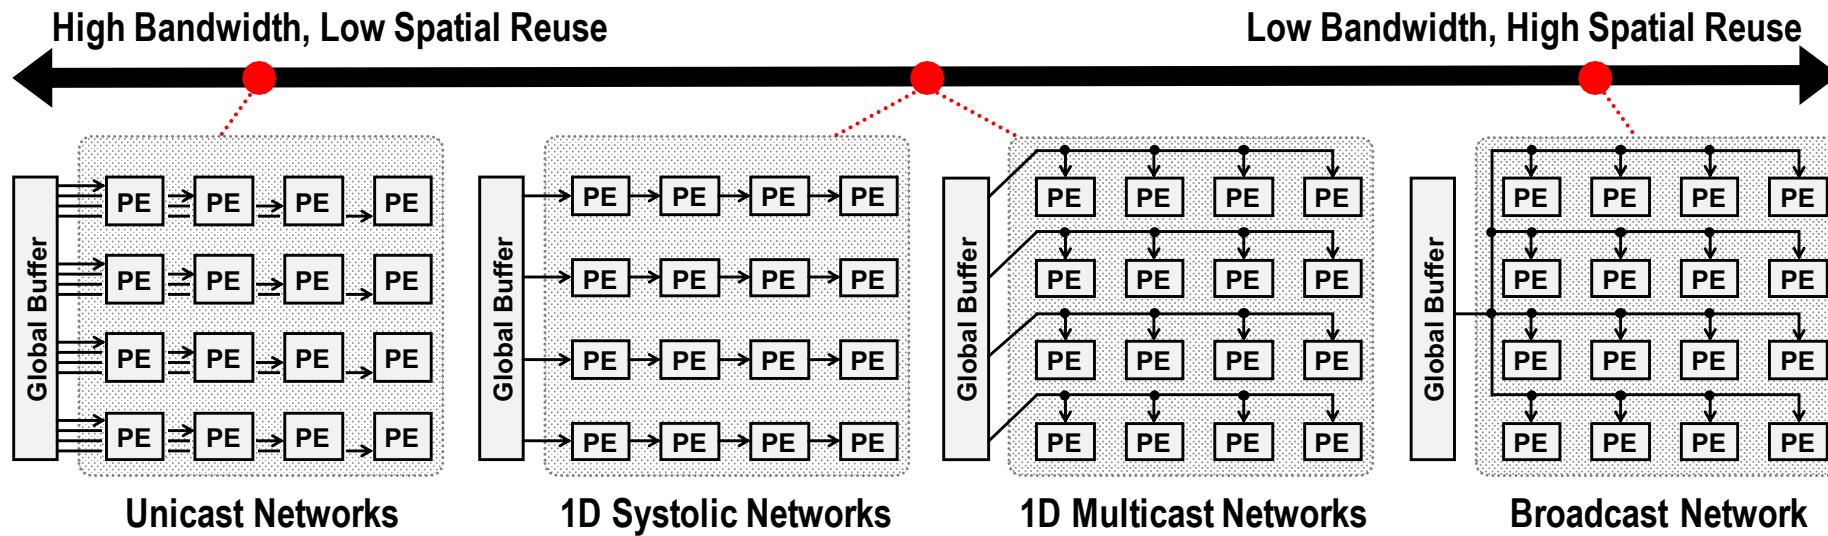
- Use flexible dataflow (**Row Stationary**) to exploit reuse in any dimension of DNN to increase energy efficiency and array utilization



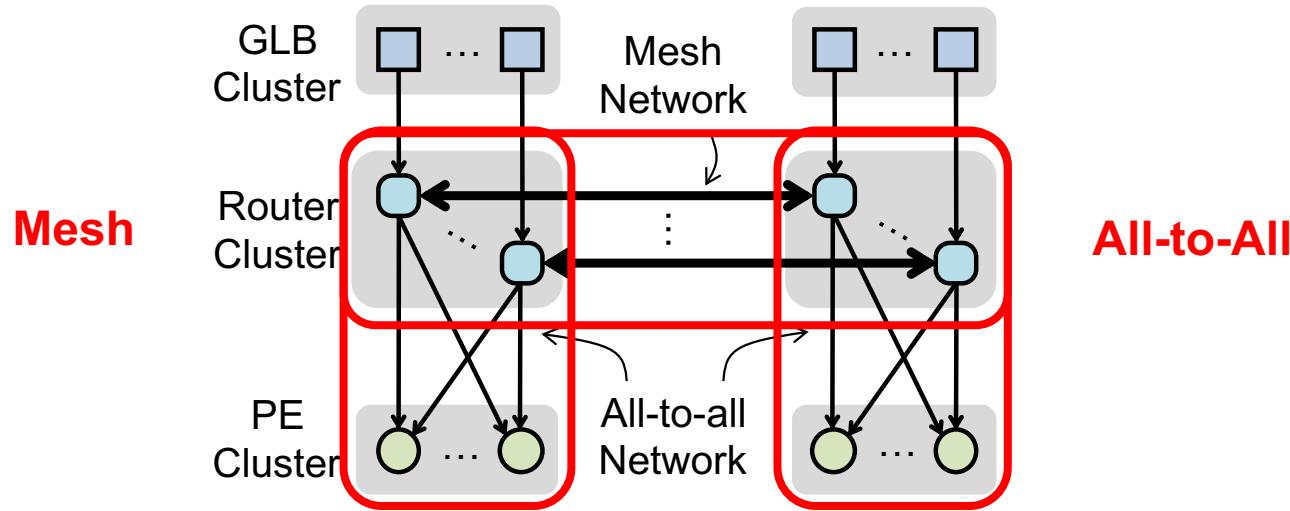
Example: Depth-wise layer

# Need Flexible NoC for Varying Reuse

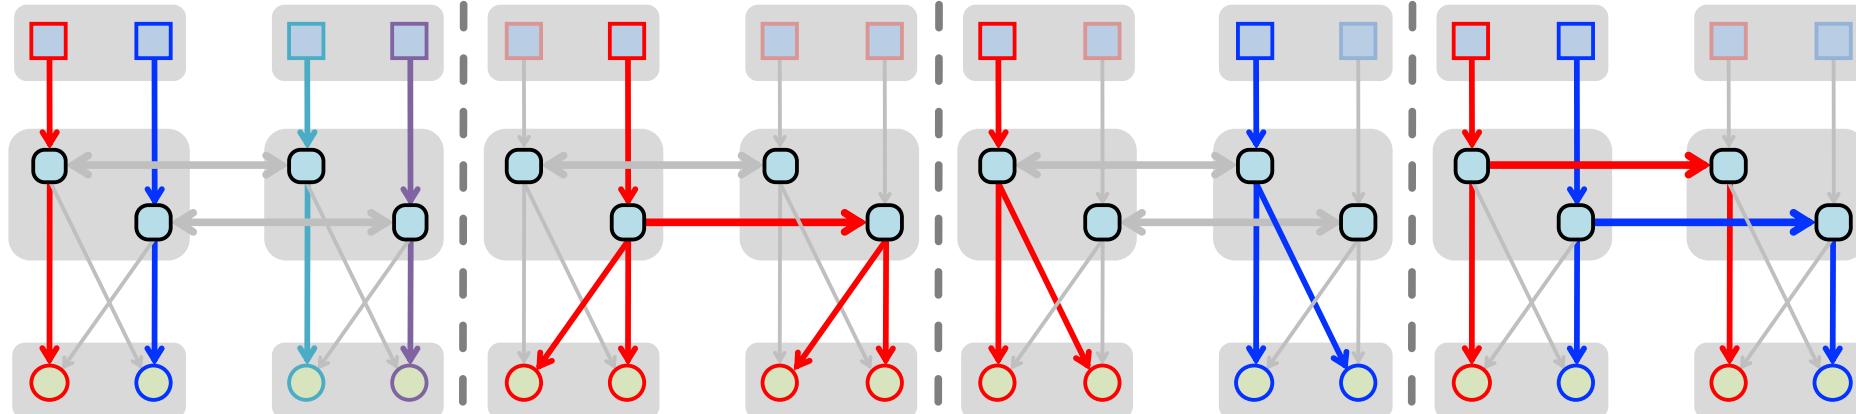
- When reuse available, need **multicast** to exploit spatial data reuse for energy efficiency and high array utilization
- When reuse not available, need **unicast** for high BW for weights for FC and weights & activations for high PE utilization
- An **all-to-all** satisfies above but too expensive and not scalable



# Hierarchical Mesh



High Bandwidth      High Reuse      Grouped Multicast      Interleaved Multicast



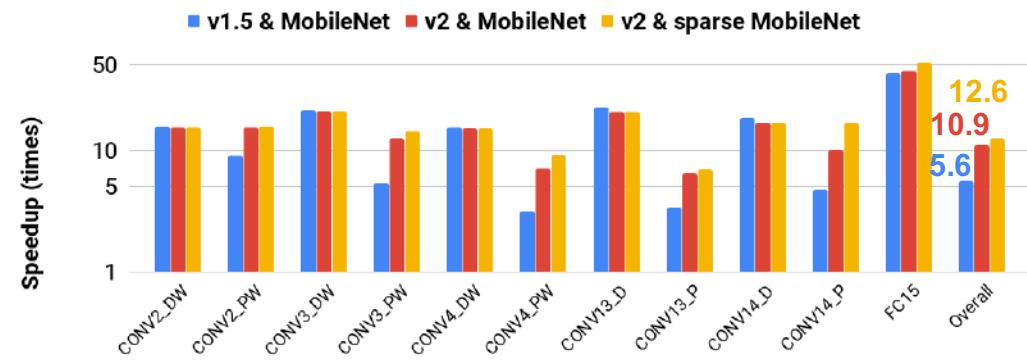
# Eyeriss v2: Balancing Flexibility and Efficiency

- Uses a flexible hierarchical mesh on-chip network to efficiently support

- Wide range of filter shapes
- Different layers
- Wide range of sparsity

- Scalable architecture

Over an order of magnitude faster and more energy efficient than Eyeriss v1



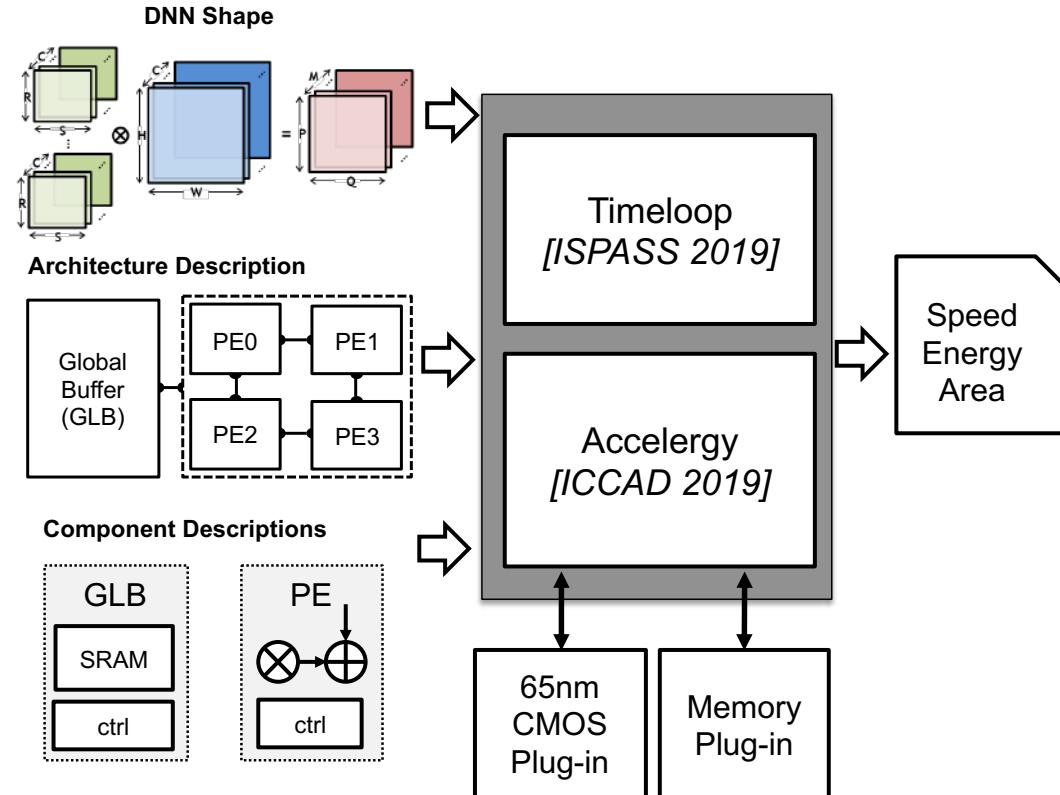
*Speed up over Eyeriss v1 scales with number of PEs*

# of PEs	256	1024	16384
AlexNet	17.9x	71.5x	1086.7x
GoogLeNet	10.4x	37.8x	448.8x
MobileNet	15.7x	57.9x	873.0x

[Joint work with Joel Emer]

# DNN Accelerator Evaluation Tools

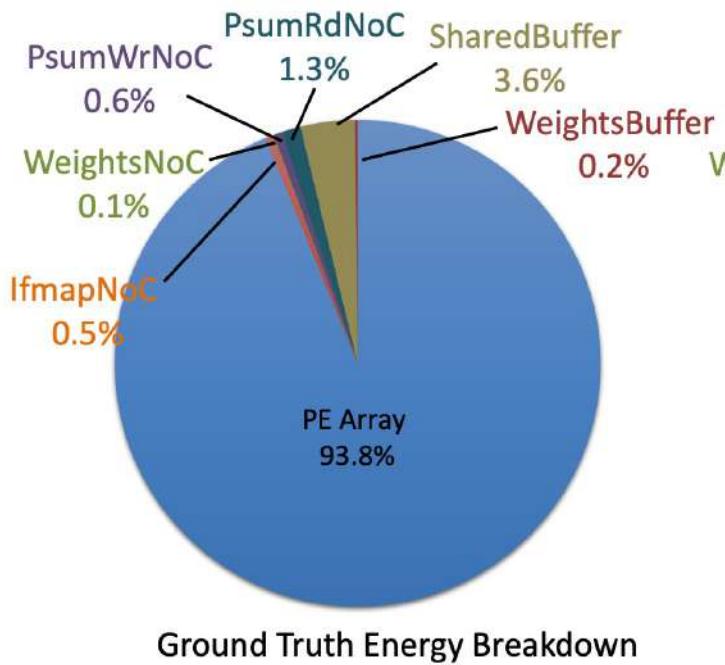
- Require systematic way to
  - Evaluate and compare DNN accelerators
  - Rapidly explore design space
- **Accelergy** [wu, ICCAD 2019]
  - Early-stage estimation tool at the architecture level
    - Estimate energy based on architecture level components (e.g., # of PEs, memory size, on-chip network)
  - Evaluate architecture level impact of emerging devices
    - Plug-ins for different technologies
- **Timeloop** [Parashar, ISPASS 2019]
  - DNN mapping tool
  - Performance Simulator → Action counts
- Bridge architecture, circuit, and device research



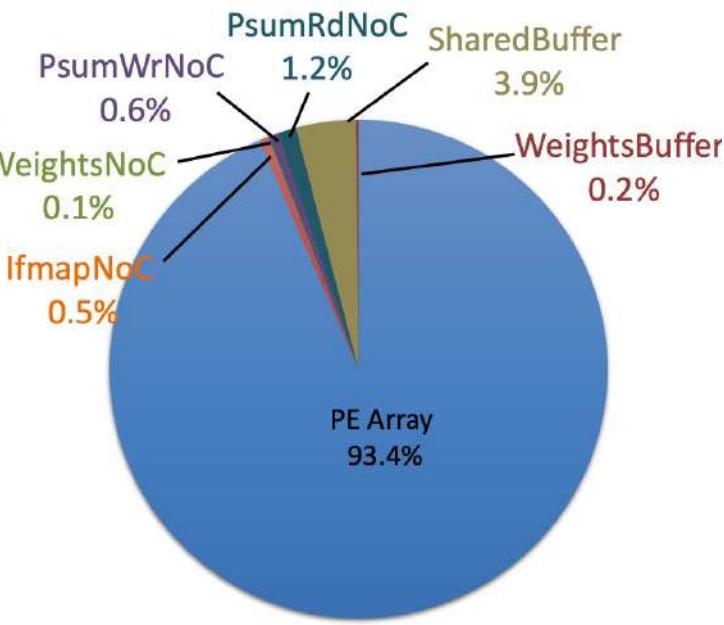
Open-source code available at:  
<http://accelergy.mit.edu>

# Accelergy Estimation Validation

- Validation on Eyeriss [Chen, ISSCC 2016]
  - Achieves 95% accuracy compared to post-layout simulations
  - Can accurately captures energy breakdown at different granularities



Ground Truth Energy Breakdown



Accelergy Energy Breakdown

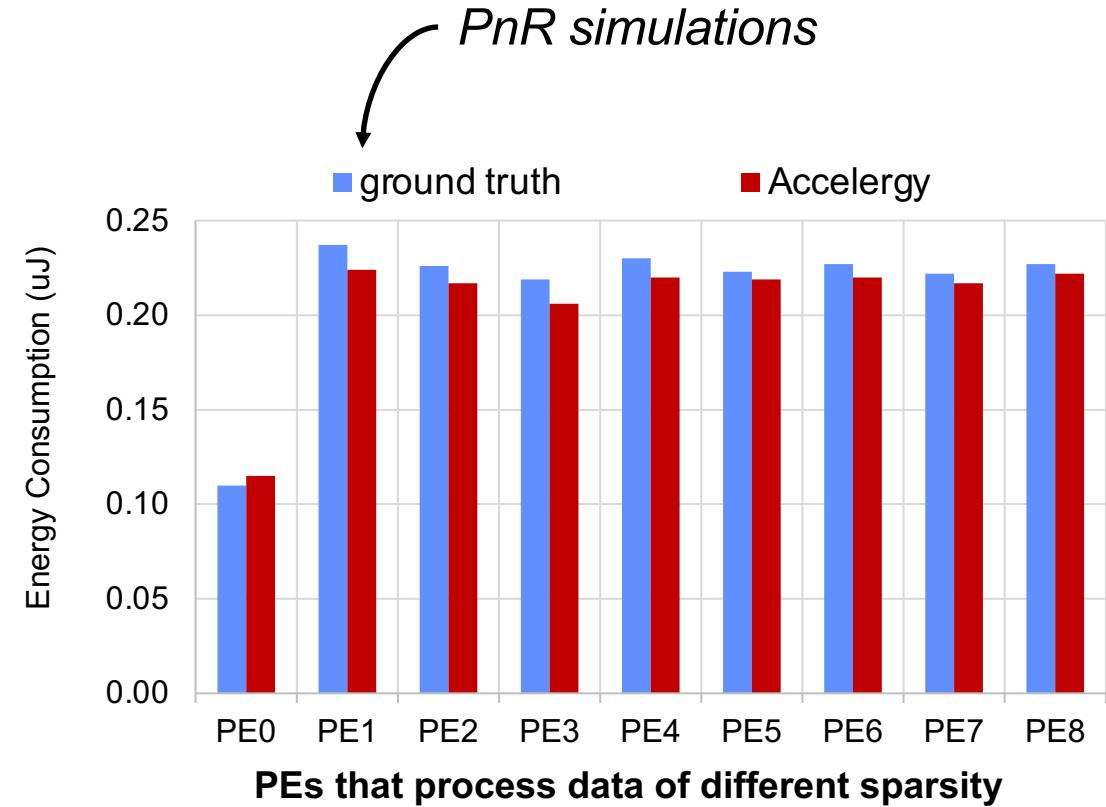
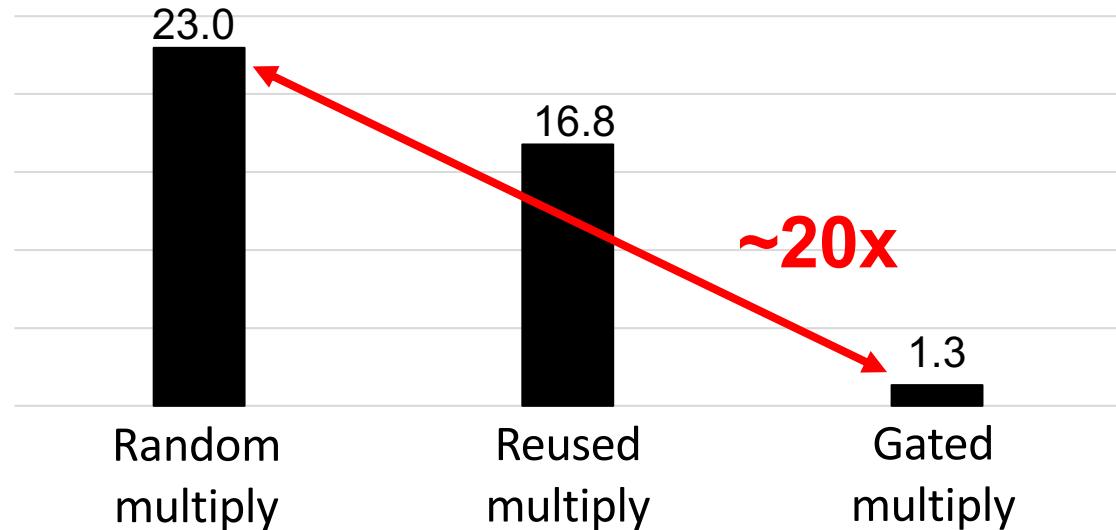
\*Total energy might not add up to exact 100.0% due to rounding

Open-source code available at: <http://accelergy.mit.edu>

# Plug-ins for Fine-Grain Action Energy Estimation

- External energy/area models that accurately reflect the properties of a macro
  - e.g., multiplier with zero-gating

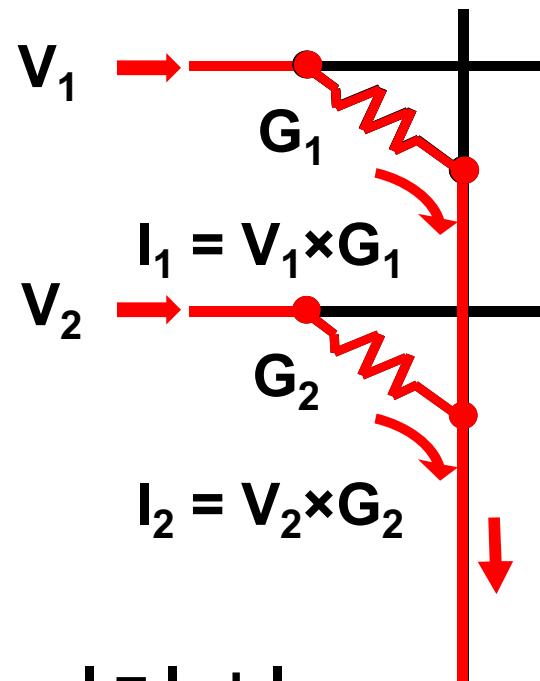
Energy characterizations of the zero-gated multiplier  
(normalized to idle)



With the characterization provided in the plug-in,  
we can capture the energy savings for sparse workloads

# In-Memory Computing (IMC\*)

Activation is input voltage ( $V_i$ )  
Weight is resistor conductance ( $G_i$ )



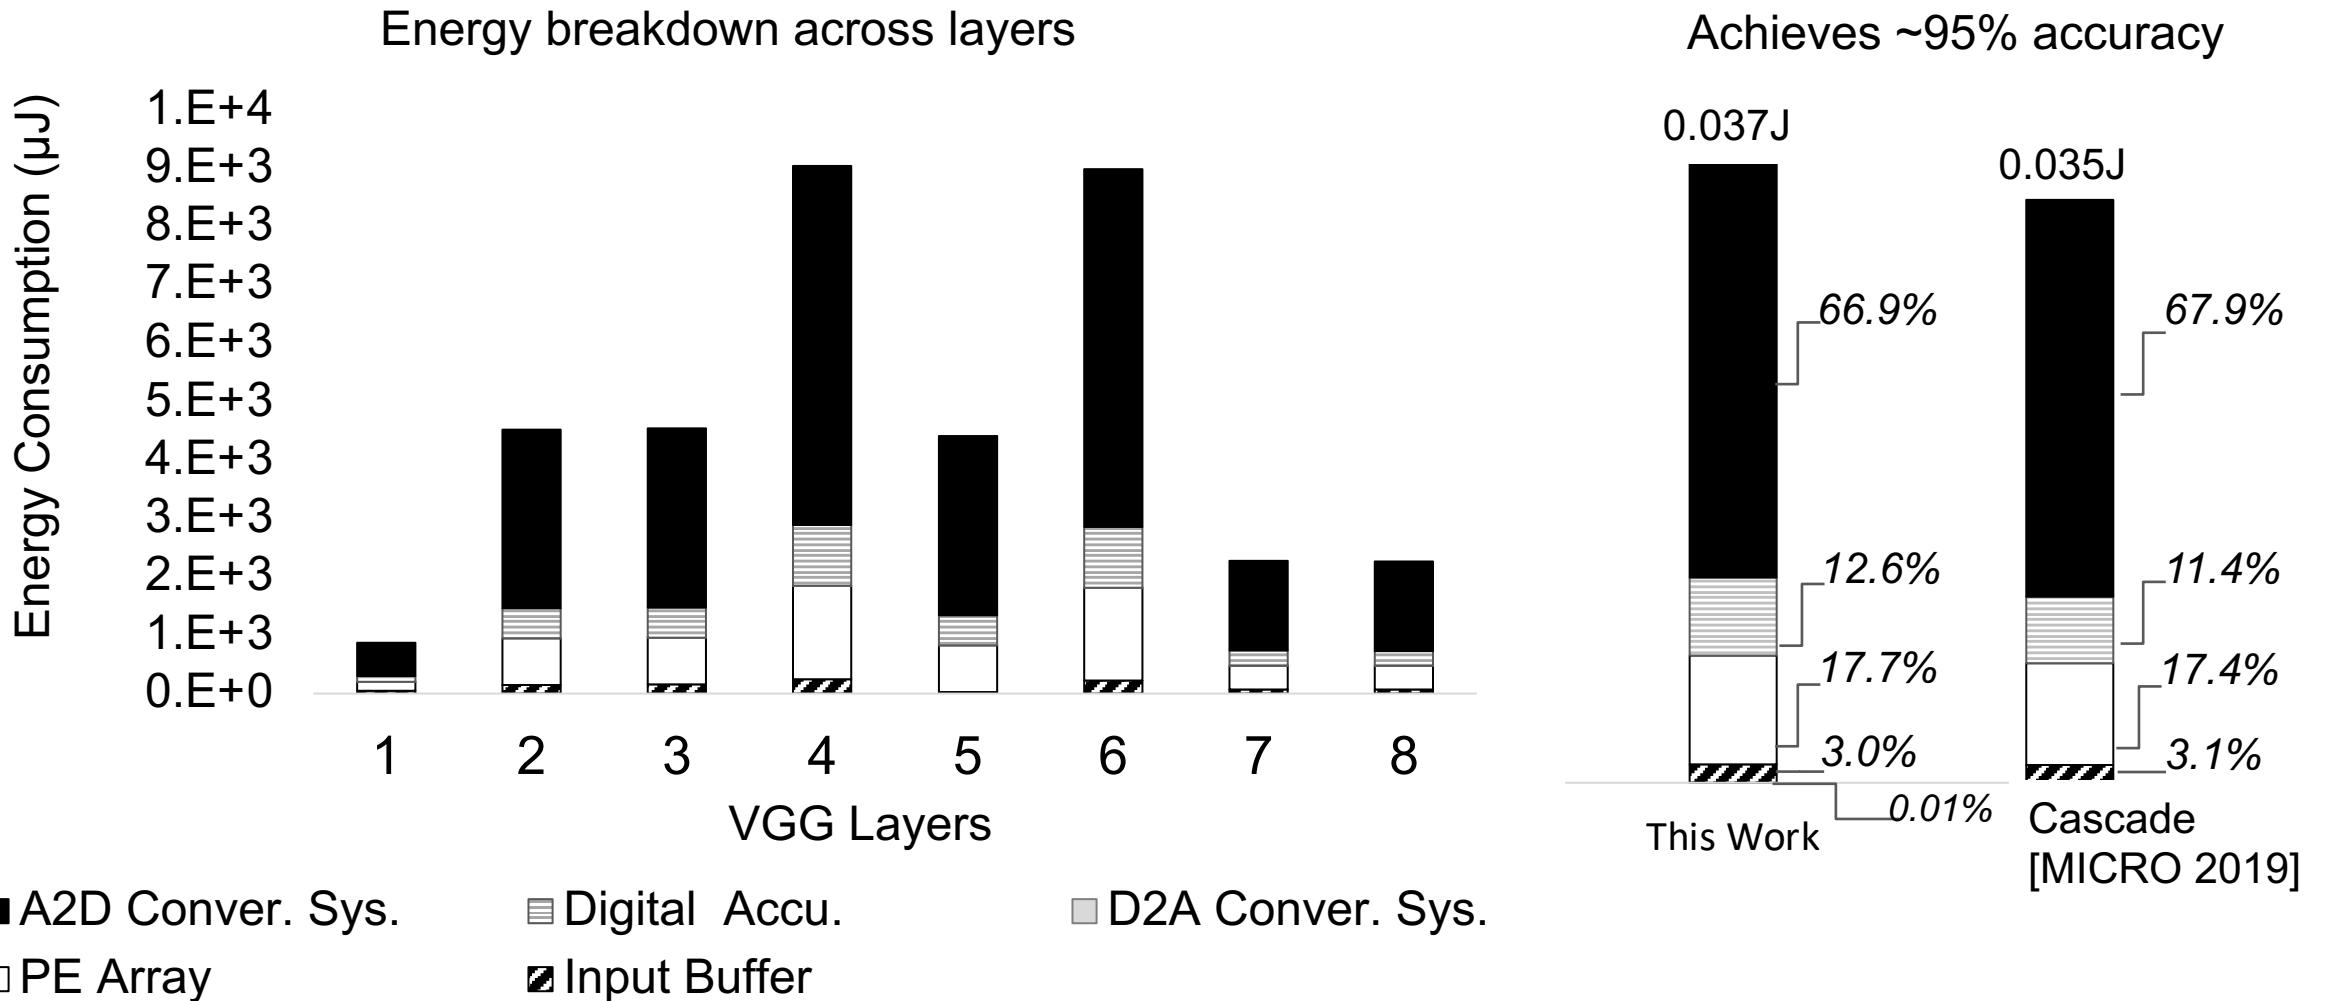
Psum  
is output  
current

Image Source: [Shafiee, ISCA 2016]

- Reduce data movement by **moving compute into memory**
- Compute MAC with memory storage element
- **Analog Compute**
  - Activations, weights and/or partial sums are encoded with analog voltage, current, or resistance
  - Increased sensitivity to circuit non-idealities
  - A/D and D/A circuits to interface with digital domain
- Leverage **emerging memory device technology**

# Accelergy for IMC

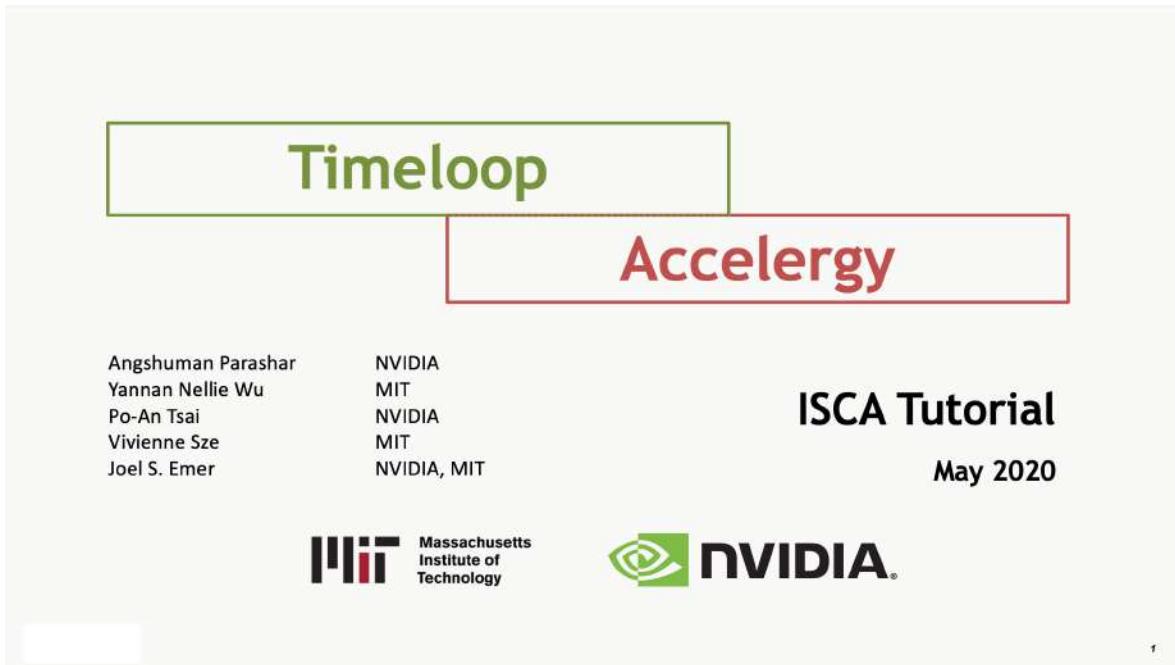
Open-source code available at:  
<http://accelergy.mit.edu>



# Accelergy + Timeloop Tutorial

Tutorial material available at <http://accelergy.mit.edu/tutorial.html>

***Includes videos and hands-on exercises***

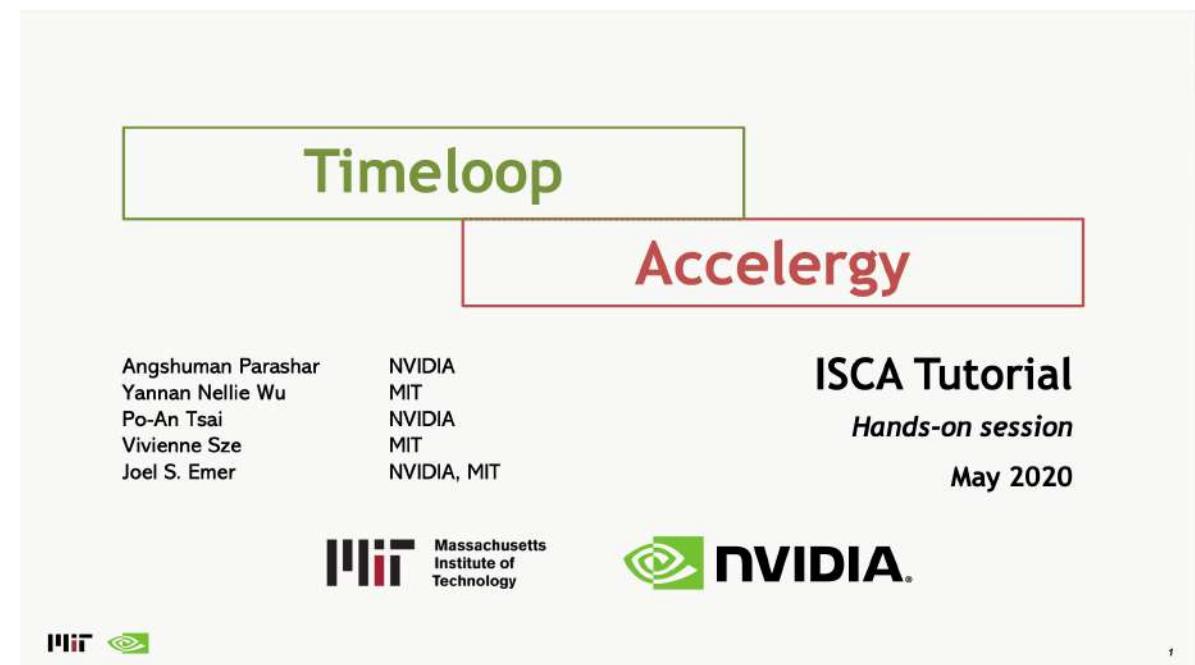


The slide features two main sections: "Timeloop" in a green box at the top left and "Accelergy" in a red box below it. Below these boxes, the text "ISCA Tutorial" is centered, followed by "May 2020". To the left, a table lists speakers and their affiliations: Angshuman Parashar (NVIDIA), Yannan Nellie Wu (MIT), Po-An Tsai (NVIDIA), Vivienne Sze (MIT), and Joel S. Emer (NVIDIA, MIT). Logos for MIT and NVIDIA are at the bottom.

Angshuman Parashar	NVIDIA
Yannan Nellie Wu	MIT
Po-An Tsai	NVIDIA
Vivienne Sze	MIT
Joel S. Emer	NVIDIA, MIT

**ISCA Tutorial**  
May 2020

MIT Massachusetts Institute of Technology NVIDIA



The slide features two main sections: "Timeloop" in a green box at the top left and "Accelergy" in a red box below it. Below these boxes, the text "ISCA Tutorial" is centered, followed by "Hands-on session" and "May 2020". To the left, a table lists speakers and their affiliations: Angshuman Parashar (NVIDIA), Yannan Nellie Wu (MIT), Po-An Tsai (NVIDIA), Vivienne Sze (MIT), and Joel S. Emer (NVIDIA, MIT). Logos for MIT and NVIDIA are at the bottom.

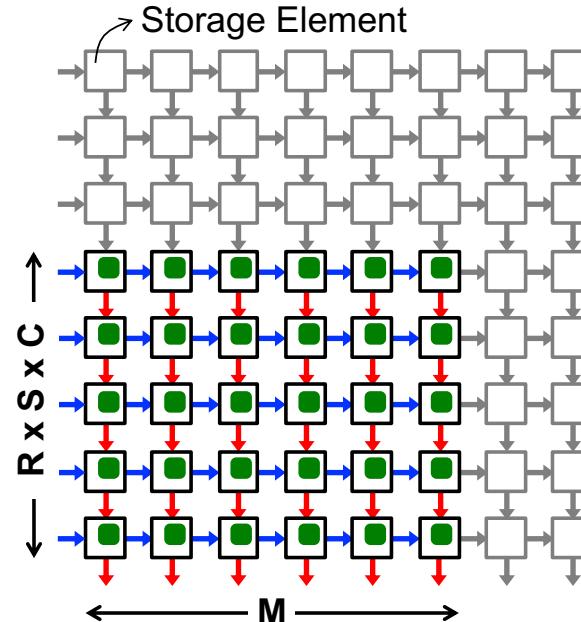
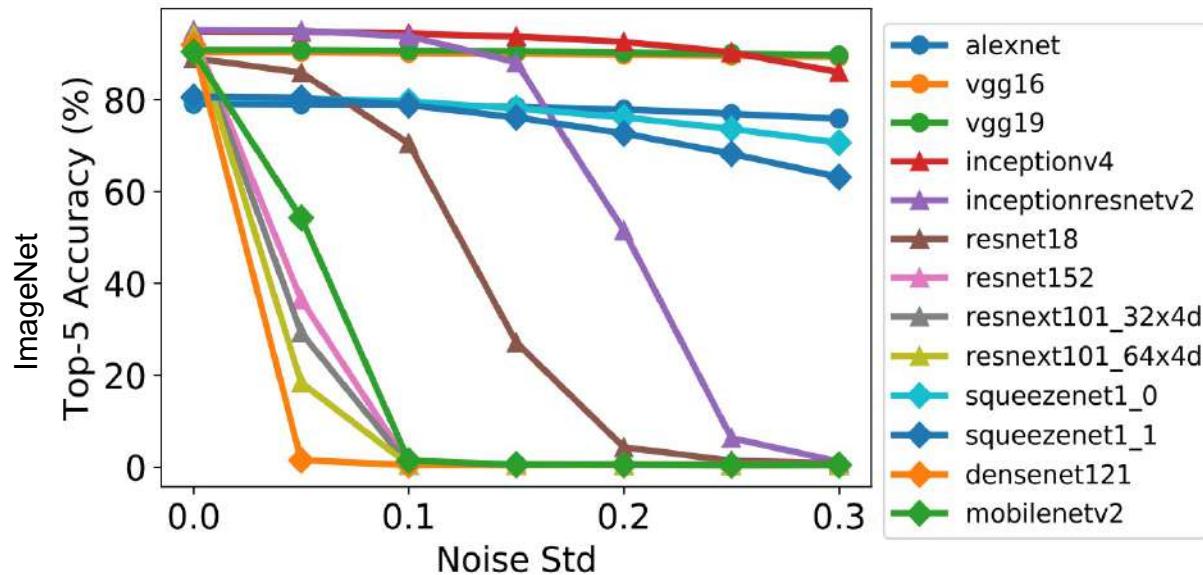
Angshuman Parashar	NVIDIA
Yannan Nellie Wu	MIT
Po-An Tsai	NVIDIA
Vivienne Sze	MIT
Joel S. Emer	NVIDIA, MIT

**ISCA Tutorial**  
*Hands-on session*  
May 2020

MIT Massachusetts Institute of Technology NVIDIA

# Designing DNNs for IMC

- Designing DNNs for IMC may differ from DNNs for digital processors
- Highest accuracy DNN on digital processor may be different on IMC
  - Accuracy drops based on robustness to non-idealities
- Reducing number of weights is less desirable
  - Since IMC is weight stationary, may be better to reduce number of activations
  - IMC tend to have larger arrays → fewer weights may lead to low utilization on IMC



# Book Chapter on In-Memory Computing

CHAPTER 10

253

## Advanced Technologies

As highlighted throughout the previous chapters, data movement dominates energy consumption. The energy is consumed both in the access to the memory as well as the transfer of the data. The associated physical factors also limit the bandwidth available to deliver data between memory and compute, and thus limits the throughput of the overall system. This is commonly referred to by computer architects as the "memory wall."<sup>1</sup>

To address the challenges associated with data movement, there have been various efforts to bring compute and memory closer together. Chapters 5 and 6 primarily focus on how to design spatial architectures that distribute the on-chip memory closer to the computation (e.g., scratch pad memory in the PE). This chapter will describe various other architectures that use *advanced memory, process, and fabrication technologies* to bring the compute and memory together.

First, we will describe efforts to bring the off-chip high-density memory (e.g., DRAM) closer to the computation. These approaches are often referred to as *processing near memory* or *near-data processing*, and include memory technologies such as embedded DRAM and 3-D stacked DRAM.

Next, we will describe efforts to integrate the computation *into* the memory itself. These approaches are often referred to as *processing in memory* or *in-memory computing*, and include memory technologies such as Static Random Access Memories (SRAM), Dynamic Random Access Memories (DRAM), and emerging non-volatile memory (NVM). Since these approaches rely on mixed-signal circuit design to enable processing in the analog domain, we will also discuss the design challenges related to handling the increased sensitivity to circuit and device non-idealities (e.g., nonlinearity, process and temperature variations), as well as the impact on area density, which is critical for memory.

Significant data movement also occurs between the sensor that collects the data and the DNN processor. The same principles that are used to bring compute near the memory, where the weights are stored, can be used to bring the compute *near* the sensor, where the input data is collected. Therefore, we will also discuss how to integrate some of the compute *into* the sensor.

Finally, since photons travel much faster than electrons and the cost of moving a photon can be *independent* of distance, processing in the optical domain using light may provide significant improvements in energy efficiency and throughput over the electrical domain. Accordingly, we will conclude this chapter by discussing the recent work that performs DNN processing in the optical domain, referred to as *Optical Neural Networks*.

<sup>1</sup>Specifically, the memory wall refers to data moving between the off-chip memory (e.g., DRAM) and the processor.

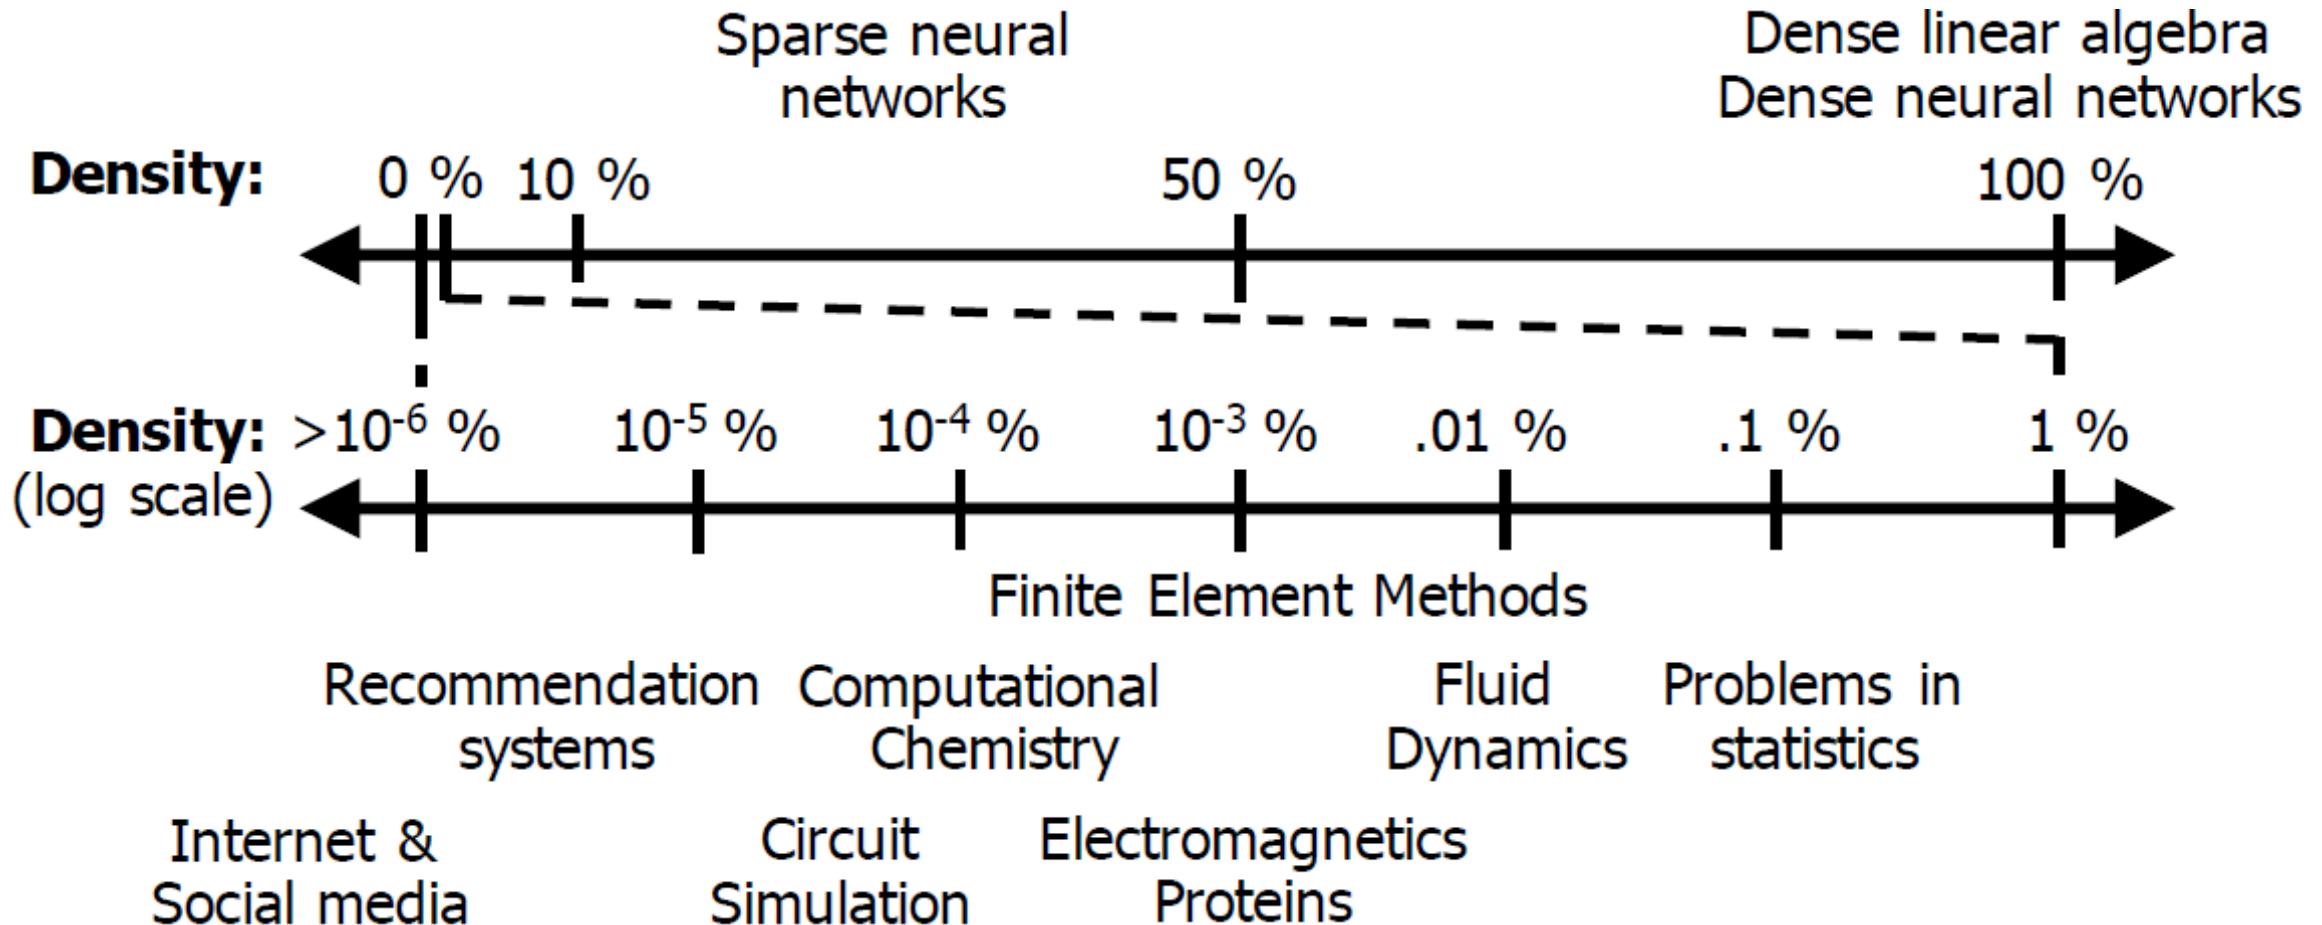
## Many Design Considerations for In-Memory Computing

- Number of Storage Elements per Weight
- Array Size
- Number of Rows Activated in Parallel
- Number of Columns Activated in Parallel
- Time to Deliver Input
- Time to Compute MAC

Tradeoffs between energy efficiency, throughput, area density, and accuracy, which *reduce the achievable gains over conventional architectures*

Available on DNN tutorial website  
<http://eyeriss.mit.edu/tutorial.html>

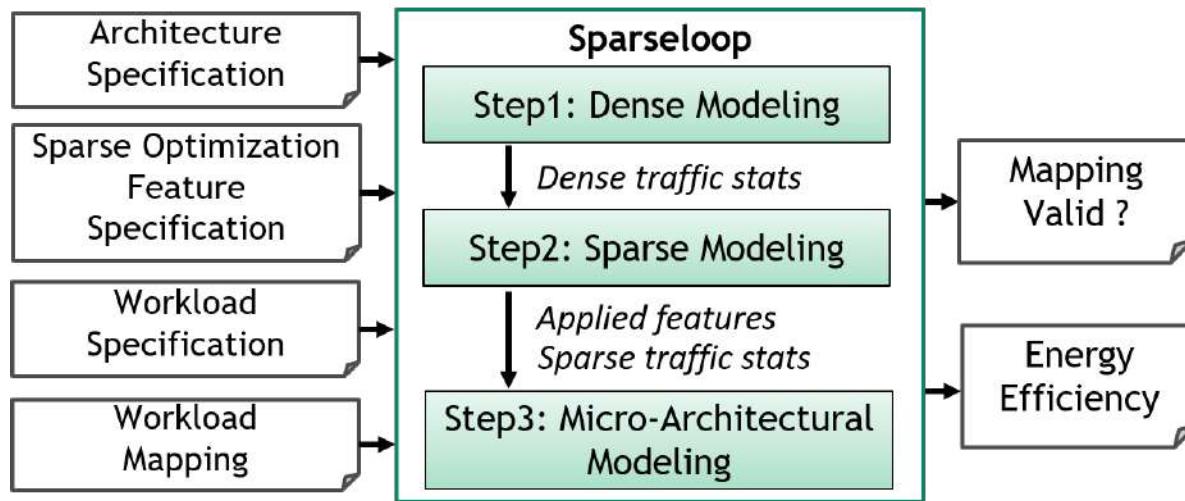
# Applications that use Sparse Tensor



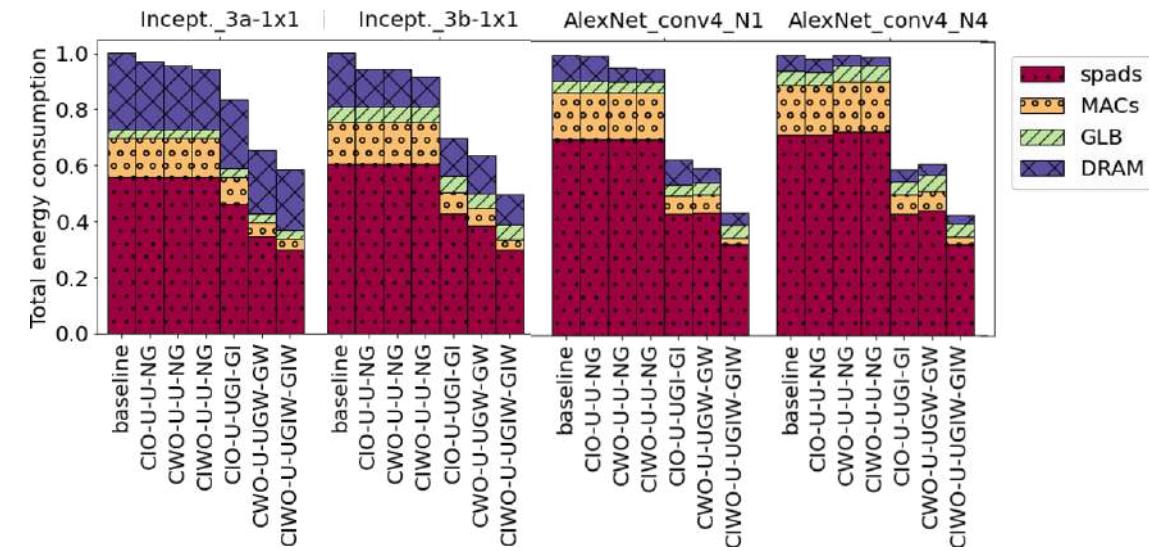
# Sparseloop: Design Space Exploration for Sparse Tensor Accelerators

- An analytical design exploration framework that comprehends a wide range of
  - Sparse optimizations (e.g., zero-gating, zero-skipping, zero-compression)
  - Data representations (e.g., uncompressed, run length coding, bitmask)

*Propose modularized three-step evaluation process*



*Energy impact of sparse optimizations at different levels of the memory hierarchy in Eyeriss-based topology*



Tutorial at ISCA 2021 (June 19): [http://accelergy.mit.edu/sparse\\_tutorial.html](http://accelergy.mit.edu/sparse_tutorial.html)

# Book Chapter on Sparse Computations

CHAPTER 8

167

## Exploiting Sparsity

A salient characteristic of the data used in DNN computations is that it is (or can be made to be) sparse. By saying that the data is sparse, we are referring to the fact that there are many repeated values in the data. Much of the time the repeated value is zero, which is what we will assume unless explicitly noted. Thus, we will talk about the sparsity or density of the data as the percentage of zeros or non-zeros, respectively in the data. The existence of sparse data leads broadly to two potential architectural benefits: (1) sparsity can reduce the footprint of the data, which provides an opportunity to reduce storage requirements and data movement. This is because sparse data is amenable to being compressed, as described in Section 8.2;<sup>1</sup> and (2) sparsity presents an opportunity for a reduction in MAC operations. The reduction in MAC operations results from the fact that  $0 \times \text{anything}$  is 0. This can result in either savings in energy or time or both. In Section 8.3, we will discuss how the dataflows for sparse data can translate sparsity into improvements in energy-efficiency and throughput. However, first in Section 8.1 we discuss the origins and ways that one can increase sparsity in the data used in DNN computations.

### 8.1 SOURCES OF SPARSITY

Efficient processing of feature map activations becomes increasingly important as the size of the input to the DNN model grows (e.g., increased image resolution), while efficient processing of filter weights becomes increasingly important as the size of the DNN model grows (e.g., increased number of layers).

This section will discuss various approaches that can exploit properties such as redundancy and correlation in the feature maps and filters to increase their activation sparsity (Section 8.1.1) and weight sparsity (Section 8.1.2), respectively. The requirements for these approaches may differ as activation sparsity is often data dependent and not known *a priori*, while weight sparsity can be known *a priori*. As a result, methods to increase sparsity for weights can be performed offline (as opposed to during inference) and can be more computationally complex than methods applied to increase activation sparsity. For instance, increasing weight sparsity can be incorporated into training.

<sup>1</sup>Note: We use the words sparsity or density to refer to a statistical property of the data, while we use the words compressed or uncompressed to describe the characteristics of a representation of the (typically sparse) data.

- Sources of Sparsity
- Compression and Sparse Tensor Representation
- Sparse Dataflows

<https://tinyurl.com/EfficientDNNBook>

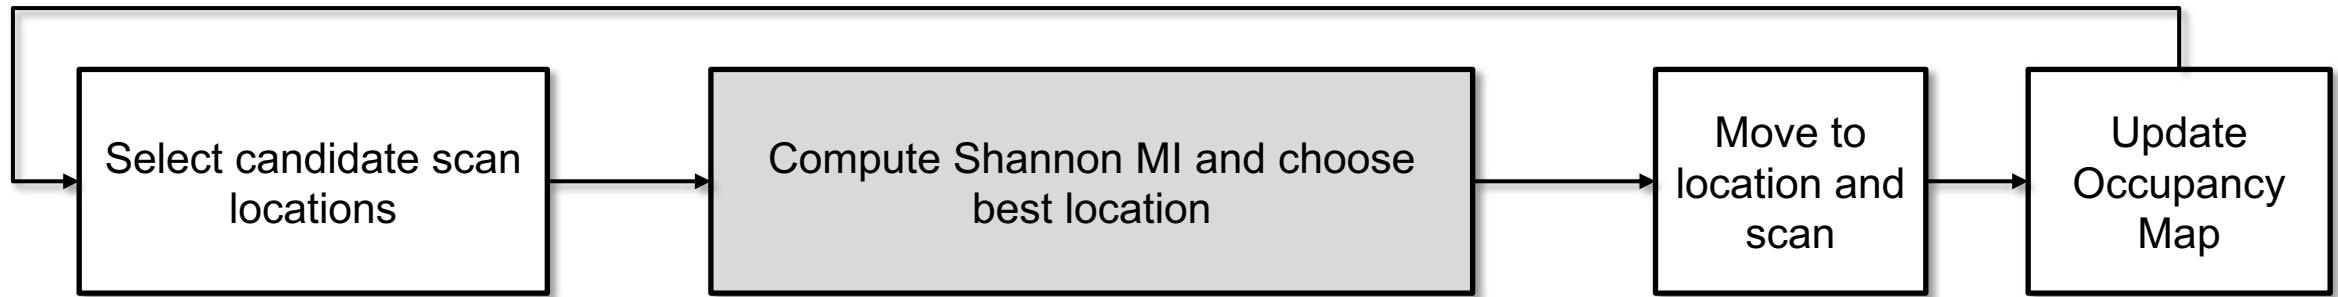
# Where to Go Next: Planning and Mapping

## Robot Exploration

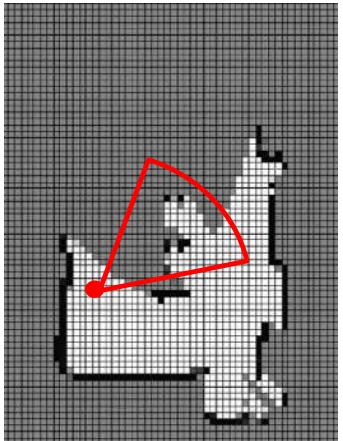


# Where to Go Next: Planning and Mapping

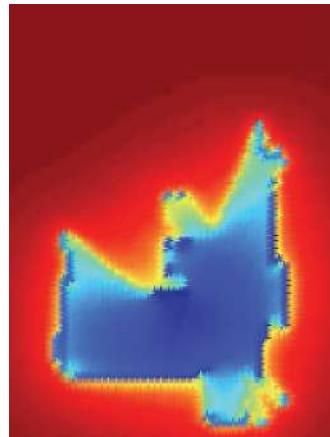
***Robot Exploration: Decide where to go by computing Shannon Mutual Information***



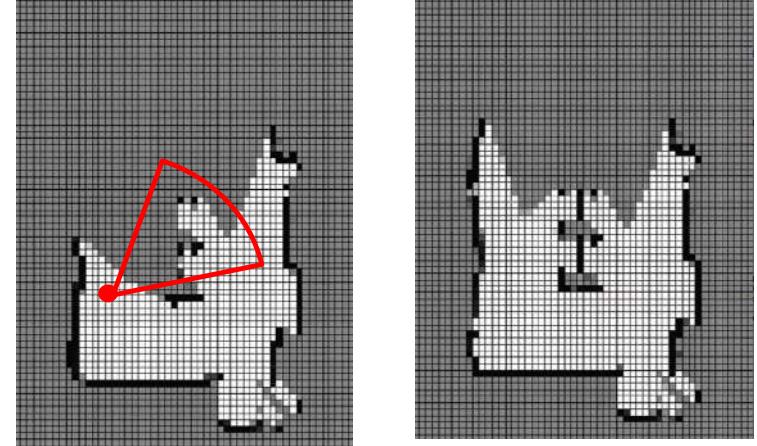
Where to scan?



Mutual Information



Updated Map



# Experimental Results (4x Real Time)



Occupancy map with  
planned path using RRT\*  
(compute MI on all possible paths)

MI surface

Exploration with a mini race car using motion capture for localization

# Building Hardware Accelerator to Compute MI

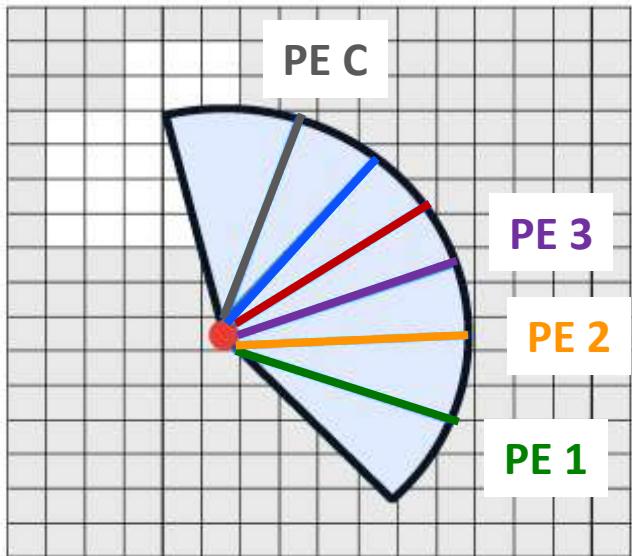
**Motivation:** Compute MI faster for faster exploration!

$$I(M; Z) = \sum_{j=1}^n \sum_{k=j-\Delta}^{j+\Delta} P(e_j) C_k G_{k,j}$$

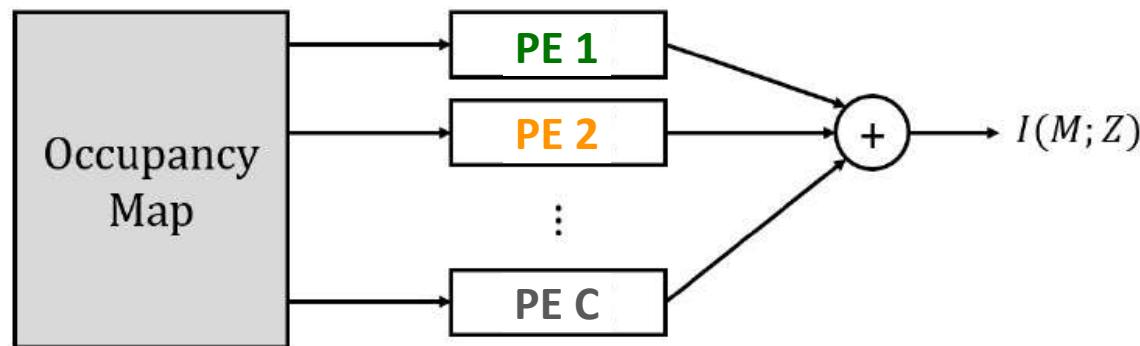
Fast Shannon  
Mutual Information (FSMI)  
[Zhang, ICRA 2019]

Algorithm is *embarrassingly* parallel!

High throughput *should* be possible with multiple processing elements (PE)

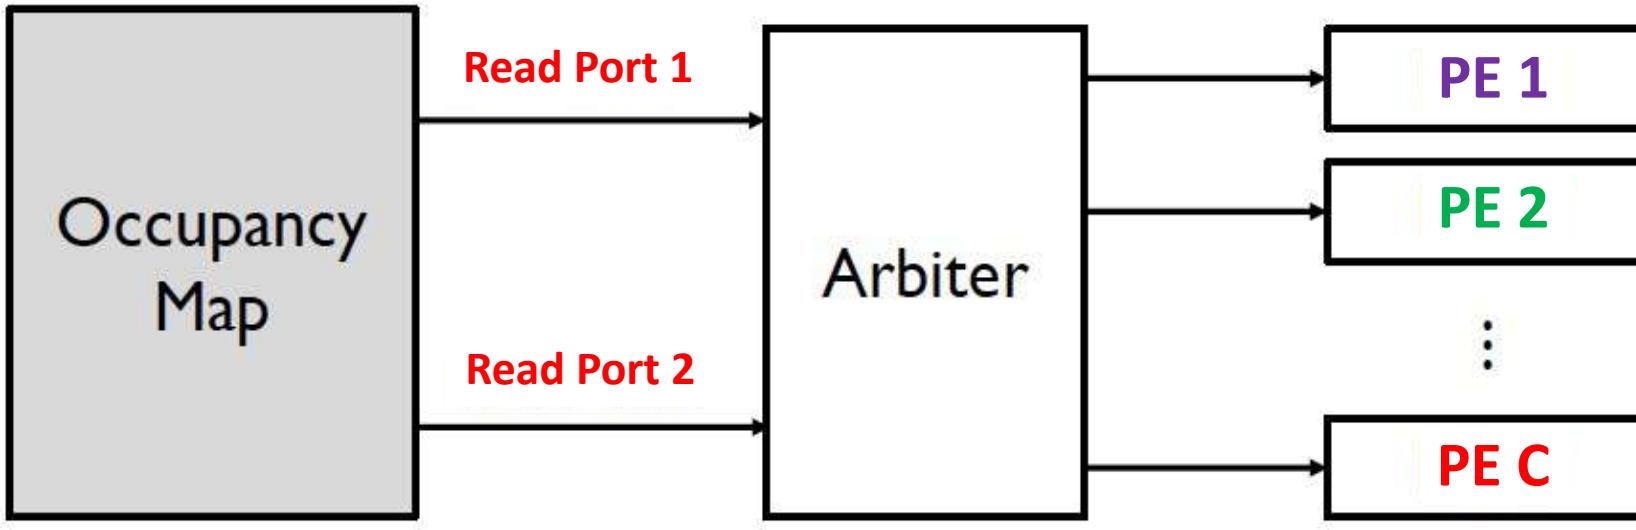


Process sensor beams in parallel with multiple PEs



# Challenge is Data Delivery to All PEs

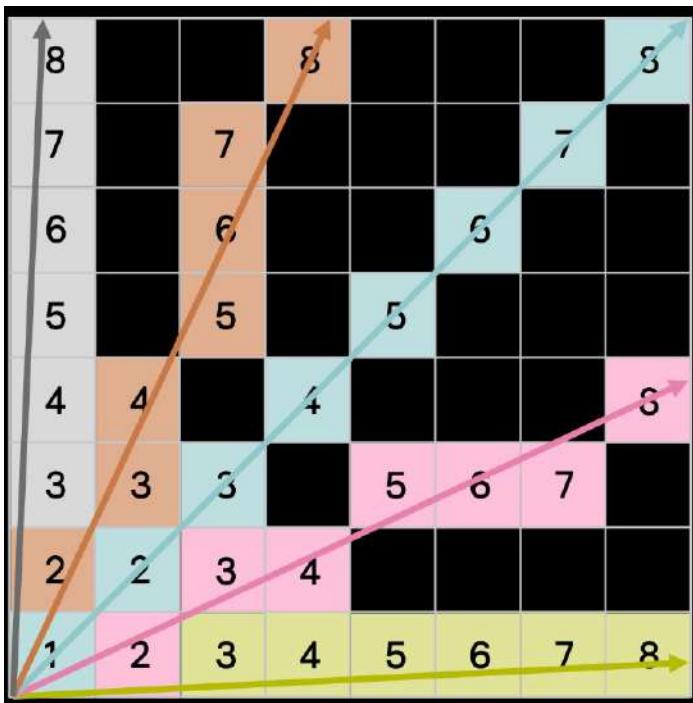
Power consumption of memory scales with number of ports.  
**Low power SRAM limited to two-ports!**



Data delivery, specifically memory bandwidth,  
limits the throughput (not compute)

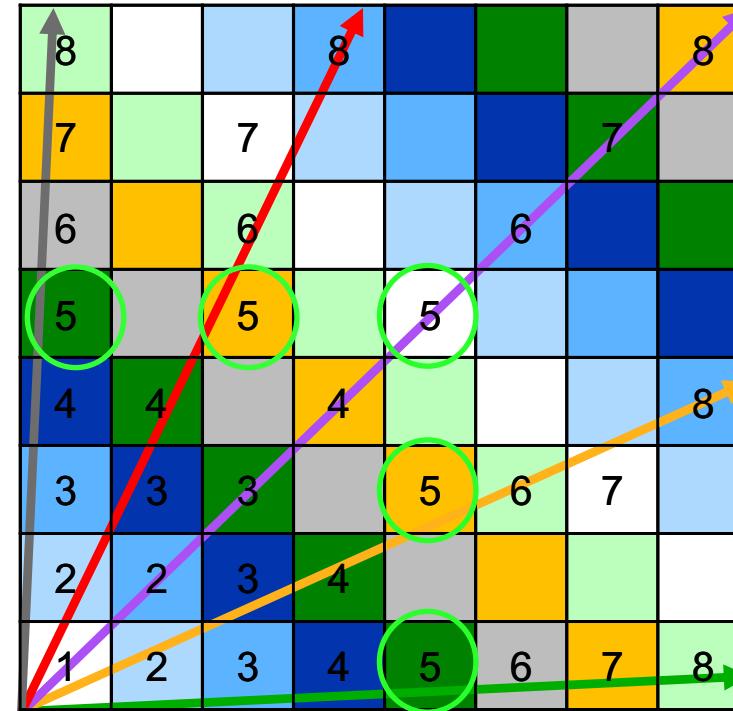
# Optimized Memory Banking Pattern

Memory Access Pattern



PEs read the map at the same row  
or column every cycle

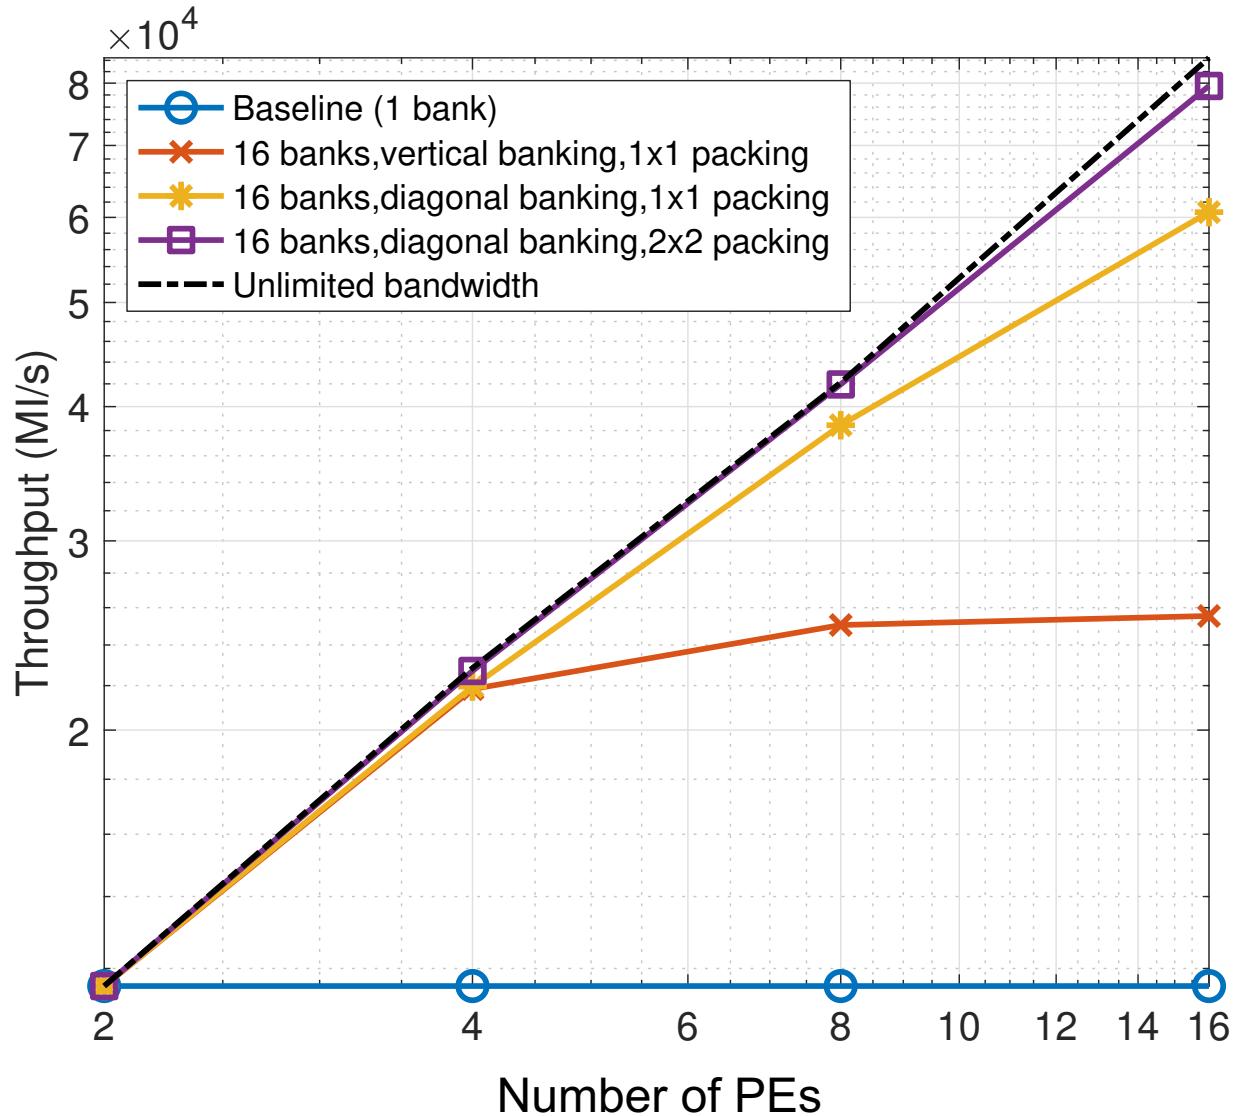
Diagonal Banking Pattern



Reduced conflicts across banks

- Bank 0
- Bank 1
- Bank 2
- Bank 3
- Bank 4
- Bank 5
- Bank 6
- Bank 7

# Experimental Results

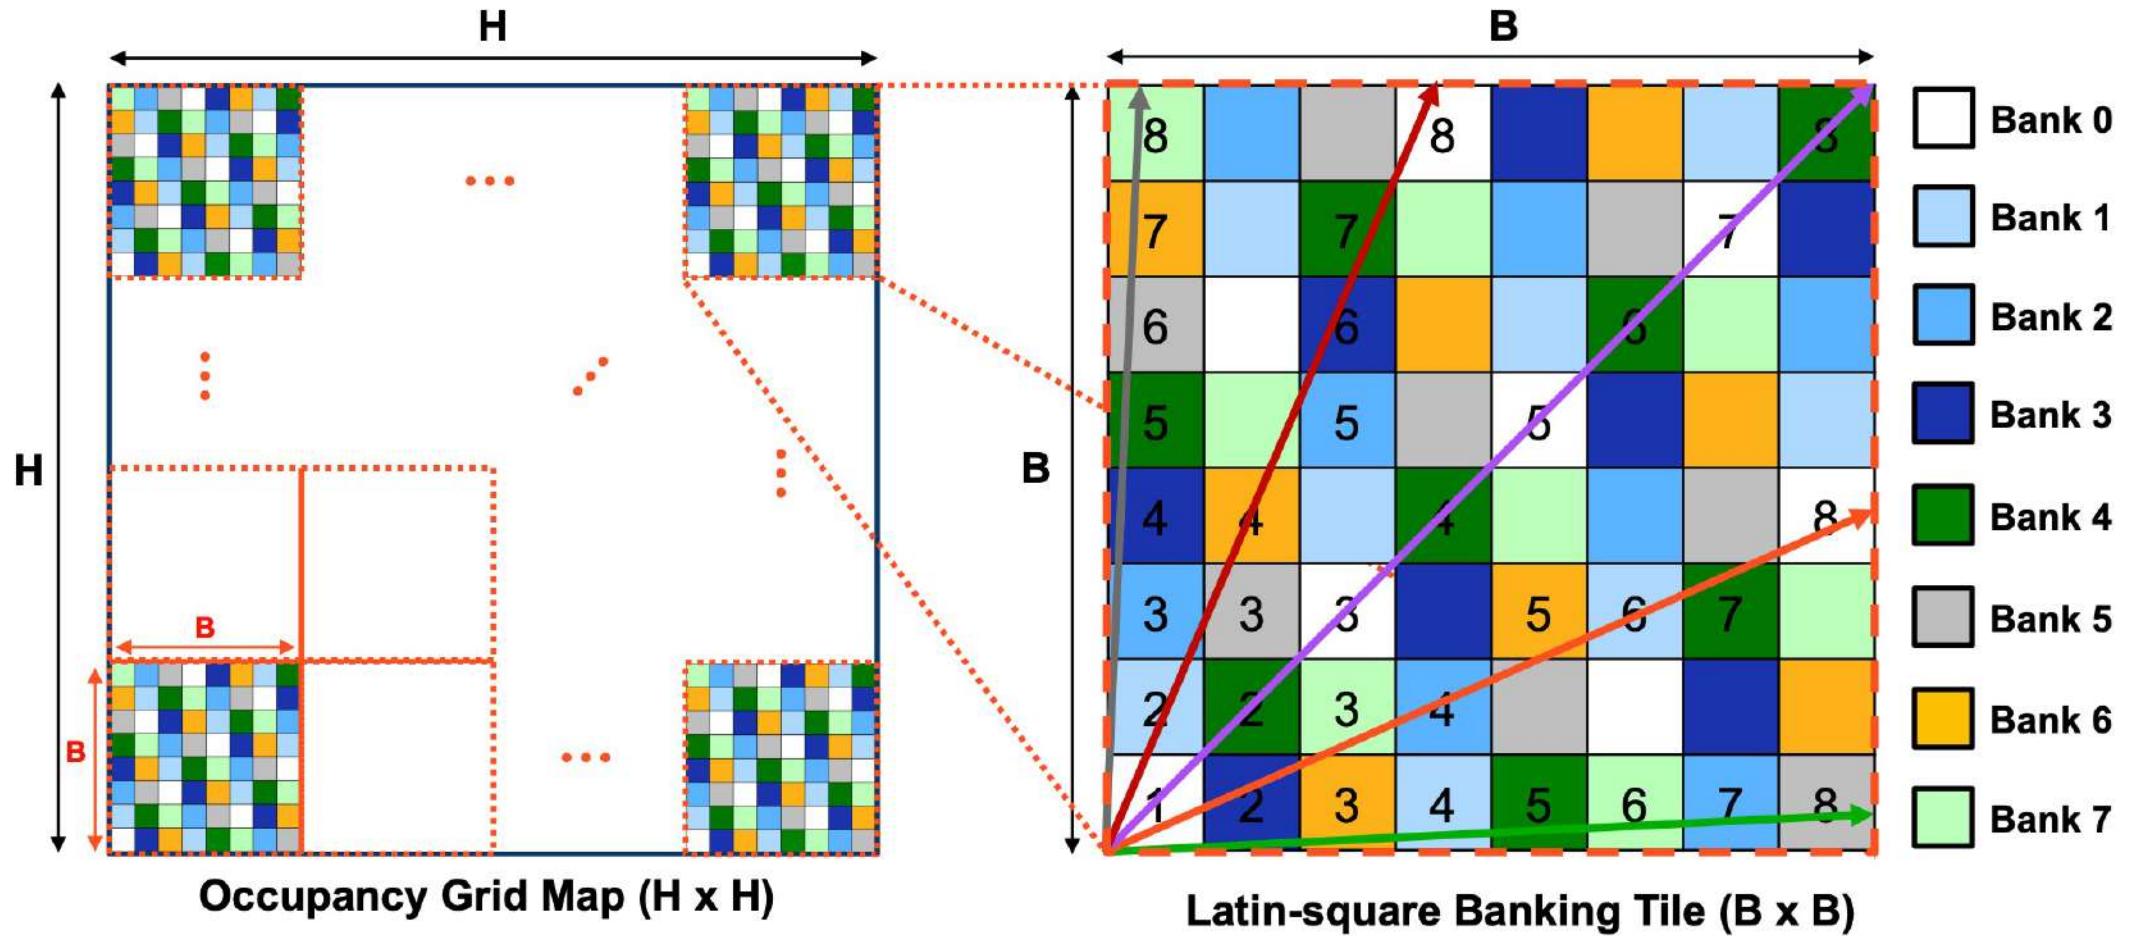


Specialized banking, efficient memory arbiter and packing multiple values at each address results in throughput **within 94% of theoretical limit** (unlimited bandwidth)

Compute MI for an **entire map** of 20m x 20m at 0.1m resolution **in under a second** on a ZC706 FPGA  
(100x faster than CPU at 10x lower power)

# Generalize to a Class of Banking Patterns

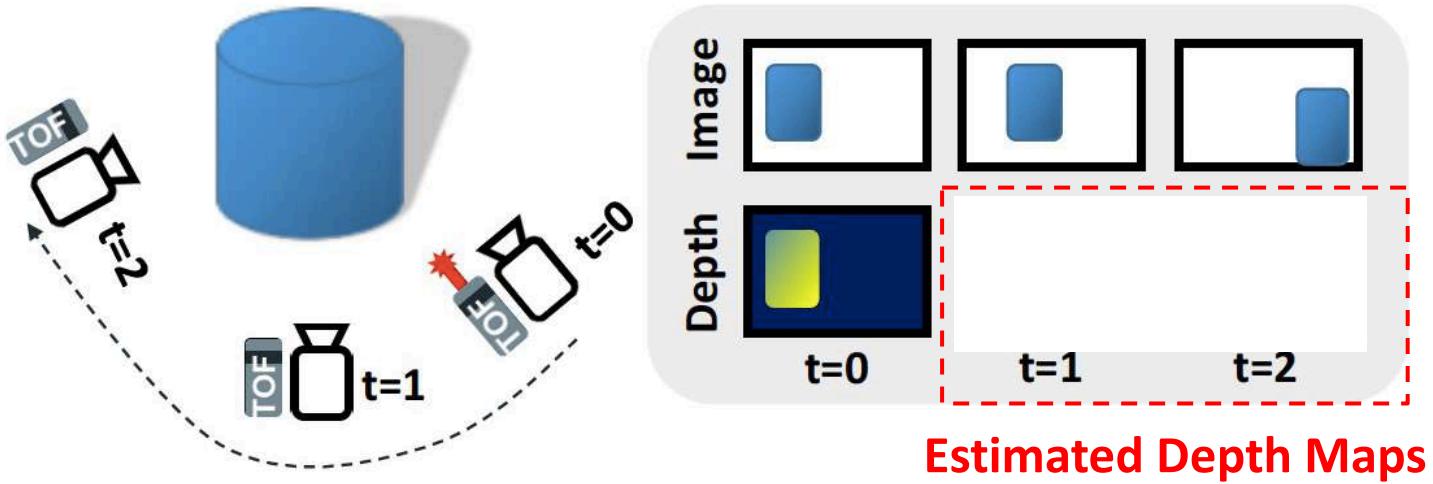
- Latin-square banking tile: cells in each column and row is assigned to different banks



We rigorously proved that Latin-square tiles usage minimizes read conflicts between PEs

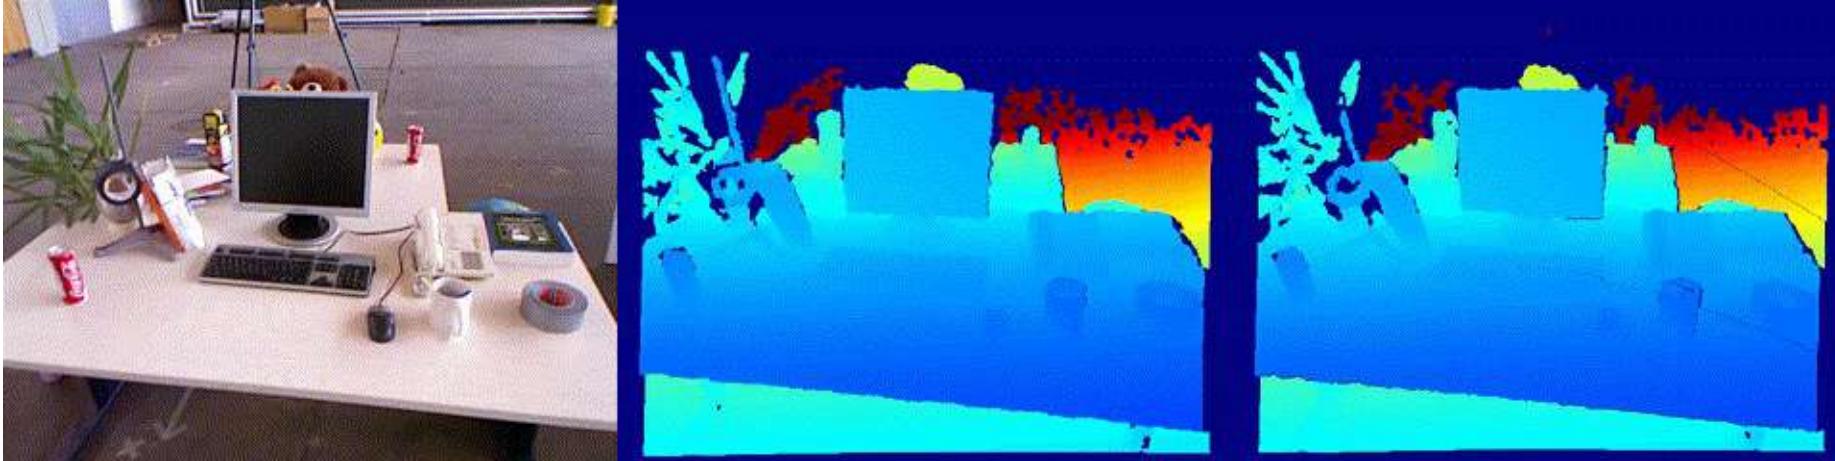
# Low Power 3D Time of Flight Imaging

- Pulsed Time of Flight: Measure distance using round trip time of laser light for each image pixel
  - Illumination + Imager Power: 2.5 – 20 W for range from 1 - 8 m
- Use computer vision techniques and passive images to estimate changes in depth without turning on laser
  - CMOS Imaging Sensor Power: < 350 mW



**Real-time Performance on Embedded Processor**  
VGA @ 30 fps on Cortex-A7 (< 0.5W active power)

# Results of Low Power Depth ToF Imaging



RGB Image

Depth Map  
Ground Truth

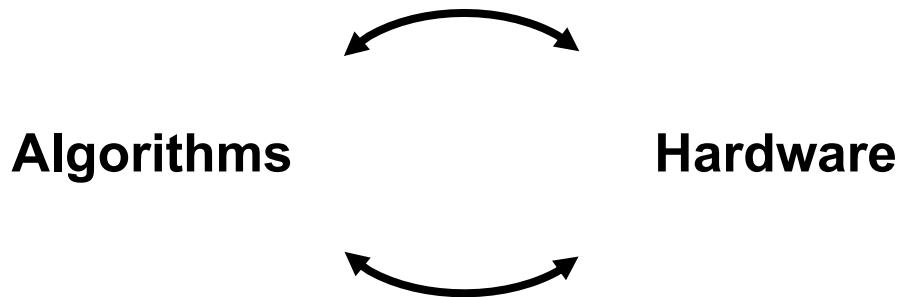
Depth Map  
Estimated

Mean Relative Error: 0.7%

Duty Cycle (on-time of laser): 11%

# Summary

- Efficient computing is critical for advancing the progress of AI & autonomous robots  
→ **Critical step to making AI & autonomy ubiquitous!**
- In order to meet computing demands in terms of power and speed, need to redesign computing hardware from the ground up → **Focus on data movement!**
- Specialized hardware creates new opportunities for the co-design of algorithms and hardware → **Innovation opportunities for the future of AI & robotics!**



# Acknowledgements



Joel Emer



Sertac Karaman

Research conducted in the **MIT Energy-Efficient Multimedia Systems Group** would not be possible without the support of the following organizations:



# Low-Energy Autonomy and Navigation (LEAN) Group



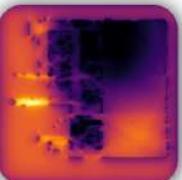
A broad range of next-generation applications will be enabled by low-energy, miniature mobile robotics including insect-size flapping wing robots that can help with search and rescue, chip-size satellites that can explore nearby stars, and blimps that can stay in the air for years to provide communication services in remote locations. While the low-energy, miniature actuation, and sensing systems have already been developed in many of these cases, the processors currently used to run the algorithms for autonomous navigation are still energy-hungry. Our research addresses this challenge as well as brings together the robotics and hardware design communities.

We enable efficient computing on various key modules of other autonomous navigation systems including perception, localization, exploration and planning. We also consider the overall system by considering the energy cost of computing in conjunction with actuation and sensing.



## Motion Planning

Many motion planning and control algorithms aim to design trajectories and controllers that minimize actuation energy. However, in low-energy robotics, computing such trajectories and controls themselves may consume a large amount of energy. We develop algorithms that optimize this trade-off.



## Mutual Information for Exploration

Computing mutual information between the map and future measurements is critical to efficient exploration. Unfortunately, mutual information computation is computationally very challenging. We develop new algorithms and hardware for efficient computation of mutual information, and demonstrate real-time computation for the whole map in a reasonably-sized map.



## Depth Sensing and Perception

Depth sensing is a critical function for robotic tasks such as localization, mapping and obstacle detection. State-of-the-art single-view depth estimation algorithms are based on fairly complex deep neural networks that are too slow for real-time inference on an embedded platform, for instance, mounted on a micro aerial vehicle. We address the problem of fast depth estimation on embedded systems.



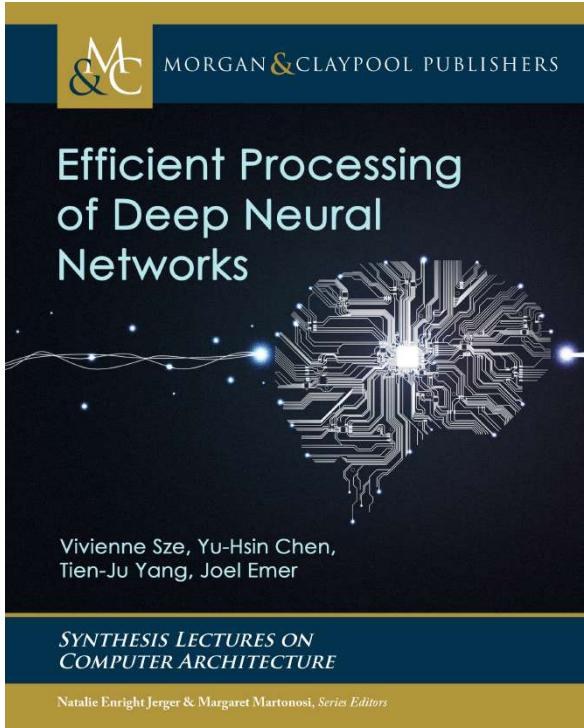
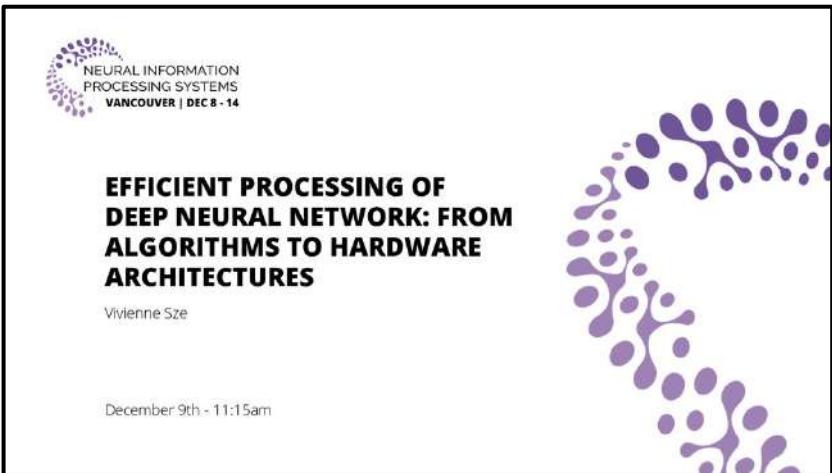
## Localization and Mapping

Autonomous navigation of miniaturized robots (e.g., nano/pico aerial vehicles) is currently a grand challenge for robotics research, due to the need for processing a large amount of sensor data (e.g., camera frames) with limited on-board computational resources. We focus on the design of a visual-inertial odometry (VIO) system in which the robot estimates its ego-motion (and a landmark-based map) from on-board camera and IMU data.



**Group Website:** <http://lean.mit.edu>

# Resources on Efficient Processing of DNNs



<http://eyeriss.mit.edu/tutorial.html>

# Excerpts of Book

## CHAPTER 3

### Key Metrics and Design Objectives

Over the past few years, there has been a significant amount of research on efficient processing of DNNs. Accordingly, it is important to discuss the key metrics that one should consider when comparing and evaluating the strengths and weaknesses of different designs and proposed techniques and that should be incorporated into design considerations. While efficiency is often only associated with the number of operations per second per Watt (e.g., floating-point operations per second per Watt as FLOPS/W or tera-operations per second per Watt as TOPS/W), it is actually composed of many more metrics including accuracy, throughput, latency, energy consumption, power consumption, cost, flexibility, and scalability. Reporting a comprehensive set of these metrics is important in order to provide a complete picture of the trade-offs made by a proposed design or technique.

In this chapter, we will

- discuss the importance of each of these metrics;
- breakdown the factors that affect each metric. When feasible, present equations that describe the relationship between the factors and the metrics;
- describe how these metrics can be incorporated into design considerations for both the DNN hardware and the DNN model (i.e., workload); and
- specify what should be reported for a given metric to enable proper evaluation.

Finally, we will provide a case study on how one might bring all these metrics together for a holistic evaluation of a given approach. But first, we will discuss each of the metrics.

#### 3.1 ACCURACY

*Accuracy* is used to indicate the quality of the result for a given task. The fact that DNNs can achieve state-of-the-art accuracy on a wide range of tasks is one of the key reasons for the popularity and wide use of DNNs today. The units used to measure accuracy vary by task. For instance, for image classification, accuracy is reported as the percent of correctly classified images, while for object detection, accuracy is reported as the mean average precision (mAP), which is related to the trade off between the true positive rate and false

43

## CHAPTER 10

### Advanced Technologies

As highlighted throughout the previous chapters, data movement dominates energy consumption. The energy is consumed both in the access to the memory as well as the transfer of the data. The associated physical factors also limit the bandwidth available to deliver data between memory and compute, and thus limits the throughput of the overall system. This is commonly referred to by computer architects as the “memory wall.”<sup>1</sup>

To address the challenges associated with data movement, there have been various efforts to bring compute and memory closer together. Chapters 5 and 6 primarily focus on how to design spatial architectures that distribute the on-chip memory closer to the computation (e.g., scratch pad memory in the PE). This chapter will describe various other architectures that use advanced memory, process, and fabrication technologies to bring the compute and memory together.

First, we will describe efforts to bring the off-chip high-density memory (e.g., DRAM) closer to the computation. These approaches are often referred to as *processing near memory* or *near-data processing*, and include memory technologies such as embedded DRAM and 3-D stacked DRAM.

Next, we will describe efforts to integrate the computation *into* the memory itself. These approaches are often referred to as *processing in memory* or *in-memory computing*, and include memory technologies such as Static Random Access Memories (SRAM), Dynamic Random Access Memories (DRAM), and emerging non-volatile memory (NVM). Since these approaches rely on mixed-signal circuit design to enable processing in the analog domain, we will also discuss the design challenges related to handling the increased sensitivity to circuit and device non-idealities (e.g., nonlinearity, process and temperature variations), as well as the impact on area density, which is critical for memory.

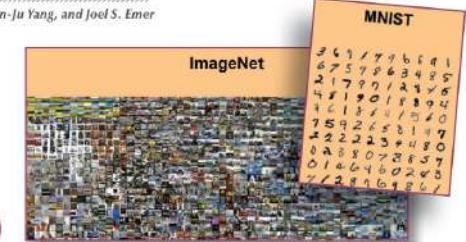
Significant data movement also occurs between the sensor that collects the data and the DNN processor. The same principles that are used to bring compute near the memory, where the weights are stored, can be used to bring the compute *near* the sensor, where the input data is collected. Therefore, we will also discuss how to integrate some of the compute *into* the sensor.

Finally, since photons travel much faster than electrons and the cost of moving a photon can be *independent of distance*, processing in the optical domain using light may provide significant

253

## ISSCC 2020 TUTORIAL

Vivienne Sze, Yu-Hsin Chen, Tien-Ju Yang, and Joel S. Emer



### How to Evaluate Deep Neural Network Processors

TOPS/W (alone) Considered Harmful



A significant amount of specialized hardware has been developed for processing deep neural networks (DNNs) in both academia and industry. This article aims to highlight the key concepts required to evaluate and compare these DNN processors.

International Solid-State Circuits Conference, as well as excerpts from the book, *Efficient Processing of Deep Neural Networks* [36].

#### Motivation and Background

Over the past few years, there has been a significant amount of research on enabling the efficient processing of DNNs. The challenge of efficient DNN processing depends on balancing multiple objectives:

- high performance (including accuracy) and efficiency (including cost)

- enough flexibility to cater to a wide and rapidly changing range of workloads

- good integration with existing software frameworks.

DNN computations are composed of several processing layers (Figure 1), where, for many layers, the main computation is a weighted sum; in other words, the main computation for DNN processing is often a

multiply-accumulate (MAC) operation. The arrangement of the MAC operations within a layer is defined by the layer shape; for instance, Table 1 and Figure 2 highlight the shape parameters for layers used in convolutional neural networks (CNNs), a popular type of DNN. Because the shape parameters can vary across layers, DNNs come in a wide variety of shapes and sizes, depending on the application. (The DNN research community often refers to the shape and size of a DNN as its *network architecture*. However, to avoid confusion with the use of the word *architecture* by the hardware community, we talk about *DNN models* and their shape and size in this article.) This variety is one of the motivations for flexibility, and it causes the objectives listed previously to be highly interrelated.

Figure 3 illustrates the hardware architecture of a typical DNN processor, which is composed of an array

Available on DNN tutorial website  
<http://eyeriss.mit.edu/tutorial.html>

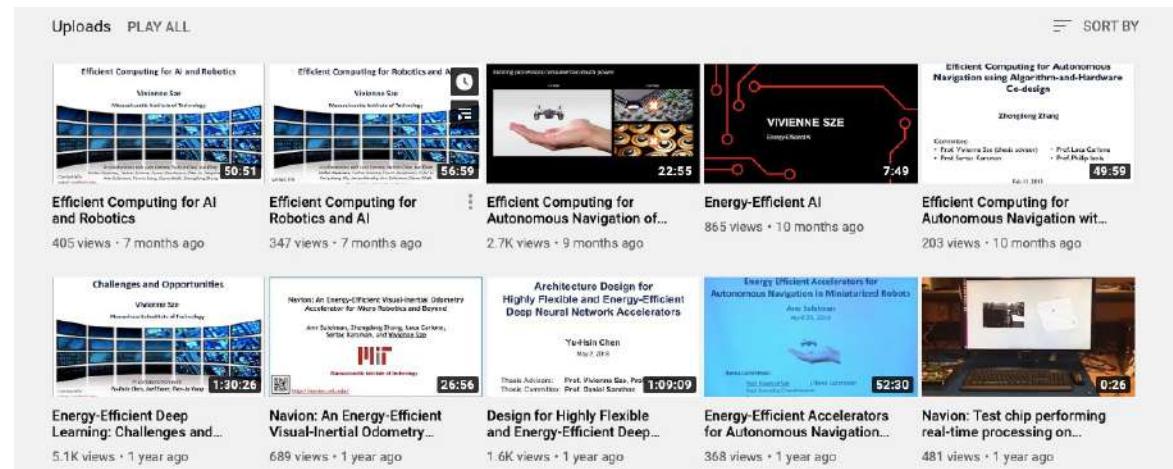
After 10.1109/ISSCC.2020.9002140  
 Date: 25 August 2020

# Additional Resources

**Talks and Tutorial Available Online**  
<https://tinyurl.com/ISCAS2021Sze>



YouTube Channel  
**EEMS Group – PI: Vivienne Sze**



# References

- Efficient Processing for Deep Neural Networks

- Project website: <http://eyeriss.mit.edu>
- Y. N. Wu, P.-A. Tsai, A. Parashar, V. Sze, J. S. Emer, “Sparseloop: An Analytical, Energy-Focused Design Space Exploration Methodology for Sparse Tensor Accelerators,” IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS), March 2021
- Y.-H. Chen, T.-J Yang, J. Emer, V. Sze, “Eyeriss v2: A Flexible Accelerator for Emerging Deep Neural Networks on Mobile Devices,” IEEE Journal on Emerging and Selected Topics in Circuits and Systems (JETCAS), Vol. 9, No. 2, pp. 292-308, June 2019.
- Y.-H. Chen, T. Krishna, J. Emer, V. Sze, “Eyeriss: An Energy-Efficient Reconfigurable Accelerator for Deep Convolutional Neural Networks,” IEEE Journal of Solid-State Circuits (JSSC), ISSCC Special Issue, Vol. 52, No. 1, pp. 127-138, January 2017.
- Y.-H. Chen, J. Emer, V. Sze, “Eyeriss: A Spatial Architecture for Energy-Efficient Dataflow for Convolutional Neural Networks,” International Symposium on Computer Architecture (ISCA), pp. 367-379, June 2016.
- Y.-H. Chen\*, T.-J. Yang\*, J. Emer, V. Sze, “Understanding the Limitations of Existing Energy-Efficient Design Approaches for Deep Neural Networks,” SysML Conference, February 2018.
- V. Sze, Y.-H. Chen, T.-J. Yang, J. Emer, “Efficient Processing of Deep Neural Networks: A Tutorial and Survey,” Proceedings of the IEEE, vol. 105, no. 12, pp. 2295-2329, December 2017.
- Y. N. Wu, J. S. Emer, V. Sze, “Accelergy: An Architecture-Level Energy Estimation Methodology for Accelerator Designs,” International Conference on Computer Aided Design (ICCAD), November 2019. <http://accelergy.mit.edu/>
- Y. N. Wu, V. Sze, J. S. Emer, “An Architecture-Level Energy and Area Estimator for Processing-In-Memory Accelerator Designs,” to appear in IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS), April 2020.
- A. Suleiman\*, Y.-H. Chen\*, J. Emer, V. Sze, “Towards Closing the Energy Gap Between HOG and CNN Features for Embedded Vision,” IEEE International Symposium of Circuits and Systems (ISCAS), Invited Paper, May 2017.
- Hardware Architecture for Deep Neural Networks: <http://eyeriss.mit.edu/tutorial.html>

# References

- **Co-Design of Algorithms and Hardware for Deep Neural Networks**

- T.-J. Yang, Y.-H. Chen, V. Sze, “Designing Energy-Efficient Convolutional Neural Networks using Energy-Aware Pruning,” IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2017.
- Energy estimation tool: <http://eyeriss.mit.edu/energy.html>
- T.-J. Yang, A. Howard, B. Chen, X. Zhang, A. Go, V. Sze, H. Adam, “NetAdapt: Platform-Aware Neural Network Adaptation for Mobile Applications,” European Conference on Computer Vision (ECCV), 2018. <http://netadapt.mit.edu>
- D. Wofk\*, F. Ma\*, T.-J. Yang, S. Karaman, V. Sze, “FastDepth: Fast Monocular Depth Estimation on Embedded Systems,” IEEE International Conference on Robotics and Automation (ICRA), May 2019. <http://fastdepth.mit.edu/>
- T.-J. Yang, V. Sze, “Design Considerations for Efficient Deep Neural Networks on Processing-in-Memory Accelerators,” IEEE International Electron Devices Meeting (IEDM), Invited Paper, December 2019.
- T.-J. Yang, Y.-L. Liao, V. Sze, “NetAdaptV2: Efficient Neural Architecture Search with Fast Super-Network Training and Architecture Optimization ,” IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2021.

- **Low Power Time of Flight Imaging**

- J. Noraky, V. Sze, “Low Power Depth Estimation of Rigid Objects for Time-of-Flight Imaging,” IEEE Transactions on Circuits and Systems for Video Technology (TCSVT), 2020.
- J. Noraky, V. Sze, “Depth Map Estimation of Dynamic Scenes Using Prior Depth Information,” arXiv, February 2020. <https://arxiv.org/abs/2002.00297>
- J. Noraky, V. Sze, “Depth Estimation of Non-Rigid Objects For Time-Of-Flight Imaging,” IEEE International Conference on Image Processing (ICIP), October 2018.

# References

- **Energy-Efficient Visual Inertial Localization**
  - Project website: <http://navion.mit.edu>
  - A. Suleiman, Z. Zhang, L. Carlone, S. Karaman, V. Sze, “Navion: A Fully Integrated Energy-Efficient Visual-Inertial Odometry Accelerator for Autonomous Navigation of Nano Drones,” IEEE Symposium on VLSI Circuits (VLSI-Circuits), June 2018.
  - Z. Zhang\*, A. Suleiman\*, L. Carlone, V. Sze, S. Karaman, “Visual-Inertial Odometry on Chip: An Algorithm-and-Hardware Co-design Approach,” Robotics: Science and Systems (RSS), July 2017.
  - A. Suleiman, Z. Zhang, L. Carlone, S. Karaman, V. Sze, “Navion: A 2mW Fully Integrated Real-Time Visual-Inertial Odometry Accelerator for Autonomous Navigation of Nano Drones,” IEEE Journal of Solid State Circuits (JSSC), VLSI Symposia Special Issue, Vol. 54, No. 4, pp. 1106-1119, April 2019.

# References

- **Fast Shannon Mutual Information for Robot Exploration**
  - Project website: <http://lean.mit.edu>
  - Z. Zhang, T. Henderson, V. Sze, S. Karaman, “FSMI: Fast computation of Shannon Mutual Information for information-theoretic mapping,” IEEE International Conference on Robotics and Automation (ICRA), May 2019.
  - P. Li\*, Z. Zhang\*, S. Karaman, V. Sze, “High-throughput Computation of Shannon Mutual Information on Chip,” Robotics: Science and Systems (RSS), June 2019
  - Z. Zhang, T. Henderson, S. Karaman, V. Sze, “FSMI: Fast computation of Shannon Mutual Information for information-theoretic mapping,” to appear in International Journal of Robotics Research (IJRR). <http://arxiv.org/abs/1905.02238>
  - T. Henderson, V. Sze, S. Karaman, “An Efficient and Continuous Approach to Information-Theoretic Exploration,” IEEE International Conference on Robotics and Automation (ICRA), May 2020
- **Balancing Actuation and Computation**
  - Project website: <http://lean.mit.edu>
  - S. Sudhakar, S. Karaman, V. Sze, “Balancing Actuation and Computing Energy in Motion Planning,” IEEE International Conference on Robotics and Automation (ICRA), May 2020