

АГЕНТЫ ИИ / ИНФРАСТРУКТУРА ИИ / СТРАТЕГИЯ ИИ

# 10 событий, определивших бурную первую половину 2026 года в сфере ИИ

Десять определяющих событий в сфере ИИ в 2026 году: закрытие Fable 5 от Anthropic, столкновения в Пентагоне, гонка за внедрение ИИ на Уолл-стрит и появление «пояса ИИ».

5 июля 2026 г. в 10:00 [Дэррил К. Тафт](#)




Фото [vackground.com](#) на [Unsplash](#).

**В середине 2026 года искусственный интеллект был в центре внимания всех крупных событий в мире разработки программного обеспечения и практически всех крупных событий за его пределами.**

**В прошлом месяце Министерство торговли США приказало компании Anthropic отключить Fable 5 и Mythos 5 по всему миру — но через 18 дней запрет был снят. Это был не единственный конфликт Anthropic с**

В других местах передовые лаборатории в мае закрепились на Уолл-стрит, создав новые подразделения и наладив партнерские отношения, в то время как Anthropic и OpenAI вышли на IPO с оценочной стоимостью более 800 миллиардов долларов. За этими оценками стоит наращивание инфраструктуры в виде чипов, центров обработки данных и сделок, призванных идти в ногу с релизами моделей, которые появляются каждые несколько недель.

Тем временем модели с открытым исходным кодом сокращают разрыв между закрытыми и загружаемыми моделями, а «обвязка» — инструменты, память и управление моделью — приобретает все большее значение по мере внедрения агентного ИИ в корпоративную среду. Добавьте сюда токеномику (реальную стоимость масштабирования рассуждений) и нетехнических руководителей, которые создают собственные инструменты, и вы получите десять важнейших событий в сфере ИИ в 2026 году — пока что.

Вот десять событий, которые определили важнейшую первую половину 2026 года в мире ИИ по мнению редакции *The New Stack*.

## 10. Указ президента Трампа об искусственном интеллекте

Президент Дональд Трамп довольно благосклонно относится к искусственному интеллекту, учитывая его поддержку технологических компаний Кремниевой долины. 2 июня 2026 года Трамп подписал **указ**, направленный на защиту американских систем от угроз, связанных с искусственным интеллектом, и в то же время отвергающий «чрезмерно обременительное регулирование». Указ предписывает Комитету по безопасности национальных систем уделять приоритетное внимание киберзащите, а Министерству финансов, Агентству национальной безопасности и Агентству по кибербезопасности и защите инфраструктуры — создать центр по обеспечению безопасности ИИ для координации сканирования уязвимостей и установки исправлений в критически важной инфраструктуре. Администрация поддерживает дерегулирование наряду с надзором за ИИ в интересах национальной безопасности.

## 9. Создание инфраструктуры ИИ

В 2026 году производители чипов и лаборатории, занимающиеся разработкой ИИ, укрепили связи, чтобы успевать за выпуском моделей. Мы также стали

процессоры vcpu и память нового поколения. Тем временем мощности центров обработки данных по всему миру увеличились. Это свидетельствует о том, что вычислительные ресурсы, энергопотребление и аппаратное обеспечение теперь являются сдерживающими факторами для развития ИИ.

#### TRENDING STORIES

1. Microsoft, AWS and Anthropic are spending billions – and not on better models
2. Taming the agentic influx: a blueprint for AI business observability
3. Forward deployed engineer is AI's hottest job as OpenAI and Google race to hire. Here's how to become one.
4. 10 moments that defined AI's turbulent first half of 2026
5. Why Anthropic just doubled Claude Cowork limits at no charge

## 8. The rise of the harness

“The harness is where the hard work is,” Harness CEO and founder Jyoti Bansal **told The New Stack** last month. As base models are performing closely on benchmarks, the harness becomes the differentiator. The harness is the tools, memory, orchestration, and guardrails wrapped around a model. The harness determines whether an agent stays on task, recovers from errors, and performs safely. The harness helps shift competitive advantage from model capability to system design.

## 7. Tokenomics

Spend is now the battleground for both AI producers and consumers. AI labs are restructuring pricing around compute consumption rather than flat subscription, and companies are looking for ways to cut their token spend. Last month, the **Linux Foundation launched the Tokenomics Foundation** with support from Google, Microsoft, IBM, JPMorgan Chase, KPMG, Oracle, and Salesforce. The organization is tasked with establishing open standards, benchmarks, and best practices across the entire AI token economy.

## 6. Agentic AI goes mainstream

agent, Claude's tool use and multistep coding runs, and Google's autonomous information agents now run continuously in the background rather than on command. Meanwhile, enterprises are adding agents into real workflows such as monitoring, code review, procurement, and customer support. However, this shift may carry security risks as agents gain access to data and systems.

## 5. The Pentagon goes to war with Anthropic

In February, Department of War Secretary Pete Hegseth summoned Anthropic CEO **Dario Amodei** to his office to **demand that the military be granted unrestricted use** of the company's technology. But Amodei held his ground and refused to allow the military to use Anthropic technology for mass surveillance on citizens or for autonomous weapons. Following that, President Trump ordered federal agencies to phase out the use of Anthropic, and Hegseth designated the company a "supply chain risk," which is a label previously used for foreign adversaries – basically blocking the company. Anthropic sued in federal court, claiming the government's move was unfounded and retaliatory. The company received a preliminary injunction from a San Francisco court, holding that the government's actions constituted unconstitutional First Amendment retaliation.

Meanwhile, amid the initial fallout from Anthropic's battle with the Pentagon, **OpenAI made its own deal with the military.**

## 4. AI titans plant their flags on Wall Street

Within 72 hours in May, Anthropic and OpenAI each launched enterprise deployment arms, announced major **financial services partnerships**, and shipped agent tooling targeting Wall Street workflows. The message was the same – the next phase of frontier AI is not about models. It's about deployment.

Anthropic's new services firm – backed by Blackstone and Hellman & Friedman alongside General Atlantic, Apollo, Goldman Sachs, and Sequoia Capital – targets mid-sized enterprises that the large consulting and systems integration firms don't prioritize. These include community banks, regional health systems, and mid-market manufacturers. Applied AI engineers from Anthropic embed directly with clients alongside the new firm's own engineering staff, doing workflow discovery, building custom Claude-powered solutions, and supporting clients long-term.

OpenAI's Deployment Company – "DeployCo" – operates one market segment up, targeting large enterprises with the same forward-deployed engineering model. Its

---

initial investment and a partner roster that includes McKinsey, Bain & Company, and Capgemini.

Meanwhile, both companies are considering IPOs with valuations exceeding \$800 billion.

And both companies are betting on the same thesis: that the deployment gap — the widening distance between what frontier AI can do and what enterprises have actually shipped — is the next major revenue opportunity. And both moved on it in the same week.

### 3. Open-weight models are coming

Chinese labs continued to close the gap with Western frontier labs in 2026. Alibaba's Qwen, Zai's GLM and Moonshot's Kimi, delivering open-weight releases that rivaled closed models on standard benchmarks. Zai's **GLM-5.2**, released June 13, beat Anthropic's Claude Opus 4.8 in some benchmarks — showing the highest marks for open-weight models. GLM-5.2 is also one-fifth the price of comparable closed models.

"The industry is all focused on which lab has the smartest model, but that focus looks to the past," David Mytton, CEO of Arcjet, told *The New Stack*. "GLM-5.2's capability indicates that usage of open-source models is about to explode. "This all seems obvious in retrospect: models became more capable at the start of the year, agents started taking real actions (particularly since most of the work happens following the chat prompt), and legal restrictions on using frontier models are causing people to look elsewhere. This will cause all sorts of security issues because managing so many model capabilities will become difficult."

Paul Sawers, contributing writer at *The New Stack*, tells us for this countdown: "Budget open-weight model panels are now matching frontier proprietary benchmarks at a fraction of the cost, undercutting the case for paying top dollar for a single closed model. Example of startup ditching Anthropic for DeepSeek."

Indeed, the European AI agent startup Lindy AI migrated 100% of its production traffic from Anthropic to DeepSeek, citing millions of dollars in savings,

### CEOs vibe coding their own tools



---

chain of command — is adopting these tools to **vibe code** a variety of agents and productivity applications.

The trend runs from simple workflow automations to full production systems serving hundreds of users. The tools are **Claude**, **Cursor**, and, increasingly, the AI features embedded in the platforms these executives already run. The motivations range from impatience with IT queues to genuine curiosity about what the technology can do. And the results are more varied than the enthusiasm surrounding them might suggest.

**Woodson Martin**, CEO of **OutSystems**, took a more structured approach to his own vibe coding experiment. He built a personal mobile app wrapper on top of **MCP services** his team had created — and he built it twice in parallel, once using OutSystems' own AI coding tool, **Mentor**, and once using Claude, connecting to the same backend both times.

"I was tired of explaining it to somebody who was supposed to build it for me," **Martin told *The New Stack* in April**. "I was just like, 'I'll do this myself.'"

The app is a personal chief-of-staff system that consolidates customer account intelligence — buying signals, website activity, internal data — into a pre-meeting briefing he can pull up on his phone. It replaced what had been a 45-minute PowerPoint session plus multiple prep meetings from his sales team.

## 1. The government cracks down (and later relents) on Anthropic Fable 5 and Mythos 5

The **Fable/Mythos takedown** illustrated how unpredictable AI policy has become. Anthropic launched Fable 5, its only Mythos-tier model, on Jun 9, 2026, with the fuller Mythos 5 reserved for a small set of trusted customers under what the company called Project Glasswing. The rollout lasted only three days.

On June 12, Commerce Secretary Howard Lutnick sent Anthropic CEO Dario Amodei a directive ordering the immediate worldwide suspension of both models for all foreign nationals, including Anthropic's own non-citizen employees.

Apparently, the trigger was a jailbreak that Amazon researchers found, which could use the models' cybersecurity capabilities. This raised concerns at Commerce. Anthropic said it didn't have a way to restrict access by nationality in real time, so it

weeks, but the freeze didn't hit until June 30 — after an 18-day period of limbo.

Meanwhile, Anthropic added extra cybersecurity safeguards to Fable 5 and began restoring global access on July 1.

Having mentioned Anthropic's run-in with the Pentagon, Frederic Lardinois of *The New Stack* tells us: "While the two situations are not directly linked, it's hard not to read the Fable controls as an extension of this existing animosity between Anthropic and the Trump administration — and in part, this seems **personal as well**."

## What's next

What's in store for the second half of 2026? We'll be tracking how much AI-generated code actually makes it into production — and the tools designed to close that gap; the expanding autonomy of agentic AI and the guardrails keeping pace with it; the fast-moving regulatory environment around the frontier labs; and enterprise adoption of open-weight models.

We're also watching AI's spread among knowledge workers, the ever-longer stretches agents run unsupervised, and the hiring and performance of forward deployed engineers.

Whatever happens this year, *The New Stack* will cover it. And if you haven't already, **subscribe to *The New Stack Daily newsletter*** for timely, thoughtful updates on how AI is reshaping software development.

TNS



Darryl K. Taft covers DevOps, software development tools and developer-related issues from his office in the Baltimore area. He has more than 25 years of experience in the business and is always looking for the next scoop. He has worked...

Read more from Darryl K. Taft →

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY



---

models

2. Taming the agentic influx: a blueprint for AI business observability
3. Forward deployed engineer is AI's hottest job as OpenAI and Google race to hire. Here's how to become one.
4. 10 moments that defined AI's turbulent first half of 2026
5. Why Anthropic just doubled Claude Cowork limits at no charge

INSIGHTS FROM OUR SPONSORS



## Build Azure Infrastructure Using AI and Terraform

11 June 2025

## GenAI Made Terraform More Relevant Than Ever

1 April 2025

## The Deployment Bottleneck No One Talks About

26 March 2025



VMware Tanzu

## From Instance to Fleet: Enterprise-Scale Service Management in Tanzu Platform 10.4

26 June 2026

## USSFCU Secures Financial Data on a Sovereign Private Cloud

25 June 2026

## July 22, 2026: Demystifying the Agentic Loop: Building Safe, Server-Driven Agents with Spring

22 June 2026

## PagerDuty

## ServiceNow Runs Your IT. PagerDuty Makes Sure It Never Stops. by Laura Chu

4 July 2026

## Make the most of shift-based schedules by Mandi Walls

3 July 2026

## Why Faster Recovery Beats Faster Shipping in the AI Era by PagerDuty

2 July 2026

# Receive a free roundup of the most recent TNS articles in your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

## ARCHITECTURE

Cloud Native Ecosystem  
Containers  
Databases  
Edge Computing  
Infrastructure as Code  
Linux  
Microservices  
Open Source  
Networking  
Storage

## OPERATIONS

DevOps  
CI/CD

## ENGINEERING

AI  
AI Engineering  
API Management  
Backend development  
Data  
Frontend Development  
Large Language Models  
Security  
Software Development  
WebAssembly

## CHANNELS

Podcasts  
Ebooks

Kubernetes

Observability

Operations

Platform Engineering

Newsletter

TNS RSS Feeds

THE NEW STACK

About / Contact

Sponsors

Advertise With Us

Contributions

 roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap

Backend Developer Roadmap

Devops Roadmap

FOLLOW TNS



© The New Stack 2026

Disclosures

Terms of Use

Advertising Terms & Conditions

Privacy Policy

Cookie Policy