

Независимый анализ ИИ

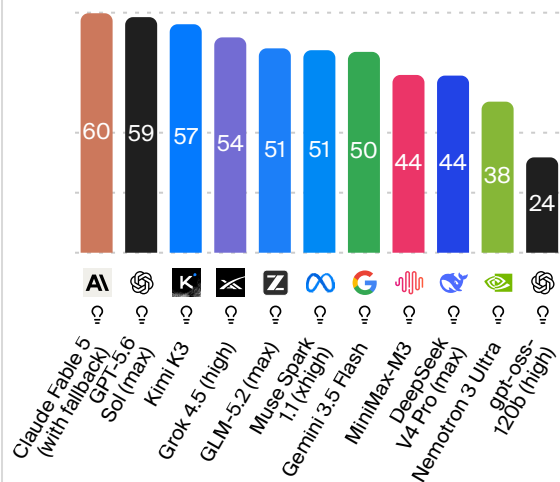
Изучите ландшафт ИИ, чтобы выбрать лучшую модель и поставщика для вашего варианта использования

Основные моменты

Интеллект

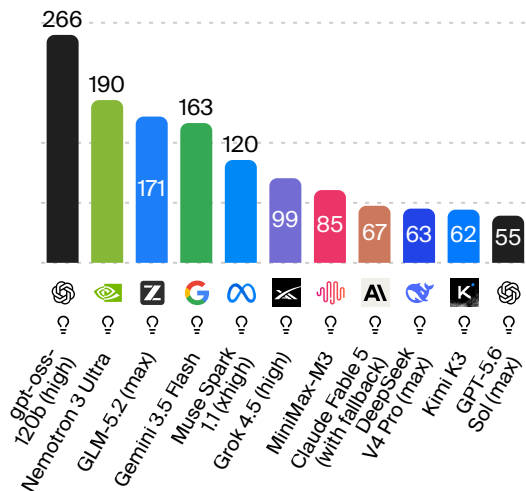
Обновлено

Индекс интеллекта искусственного анализа · Чем выше, тем лучше



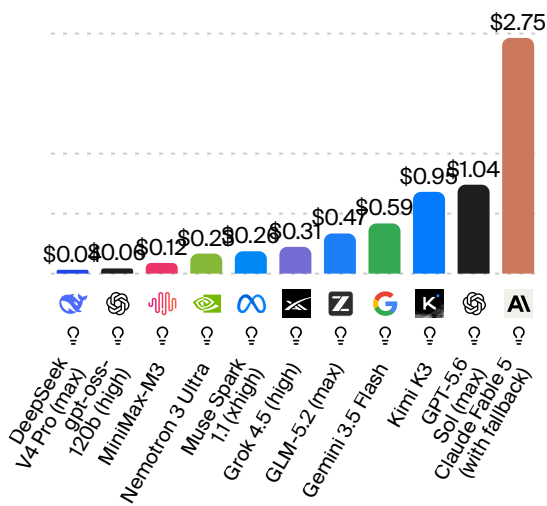
Скорость

Вывод токенов в секунду · Чем выше, тем лучше



Стоимость за задачу

Средневзвешенная стоимость (USD) за задачу с индексом интеллекта · Чем ниже, тем лучше



Персонализированный рекомендатель модели

Получайте персонализированные рекомендации, основанные на ваших приоритетах в отношении интеллекта, скорости и стоимости

Изучите агентов для общей работы, программирования, поддержки клиентов и многого другого

Сравните агентов искусственного интеллекта по возможностям, ценам и поддержке платформы

Изучите премиальные планы

Получите доступ к расширенным контрольным данным, пользовательским визуализациям, отраслевым отчетам и многому другому

Журнал изменений



Оценка новой языковой модели · 16 июля

 Кими К3

Опубликована новая статья · 15 июля

 Компания Thinking Machines выпустила Inkling, новую ведущую весовую модель Открытого чемпионата США по теннису

Интеллект

Обновлено

Интеллект ведущих моделей ИИ по результатам наших независимых оценок

Индекс искусственного интеллекта

Обновлено

Индекс кодирования

Индекс агентности

Индекс искусственного интеллекта

Индекс искусственного аналитического интеллекта v4.1 включает 9 оценок: GDPval-AA v2, τ3-Банковское дело, терминал-Бенч v2.1, SciCode, Последний экзамен человечества, GPQA Diamond, CritPt, AA-Всеведение, AA-LCR

↶

📌


📊

28 из 576 моделей

⬆

⌵

+ Добавьте модель от конкретного поставщика



💡 Обосновывающие модели обозначены значком лампочки

Индекс интеллекта искусственного анализа по Открытым Весам / Проприетарный

Индекс искусственного аналитического интеллекта v4.1 включает 9 оценок: GDPval-AA v2, τ 3-Банковское дело, терминал-Бенч v2.1, SciCode, Последний экзамен человечества, GPQA Diamond, CritPt, AA-Всеведение, AA-LCR

↩

🔍

📊

28 из 576 моделей

⬆


⚙

+

Добавьте модель от определенного поставщика

Запатентованные открытые веса

Открытые веса (коммерческое использование ограничено)



💡 Модели рассуждений обозначены значком лампочки

- 📘

Индекс интеллекта искусственного анализа

+
- 📘

Open Weights

+

Стоимость задачи по индексу интеллектуальности

Средневзвешенная стоимость (в долларах США) задачи по индексу интеллектуальности искусственного анализа, сегментированная по типу токенов.

Чем ниже, тем лучше




28 из 576 моделей



Ответ Обоснование Запись в кэш Попадание в кэш Вход

Artificial Analysis



Модели рассуждений обозначены значком лампочки

Стоимость задачи по индексу интеллектуальности



Интеллект языковой модели Frontier в динамике


Индекс искусственного аналитического интеллекта v4.1 включает 9 оценок: GDPval-AA v2, τ 3-Банковское дело, терминал-Бенч v2.1, SciCode, Последний экзамен человечества, GPQA Diamond, CritPt, AA-Всеведение, AA-LCR



15 of 54 model creators



Anthropic OpenAI Kimi SpaceXAI Z AI Meta Google Alibaba
MiniMax DeepSeek Xiaomi Thinking Machines Mistral Cohere
MBZUAI Institute of Foundation Models



Reasoning models are indicated by a lightbulb icon.

Artificial Analysis Intelligence Index



Artificial Analysis Coding Agent Index

Performance, cost, and execution time for leading coding agents on end-to-end software engineering tasks

Explore Artificial Analysis Coding Agent Index

Index

Cost

Execution Time

Artificial Analysis Coding Agent Index

Composite average pass@1 across DeepSWE, Terminal-Bench v2, and SWE-Atlas-QnA · Higher is better



Color by

Model

Agent

14 of 43 models






Image & Video Leaderboards

Top models from our Image Arena and Video Arena leaderboards, with 95% confidence intervals

Text to Image

Image Editing

Text to Video


Image to Video

Video Editing

Text to Image Leaderboard

Elo scores from blind preference votes in our Image Arena. [See the full leaderboard here.](#)

15 of 151 models



Speech Leaderboards

Top models from our Text to Speech Arena, Speech to Text and Speech to Speech evaluations

Text to Speech Arena Elo

AA-WER Index (Non-Streaming)

AA-WER Streaming Index (Fi

Искусственный Анализ

Премиум

Войти

Q

≡

Text to Speech Arena Leaderboard

Elo scores from blind preference votes in our Text to Speech Arena · [See the full leaderboard here.](#)


↶

↷

📊

15 of 85 models

⌵



Quality Elo



Capability Indices

New

See more ↗


Measures the performance of models on specific capabilities and industries

- Agentic
- Coding
- Finance & Accounting
- Strategy & Ops
- Legal
- Healthcare & Medical
- Engineering
- Economics

Artificial Analysis Agentic Index ↗

Measures performance in agentic workflows, focusing on behaviors like tool use, planning, autonomy, and complex problem solving.

27 of 149 models



💡 Reasoning models are indicated by a lightbulb icon

Intelligence Breakdown

Intelligence Evaluations

Intelligence evaluations measured independently by Artificial Analysis · Higher is better

↶

↷

📊

18 of 22 evaluations

⌵

28 of 576 models


⌵

+

 Add model from specific provider


GDPval-AA v2 ↗

Agentic real-world work tasks, (Elo-500)/2000




Terminal-Bench v2.1 ↗

Agentic coding & terminal use


SciCode

Coding




GPQA Diamond ↗


Scientific reasoning



CritPt ↗


Physics reasoning

AA-Omniscience Accuracy [↗](#)




AA-Omniscience Non-Hallucination Rate ↗

1 - hallucination rate





AA-LCR ↗

Long context reasoning




AA-Briefcase ↗ NEW

Agentic knowledge work, Elo




Harvey LAB-AA NEW

Legal agentic work, task all-pass rate




EnterpriseOps-Gym-AA NEW

Agentic business operations




APEX-Agents-AA [↗](#)

Long-horizon agentic tasks



ITBench-AA [↗](#)

Kubernetes incident root-cause analysis



💡 Reasoning models are indicated by a lightbulb icon.

📄 Intelligence Evaluation Relevance +

📄 Artificial Analysis Intelligence Index +

AA-Briefcase New

AA-Briefcase is a frontier agentic evaluation for long-horizon knowledge work, testing agents on realistic business workflows that require deliverables such as spreadsheets, presentations, and memos

AA-Briefcase Elo

AA-Briefcase Rubric Score (%)

Analytical Quality & Presentation Elo

AA-Briefcase Elo

AA-Briefcase is an agentic knowledge work benchmark developed by Artificial Analysis. AA-Briefcase Elo is a combined metric that aggregates rubric pass rate, analytical quality Elo and presentation Elo · Higher is better

↶

↷

📄


25 of 49 models

⬆

⬆

+ Add model from specific provider

Artificial Analysis



Reasoning models are indicated by a lightbulb icon

i AA-Briefcase Elo

+

AA-Omniscience

AA-Omniscience is a knowledge and hallucination benchmark that rewards accuracy, punishes bad guesses and provides a comprehensive view of which models produce factually reliable outputs across different domains

AA-Omniscience Index

AA-Omniscience Accuracy

AA-Omniscience Hallucination Rat

AA-Omniscience Index

AA-Omniscience Index (higher is better) measures knowledge reliability and hallucination. It rewards correct answers, penalizes hallucinations, and has no penalty for refusing to answer. Scores range from -100 to 100, where 0 means as many correct as incorrect answers, and negative scores mean more incorrect than correct.




27 of 446 models



+ Add model from specific provider

Artificial Analysis



Reasoning models are indicated by a lightbulb icon

AA-Omniscience Index

+

GDPval-AA v2

GDPval-AA v2 evaluates AI models on real-world, economically valuable tasks across a wide range of occupations

GDPval-AA v2 Leaderboard

Elo rating for performance on real-world work tasks · Anchored to a human baseline of 1,000 · Higher is better




27 of 162 models



+ Add model from specific provider

----- Human Baseline (1,000)



💡 Reasoning models are indicated by a lightbulb icon

Artificial Analysis Openness Index

Artificial Analysis Openness Index assesses how 'open' models are on the basis of their availability and transparency across different components.

Openness Index Components

Openness Index

Artificial Analysis Openness Index: Components

Openness Index underlying score contribution by components, up to a maximum of 18 (higher is more open)



13 of 267 models




+ Add model from specific provider

Transparency - Pre-training Data

Transparency - Post-training Data

Transparency - Methodology

Model Availability



💡 Reasoning models are indicated by a lightbulb icon

Artificial Analysis Openness Index vs. Artificial Analysis Intelligence Index

↔️ ↺️ 📊


13 of 267 models ⬆️ ⚙️

+

Add model from specific provider

Most attractive quadrant

Z AI MiniMax DeepSeek Kimi Xiaomi NVIDIA Mistral Anthropic
Google OpenAI Cohere MBZUAI Institute of Foundation Models



Output Tokens

Output tokens of leading AI models based on our independent evaluations

Output Tokens per Task

Intelligence vs. Output Tokens per Task

Output Tokens

Output Tokens per Intelligence Index Task

Weighted average number of output tokens used to run one task in the Artificial Analysis Intelligence Index




28 of 576 models



Answer

Reasoning



Reasoning models are indicated by a lightbulb icon

Output Tokens per Intelligence Index Task

Price and Cost

Price and real-world costs of leading AI models based on our independent evaluations

Cost per Task

Intelligence vs. Cost per Task

Evaluation Breakdown

Cost per Intelligence Index Task

Weighted average cost (USD) per Artificial Analysis Intelligence Index task, segmented by token type. Lower is better

28 of 576 models




Answer

Reasoning

Cache Write

Cache Hit

Input



Reasoning models are indicated by a lightbulb icon

Cost per Intelligence Index Task

Total Cost Intelligence vs. Total Cost Intelligence vs. Compute

Cost to Run Artificial Analysis Intelligence Index

Cost (USD) to run all evaluations in the Artificial Analysis Intelligence Index

↔

📷


📊

28 of 576 models

+

Add model from specific provider

Output Reasoning Cache Write Cache Read Non-Cache Input



💡 Reasoning models are indicated by a lightbulb icon

Cost to Run Artificial Analysis Intelligence Index +

Cache Hit, Input, and Output Pricing Blended Price Blended Price (Stacked) CAC

Pricing: Cache Hit, Input, and Output

Price (USD per M Tokens)

↔

📷

📊


28 of 576 models

⬆

⬆

+ Add model from specific provider

Cache Hit Input Output



Reasoning models are indicated by a lightbulb icon

- Cache Hit
- Input Price
- Cache Pricing by Provider
- Output Price
- Model Performance Representation

Speed & Latency

Comparison of first-party API performance


Speed	Latency	End-to-End Response Time	Intelligence vs. Speed
Output Speed Output tokens per second · Higher is better			

28 of 576 models



+ Add model from specific provider

Artificial Analysis



💡 Reasoning models are indicated by a lightbulb icon

ⓘ Output Speed



ⓘ Model Performance Representation



Time per Task

Intelligence vs. Time per Task

Cost vs. Time per Task


Time per Intelligence Index Task

Weighted average decode time (minutes) per task; excludes TTFT and overhead time · Lower is better



28 of 576 models





Reasoning models are indicated by a lightbulb icon

Time per Intelligence Index Task

API Provider Performance

Output Speed vs. Price: gpt-oss-120b (high)


Output tokens per second · USD per 1M tokens (blended) · 10,000 input tokens

19 of 19 providers

⌵

Most attractive quadrant

- Amazon Azure Baseten Cerebras Cloudflare CoreWeave Databricks
- DeepInfra DeepInfra (Turbo) Fireworks Google Vertex Groq Makora
- Nebius (Base) Novita Parasail SambaNova Scaleway Together AI



Reasoning models are indicated by a lightbulb icon.

i	Emerging Provider Competition	+
i	Price	+
i	Output Speed	+
i	Median	+


Pricing (Cache Hit, Input, and Output): gpt-oss-120b (high)

Price (USD per M Tokens) · Lower is better · 10,000 input tokens



19 of 19 providers

Cache Hit Input Output



- Cache Hit +
- Input Price +
- Cache Pricing by Provider +
- Output Price +

Speed Speed Over Time


Output Speed: gpt-oss-120b (high)

Output speed: output tokens per second · 10,000 input tokens



19 of 19 providers






- ⓘ

Output Speed

+
- ⓘ

Model Performance Representation

+



Artificial Analysis

Get notified about new articles

What's your email?

Explore

- LLM Leaderboard
- Image Arena
- Video Arena
- AI Agents
- Evaluations

Company

- Methodology
- Services
- Contact
- Articles
- FAQ

