

[← Вернуться к блогам](#)
[Новости ИИ](#)

Новости ИИ на сегодня, 16 июля 2026 года: 15 главных событий

📅 16 июля 2026 года ⌚ 25 мин на чтение



Поделиться:



Компания Apple только что преодолела последнее нормативное препятствие на пути к запуску Apple Intelligence в Китае, и теперь за дело берутся модели Qwen от Alibaba. В тот же день один стартап уменьшил модель с 27 миллиардами параметров до 3,9 гигабайта, чтобы она работала на iPhone, компания OpenAI выставила на продажу физическую клавиатуру для программирования, а рабочие завода в Корее устроили забастовку из-за роботов. Все это

произошло за день до того, как 17 июля должны были состояться Gemini 3.5 Pro и Шанхайская всемирная конференция по искусственному интеллекту.

Вот 15 важных событий, произошедших 16 июля 2026 года, с указанием цифр, дат и честных оговорок. Чтобы быть в курсе всех релизов этого месяца, добавьте в закладки наш [раздел с новостями и тенденциями в сфере искусственного интеллекта](#).

1. Apple Intelligence Очищает Китай с помощью Qwen и Baidu от Alibaba

Apple Intelligence была зарегистрирована в Администрации киберпространства Китая - нормативное одобрение, необходимое Apple для внедрения своих функций искусственного интеллекта в материковый Китай, и она будет работать на моделях Qwen от Alibaba. Apple также работает с Baidu над дополнительными функциями для китайских пользователей iPhone. Одобрение, о котором сообщалось 15 июля, устраняет барьер, из-за которого функции искусственного интеллекта Apple были недоступны сотням миллионов китайских клиентов, в то время как конкуренты осуществляли поставки.

Стратегическая реальность такова: американский технологический гигант не может использовать собственные модели искусственного интеллекта в Китае, поэтому ему приходится арендовать китайскую модель. В Китае каждая крупная языковая модель должна быть зарегистрирована и одобрена перед публичным выпуском, а иностранные модели не проходят эту проверку. Именно поэтому Apple, самая дорогая компания в мире, использует Qwen от Alibaba в рамках Apple Intelligence для своего второго по величине рынка. Это самый явный признак того, что мир искусственного интеллекта разделяется на два лагеря — западный и китайский, и даже Apple вынуждена использовать местные модели для работы на местном рынке.

Дата запуска пока не определена, поскольку регистрация — это скорее этап, чем релиз. Но направление задано, и оно меняет конкурентную среду. Alibaba получает огромное признание, став поставщиком комплектующих для iPhone в Китае, и эта сделка показывает, насколько далеко продвинулись китайские модели. Мы следили за этой темой, когда Goldman Sachs начал рекомендовать их клиентам в нашем [обзоре новостей об искусственном интеллекте от 13 июля](#). Мое мнение: это самая важная бизнес-новость недели, потому что она показывает, что решающим фактором в том, кто выиграет в той или иной стране, становится доступ к рынку, а не качество моделей.

2. PrismML's Bonsai 27B Runs a 27B Model on an iPhone

PrismML released Bonsai 27B on July 14, a 27-billion-parameter model compressed to just 3.9 gigabytes that runs locally on an iPhone 17 Pro at 11 tokens per second, and many are calling it a DeepSeek moment for on-device AI. Built from Qwen3.6-27B using 1-bit and 1.58-bit ternary compression, the model is multimodal, handles reasoning, coding, and agentic tasks, and keeps more than 90 percent of full-precision performance in the 1-bit version and over 95 percent in the ternary build. It is free under the Apache 2.0 license.

The reason this matters is that a capable 27B-class model has never fit comfortably on a phone before. Running frontier-adjacent intelligence entirely on-device, with no cloud call, means privacy by default, zero per-token cost, and no internet dependency, which changes what a mobile app can do. PrismML CEO Babak Hassibi told CNBC that Apple and others have been evaluating the models for speed and energy efficiency on their hardware, which lands with extra weight the same week Apple secured its China AI approval. If a phone can run a 27B model well, a large slice of everyday AI work stops needing a data center at all.

The honest caveat is that 11 tokens per second is usable but not fast, and 1-bit compression does lose real capability on the hardest tasks, so this is

not a frontier-model replacement. But it does not need to be. Most on-device jobs, summarizing, drafting, classifying, answering, do not require Sol-level power, and Bonsai 27B shows those jobs can run locally today. For developers, the shift toward efficient local models is one we track alongside the cloud tools in our [AI coding tools hub](#).

3. Liquid AI's Antidoom Cuts Small-Model Failures From 22.9% to 1%

Liquid AI introduced a method called Antidoom that, when applied to the small Qwen3.5-4B model, cut its failure rate on a target task from 22.9 percent to 1 percent. The technique targets the biggest weakness of small on-device models, which is reliability: they are fast and cheap but tend to break on edge cases in ways that large cloud models do not. Antidoom closes most of that gap on the tested workload.

Pair this with Bonsai 27B from story 2 and a clear theme emerges: the frontier of AI in mid-July 2026 is not just getting bigger, it is getting smaller and more reliable at the edge. A 4-billion-parameter model that fails one time in a hundred instead of one time in four is genuinely deployable in production, and that is the threshold that turns a cheap local model from a toy into a tool. The economics are hard to ignore, since a reliable small model running on a user's own device costs the developer nothing per query, compared with the \$30 per million output tokens of a flagship like GPT-5.6 Sol.

The caveat, as always with a single benchmark, is that a 22.9-to-1 improvement on one task does not guarantee the same gain everywhere, and independent testing across varied workloads is what will prove the method. Still, reliability has been the wall keeping small models out of serious products, and any credible progress on it matters more than another point of benchmark score on a giant model. This is the unglamorous engineering that quietly reshapes what ships.

4. The On-Device AI Shift Becomes the Real Story of the Week

Taken together, Bonsai 27B and Antidoom mark a turning point: July 2026 is when running capable AI on your own phone or laptop stopped being a demo and started being practical. The pattern is bigger than two releases. It represents a migration of a real share of AI work from the cloud back to the device, driven by three forces at once, which are cost, privacy, and independence from a network connection.

Why this reshapes the industry is a matter of who pays and who controls. Cloud AI means every query flows through a frontier lab that meters it, logs it, and bills for it. On-device AI means the query never leaves your hardware, which is better for privacy and free at the margin. For the labs, that is a threat to the token-metering business model that funds everything; for users and app developers, it is liberation. Apple's China deal in story 1 is the same theme from a different angle, since a model that runs locally is also a model that satisfies a regulator worried about data leaving the country.

My honest read is that the future is hybrid, not either-or. Frontier reasoning that demands maximum capability stays in the cloud, while the enormous volume of routine AI work, the summaries, the drafts, the classifications, moves on-device where it is cheaper and more private. The companies that win the next phase will be the ones that route each task to the right place automatically. If you are building this kind of routing, the patterns in our [open-source Gen AI cookbooks](#) are a practical starting point.

5. OpenAI's Codex Micro Puts a \$230 Coding Keypad on Your Desk

OpenAI, in collaboration with hardware maker Work Louder, launched Codex Micro, a \$230 desktop keypad built for its Codex coding agent, featuring backlit keys, a rotary knob, and a tiny joystick. It is OpenAI's first real consumer hardware accessory, a physical control surface

designed to make working with an AI coding agent feel more tactile and immediate than a chat box in a browser tab.

A dedicated keypad for an AI agent is a small product with a big signal inside it. It says OpenAI believes AI coding is becoming its own workflow worthy of dedicated hardware, the way video editing or music production earned their own control surfaces. It also fits OpenAI's broader hardware ambitions, from its acquisition of Jony Ive's design startup to the consumer devices at the center of Apple's lawsuit against it. A \$230 keypad is not the device that ambition is ultimately about, but it is a toe in the water of selling physical objects, not just tokens.

The skeptical view is that this is a niche gadget that most developers will happily skip, and that is probably right. But niche accessories are how hardware companies learn to make hardware, and OpenAI is clearly practicing. Whether Codex Micro sells well matters far less than what it reveals about a software company steadily building the muscles to ship physical products. Keep an eye on where this leads rather than on the keypad itself.

 **COHORT WAITLIST OPEN**

Go From AI User to AI Builder

Don't just use ChatGPT. Learn to build custom LLM agents, RAG pipelines, and full-stack Agentic AI apps in our intensive 6-week program.

-
- | | |
|-------------------------------|---------------------------------|
| ✓ 6 Weeks Live Mentorship | ✓ Deploy 5+ Real-world Apps |
| ✓ Weekly App Templates & Code | ✓ No Coding Experience Required |
-

[Explore Program](#)

Join 1,000+ graduates • Free Registration

6. ASML Moves to Raise Chip-Tool Prices as TSMC Pushes Back

ASML, the Dutch company that makes the extreme ultraviolet lithography machines required to build the most advanced chips, has discussed raising prices for its EUV systems with TSMC and plans to charge about 10 percent more for its older DUV systems, with TSMC resisting the increases. ASML is the only company in the world that makes EUV machines, which makes this a negotiation between two irreplaceable links in the AI supply chain.

This is a rare glimpse of pricing power at the very bottom of the AI stack. Every advanced AI chip, from Nvidia's accelerators to Apple's silicon, is etched by an ASML machine, and every one of those machines costs hundreds of millions of dollars. When ASML raises prices and even TSMC, the world's most important chipmaker, has to push back, it shows how much leverage sits with the handful of companies that own irreplaceable technology. The cost eventually flows downstream, since pricier tools mean pricier chips, which mean pricier AI compute for everyone building on top.

The bigger point is that the AI boom rests on an extraordinarily narrow foundation. One company makes the EUV machines, one company dominates leading-edge fabrication, one company leads high-bandwidth memory. Each is a potential single point of failure and a holder of real pricing power, which is exactly why nations are pouring hundreds of billions into building alternatives. This quiet supplier negotiation is the AI economy showing its plumbing.

7. Apple Hunts Chip Acquisitions for Its Own AI Server Silicon

Apple is reportedly seeking to acquire chip companies to accelerate its effort to build custom AI server chips, a move that would extend its industry-leading silicon design from phones and laptops into the data center. Apple already designs the best mobile chips in the world through its A-series and M-series lines, and turning that expertise toward AI

server silicon would reduce its dependence on Nvidia while giving it hardware tuned to its own models.

The timing connects several threads from this week. Apple just secured its China AI approval in story 1, it is being evaluated as a customer for on-device models in story 2, and now it wants its own server chips too. The strategy is vertical integration end to end, from the phone in a customer's hand to the data center behind Apple Intelligence, all running Apple-designed silicon. It puts Apple on the same path as Google, Amazon, Meta, and OpenAI, each of which is building custom chips to escape Nvidia's pricing and tune hardware to their software.

The catch is that data-center chips are a different discipline from mobile chips, and buying expertise through acquisition is faster but riskier than building it. Apple has the balance sheet to buy almost anyone and the design pedigree to make it work, so this is a credible threat to the merchant-silicon status quo. If Apple ships competitive AI server chips, it changes the economics of running Apple Intelligence and pressures Nvidia at the same time. This is a multi-year bet worth watching.

8. Hyundai Workers Strike Over Wages, AI, and Humanoid Robots

Hyundai auto workers launched a partial strike over wages, AI deployment, and the introduction of humanoid robots on the factory floor, making it one of the first major labor actions where AI and robots are named explicitly as grievances. The walkout signals that the tension between automation and employment, long discussed in white-collar terms, has arrived on the industrial shop floor where robots directly replace physical labor.

This connects to a growing pattern of workers responding to AI, from the Mews layoffs blamed openly on AI efficiency to the survey showing most US workers want AI profits shared, both covered in our [July 15 AI news recap](#). What makes the Hyundai action distinct is that it is blue-collar and it is specifically about physical robots, the same humanoid machines

that South Korea just pledged to grow from 1 percent to 20 percent of its market by 2028. When a government subsidizes robots and a company deploys them, the workers whose jobs are on the line are the ones left to negotiate, and increasingly they are choosing to strike.

The honest tension is that automation genuinely raises productivity and genuinely displaces workers, and both facts are true at once. Korea's national bet on humanoid robots and its workers' strike against them are two sides of the same policy, and how that gets resolved, through retraining, profit-sharing, slower deployment, or conflict, will be a template other industrial economies watch closely. The robot question stopped being theoretical the moment the machines showed up on the line.

9. Spectro Cloud Raises \$100 Million at a \$1 Billion-Plus Valuation

AI infrastructure company Spectro Cloud raised a \$100 million Series D at a valuation above \$1 billion, up from \$750 million in 2024, with backing from AMD, LG, and others. Spectro Cloud helps companies manage Kubernetes and edge infrastructure, the unglamorous plumbing that runs AI workloads across data centers and devices, and the round reflects strong investor appetite for the layer beneath the models.

The interesting detail is who is investing. AMD backing an infrastructure company is a chipmaker betting on the software that will orchestrate its hardware, and LG signals interest from a major electronics manufacturer eyeing edge AI. The round fits the week's on-device theme, since managing AI that runs across both cloud and edge is exactly the problem Spectro Cloud addresses, and it is a problem that gets harder as models like Bonsai 27B push more work onto local devices.

Infrastructure that spans cloud and edge is quietly becoming essential.

A \$100 million round for infrastructure rarely makes headlines next to a model launch, but it is a useful signal of where smart money sees durable value. The models get the attention while the plumbing gets the

recurring revenue, and investors know it. As AI deployments sprawl across clouds, data centers, and now phones, the companies that manage that sprawl are positioned to collect regardless of which model wins.

 **COHORT PROGRAM OPEN**

Claude Mastery: *Cowork & Code*

The only comprehensive program designed to take you from basic prompting to building interactive Artifacts, custom integrations, and deploying production-ready code with Claude Code.

✓ No coding experience needed

✓ Build interactive Artifacts & Agents

✓ Deploy apps with Claude Code

✓ Cohort-based learning & mentorship

Explore Program

Join Waitlist

Cohort-based training

• Register Now

10. Zeroth Raises \$73.6 Million for Humanoid Robots, Led by Ant Group

Humanoid robotics startup Zeroth raised approximately \$73.6 million in Series A funding led by Ant Group, the Alibaba-affiliated fintech giant, adding to a torrent of capital flowing into physical AI. The round places another well-funded competitor into a humanoid robotics field that has become one of the hottest and most crowded corners of AI investment in 2026.

Ant Group leading the round is the notable part. A major Chinese fintech backing humanoid robots shows how broadly the physical-AI thesis has spread beyond the obvious robotics specialists, and it fits the pattern of Chinese capital and companies moving aggressively into embodied AI. Between Unitree's IPO, Zeroth's raise, and South Korea's national robot target, the humanoid race is increasingly an Asian-led contest, with

Western players like Tesla and Boston Dynamics competing against a wall of well-capitalized regional rivals.

The same caveat that shadows every humanoid story applies here: no company has yet proven a humanoid robot pays for itself at scale, and a \$73.6 million Series A buys development runway, not a viable business. But capital at this volume, from backers this serious, is a bet that the economics will eventually close. Whether it does is the multibillion-dollar question hanging over the entire sector, and rounds like Zeroth's are the market wagering that the answer is yes.

11. Hemispheric Raises \$52 Million for Brain-Activity AI

Israel-based startup Hemispheric raised \$52 million for AI that analyzes brain activity, pushing artificial intelligence into neuroscience and medical diagnostics. The company applies machine learning to interpret neural signals, a field with applications spanning medical diagnosis, mental health monitoring, and eventually brain-computer interfaces, and the round reflects growing investor interest in AI applied to the hardest scientific domains.

This is a reminder that AI in 2026 is not only chatbots and coding agents. Applying machine learning to brain activity sits at the frontier of medical AI, where the potential upside, earlier diagnosis of neurological conditions, better understanding of mental health, is enormous, and where the stakes and the caution required are equally high. It fits a broader move of AI capital into science and health, alongside efforts like Anthropic's Claude Science push and the wave of AI-for-biology startups.

The responsible note is that brain-activity AI carries real privacy and ethical weight, since neural data is about as personal as data gets, and the field deserves scrutiny proportional to its sensitivity. A \$52 million round is early-stage validation, not a proven product, and medical AI faces a long road of regulation and clinical evidence before it reaches patients. But the direction is meaningful: AI is increasingly being pointed

at problems that matter far beyond productivity software, and this is one of them.

12. TikTok Shop's AI Creators Drive Toward \$23.4 Billion in US Sales

AI-generated creators and content are driving sales growth on TikTok Shop, which is projected to reach \$23.4 billion in US sales in 2026, as synthetic influencers and AI-assisted content reshape social commerce. The trend points to a future where a meaningful share of the creators selling products online are partly or entirely AI-generated, blurring the line between human influencer and software.

The commercial logic is powerful and a little unsettling. An AI creator never sleeps, never demands a fee, can be generated in any style, and can produce endless content optimized for engagement, which is a compelling proposition for brands chasing conversions. At the same time, it raises real questions about disclosure and trust, since a shopper may not know whether the enthusiastic reviewer recommending a product is a person or a generated character. A \$23.4 billion market moving in this direction makes those questions urgent rather than academic.

My take is that this is one of the most consequential and least discussed AI trends of 2026, because it touches how ordinary people are persuaded to spend money. The technology to generate convincing synthetic creators is here, the economic incentive to use them is overwhelming, and the rules around disclosure are barely written. However the industry handles transparency will shape whether AI creators become a useful tool or an erosion of trust in everything people see online.

13. KredosAI's Round and the Quiet Boom in Vertical AI Funding

KredosAI, an AI-powered debt-collections platform that uses behavioral intelligence to improve revenue recovery, closed a \$7 million Series A led

by BMW i Ventures, one of a steady stream of smaller vertical-AI rounds that rarely make headlines but collectively show where applied AI is taking hold. Alongside larger raises like Even Realities at \$150 million and Bespoke Labs at \$40 million earlier in July, the pattern is unmistakable: AI is being funded not just at the frontier but deep inside specific industries.

Vertical AI is where the technology quietly becomes a business. A collections platform, an insurance underwriter, a warehouse-inspection tool, these are not glamorous, but they solve concrete problems for customers willing to pay, and they are where a large share of AI's real economic value will be created. BMW i Ventures backing a collections startup is a corporate venture arm betting that behavioral AI improves a boring, essential process, and that is exactly the kind of unshowy application that adds up. In 2026, AI startups are attracting roughly a third of all venture funding, and much of that is flowing into these focused, industry-specific tools.

The contrast with the frontier is instructive. While a handful of labs soak up tens of billions, hundreds of vertical startups are quietly raising single-digit and double-digit millions to apply AI to one problem well. Both layers matter, but the vertical layer is where most businesses will actually encounter AI, embedded in the software they already use. The headline rounds get the attention, but the vertical boom is where the technology touches the real economy.

14. China's AI Map Redraws Around Apple, Alibaba, and Baidu

Apple's China approval is not an isolated deal but a marker of how thoroughly the Chinese AI market now runs on Chinese models, with Alibaba's Qwen and Baidu's systems becoming the mandatory partners for any foreign company that wants to offer AI there. The pattern is now the rule rather than the exception: to operate AI in China, you use a Chinese model, full stop, and the most powerful Western technology companies are no exception.

This bifurcation has been building all year, from the US restricting Chinese access to frontier models to China requiring domestic registration for every model, and Apple's deal crystallizes it. The world is settling into two AI ecosystems, one built on OpenAI, Anthropic, and Google, the other on Alibaba, Baidu, DeepSeek, and their peers, with a hard regulatory border between them. Companies that operate globally, like Apple, increasingly have to run one stack in the West and a different stack in China, which is a profound shift from the single global internet that technology grew up on.

The long-term implication is that AI is becoming an instrument and a reflection of geopolitics, not a borderless technology. Model quality still matters, but market access, regulatory approval, and national alignment increasingly decide who serves which billion people. For the companies caught in the middle, the cost is complexity and duplicated stacks; for the two blocs, it is a deepening technological divide. This is the quiet structural story underneath every flashy model launch, and Apple just made it impossible to ignore.

15. One Day to July 17: Gemini 3.5 Pro Meets the World AI Conference

July 17 is now one day away, and it remains the single biggest date on the AI calendar this summer, with Google's Gemini 3.5 Pro expected to launch the same day Shanghai's 2026 World Artificial Intelligence Conference opens with President Xi Jinping attending in person for the first time since the event began in 2018. One date, two hemispheres: the West's most anticipated model going live while the East's most powerful leader takes the world's biggest AI-governance stage.

The Gemini stakes are unchanged and unforgiving. The model is six weeks late, lands a week after GPT-5.6 and nine days after Grok 4.5, and carries a formidable leaked spec of a 2-million-token context window, Deep Think reasoning on the \$250 Ultra tier, and pricing near \$1.25 input and \$10 output per million tokens. It has to beat GPT-5.6 Sol on at least one headline benchmark, hold long-context recall at full length, and

actually ship on time after a bruising run of talent departures. Grok 4.5 already reset the value floor at \$2 and \$6, as our [Grok 4.5 hands-on review](#) details, so matching the field on price will not be enough.

The World AI Conference half signals the bigger shift the whole week has pointed toward. Between Apple running Chinese models in China, South Korea's \$880 billion plan, and Xi personally opening the Shanghai conference, the through-line of July 2026 is that AI has become a multi-superpower contest fought with models, chips, capital, and policy at once. When a frontier launch and a head-of-state AI summit share one calendar square on opposite sides of the planet, that is the whole year compressed into a day. We will have the full recap on our [best AI models July 2026 leaderboard](#) once the dust settles.

The July 16 On-Device vs Cloud Model Snapshot

July 16, 2026 is defined by AI moving in two directions at once: smaller and local, or larger and cloud-bound. Here is where the field sits, one day before Gemini 3.5 Pro.

On-device benchmarks and launch-week pricing claims both deserve independent verification, and this snapshot will shift the moment Gemini 3.5 Pro ships.

Frequently Asked Questions

Is Apple Intelligence available in China?

As of July 15, 2026, Apple Intelligence has been registered and approved by the Cyberspace Administration of China, using Alibaba's Qwen models, with Baidu also involved. This clears the main regulatory barrier, though Apple and regulators have not yet announced a public launch date for the features in China.

Can a 27B AI model run on a phone?

Yes. PrismML's Bonsai 27B, released July 14, 2026, compresses a 27-billion-parameter model to 3.9 gigabytes and runs on an iPhone 17 Pro at about 11 tokens per second. The 1-bit version keeps over 90 percent of full-precision performance, and the model is free under the Apache 2.0 license.

What is Bonsai 27B?

Bonsai 27B is a compressed, multimodal AI model from PrismML based on Qwen3.6-27B, built in 1-bit and 1.58-bit ternary versions to run locally on phones and laptops. At 3.9 gigabytes it fits on high-end mobile devices while handling reasoning, coding, and agentic tasks, and many observers called it a DeepSeek moment for on-device AI.

Which AI models does Apple use in China?

For Apple Intelligence in China, Apple is using Alibaba's Qwen models as the core, with Baidu contributing additional features. China requires all large language models to be registered domestically, which is why Apple relies on Chinese models there rather than its own or other Western systems.

Why is ASML raising prices?

ASML, the sole maker of EUV chip-manufacturing machines, has discussed raising EUV prices with TSMC and plans to charge about 10 percent more for its DUV systems, reflecting the pricing power of a company that has no competitor for its most advanced tools. TSMC is resisting the increases, and any rise eventually flows into the cost of advanced AI chips.

When will Gemini 3.5 Pro launch?

Leaked plans point to July 17, 2026, one day after this post and the same day Shanghai's World AI Conference opens. Expected specs include a 2-million-token context window, Deep Think reasoning on the \$250 per

month Ultra tier, and API pricing near \$1.25 input and \$10 output per million tokens. Google has not officially confirmed the date.

Recommended Blogs

- [AI News Today July 15 2026: 15 Biggest Stories](#)
- [AI News Today July 14 2026: 15 Biggest Stories](#)
- [AI News Today July 13 2026: 15 Biggest Stories](#)
- [Best AI Models July 2026: Full Ranked Leaderboard](#)
- [Grok 4.5 Review: xAI's Coding Model Tested](#)
- [AI Coding Tools 2026: The Complete Hub](#)

Resources & Community

Join our community of **70,000+** AI enthusiasts and learn to build powerful AI applications! Whether you're a beginner or an experienced developer, Build Fast with AI helps you understand and implement AI in your projects.

- [Website — buildfastwithai.com](#)
- [LinkedIn — Build Fast with AI](#)
- [Instagram — @buildfastwithai](#)
- [Founder Twitter — @satvikps](#)
- [Twitter — @BuildFastWithAI](#)

Agentic AI Launchpad 2026

A structured **6-week cohort program** that takes you from AI basics to building and deploying real-world agentic AI systems. Includes live sessions, expert mentorship, project reviews, and a builder community network.

Ready to go from learning to building? Join the next cohort → [Agentic AI Launchpad 2026](#)

Free AI Resources

Access free tools, workshops, and micro-learning to keep building:

- [AI Workshops - Free resources, upcoming events & past recordings](#)
- [Unrot - Learn AI in 5 minutes a day \(free micro-learning app\)](#)

Tomorrow is July 17, the biggest AI day of the year, with Gemini 3.5 Pro and the World AI Conference colliding. Follow Build Fast with AI and subscribe so the recap lands before your standup.

References

- [TechCrunch - Apple Intelligence Approved for China With Alibaba's Qwen](#)
- [9to5Mac - PrismML Releases Bonsai 27B Fit for iPhone](#)
- [MarkTechPost - Bonsai 27B 1-bit and Ternary Builds of Qwen3.6-27B](#)
- [LLM Stats - LLM News Today, July 2026](#)
- [The AI Insider - KredosAI Raises \\$7M Series A](#)
- [Crescendo AI - Latest VC Investment Deals in AI Startups](#)
- [TechCrunch - OpenAI Launches the GPT-5.6 Family](#)

- [Fortune - Anthropic Overtakes OpenAI on Revenue](#)

Enjoyed this article? Share it →

Share:



Ship Your First AI App

From zero to deployed app with our Agentic AI Launchpad

Start Building Today

0 Comments

Sort by



Sign in to join the conversation...

Ctrl+Enter to submit

Sign in to comment



Be the first to comment
Start the conversation below



AI That Keeps You Ahead

Get the latest AI insights, tools, and frameworks delivered to your inbox. Join builders who stay ahead of the curve.

your@email.com

Join 50,000+ builders

You Might Also Like



AI News

AI News Today - June 8, 2026: 16 Biggest Stories

16 AI stories: Apple WWDC 2026 recap, Claude becomes an iPhone option, Microsoft Foundry, EU AI Act countdown, Pentagon AI race, SpaceX IPO, and more.





AI News

AI News Today June 22 2026: 15 Biggest Stories

Fable 5 free trial closes, Noam Shazeer joins OpenAI, GLM-5.2 beats GPT-5.5, GPT-5.6 preview, Agentjacking attacks, and 10 more AI stories for June 22.

✦ Personalized Growth Engine

What's your AI Score?

Measure your AI readiness and unlock a personalized roadmap with curated tools, frameworks, and resources tailored to your role.

Check My AI Score →

✓ Takes 2 minutes ✓ Free forever ✓ Actionable advice