

ИИ / АГЕНТЫ ИСКУССТВЕННОГО ИНТЕЛЛЕКТА

Anthropic представляет Opus 4.6 с выдающимися результатами в решении сложных задач, которые не под силу другим ИИ

Новая флагманская модель Anthropic также используется в агентных командах Claude Code и теперь имеет контекстное окно на миллион токенов.

5 февраля 2026 г. в 9:45 [Фредерик Лардинуа](#)



В четверг компания Anthropic **выпустила** Opus 4.6 — обновленную версию своей флагманской модели Opus, которая значительно превосходит своего предшественника и многих конкурентов практически по всем показателям.

токенов и команды агентов в коде Claude, которые могут работать над задачами параллельно.

Цены остались прежними: 5 долларов за миллион входных/выходных токенов.

Качественный скачок для корпоративных пользователей

Компания утверждает, что Opus 4.6 — это качественный скачок в использовании больших языковых моделей (БЯМ) для корпоративных рабочих процессов благодаря их способности решать более сложные задачи и быстрее выдавать результаты.

Как сообщил представитель Anthropic изданию *The New Stack*, «с первой попытки качество становится гораздо ближе к готовому к использованию, чем у любой другой модели. Документы, электронные таблицы и презентации потребуют меньше итераций».

Anthropic отмечает, что Claude в Excel, например, теперь может выполнять более длительные и сложные задачи, а также вносить многоэтапные изменения за один проход.

The screenshot shows a PowerPoint presentation slide titled "European EV charging market value projected to reach \$15.1B by 2030". The slide features a stacked bar chart showing market size growth from 2024 to 2030, broken down by charger type: AC Charging and DC Fast Charging. The chart shows a significant increase in market value over the period.

On the right side of the slide, there is a sidebar with the Opus 4.6 interface. It displays a summary of the chart data and key market facts.

Chart: A stacked bar chart showing market size growth from 2024 to 2030, broken down by charger type:

- AC Charging:** Growing from \$3.4B (2024) to \$5.9B (2030)
- DC Fast Charging:** Growing from \$2.3B (2024) to \$9.2B (2030)
- Total market:** \$5.7B (2024) → \$15.1B (2030)

Key Market Facts:

- 21.5% CAGR expected 2025-2030
- Fast chargers driving highest growth
- Netherlands projected fastest CAGR
- Europe accounts for 14.7% of global market

The sidebar also includes a search bar with the text "What can I do for you?" and a button to "Accept all edits".

Opus 4.6 в PowerPoint (фото: Anthropic).


Zoom

результаты в бенчмарках по написанию кода, за исключением тестов, проверенных SWE-bench, и **бенчмарка MCP Atlas** для тестирования использования инструментов, в которых наблюдается незначительная регрессия. Это своего рода аномалия, особенно если учесть, что модель показывает чрезвычайно высокие результаты в аналогичных бенчмарках, изучающих агентное кодирование в терминале (**Terminal Bench 2.0**) и агентное использование инструментов (**t2-bench**).

TRENDING STORIES

1. **OpenAI's first gadget is the \$230 Codex Micro macropad**
2. **Amazon, Microsoft, and Google are converging on the same enterprise agent architecture**
3. **Anaconda buys Kilo, the open source coding agent that answers to no single model maker**
4. **The bottleneck for AI agents isn't the model anymore. It's the context layer.**
5. **OpenAI hits 8 million Codex users — what developers need to know**

On Terminal Bench, Opus 4.6 scores 65.4%, up from 59.8% for Opus 4.5, and on the OSWorld agentic computer use benchmark, its score rose from 66.3% to 72.7%. This now puts it ahead of OpenAI's GPT-5.2 and Google's Gemini 3 Pro, and according to Anthropic, the new model performs especially well at diagnosing more complex bugs.



Claude Opus 4.6 Terminal Bench benchmark (credit: Anthropic).

Zoom

Anthropic reports similar gains across benchmarks. The standout, though, is its score of 68.8% on the **ARC AGI 2** benchmark, which is less about achieving PhD-level performance in specialized tasks and more about solving problems that are easy for humans and very hard for AI systems. Opus 4.5 scored only 37.6%, while Gemini 3 Pro scored 45.1% and GPT-5.2 scored 54.2%.

Benchmarks, of course, only tell some of the story and don't always reflect how these models work in practice. Anthropic argues that, in its internal use, Opus 4.6 has handled far more challenging tasks, even without explicit instructions, and has done so more quickly and with better results.

In its safety evaluations, Anthropic found that Opus 4.6 matches Opus 4.5 in terms of misalignments such as deception, sycophancy, and encouraging user delusions.

Agentic terminal coding Terminal-Bench 2.0	65.4%	59.8%	51.0%	56.2% (54.2% self-reported)	64.7% (64% self-reported) (Codex CLI)
Agentic coding SWE-bench Verified	80.8%	80.9%	77.2%	76.2%	80.0%
Agentic computer use OSWorld	72.7%	66.3%	61.4%	—	—
Agentic tool use t2-bench	Retail 91.9%	Retail 88.9%	Retail 86.2%	Retail 85.3%	Retail 82.0%
	Telecom 99.3%	Telecom 98.2%	Telecom 98.0%	Telecom 98.0%	Telecom 98.7%
Scaled tool use MCP Atlas	59.5%	62.3%	43.8%	54.1%	60.6%
Agentic search BrowseComp	84.0%	67.8%	43.9%	59.2% (Deep Research)	77.9% (Pro)
Multi-discipline reasoning Humanity's Last Exam	40.0% without tools	30.8% without tools	17.7% without tools	37.5% without tools	36.6% without tools
	53.1% with tools	43.4% with tools	33.6% with tools	45.8% with tools	50.0% with tools (Pro)
Agentic financial analysis Finance Agent	60.7%	55.9%	54.2%	44.1%	56.6% (5.1)
Office tasks GDPVal-AA Elo	1606	1416	1277	1195	1462
Novel problem solving ARC AGI 2	68.8%	37.6%	13.6%	45.1% (Deep Thinking)	54.2% (Pro)
Graduate-level reasoning GPQA Diamond	91.3%	87.0%	83.4%	91.9%	93.2% (Pro)
Visual reasoning MMMU Pro	73.9% without tools	70.6% without tools	63.4% without tools	81.0% without tools	79.5% without tools
	77.3% with tools	73.9% with tools	68.9% with tools	— with tools	80.4% with tools
Multilingual Q&A MMMLU	91.1%	90.8%	89.5%	91.8%	89.6%

 Zoom

More features

Even though the version number change is small, there are other updates here beyond improved reasoning performance. Opus 4.6 is the first model in the Opus family to feature a one-million-token context window, for example.

It's also the first Anthropic model to use adaptive thinking, which allows it to consider contextual clues to determine how much effort to invest in a prompt.

Developers still get more control over this with the `/effort` parameter to make explicit tradeoffs between quality, inference speed, and cost. Previously, though,

For API users, Claude can now use compaction to summarize context, allowing it to handle longer-running tasks without hitting its context limits.

There's also a nod to digital sovereignty in this update. If your workloads can only run in the United States, that's now an option, but you will pay 10% more for it. More typically, we see companies offer this for workloads that *can't* run in the United States,

Agent teams

For developers, the most interesting new feature may be agent teams, though. While developers have found ways around this, by default, Claude Code has only run one agent at a time until now. Now, Anthropic is introducing agent teams, which allow developers to split work across agents. Those agents can then work in parallel and coordinate their efforts autonomously.

Anthropic notes that this is especially useful for read-heavy work, such as codebase reviews.

TNS



Before joining The New Stack as its senior editor for AI, Frederic was the enterprise editor at TechCrunch, where he covered everything from the rise of the cloud and the earliest days of Kubernetes to the advent of quantum computing....

[Read more from Frederic Lardinois →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY



TRENDING STORIES

1. OpenAI's first gadget is the \$230 Codex Micro macropad

3. Anaconda buys Kilo, the open source coding agent that answers to no single model maker
4. The bottleneck for AI agents isn't the model anymore. It's the context layer.
5. OpenAI hits 8 million Codex users — what developers need to know

INSIGHTS FROM OUR SPONSORS



Snowflake context and Meta campaigns: close the signal loop

21 July 2026

What's New in Snowflake Cortex Agents for Enterprise AI

21 July 2026

Snowflake CoCo: Built to Scale Enterprise AI with Trust

21 July 2026



Build Azure Infrastructure Using AI and Terraform

11 June 2025

GenAI Made Terraform More Relevant Than Ever

1 April 2025

The Deployment Bottleneck No One Talks About

26 March 2025



21 July 2026

Microsoft and Mistral expand strategic partnership to give enterprises and regulated industries frontier AI they can control

21 July 2026

Microsoft expands Azure AI and HPC infrastructure with AMD

20 July 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the most recent TNS articles in your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

Databases

Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

API Management

Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

 roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



[Privacy Policy](#)

[Cookie Policy](#)