

Избранное

Китайский Moonshot AI выпускает Kimi K3, крупнейшую в истории модель с открытым исходным кодом, конкурирующую с ведущими системами США

Michael Nuñez

12: 42 по североамериканскому времени, 16 июля 2026 г.



Фото: VentureBeat, сделанное совместно с Midjourney



[Moonshot AI](#), пекинский стартап в области искусственного интеллекта, поддерживаемый Alibaba, в четверг выпустил [Kimi K3](#) – модель с 2,8 триллионами параметров, которая, по словам компании, в настоящее время является крупнейшей моделью искусственного интеллекта с открытым исходным кодом в мире и которая, как показывают тесты, работает наравне с самыми мощными проприетарными системами от [Anthropic](#) и [OpenAI](#).

Релиз, приуроченный к началу [Всемирной конференции по искусственному интеллекту 2026](#) в Шанхае, является резким обострением глобальной гонки вооружений в области ИИ и переломным моментом для движения за ИИ с открытым исходным кодом. Он также знаменует собой впечатляющее возвращение компании, чьи позиции на рынке значительно ослабли за последние полтора года после стремительного взлёта DeepSeek.

Согласно информации, предоставленной исследователями, которые ознакомились с технической документацией компании, полные веса модели планируется опубликовать 27 июля. Если вы хотите опробовать [Kimi K3](#) прямо сейчас, то можете это сделать — просто зайдите на kimi.com, зарегистрируйтесь, указав аккаунт Google или номер телефона (кредитная карта не требуется), и начните общаться с, возможно, самой мощной моделью с открытым исходным кодом из когда-либо созданных.



В основе архитектуры, на которой работает крупнейшая в мире модель искусственного интеллекта с открытым исходным кодом

[Kimi K3](#), лежит большая языковая модель передового уровня с 2,8 триллионами параметров — примерно на 75 процентов больше, чем у [V4 Pro от DeepSeek](#), которой, согласно собственной временной шкале компании, около 1,6 триллионов параметров.

параметров. Модель имеет контекстное окно в 1 миллион токенов, встроенные возможности визуального восприятия и постоянно включенный режим логического вывода, который в компании называют «режимом мышления».

The model is built on two key architectural innovations developed internally at Moonshot AI: [Kimi Delta Attention](#), a hybrid linear attention mechanism, and [Attention Residuals](#), which the company describes as a drop-in replacement for residual connections that delivers consistent scaling gains. Both techniques were previously published as open research by the Moonshot team on [GitHub](#).

On the [API side](#), Kimi K3 is compatible with the [OpenAI SDK](#), lowering the integration barrier for developers already building on OpenAI or Anthropic toolchains. The model is priced at \$3 per million input tokens and \$15 per million output tokens, with cached input tokens dropping to just \$0.30 per million — pricing that positions it roughly in line with mid-tier offerings from Western labs, but at a performance level the company claims approaches the top of the market. A promotional top-up rebate running through August 12 offers up to 30 percent back in vouchers for API credits of \$1,000 or more.

As [Xinhua reported](#), a Moonshot AI executive explained the significance of the parameter count in simple terms: parameters are like neural connections in the human brain, and nearly 3 trillion of them means the model can "store more knowledge and patterns in its brain, understand more, think deeper, and answer more accurately."



Benchmark results show Kimi K3 trading blows with Claude and GPT at the top of the leaderboard

The benchmark results, drawn from public leaderboard data and a private evaluation by analytics firm Artificial Analysis, tell a striking story.

On [GDPval-AA v2](#), a benchmark measuring real-world tasks across 44 occupations and 9 major industries, Kimi K3 scored 1,687 — placing it third overall, behind only Claude Fable 5 Max (1,815) and GPT-5.6 Sol Max (1,747.8), and ahead of Claude Opus 4.8 (1,600).

On [AA-Briefcase](#), a private agentic benchmark from Artificial Analysis designed to test long-horizon knowledge work, K3 climbed to second place with a score of 1,527 — beating GPT-5.6 Sol Max (1,495) and trailing only Fable 5 Max (1,587).

Perhaps most impressively, K3 achieved a state-of-the-art score of 91.2 out of 100 on [BrowseComp](#), a benchmark for long-horizon, high-difficulty information seeking.

General Agents All maxed out on thinking effort: max or xhigh.


GDPval-AA v2 Elo



AA-Briefcase Elo



Automation Bench




JobBench



SpreadsheetBench 2



BrowseComp



Visual Agents All maxed out on thinking effort: max or xhigh.

CharXiv (RQ) w/ tool



Zerobench w/ tool (Pass@5)



In tests of real-world task automation, Kimi K3 ranked first in four out of eight benchmarks — including Automation Bench, SpreadsheetBench 2 and BrowseComp — while finishing second to Fable 5 in most others. Fable 5 and GPT-5.6 Sol were its closest competitors overall. (Source: Moonshot AI)


The company says it accomplished this in a single-agent setup using its 1-million-token context window, without any context compression or additional context management techniques — a feat that suggests raw context length, when paired with strong retrieval capabilities, may be more powerful than elaborate multi-agent workarounds.

As [one widely followed AI commentator](#) put it on social media: "Open source is no longer lagging six months behind Western closed-source models. Read that again, and think about what it all means."

That observation captures the significance of the moment. For much of the past three years, open-source models have typically trailed their proprietary counterparts by a meaningful margin. Kimi K3 appears to have closed that gap almost entirely.

Frontend Code Arena

Kimi-K3: Ranked #1



Kimi K3 claimed the No. 1 spot on Arena.AI's Frontend Code Arena with a score of 1,679, outpacing Claude Fable 5 and GPT-5.6 Sol by a significant margin. The leaderboard ranks models by human preference in head-to-head frontend coding comparisons. (Source: arena.ai)

How a 48-hour autonomous chip design demo reveals Moonshot's real ambitions

Beyond raw benchmarks, [Moonshot AI](#) showcased a proof-of-concept that may be even more revealing of K3's capabilities and the company's strategic direction.

In a demonstration documented in the company's technical materials, [Kimi K3](#) was tasked with designing a physical chip to run a nano-scale version of itself. Over 48

hours of continuous autonomous agent operation, K3 independently completed the chip's full construction pipeline — from architectural design through optimization and verification — using open-source electronic design automation tools. The result was a tiny but functional chip design, just 4 square millimeters, that achieved timing convergence at 100 MHz and could decode more than 8,700 tokens per second in simulation.

This is not a production chip. It is a demonstration of what [Moonshot AI](#) clearly views as the next competitive frontier: long-range autonomous agent capabilities. The ability to sustain coherent, multi-step technical work over a 48-hour window — reading documentation, making design decisions, running verification loops, and iterating on failures — represents a qualitative leap beyond the kind of single-turn question-answering that defined the first generation of large language models.

The company also highlighted a case in computational astrophysics, where K3 reportedly reproduced the universal [I-Love-Q relation](#) — a complex calculation that typically takes a senior researcher one to two weeks — in approximately two hours, reading and cross-validating more than 20 papers and implementing a complete numerical pipeline along the way.

Moonshot AI's fall and rise tells the story of China's brutal AI market

To understand why [Kimi K3](#) matters, you need to understand where Moonshot AI was 18 months ago — and how far it fell.

Founded in 2023 by [Yang Zhilin](#), a Tsinghua University graduate who previously conducted research at Google and Meta, Moonshot AI quickly became one of China's most prominent AI startups. The company gained early traction in 2024 when users flocked to its [Kimi platform](#) for its long-text analysis capabilities and AI search functions. By early 2026, it had raised roughly [\\$1.5 billion](#) across multiple rounds, with its valuation climbing from \$2.5 billion to \$4.3 billion and the company reportedly [seeking a new round at \\$5 billion](#).

Then DeepSeek happened. The release of DeepSeek's low-cost R1 model in January 2025 disrupted the entire Chinese AI landscape, and Moonshot AI was among the hardest hit. Kimi, which had ranked third in monthly active users in China, slid to seventh. The company's strategic pivot to open-source models — beginning with Kimi K2 in July 2025 and accelerating with K2.5 in January 2026 — was in large part an effort to reclaim relevance.

[Kimi K3](#) is the culmination of that effort — and the sheer scale of the model suggests that Moonshot AI has been planning this move for some time. Training a 2.8-trillion-

parameter model requires enormous computational resources and months of preparation, which means the architectural and infrastructure decisions behind K3 were likely locked in well before the model reached the public.

Why open-sourcing the world's biggest model is a geopolitical chess move

The decision to release K3's full weights on July 27 is strategically significant and worth parsing carefully.

The company's own timeline chart of open-source frontier model scale positions K3 as a dramatic outlier, towering above competitors like [DeepSeek](#) (1.6T), [Xiaomi](#) (1.02T), and [Alibaba](#) (397B). By releasing the world's largest open-source model, Moonshot AI is making a bid to become the center of gravity for the global open-source AI developer community.

This follows a broader trend among Chinese AI companies. As [Reuters noted](#), open-sourcing allows companies to "showcase their technological capabilities and expand developer communities as well as their global influence, a strategy likely to help China counter U.S. efforts to limit Beijing's tech progress." DeepSeek, Alibaba, Tencent, and Baidu have all released open-source models. But none have released anything at this parameter count.


For enterprise technology leaders, the implications are concrete. A 2.8-trillion-parameter open-source model that performs at near-frontier levels creates new options for companies that want to fine-tune, self-host, or build proprietary systems on top of a capable base model — without being locked into API contracts with OpenAI or Anthropic. The trade-off, of course, is that running a model of this size requires substantial GPU infrastructure. Inference at 2.8 trillion parameters is not something that runs on a single server rack.

That said, [Moonshot AI](#) has signaled awareness of this challenge. Its Mooncake project, which won the Best Paper award at FAST 2025, pioneered KV-cache-centric disaggregated serving for large language models — an architecture designed specifically to make inference at extreme scale more practical and cost-efficient.

Kimi Code and a three-tier model lineup form the foundation of Moonshot's enterprise play



Alongside K3, Moonshot AI continues to invest heavily in its coding agent ecosystem. [Kimi Code](#), the company's open-source coding tool that competes with Anthropic's Claude Code and Google's Gemini CLI, received two major updates on the same day as K3's launch — versions 0.25.0 and 0.26.0 — adding features like expanded subagent tooling, background task management, and security fixes.



Kimi K3 consistently placed among the top three models across six coding benchmarks, leading all competitors in SWE Marathon and Program Bench, and trailing only GPT-5.6 Sol in Terminal Bench 2.1 by half a point. All models were tested at maximum thinking effort. (Credit: Moonshot AI)

The [Kimi Code CLI](#) has accumulated over 3,100 stars on GitHub and features integration with VSCode, Cursor, and Zed. The latest release expanded the "coder subagent" tool set to include background tasks, todo lists, plan mode, skill invocation, and nested agents — effectively turning the coding agent into a multi-layered autonomous system capable of managing complex software engineering projects with minimal human intervention.

This is not incidental. Coding tools have become a critical revenue driver for AI labs. As Anthropic disclosed in January, [Claude Code reached \\$1 billion in annualized recurring revenue](#). By building Kimi Code as an open-source alternative that defaults to Kimi's own models — but supports other providers — Moonshot AI is positioning itself to capture developer workflows and, eventually, enterprise contracts.

The company's model lineup now includes three tiers: [K3](#) as the flagship (\$3/\$15 per million tokens for input/output), [K2.7 Code](#) as a specialized coding model (\$0.95/\$4) and [K2.6](#) as a general-purpose option (\$0.95/\$4). All three support context windows of

256,000 tokens or above, with K3 offering the full 1-million-token window. Context caching is automatic — no cache ID, TTL, or extra parameter is required — a small but meaningful developer-experience advantage over competitors that require explicit cache management.

What Kimi K3 means for the future of enterprise AI and the global model landscape

Kimi K3's release forces a recalibration of several assumptions that have guided enterprise AI strategy.

The performance gap between open-source and proprietary models has functionally closed at the frontier. If K3's benchmark numbers hold up under independent evaluation — and particularly once the open weights are available for community testing on July 27 — it will be difficult for closed-source providers to justify premium pricing purely on the basis of capability.

The locus of AI innovation, meanwhile, continues to shift. China's AI ecosystem, which many Western observers questioned after early struggles with chip export restrictions, has now produced a model that competes with the best systems from companies with direct access to Nvidia's most advanced hardware. The architectural innovations behind K3 — particularly the hybrid linear attention mechanism — suggest that algorithmic efficiency may matter as much as raw compute.

And the agentic capabilities demonstrated by K3 — chip design, multi-week research compression, long-horizon information seeking — point toward a future where AI models are not just answering questions but autonomously executing complex, multi-day projects. For enterprises evaluating AI investments, this shifts the value proposition from "productivity copilot" to "autonomous technical workforce."

[Xinhua](#), China's state news agency, framed the release as a national milestone, reporting that K3 "marks a new step forward in the development of China's artificial intelligence models." Liu Tieyan, dean of the Zhongguancun Academy in Beijing, was quoted as saying that a wave of Chinese open-source models has moved from isolated breakthroughs to collective advancement, providing "new solutions and new paths" for global AI development.

Just two years ago, [Moonshot AI](#) was a scrappy startup named for the audacious problems it hoped to solve. Eighteen months ago, it was a cautionary tale about how quickly a market darling can lose its footing. Today, it is the maker of the world's largest open-source AI model — one that can, given 48 hours and an internet

connection, design a chip to run itself. The frontier, it turns out, is not a place. It is a race. And the field just got a lot more crowded.

More





Thinking Machines open sources first multimodal language model, Inkling, focused on low cost and 'resistance to censorship'

An Apache 2.0 designation makes Inkling a true open-source foundation. This gives developers the legal freedom to download, modify, integrate, and commercialize the model weights.

Carl Franzen

July 16, 2026





Cohere VP says enterprise AI sovereignty requires control of the full agent stack at VB Transform 2026

Asked what would persuade enterprises to move beyond bundled AI services from existing cloud providers, Alao returned to data control and portability.

Carl Franzen

July 15, 2026



1Password moves into AI cost management, betting that token spend is the next enterprise budget crisis

The move marks the latest strategic expansion for a company that built its reputation on password management for consumers and, over the past three years, has aggressively repositioned itself as a broader identity security and SaaS governance...

Michael Nuñez

July 14, 2026



Canva launches Code 2.0, offering AI website building to every user — including free accounts

The move is Canva's most aggressive push yet into the fast-growing "vibe coding" market, a category that barely existed 18 months ago but has already minted billion-dollar startups and reshaped how non-developers think about building software. But...

Michael Nuñez

July 14, 2026






PARTNER CONTENT

The desktop infrastructure problem that Kubernetes finally solves

Presented by Kasm Technologies

Daniel Ben-Chitrit, Kasm Technologies

July 13, 2026




Google's TabFM skips per-dataset training and still predicts on tables it's never seen

Data scientists spend weeks tuning hyperparameters and rebuilding pipelines for every new dataset. Google's latest model does the job in one API call instead.

Ben Dickson

July 10, 2026



SpaceX's Grok 4.5 launches at half the price of rivals — here's why that could rattle Anthropic and OpenAI

Elon Musk's SpaceX released Grok 4.5 on Wednesday, the first artificial intelligence model the company has trained specifically for coding and autonomous agents — and the first tangible product of its \$60 billion acquisition of the AI coding startup Cursor,...

Michael Nuñez

July 9, 2026






OpenAI launches GPT-Live, a full-duplex voice upgrade that lets ChatGPT talk more like a person

The two models, GPT-Live-1 and GPT-Live-1 mini, are rolling out globally starting today across iOS, Android, and ChatGPT.com. GPT-Live-1 becomes the default voice model for paid ChatGPT users on the Go, Plus, and Pro tiers, while GPT-Live-1 mini...

Michael Nuñez

July 8, 2026




Slack's Slackbot can now pull your CRM data, generate charts, and send DocuSigns — all from a chat message.

Five years and \$27.7 billion after Salesforce acquired Slack, the two products are finally starting to function as a single system. On Wednesday, Slack launched an integration that connects Slackbot — the personal AI agent built into every workspace — to the...

Michael Nuñez

July 8, 2026




Anthropic brings Claude Cowork to mobile and web as usage data shows most users aren't coding

The rollout, which begins in beta with Max subscribers before expanding to additional plans, marks a strategic inflection for Anthropic. It transforms Cowork from a desktop-only agent into a cross-device platform where tasks can start on a laptop, continue...

Michael Nuñez

July 7, 2026






Anthropic's new "J-lens" reveals a silent workspace inside Claude that mirrors a leading theory of consciousness

The 16-author study, titled "*Verbalizable Representations Form a Global Workspace in Language Models*," describes how Anthropic's researchers used a new mathematical technique to peer inside Claude's neural network and discovered what they call a "J-...

Michael Nuñez

July 7, 2026



Tencent's Apache-licensed Hy3 takes on GLM-5.2 at half the size — and wins everywhere except coding

Tencent's Hy3 ditches restrictive licensing, halves its hallucination rate, and is sized to run on export-compliant GPUs — a serious option for IT teams evaluating open models.

Sam Witteveen

July 6, 2026



[Press Releases](#)

[Contact Us](#)

[Advertise](#)



[Share a News Tip](#)

[Contribute](#)

[Privacy Policy](#)

[Terms of Service](#)

[Consent Preferences](#)

[Do Not Sell or Share My Personal Information](#)

[Limit the Use Of My Sensitive Personal Information](#)

© 2026 VentureBeat. All rights reserved.

