
[ГЛАВНАЯ](#) > [ТЕХНОЛОГИИ](#)

Gemini 3.5 Pro рассчитана на 17 июля, поскольку разработчики DeepSeek истекают к 24 июля

Google перестроила свою базовую модель с нуля; крайний срок предоставления API-псевдонима DeepSeek истекает 24 июля.

Автор: [Ричард Л. Уэллс](#)

Опубликовано: 08 июля 2026, 1:18 EDT



Google Gemini (PHOTO: GOOGLE.COM)

Три самых ожидаемых события, связанных с моделями искусственного интеллекта, в 2026 году приходятся на одну неделю, и одно из них не терпит отлагательств. Google DeepMind планирует выпустить общедоступную версию Gemini 3.5 Pro 17 июля — модель, которую компания решила создать с нуля, а не выпускать в том виде, в котором она была создана ранее. В тот же день DeepSeek планирует перевести семейство V4 из стадии

появились строки с названием модели, а в любой момент может начаться подписка. Независимо от того, выйдет ли обновление 17 июля для всех трех моделей, практический срок, который больше всего важен для любого разработчика, использующего производственную систему на базе API DeepSeek, — это 24 июля, когда перестанут работать устаревшие псевдонимы deepseek-chat и deepseek-reasoner, и никаких продлений не будет.

Почему Google отказалась от прежней архитектуры и начала все сначала

Gemini 3.5 Pro должен был выйти в июне. 19 мая на конференции Google I/O Сундар Пичаи попросил аудиторию «подождать до следующего месяца». Но этого не произошло. Промах пришелся на то же двухнедельное окно, что и волна увольнений исследователей, пошатнувшая рыночную оценку Alphabet: Ноам Шапиро — соруководитель Gemini и соавтор статьи 2017 года "Внимание — это все, что вам нужно", в которой была представлена архитектура transformer, лежащая в основе практически всех используемых сегодня крупных языковых моделей, — объявил о своем уходе из OpenAI 18 июня. Джон Джампер, нобелевский лауреат AlphaFold и девятилетний ветеран DeepMind, объявил о своем переходе в Anthropic 19 июня. Вместе с уходом двух других старших научных сотрудников эти события привели к падению акций Alphabet на 5% за одну торговую сессию 22 июня, что привело к снижению [рыночной стоимости](#) примерно на 225 миллиардов долларов.

Сама по себе задержка оказалась не самым удивительным фактом. Что привлекло всеобщее внимание, так это [заявленное решение Google DeepMind](#) полностью отказаться от базовой модели Gemini 2.5 Pro и запустить совершенно новый цикл предварительного обучения с нуля. Заявленное обоснование связано с тремя пробелами в производительности, которые, как сообщается, не удалось устранить с помощью существующей архитектуры: математические вычисления, генерация сцен с помощью масштабируемой векторной графики (SVG) и общее качество изображения. Инкрементальная тонкая настройка достигла предела в каждом из этих направлений. Запуск совершенно нового цикла предварительного обучения в передовых масштабах обходится в сотни миллионов долларов и требует нескольких месяцев работы на графических процессорах. Google все равно выбрала этот путь — это сигнал о том, насколько не оправдал ожиданий предыдущий кандидат, или о том, насколько высокую планку установили конкуренты, вышедшие на рынок в тот же день.

Сообщается, что обновленная модель будет иметь контекстное окно на 2 миллиона токенов, что вдвое больше, чем 1 миллион в Gemini 2.5 Pro, а также уровень рассуждений Deep Think для многоэтапной логики и возможности автономного рабочего процесса для выполнения сложных задач, связанных с кодированием и использованием инструментов. Эти сведения взяты из сторонних источников и утечек, а не из официальной документации Google. По состоянию на 7 июля 2026 года в общедоступном API Gemini указаны только версии gemini-3.5-flash и gemini-3.1-pro-preview. Не было опубликовано ни карточки модели, ни подтверждения цен, ни официальных сравнительных тестов. 17 июля — это дата, о которой много пишут, но это не анонс с подписанным пост-релизом.

A context window is the total number of tokens — input plus output combined — that a model can hold in a single inference pass. At 2 million tokens, Gemini 3.5 Pro would be able to process roughly 1.5 million words in a single prompt: a full large codebase, a year's worth of meeting transcripts, or a multi-volume research dataset. That represents a genuine engineering advance and a real [capability gap versus most current alternatives](#).

The engineering complexity behind it matters for evaluating the claim. Transformer attention scales quadratically with sequence length, which means processing 2 million tokens demands orders of magnitude more compute than processing 100,000. Extending context to that scale requires significant architectural work — researchers at Microsoft demonstrated a technique called LongRoPE that extends context windows to 2 million tokens, but achieving reliable, accurate retrieval across the full span is a separate problem from technically accepting that many tokens. Researchers at Stanford and other institutions have documented a phenomenon where model performance degrades for information located in the middle 50% of a very long context, regardless of whether the model technically fits it. Effective context window — the range where the model reliably uses information — often falls well short of the advertised limit.

Until independent evaluators run long-context retrieval benchmarks on Gemini 3.5 Pro, the 2 million token headline is a capability claim, not a verified specification. The evaluation to watch for is not whether the model accepts a 2 million token prompt but whether reasoning quality holds across the full range.

DeepSeek V4: Cheap, Capable, and Carrying Legal Obligations No Contract Can Override

[DeepSeek V4-Pro went live as a preview on April 24, 2026](#) — the same day OpenAI shipped GPT-5.5, a timing that appears deliberate. The model uses a Mixture-of-Experts (MoE) architecture: rather than activating all of its 1.6 trillion total parameters for every token, a routing network selects the 49 billion most relevant parameters per token. The rest stay dormant. That selective activation is what makes a 1.6-trillion-parameter model economically viable to serve — each token costs roughly the same compute as a much smaller dense model, while the full parameter count is available in aggregate across diverse tasks. V4-Flash (284 billion total, 13 billion active) takes that principle further, delivering much of the Pro tier's benchmark performance at a fraction of the cost.

The [price gap versus Western frontier APIs](#) is substantial. V4-Pro's permanent list price of \$0.87 per million output tokens compares with roughly \$25 per million for Claude Opus 4.7 and \$30 per million for GPT-5.5. For high-volume coding agent workloads, that differential can mean the difference between a viable product economics model and one that is not. Developers who self-host the open-weight MIT-licensed model on their own infrastructure eliminate the per-token API cost entirely — and V4-Flash, at 160GB, is within reach of a high-spec local setup with light quantization.

Before treating that price gap as the deciding factor, four dimensions require equal weight.

First, independent benchmarks show real performance gaps. On DeepSWE — a contamination-free, independently run benchmark from the [yage.ai lab](#) that uses a higher-fidelity verifier than the widely cited SWE-bench Verified — DeepSeek V4-Pro scores 8% pass@1 against GPT-5.5 at 70% and Claude Opus 4.7 at 54%. On the vendor's own technical report, the model is described as "marginally short" of GPT-5.4 and Gemini 3.1 Pro on reasoning tasks, placing it approximately three to six months behind the frontier. V4 is also text-only: it does not process audio, video, or images, a limitation that closed-source peers have closed. [Scale AI's SEAL leaderboard](#), as of this writing, lists no DeepSeek V4 entry at all.

Second, published benchmark scores from a company headquartered in China require independent replication before they should inform production decisions. The vendor-reported SWE-bench Verified score of 80.6% uses a verifier that the yage.ai audit found accepts approximately 8.5% of incorrect solutions. The DeepSWE result, which uses a verifier with 0.3% false positive rate, tells a materially different story about long-horizon agentic capability. [BenchLM](#), an independent model tracking service, ranks V4-Pro 29th out of 33 on its verified leaderboard — and explicitly classifies it as "not a frontier model" by their methodology.

Third, enterprise teams should factor in ecosystem costs beyond token pricing. DeepSeek V4 runs on Huawei Ascend 950 chips — fully outside U.S. export controls — which matters if your infrastructure planning or supply-chain policy requires Western semiconductor provenance. Peak-hour pricing at 2× the baseline rate takes effect with the July stable release, applying during Beijing business hours (9 AM–12 PM and 2–6 PM China Standard Time). Workloads running during those windows — which overlap partially with morning hours on the U.S. East Coast — will face double the off-peak rate for API calls.

Fourth, and this is not a variable risk to weigh against price, it is a fixed legal condition: DeepSeek is operated by Hangzhou DeepSeek Artificial Intelligence Co., Ltd., a Chinese company. [China's National Intelligence Law \(2017\), Article 7](#), requires all Chinese organizations and citizens to "support, assist, and cooperate with national intelligence work" — with no exceptions and no court order required. This obligation applies regardless of where DeepSeek's servers are physically located, regardless of the company's stated privacy policy, and regardless of any contractual terms between DeepSeek and its users. The Data Security Law (2021) and Cybersecurity Law (2017) add data localization and government access requirements on top of that baseline.

SecurityScorecard's research team found DeepSeek code integrating with ByteDance services and data flowing to Chinese state-linked entities. [Feroot Security](#), a cybersecurity firm, reported in late 2024 that DeepSeek code contained lines transmitting user data to CMPassport.com, a portal operated by China Mobile — a telecom operator that the U.S.

prevents unencrypted data transmission — and used 3DES encryption with a key hardcoded in the app code. [Wiz Research](#) found an unauthenticated DeepSeek database exposing more than 1 million log entries, including chat history and API keys, before the company secured it after disclosure.

The self-hosting path changes the calculation meaningfully. Because DeepSeek V4 is released under the MIT license, any developer can download the model weights and run inference on their own infrastructure using tools like vLLM. When you do that, prompts never reach DeepSeek's servers in China, and the cross-border data-flow concern disappears. The legal obligations that run with operating under Chinese jurisdiction do not attach to a self-hosted deployment on Western infrastructure. The censorship behavior baked into the model's weights — [CrowdStrike documented](#) that R1 refused to generate code associated with certain politically sensitive topics in approximately 45% of tested cases — does carry over regardless of deployment environment; that is trained behavior, not a server-side filter. For regulated industries, client data, or anything sensitive, the right answer is self-hosting on compliant infrastructure, not the hosted API.

DeepSeek API Migration: The Deadline That Is Already Here

Regardless of what ships on July 17, every development team using DeepSeek's hosted API has one non-optional task: update any code that calls `deepseek-chat` or `deepseek-reasoner` before July 24, 2026, at 15:59 UTC. After that timestamp, those aliases return errors. [No extension has been announced.](#)

The migration is a one-line code change per call — update the model parameter to `deepseek-v4-pro` or `deepseek-v4-flash` on the same base URL, same API key. The important catch: `deepseek-reasoner` maps to `deepseek-v4-flash` (thinking mode), not to V4-Pro. A team that was using `deepseek-reasoner` for heavy reasoning workloads and assumes the alias migration preserves capability parity will end up on Flash-tier reasoning at Flash prices, which may be a mismatch. Anyone who relied on `deepseek-reasoner` for complex tasks should explicitly evaluate V4-Pro — not just swap the alias name.

Batch processing, evaluation runs, and non-latency-sensitive inference should also be scheduled outside Beijing peak hours once the stable release launches to avoid the 2× peak rate. The windows most affected for U.S.-based teams are Beijing 9–12 AM (1–4 AM ET) and 2–6 PM (2–6 AM ET the following day), meaning overnight batch jobs on a U.S. schedule will largely miss peak-rate windows — which is useful to know for cost planning.

Grok 4.5: Interesting Architecture, No Independent Numbers

SpaceXAI — the entity formed when xAI was folded into SpaceX in February 2026 — confirmed on June 28 that Grok 4.5 entered private beta at SpaceX and Tesla, built on a 1.5-trillion-parameter V9 foundation model. Reinforcement learning was still running at the time of the announcement. Canary strings reading "Unlock the full power of Chat with Grok 4.5" surfaced in the [Grok web UI on July 6](#), and the model appears in the Grok Build coding agent's upgrade flow — both signals that a subscriber release is approaching, though no public date has been confirmed.

and Tesla engineers, organizations that now sit under the same corporate umbrella as xAI. No independent benchmark — SWE-bench, Humanity's Last Exam, Artificial Analysis, LMArena — has produced a public score for Grok 4.5. As of mid-June 2026, [no publicly verified numeric scores existed even for Grok 4.3](#), the current production model. The AI industry has documented cases where a developer's self-reported benchmark score diverged from independently measured results by more than 20 percentage points; a claim about an inaccessible model backed by no external evaluators carries significant epistemic uncertainty.

The V9 architecture incorporates Cursor's developer-workflow data, but as supplemental training — added after the initial pre-training run rather than from the start. An xAI engineer acknowledged that supplemental inclusion is "not quite as good as having it in initial training." SpaceX has agreed to acquire Cursor for \$60 billion, with the deal expected to close in Q3 2026; the next model in training is being built to incorporate Cursor data from the beginning of pre-training, which xAI expects to produce meaningfully stronger coding performance.

Read more: [Gemini 3.5 Pro Cleared for July Launch as Fable 5 Nears Return, GPT-5.6 Stays Locked](#)

What the Rebuild Signals About the Competitive State of Play

Three major model events clustering around the same week is not an accident. It reflects where the frontier has moved: labs are now shipping at a pace that makes the same-week collision of a Google flagship launch, a DeepSeek stable release, and a Grok rollout feel routine rather than extraordinary.

What is not routine is Google's decision to scrap a pre-training run at the frontier. Rebuilding from scratch at this scale costs hundreds of millions of dollars and forfeits months of development time. The decision signals one of two things: the prior candidate genuinely failed internal quality thresholds, or the competitive bar set by the models arriving on the same date made the prior candidate unacceptable even if it met prior thresholds. Either way, the gamble is that a ground-up rebuild can deliver substantially better reasoning and coding performance than iterative improvement on the 2.5 Pro base could have. That bet will be evaluated against independent agentic benchmarks — SWE-bench, Terminal-Bench 2.0, DeepSWE — not against Google's own published launch materials.

The architectural diversity converging on this week is also notable. Google's dense transformer rebuild prioritizes reasoning quality and visual precision. DeepSeek's sparse Mixture-of-Experts design targets enterprise-scale cost efficiency. SpaceXAI's vertically integrated approach — where a single entity controls the compute cluster, the model, the coding tool that feeds training data, and the companies serving as production stress-test environments — has no direct precedent among frontier labs. Those are three genuinely different technical bets on what the next generation of capable AI looks like.

Whether all three actually arrive on July 17 is secondary to the underlying reality: the decision window for any development team choosing or upgrading an AI stack in the second half of 2026 has compressed to days, not quarters. For the one deadline that is already firm, July 24 is seventeen days away.

Frequently Asked Questions

When is Gemini 3.5 Pro actually releasing?

Multiple third-party outlets, including Business Insider and Geeky Gadgets, report July 17, 2026 as the target date. Google has not officially confirmed this date with a model card, API documentation, or pricing announcement. As of July 7, 2026, the public Gemini API does not list gemini-3.5-pro as a generally available model. Treat July 17 as a reported target: worth planning around, but not a commitment you should build a production dependency on.

Do I need to update my DeepSeek API integration before July 24?

Yes, if your code calls `deepseek-chat` or `deepseek-reasoner`. Both aliases stop working on July 24, 2026, at 15:59 UTC with no announced extension. Migrate to `deepseek-v4-pro` or `deepseek-v4-flash` — same base URL, same API key, one parameter change. Important: `deepseek-reasoner` maps to V4-Flash (thinking mode), not V4-Pro. If you were using `deepseek-reasoner` for heavy reasoning tasks and need equivalent capability, you must explicitly call `deepseek-v4-pro`, not just replace the alias name.

Is DeepSeek V4 safe for enterprise use?

It depends on how you deploy it. The hosted cloud API routes all data through DeepSeek's servers in China, where China's National Intelligence Law (2017, Article 7) legally requires the company to cooperate with government intelligence requests on demand — no court order, no exceptions. This applies regardless of DeepSeek's privacy policy or the physical location of any given server. For sensitive data, regulated industries, or environments covered by GDPR or U.S. data protection frameworks, the hosted API is not appropriate. The self-hosted path — downloading V4's MIT-licensed open weights and running them on your own infrastructure — eliminates cross-border data flow risk, though behavioral censorship baked into the model weights remains regardless of deployment environment. Security research testing earlier DeepSeek model generations (not V4 specifically) found significant jailbreak vulnerabilities, plaintext credential storage in the mobile app, and a database exposure affecting over one million log entries.

in practical terms, 2 million tokens is roughly 1.5 million words — enough to hold an entire large software codebase, a year of business correspondence, or a multi-volume research dataset in a single prompt. The engineering challenge is that transformer attention scales quadratically with sequence length, so doubling the context window does not cost twice as much compute — it costs orders of magnitude more. Achieving accurate, reliable performance across the full span is also a separate problem from technically accepting that many tokens: research has documented performance degradation for information in the middle of very long contexts even when it technically fits. Until independent evaluators publish long-context retrieval benchmarks for Gemini 3.5 Pro, the 2 million token claim should be treated as a headline specification, not a guarantee of reliable performance at that range.

© 2026 TECHTIMES.com All rights reserved. Do not reproduce without permission.

Tags: [Gemini](#) [Google](#)

Tech Times

Find the latest tech news!


MOST POPULAR

1.



China AI Companion Law Takes Effect: Doubao and Qwen Shut Down, Millions Lose Chat Data

2.



Canadian Wildfire Smoke Grounds Flights, Cancels FIFA Events as PM2.5 Hits 100 Million Americans

3.



Rebuilt Gemini 3.5 Pro Misses Third Deadline: Google Eyes Stopgap Release

4.



Quantum Computing Stocks Drop 35% as Federal Mandate Reshapes Investment Case



Tech Times

Find the latest
tech news!

Subscribe

[TECH](#) [SCIENCE](#) [BUSINESS](#) [HEALTH](#) [CULTURE](#) [FEATURES](#) [BUZZ](#)

[ABOUT US](#) [CONTACT US](#) [TERMS & CONDITIONS](#) [PRIVACY POLICY](#) [PRIVACY SETTINGS](#)

© 2026 Tech Times Media Inc. All rights reserved.