



OpenAI • Собственная модель • Выпущена в июле 2026 года

GPT-5.6 Terra (max)

Тестирование и анализ API-провайдеров

Сравнение моделей

Анализ API-провайдеров для GPT-5.6 Terra (max) по таким показателям производительности, как задержка (время до выдачи первого токена), скорость вывода (количество токенов в секунду), цена и другие. Среди протестированных провайдеров API — OpenAI.

⚡ Самый быстрый

#1 OpenAI 145.2 t/s

Скорость вывода Всего 1 провайдеров

🕒 Самая низкая задержка

#1 OpenAI 137.48 s

Время до первого токена Всего 1 провайдеров
ответа

💰 Самая низкая цена

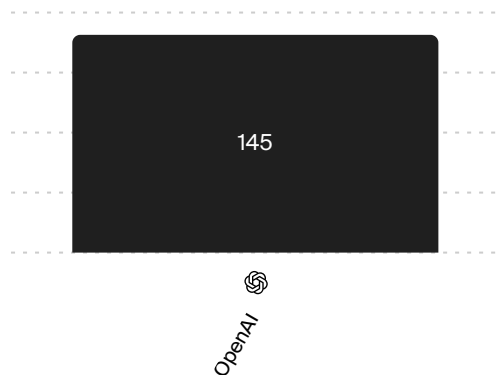
#1 OpenAI \$2.17

Смешанная цена (за 1 млн токенов) Всего 1 провайдеров

GPT-5.6 Terra (макс.) доступен через OpenAI.

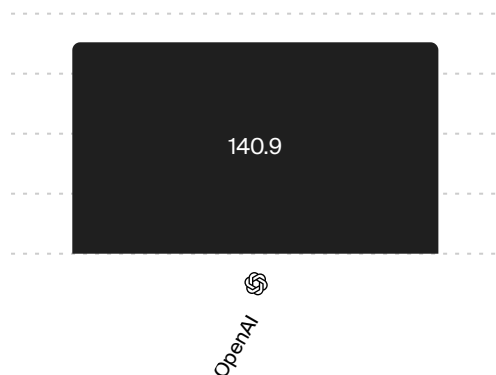
Скорость

Токены на выходе в секунду · Чем выше, тем лучше



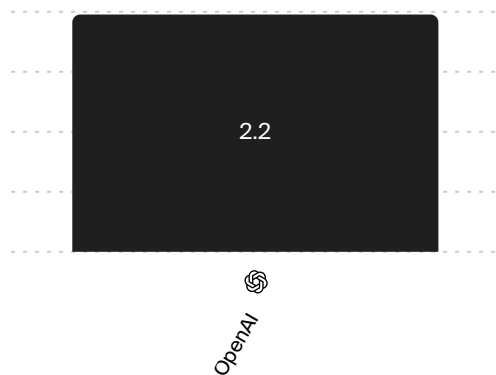
Сквозное время отклика

Секунды · Чем меньше, тем лучше



Цена

USD за 1 млн токенов (смешанная модель) · Чем ниже, тем лучше



Обновление: рабочая нагрузка для тестирования производительности по умолчанию обновлена до 10 000 входных токенов, чтобы лучше отражать производственные сценарии использования. Вы по-прежнему можете выбрать другие рабочие нагрузки выше.

Цены

Цены: Hit, Input и Output


Цена (долл. США за млн токенов) · Чем ниже, тем лучше · 10 000 входных токенов



1 из 1 поставщиков

Hit Input Output

Artificial Analysis



Cache Hit	+
Input Price	+
Cache Pricing by Provider	+


Pricing: Blended Price

Blended at 7:2:1 (cache-input-output) · USD per 1M tokens (blended) · Lower is better

7:2:1 General agentic (recommended)

1 of 1 providers

Artificial Analysis




Pricing: Cache Discount

1 - (cache hit price / input price) · Higher is better

1 of 1 providers

⌵

 Artificial Analysis



i Cache Price Discount	+
i Cache Hit	+
i Cache Pricing by Provider	+
i Median	+

Output Speed vs. Price

Blended at 7:2:1 (cache-input-output) · Output speed: output tokens per second · USD per 1M tokens (blended)

7:2:1 General agentic (recommended)


⌵

1 of 1 providers

⌵

☐ Most attractive quadrant

OpenAI



i Output Speed	+
i Price	+
i Cache Hit	+
i Cache Pricing by Provider	+
i Median	+

Speed


Measured by Output Speed (tokens per second)

Output Speed	Output Speed by Prompt Type	Output Speed Variance	Output Sp
Output Speed: GPT-5.6 Terra (max) Output speed: output tokens per second · 10,000 input tokens			

1 of 1 providers

⌵

 Artificial Analysis



-  Output Speed

+
-  Model Performance Representation

+

Latency vs. Output Speed


Latency: seconds to first token received · Output speed: output tokens per second · 10,000 input tokens

1 of 1 providers

⌵

☐ Most attractive quadrant

OpenAI



i	Output Speed	+
i	Latency (Time to First Token)	+
i	Price	+
i	Median	+

Latency

Measured by Time (seconds) to First Token

Time to First Answer Token

Time to First Token

Latency by Prompt Type

Latency

Time to First Answer Token

Seconds to first token received · Lower is better · 10,000 input tokens

↔

📄

📊


1 of 1 providers



'Thinking' time (reasoning models, where applicable)

Input processing time

 Artificial Analysis



 Time to First Answer Token



 Median



End-to-End Response Time

Seconds to output 500 tokens, calculated based on time to first token, 'thinking' time for reasoning models, and output speed

End-to-End Response Time

End-to-End Response Time by Prompt Type

End-to-End R

End-to-End Response Time

Seconds to output 500 tokens, including reasoning model 'thinking' time · Lower is better · 10,000 input tokens



1 of 1 providers




Outputting time

'Thinking' time (reasoning models)

Input processing time

Artificial Analysis




End-to-End Response Time	+
Standardized Reasoning Tokens	+
Median	+

Key Comparison Metrics & API Features

Filter, e.g. OpenAI, GPT, Meta					Expand Columns	
API Provider	Features				Price	Speed
	Context Window	Function Calling	JSON Mode	License	Blended USD/1M Tokens	Median Tokens
OpenAI	1M			Proprietary	\$2.17	14

Frequently Asked Questions

Where can I access GPT-5.6 Terra (max)?	+
How many API providers offer GPT-5.6 Terra (max)?	+
Which provider is fastest for GPT-5.6 Terra (max)?	+
Which provider has the lowest latency for GPT-5.6 Terra (max)?	+
Which provider is cheapest for GPT-5.6 Terra (max)?	+
Which provider has the lowest input price for GPT-5.6 Terra (max)?	+
Which provider has the lowest output price for GPT-5.6 Terra (max)?	+
Which GPT-5.6 Terra (max) providers support JSON mode?	+
Which GPT-5.6 Terra (max) providers support function calling?	+
Which is the best provider for GPT-5.6 Terra (max)?	+
How do I choose a provider for GPT-5.6 Terra (max)?	+
Does provider performance for GPT-5.6 Terra (max) change over time?	+
What are the overall capabilities of GPT-5.6 Terra (max)?	+



Artificial Analysis

Get notified about new articles

Subscribe

Explore

[LLM Leaderboard](#)

[Image Arena](#)

[Video Arena](#)

[AI Agents](#)

[Evaluations](#)

Company

[Methodology](#)

[Services](#)

[Contact](#)

[Articles](#)

[FAQ](#)

