



Kimi • Собственная модель • Выпущена в июле 2026 года

Kimi K3 Анализ интеллекта, производительности и цены

↔ Сравнить

Сравнительные тесты API

Интеллект обновлено

#4 / 189

Скорость

#91 / 189



57

Индекс интеллекта
искусственного анализа

62.0

Токены на выходе в
секунду

Цена

#127 / 189

Цена за попадание в кэш

#77 / 189



Вход

\$3.00

за 1 млн токенов

Выход

\$15.00

за 1 млн токенов

\$0.30

(-90%)

USD за 1 млн токенов

Подробность

#63 / 189



130 млн

Выходных токенов из
индекса интеллекта

Kimi K3 — одна из ведущих моделей в области искусственного интеллекта, но она несколько дороже других моделей аналогичной стоимости. Кроме того, она медленнее, чем в среднем, и очень многословна. Модель поддерживает ввод текста и изображений, выводит текст и имеет контекстное окно в 1 млн токенов.

Kimi K3 набирает 57 баллов по индексу искусственного интеллекта, что значительно выше среднего среди сопоставимых моделей (в среднем 30). При оценке индекса интеллекта было сгенерировано 130 млн токенов, что очень много по сравнению со средним показателем 63 млн.


Цена Kimi K3 составляет \$3,00 за 1 млн входных токенов (довольно дорого, в среднем: \$1,75) и \$15,00 за 1 млн выходных токенов (довольно дорого, в среднем: \$8,40). В общей сложности оценка Kimi K3 по индексу интеллекта обошлась в 2709,75 \$.

При 62 токенах в секунду Kimi K3 работает медленнее среднего (73).

💡 Обоснование	Да ⓘ
⬇️ Режим ввода	T 🖼️ 🗣️ 📺
⬆️ Режим вывода	T 🖼️ 🗣️ 📺
📄 Контекстное окно	1 м ⓘ


189модели этого класса

+




Speed

Output tokens per second · Higher is better



Cost per Task

Weighted average cost (USD) per Intelligence Index task · Lower is better



Prompt Options



Artificial Analysis Intelligence Index

Artificial Analysis Intelligence Index v4.1 incorporates 9 evaluations: GDPval-AA v2, τ^3 -Banking, Terminal-Bench v2.1, SciCode, Humanity's Last Exam, GPQA Diamond, CritPt, AA-Omniscience, AA-LCR




28 of 576 models



+ Add model from specific provider

Artificial Analysis



💡 Reasoning models are indicated by a lightbulb icon

📘 Artificial Analysis Intelligence Index



Artificial Analysis Intelligence Index by Open Weights / Proprietary

Artificial Analysis Intelligence Index v4.1 incorporates 9 evaluations: GDPval-AA v2, τ^3 -Banking, Terminal-Bench v2.1, SciCode, Humanity's Last Exam, GPQA Diamond, CritPt, AA-Omniscience, AA-LCR



28 of 576 models




+ Add model from specific provider

Proprietary

Open Weights

Open Weights (Commercial Use Restricted)

Artificial Analysis



Reasoning models are indicated by a lightbulb icon

Artificial Analysis Intelligence Index

+

Open Weights

+

Intelligence Breakdown

Intelligence Evaluations

Intelligence evaluations measured independently by Artificial Analysis · Higher is better



18 of 22 evaluations




28 of 576 models




+ Add model from specific provider

GDPval-AA v2 ↗




τ^3 -Banking ↗

Agentic tool use




Terminal-Bench v2.1 ↗

Agentic coding & terminal use




SciCode ↗


Coding



Scientific reasoning




Physics reasoning



AA-Omniscience Non-Hallucination Rate

1 - hallucination rate



AA-LCR

Long context reasoning

Artificial Analysis



Agentic SaaS workflows




Legal agentic work, task all-pass rate



Agentic business operations


Artificial Analysis



IFBench

Instruction following


Artificial Analysis



APEX-Agents-AA


Long-horizon agentic tasks

Artificial Analysis



ITBench-AA

Kubernetes incident root-cause analysis




① Intelligence Evaluation Relevance

 $+$

① **Artificial Analysis Intelligence Index**

+

AA-Briefcase is an agentic knowledge work benchmark developed by Artificial Analysis. AA-Briefcase Elo is a combined metric that aggregates rubric pass rate, analytical quality Elo and presentation Elo · Higher is better



AA-Omniscience

AA-Omniscience Index

AA-Omniscience Accuracy

AA-Omniscience Hallucination Rat

AA-Omniscience Index

AA-Omniscience Index (higher is better) measures knowledge reliability and hallucination. It rewards correct answers, penalizes hallucinations, and has no penalty for refusing to answer. Scores range from -100 to 100, where 0 means as many correct as incorrect answers, and negative scores mean more incorrect than correct.

↶

↷


📊

28 of 446 models

⌵

⌵

+ Add model from specific provider



Reasoning models are indicated by a lightbulb icon

AA-Omniscience Index

Intelligence Index Comparisons

Intelligence vs. Cost per Task Intelligence vs. Time per Task Intelligence vs. Output Speed

Intelligence vs. Cost per Intelligence Index Task

Artificial Analysis Intelligence Index · Weighted average cost (USD) per Artificial Analysis Intelligence Index task

↶

📄

📊

28 of 576 models

⬆

⌵

Most attractive quadrant

Kimi

Google

Anthropic

Z AI

OpenAI


DeepSeek

SpaceXAI

NVIDIA

MiniMax

Meta



Cost per Intelligence Index Task	+
Artificial Analysis Intelligence Index	+

Token Use

Output Tokens per Task

Intelligence vs. Output Tokens per Task

Intelligence Index Token


Output Tokens per Intelligence Index Task

Weighted average number of output tokens used to run one task in the Artificial Analysis Intelligence Index

28 of 576 models

Answer

Reasoning



Reasoning models are indicated by a lightbulb icon

Output Tokens per Intelligence Index Task

Price and Cost

Cost per Task

Intelligence vs. Cost per Task

Evaluation Breakdown

Cost per Intelligence Index Task

Weighted average cost (USD) per Artificial Analysis Intelligence Index task, segmented by token type. Lower is better

28 of 576 models




Answer

Reasoning

Cache Write

Cache Hit

Input



💡 Reasoning models are indicated by a lightbulb icon

Cost per Intelligence Index Task

Intelligence Index Total Cost

Intelligence vs. Total Cost

Cost to Run Artificial Analysis Intelligence Index

Cost (USD) to run all evaluations in the Artificial Analysis Intelligence Index

↔

📷

📊


28 of 576 models

⌵

⌵

+ Add model from specific provider

Output Reasoning Cache Write Cache Read Non-Cache Input



Reasoning models are indicated by a lightbulb icon

Cost to Run Artificial Analysis Intelligence Index

Cache Hit, Input, and Output PricingBlended PriceBlended Price (Stacked)Cac


Pricing: Cache Hit, Input, and Output

Price (USD per M Tokens)

28 of 576 models

+ Add model from specific provider

Cache HitInputOutput



Reasoning models are indicated by a lightbulb icon

- Cache Hit +
- Input Price +
- Cache Pricing by Provider +
- Output Price +
- Model Performance Representation +

Context Window

Context Window Intelligence vs. Context Window

Context Window

Context window: tokens limit · Higher is better

28 of 576 models

↶


↷

⌵

⌶

+

Add model from specific provider



Speed

Measured by Output Speed (tokens per second)

Output Speed

Output Speed by Prompt Type

Output Speed Variance

Output Sp

Output Speed

Output tokens per second · Higher is better

↺

🖼️

📊


28 of 576 models

⌵

⌵

+

Add model from specific provider



Reasoning models are indicated by a lightbulb icon

- Output Speed

+
- Model Performance Representation

+

Time per Task Intelligence vs. Time per Task Cost vs. Time per Task

Time per Intelligence Index Task

Weighted average decode time (minutes) per task; excludes TTFT and overhead time .
Lower is better

28 of 576 models

↕

≡



Reasoning models are indicated by a lightbulb icon

Time per Intelligence Index Task

Latency

Measured by Time (seconds) to First Token

Time To First Answer Token

Time To First Token

Latency by Prompt Type

Latency

Latency: Time To First Answer Token

Seconds to first answer token received · Accounts for reasoning model 'thinking' time

↔

🔗

📊

28 of 576 models

⌵


⌵

+

Add model from specific provider

Thinking (reasoning models, when applicable)

Input processing



💡 Reasoning models are indicated by a lightbulb icon

Time to First Answer Token +

End-to-End Response Time

Seconds to output 500 tokens, calculated based on time to first token, 'thinking' time for reasoning models, and output speed

End-to-End Response Time

End-to-End Response Time by Prompt Type

End-to-End R

End-to-End Response Time

Seconds to output 500 tokens, including reasoning model 'thinking' time · Lower is better

↔

🔍

📊

28 of 576 models

⌵

⌵


+

Add model from specific provider

Outputting time

'Thinking' time (reasoning models)

Input processing time



💡 Reasoning models are indicated by a lightbulb icon

- 📘 End-to-End Response Time

+
- 📘 Model Performance Representation

+

Frequently Asked Questions

Common questions about Kimi K3

- When was Kimi K3 released?

+
- Who created Kimi K3?


+
- How intelligent is Kimi K3?

+
- How fast is Kimi K3?

+
- What is the latency of Kimi K3?

+

How much does Kimi K3 cost?	+
What is Kimi K3 API pricing?	+
How verbose is Kimi K3?	+
Is Kimi K3 a reasoning model?	+
What input modalities does Kimi K3 support?	+
What output modalities does Kimi K3 support?	+
Can Kimi K3 process images?	+
Is Kimi K3 multimodal?	+
What is the context window of Kimi K3?	+
Is Kimi K3 open source?	+
How many parameters does Kimi K3 have?	+
How does Kimi K3 perform on benchmarks?	+
Is Kimi K3 available via API?	+
Where can I use Kimi K3?	+



Artificial Analysis

Get notified about new articles

What's your email?

Subscribe

Explore

[LLM Leaderboard](#)

[Image Arena](#)

[Video Arena](#)

[AI Agents](#)

[Evaluations](#)

Company

[Methodology](#)

[Services](#)

[Contact](#)

[Articles](#)

[FAQ](#)

