

РАЗРАБОТКА ИИ / ЭКСПЛУАТАЦИЯ ИИ / СТРАТЕГИЯ ИИ

Microsoft, AWS и Anthropic тратят миллиарды — и не на улучшение моделей

Microsoft, AWS, OpenAI и Anthropic — все они стремятся внедрить своих инженеров в клиентские компании. Вот почему именно внедрение, а не модель, определяет будущее корпоративного ИИ.

5 июля 2026 г. в 11:00 [Джанакiram MSV](#)




Фото Игоря Омилаева на Unsplash

2 июля Джадсон Альтхофф, генеральный директор коммерческого подразделения Microsoft, **объявил** о создании компании Microsoft Frontier.

Новое подразделение будет включать в себя 6000 отраслевых и технических пертов, которые будут работать в организациях клиентов над

руководить подразделением будет Гэбриел Кедд-Лима, бывший президент Microsoft в Азии.

Это заявление было сделано всего через два дня после того, как AWS выделила 1 миллиард долларов на собственную инженерную организацию с опережающим развертыванием. В мае Anthropic и OpenAI уже запустили сервисные предприятия по тому же принципу. За две недели все крупные поставщики платформ искусственного интеллекта пришли к такому же выводу.

См. также: *AWS только что вложила 1 миллиард долларов в инженеров с опережающим развертыванием. Вот почему это важно для корпоративных команд.*

Ограничивающим фактором для корпоративного ИИ стала не сама модель, а инженерные ресурсы, необходимые для ее развертывания.

Ограничивающим фактором для корпоративного ИИ стала не сама модель, а инженерные ресурсы, необходимые для ее внедрения. Такой поворот коренным образом меняет отрасль, вынуждая компании пересматривать подходы к привлечению специалистов, конкурентную динамику глобальных системных интеграторов и базовую экономику для всех заинтересованных сторон.

Внутри компании Microsoft Frontier

Компания Frontier не является отдельным юридическим лицом. Большинство ее сотрудников уже работают в Microsoft и входят в состав существующих команд инженеров, технических консультантов, сотрудников службы поддержки и отраслевых отделов продаж. Штат будет расти за счет внутренних перемещений и найма новых сотрудников. Среди первых клиентов компании — Unilever и Novo Nordisk. Microsoft также заключила партнерские соглашения с глобальными системными интеграторами, включая Accenture, Capgemini, EY, KPMG и PwC.

Engineering and will be the largest, most capable, outcome-driven engineering organization in the industry.”

Althoff positioned the unit as more than a services arm: “This goes beyond what has been labeled as Forward Deployed Engineering and will be the largest, most capable, outcome-driven engineering organization in the industry,” he writes in the announcement.

Microsoft calls the offering Frontier Transformation, which combines industry knowledge, change management, and enterprise AI engineering. The company commits not to use customer data or intellectual property for model training.

TRENDING STORIES

1. [Microsoft, AWS and Anthropic are spending billions — and not on better models](#)
2. [Taming the agentic influx: a blueprint for AI business observability](#)
3. [Forward deployed engineer is AI's hottest job as OpenAI and Google race to hire. Here's how to become one.](#)
4. [10 moments that defined AI's turbulent first half of 2026](#)
5. [Why Anthropic just doubled Claude Cowork limits at no charge](#)

The financial details deserve a closer look than the headline number. Microsoft **declined** to say whether the \$2.5 billion is fresh capital or repurposed from existing consulting budgets. It has not specified the spending period either.

Microsoft already operates Industry Solutions Delivery, a sizeable consulting organization, alongside its FastTrack deployment programs. Whether the Frontier Company is genuinely a new capability or a rebranding with an increased focus on AI will become clear from customer engagements.

The Palantir playbook goes mainstream

2005. It placed software engineers inside intelligence agencies and enterprise clients whose data could not leave the building and whose requirements changed faster than any specification document.

The engineer worked within the customer's environment, understood the workflows firsthand, and shipped working systems rather than slide decks. The model remained unique to Palantir for over a decade, and its public filings still describe forward-deployed engineers as staff who help customers identify use cases and modernize data architectures.

The frontier labs adopted the playbook first. In May, Anthropic **launched** an AI services venture with Blackstone, Hellman & Friedman, and Goldman Sachs, aimed at midsize businesses deploying Claude. Days later, OpenAI unveiled its Deployment Company, backed by TPG, Advent International, Bain Capital, and Brookfield. Both labs chose joint ventures, which bring outside capital and partner firms into the deployment relationship rather than funding it from their own balance sheets.

The pattern goes beyond the platform giants and frontier labs. Cursor runs its own forward-deployed engineering team under Pauline Brunet. She **describes** the role at the AI Engineer World's Fair as going on-site, working inside a customer's systems, and deploying customized agent workflows across the software development lifecycle.

A startup selling a code editor and the world's largest software company have agreed on the same delivery model within a quarter. The embedded engineer closely follows the enterprise AI playbook rather than a niche tactic.

AWS vs. Microsoft: Two designs for the same job

There are quite a few similarities between the AWS and Microsoft programs. Both embed vendor engineers within customer teams to build production AI systems on the customer's own data and under the customer's governance — both fund the effort from their own balance sheets rather than through joint ventures.

Sprint versus embed

The key difference lies in the engagement design: AWS optimizes for short, self-validating deployments, whereas Microsoft builds a persistent presence within the account.

cycles, with each customer hosting a pod of five or six engineers. The unit describes itself as agentic-first, meaning the teams use purpose-built agents to build agentic systems.

Success is measured by whether the customer becomes self-sufficient, with knowledge graphs, runbooks, and trained internal staff left behind at the end of an engagement. AWS says teams are already embedded with Cox Automotive, the NBA, the NFL, Ricoh, and Southwest Airlines.

Microsoft, by contrast, describes continuous co-design and improvement with measurable business outcomes, a model that looks closer to an embedded transformation partner than a 45-day sprint. That approach is a calculated risk, since the same depth that strengthens the relationship also deepens the customer’s dependence on Azure tooling.

The four initiatives fundamentally differ in structure, capital, and intent.

Vendor	Commitment	Structure	Engagement design
Microsoft Frontier Company	\$2.5 billion, 6,000 people	Internal operating unit	Persistent embedded teams, outcome-based, multi-model by design
AWS Forward Deployed Engineering	\$1 billion, thousands of engineers	Internal organization, own balance sheet	45-day pods of 5 to 6 engineers, built to leave customers self-sufficient
OpenAI Deployment Company	Private equity-backed, reported at over \$4 billion	Joint venture with TPG, Advent, Bain Capital, and Brookfield	Deploys OpenAI models into large enterprises
Anthropic services venture	Reported at roughly \$1.5 billion	Joint venture with Blackstone, Hellman & Friedman, and Goldman Sachs	Targets midsized businesses deploying Claude

Each row carries a trade-off rather than a ranking. The joint ventures shift cost and risk to partners but dilute control over the customer relationship. The hyperscaler models keep control while absorbing the full cost of thousands of embedded engineers.

Why every vendor reached the same conclusion

enterprise AI pilots fail to produce a measurable return. Enterprises bought subscriptions to ChatGPT, Claude, Gemini, and Copilot at scale.

Deployment rather than model quality becomes a critical factor

They then realized that impressive demos rarely succeed when integrated with proprietary data and legacy workflows, and that, due to organizational inertia, deployment rather than model quality becomes the critical factor.

Althoff himself supplied the most candid admission of the quarter. He **said** that Microsoft made a mistake by binding the original Copilot exclusively to OpenAI models. Customers, according to him, care more about the combination of their data and the models than about any particular model.

Model choice becomes just a configuration while the embedded engineers become the durable relationship.

This assertion fundamentally shifts the competitive landscape, especially since it originates from the vendor with the industry's most profound single-model alliance. With the Frontier Company, Microsoft enables customers to mix OpenAI, Anthropic, Microsoft AI, and open-source models in a single engagement. Model choice becomes just a configuration while the embedded engineers become the durable relationship.

What buyers should ask

Customers can use this approach to their advantage. An enterprise negotiating an AI engagement can now demand outcome-based pricing and insist that the vendor's engineers train internal staff rather than create dependency.

The first question for enterprise buyers is ownership. Which parts of the deployed system are proprietary logic the customer keeps, and which are wrappers around the vendor's managed services? The second question for buyers is exit cost. When an AWS pod leaves after 45 days, or a Microsoft team concludes its rotation, what does it take to run and evolve the system without them?

Company competes for the same transformation budgets those firms have owned for decades.

A vendor with thousands of embedded engineers and direct access to its own product roadmap can undercut an SI on both speed and depth. The integrators that thrive will be those that move up to industry-specific process knowledge that platform vendors cannot replicate.

The forward-deployed engineering race marks a milestone in enterprise AI where distribution dominates capability as the competitive front. Model quality gaps between frontier vendors have narrowed, and enterprises now hedge across providers as a matter of policy. Proximity has now become the new focus of competition.

Vendors now fight over who sits closest to the customer's data, workflows, and budget. Microsoft's answer is compelling and convincing on paper, provided the Frontier Company turns out to be new capacity rather than a fresh coat of paint on its existing consulting business.

If the early engagements with Unilever and Novo Nordisk turn into referenceable outcomes, the Frontier Company becomes the template against which AWS, Google Cloud, and the frontier labs are measured. The next signal to watch is Google. A formal forward-deployed organization from Google Cloud would complete the hyperscaler set and confirm that the embedded engineer is now as strategic to the AI platform wars as the model itself.

TNS



Janakiram MSV (Jani) is a practicing architect, research analyst, and advisor to Silicon Valley startups. He focuses on the convergence of modern infrastructure powered by cloud-native technology and machine intelligence driven by generative AI. Before becoming an entrepreneur, he spent...

[Read more from Janakiram MSV →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

TRENDING STORIES

1. Microsoft, AWS and Anthropic are spending billions — and not on better models
2. Taming the agentic influx: a blueprint for AI business observability
3. Forward deployed engineer is AI's hottest job as OpenAI and Google race to hire. Here's how to become one.
4. 10 moments that defined AI's turbulent first half of 2026
5. Why Anthropic just doubled Claude Cowork limits at no charge

INSIGHTS FROM OUR SPONSORS



Enterprise-Grade MCP Access Control Is Here. Your Gateway Makes It Real.

3 July 2026

Kong Gateway Enterprise 3.11.x.x EOL

3 July 2026

AI Gateway vs. Direct LLM API Integration: The Architecture Decision Defining Your AI Strategy

2 July 2026



The AI Code Verification Bottleneck: Why Faster Code Generation Means Slower Reviews

3 June 2026

^' Code Review Is Still a Review

ay 2026



Deep Dive into SASL PLAIN and SCRAM in Kafka: Login Modules and Config Hot-Reload

6 July 2026

Building Secure, Resilient, and Compliant Fraud Detection With Confluent Cloud

29 June 2026

Stream Governance: Making Compliance a Property of Data in Motion

29 June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in
your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

FOLLOW TNS



ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)