


ИНФРАСТРУКТУРА ИИ / МОДЕЛИ ИИ / ОБЛАЧНЫЕ СЕРВИСЫ

Компания Moonshot запустила Kimi K3. Затем из-за высокого спроса подписки были закрыты в течение 48 часов.

Узким местом в сфере ИИ становится не обучение, а логический вывод, поскольку закрытие подписки на Kimi K3 показало, что мощность графических процессоров не может удовлетворить растущий спрос на агентные рабочие нагрузки.

21 июля 2026 г. в 16:48 от [Аманда Касвелл](#)



Moonshot AI стала последней компанией в сфере ИИ, которая обнаружила, что запуск популярной модели — это только полдела. Менее чем через два дня после выпуска **Kimi K3** компания прекратила прием новых подписчиков из-за того, что спрос исчерпал доступные мощности графических процессоров. Существующие пользователи сохраняют доступ, пока Moonshot расширяет свою инфраструктуру и постепенно возобновляет подписки.

рос на логический вывод превышает предложение

задач, связанных с кодированием и агентными системами, компании обнаруживают, что им требуется больше мощностей для логического вывода, чем они предполагали.

«Kimi K3 получила гораздо больше любви, чем мы ожидали, — пишет официальный аккаунт Moonshot в X. — За последние 48 часов спрос приблизился к пределу наших текущих возможностей. Мы наращиваем мощности так быстро, как только можем, и будем открывать новые места для подписки партиями».

«Мы наращиваем мощности так быстро, как только можем, и будем открывать новые места для подписки партиями.»

Для инженеров и разработчиков, занимающихся инфраструктурой, возникший дефицит мощностей наглядно демонстрирует, почему такие компании, как OpenAI, Anthropic и Moonshot, **ограничивают доступ** вместо того, чтобы продавать безлимитный доступ.

крытый вес, закрытая емкость

публикованное снижение веса на 27 июля. В Google Code Arena от Arena.ai KO обошла и GPT-5.6 Sol от OpenAI, и Claude Fable 5 от Anthropic. По более широкому **индексу искусственного интеллекта**, он уступает обоим: у Fable 5 60 баллов, а у Sol — 59. Но это не делает игру проще в управлении.

Open weights let anyone deploy the model, but whoever hosts it still has to pay the inference bill. Coding activities tend to tie up GPU resources far longer than a typical chatbot interaction, making it harder to keep latency low as more developers pile on.

“Agent tasks are not one-off question answering; they continuously generate, read, and process tokens during ongoing tasks,” wrote **Citigroup semiconductor analyst Peter Lee** in a **research note**.

TRENDING STORIES

1. **Kubernetes won the container decade. Google's Agent Substrate wants the next one.**
2. **Anthropic employees worked "literally around the clock" to keep Fable 5 from disappearing**
3. **Moonshot launched Kimi K3. Then demand shut down subscriptions in 48 hours.**
4. **Why most AI projects fail: It's infrastructure and people**
5. **Arm and Google offer a smarter option to run agentic AI workloads**

Lee argued that as developers build longer agentic workflows, lower inference costs are quickly “re-converted into higher total resource consumption,” shifting the bottleneck from compute to server memory.

Moonshot's subscription pause is a sign that keeping enough inference capacity online once developers start using it at scale might be just as hard as building the model.

generate, read, and process tokens during ongoing tasks.”

China's chip constraints compound crunch

For a company like Moonshot, this general industry bottleneck is compounded by regional infrastructure realities. Unlike traditional software companies, AI developers typically rent much of this computing power from cloud providers such as Alibaba Cloud, Tencent Cloud, and Huawei Cloud rather than owning extensive data-center infrastructure themselves.

The capacity crunch illustrates the mounting challenge facing Chinese AI developers as US export controls continue to restrict access to leading chip provider Nvidia's most advanced AI chips. As a result, companies such as Moonshot depend on a combination of older chips and domestically produced alternatives. These constraints have forced Chinese developers to concentrate strongly on software tuning and more efficient use of computing resources to narrow the performance gap with US rivals.

Token economics under pressure

The scramble for computing power has fueled a data-center construction boom across China. Alibaba has committed more than \$53 billion to AI and cloud infrastructure over three years, while ByteDance is reportedly considering spending as much as \$70 billion this year on AI data centers and related infrastructure.

AI companies typically charge customers based on the number of tokens, or units of text, processed by a model, making token prices a key measure of operating costs.

According to **Bernstein Research**, Moonshot charges \$3 per million input tokens and \$15 per million output tokens for Kimi K3. This makes it about 40% cheaper than Anthropic's Opus 4.8 and roughly 70% cheaper than Claude Fable 5.

“A world where there are only [two to three] dominant frontier labs with 90 percent inference margins is net negative for every other layer while being awesome for those [two to three] labs,” wrote Atreides Management founder **Gavin Baker** on X.

‘er argued that models like Kimi K3, Grok 4.5, and Muse 1.1 could shift value

For developers, Moonshot's subscription freeze serves as an **architectural warning**. The era of assuming infinite, cheap API access is ending.

"A world where there are only [two to three] dominant frontier labs with 90 percent inference margins is net negative for every other layer while being awesome for those [two to three] labs."

TNS



Amanda Caswell is an AI journalist, certified prompt engineer, and technology commentator whose work and expertise have been featured on Fox News and CBS News. She covers artificial intelligence, developer tools, foundation models, and emerging technologies, with a particular focus...

[Read more from Amanda Caswell →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY



TRENDING STORIES

1. Kubernetes won the container decade. Google's Agent Substrate wants the next one.
2. Anthropic employees worked "literally around the clock" to keep Fable 5 from disappearing
3. Moonshot launched Kimi K3. Then demand shut down subscriptions in 48 hours.
4. Why most AI projects fail: It's infrastructure and people

INSIGHTS FROM OUR SPONSORS

**When the Big AI bubble pops, we'll need Lean AI**

22 July 2026

Calero: Managing Enterprise Tech Spend

20 July 2026

Baz: Review-driven-development for agentic software lifecycles

20 July 2026

**Code review isn't your bottleneck, even with AI**

21 July 2026

Octopus Easy Mode - Argo CD Manifest Update

17 July 2026

The compliance ratchet

15 July 2026

**Scaling agent systems with Apache Kafka and Agent2Agent (Part 2): Understanding the A2A object model**

21 July 2026

Five years of OpenSearch: What we've learned building on the open source future of search

21 July 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in
your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

 roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



© The New Stack 2026

Disclosures
Terms of Use
Advertising Terms & Conditions
Privacy Policy
Cookie Policy

