


ИНФРАСТРУКТУРА ИИ / МОДЕЛИ ИИ / СТРАТЕГИЯ ИИ

Алекс Карп из Palantir и Артур Менш из Mistral сходятся во мнении: предприятия будут вынуждены использовать ИИ

Алекс Карп из Palantir и Артур Менш из Mistral с разных сторон приводят одни и те же аргументы: не позволяйте закрытым поставщикам ИИ контролировать ваши данные и внедрение.

6 июля 2026 г. в 13:17 от [Аманда Касвелл](#)



Рютаро Уодзуми, Unsplash

На прошлой неделе генеральный директор Palantir Алекс Карп принял участие в программе *Squawk Box* на канале CNBC, чтобы обсудить новое партнерство с Nvidia по внедрению моделей ИИ с открытым исходным кодом в государственных структурах. Но зрители получили почти 20-минутную тщательную речь в адрес всей индустрии передовых моделей искусственного интеллекта. Ведущий назвал ее «чертовски безумной» и

Спустя несколько дней генеральный директор Mistral Артур Менш выступил с поразительно **похожим заявлением в LinkedIn**, предупредив, что закрытые поставщики ИИ получают «огромное преимущество» перед корпоративными клиентами, поскольку организации подключают собственные рабочие процессы к размещенным моделям. Он предлагает использовать модели с открытым кодом, системы с открытыми данными и создавать собственные «маховики» для обучения.

Эти два руководителя смотрят на ситуацию с противоположных сторон рынка, но их совпадение в том, что касается одного и того же тезиса, на той же неделе подчеркивает важность архитектурного контроля.

Два предложения, один аргумент

Карп руководит компанией, которая продает приложения и онтологический уровень, призванные стать связующим звеном между предприятиями и моделями. В рамках **сделки между Palantir и Nvidia** открытые модели Nemotron от Nvidia объединяются с операционной системой Sovereign AI от Palantir, построенной на AIP, Foundry, Ontology и Apollo, что позволяет государственным учреждениям и операторам критически важной инфраструктуры развертывать, настраивать и проверять модели ИИ в собственных изолированных средах.

Palantir CEO Alex Karp says 'something has gone completely
CNBC Television

выпуска. Подписывайтесь на наш канал на YouTube, чтобы смотреть все наши подкасты, интервью, демонстрации и многое другое.

When CNBC's Becky Quick told Karp he sounded angry, he pushed back, saying, "This is the voice of American business that is being channeled through me," and urged the panelists to call any CEO privately to verify.

"This is the voice of American business that is being channeled through me."

Mensch's company sells open-weight models, and he has a custom training platform called Forge, which frames the problem differently but reached the same conclusion. He argues in the post that closed providers have a track record of going after their most successful customers once they learn what those customers are building. His program runs from open models to open data stores, strict access controls, and a continuous training flywheel that improves systems on internal interactions.

Lock-in gets an upgrade

If you've been building software at scale for any length of time, you already know that when new technology arrives, enterprises can't help but rush to adopt it. But the dependency problem becomes impossible to ignore.

TRENDING STORIES

1. **Palantir's Alex Karp and Mistral's Arthur Mensch agree: AI lock-in is coming for enterprises**
2. **AWS, Microsoft, and Google agree the session is the new unit of compute. They disagree on how to isolate it.**
3. **Microsoft just admitted its biggest AI mistake — and spent \$2.5 billion fixing it**
4. **The \$1.3 million theft that exposed AI's blind spot**

We saw the same problem with cloud computing when companies went all-in on a single hyperscaler's proprietary services, only to later discover that the cost of switching providers could exceed the cost of staying, even when staying meant overpaying. It's one of the reasons the industry spent years building abstraction layers, portability tooling, and multi-cloud strategies in response.

Foundation models are raising the same questions, but there's a twist. When an enterprise connects a model to its internal data, including customer records, proprietary processes, and domain-specific knowledge, the dependency becomes informational. Karp's argument is that model quality is converging across providers, but the operational leverage accrues to whoever controls the deployment layer and the data flowing through it.

Mensch's claim has a concrete referent that enterprise architects will recognize. In 2025, Anthropic cut off model access to coding startup Windsurf while building its competing product, Claude Code. The Brookings Institution has separately warned that model providers increasingly compete with their own customers as they chase application-layer revenue.

When access disappears overnight

When the U.S. government ordered Anthropic to suspend **access to its most advanced models** for foreign nationals, the company cut access across the board, including to enterprise customers in Europe who had built workflows on top of those models.

turned compute resources without warning. It's probably why the incident sent European policymakers into overdrive. Mensch, whose company had open-weight alternatives ready, seized the moment.

Enterprise teams need a plan for when a critical dependency is modified, repriced, or revoked by a provider whose incentives may not always align with their own.

Enterprise teams need a plan for when a critical dependency is modified, repriced, or revoked by a provider whose incentives may not always align with their own.

Architecture shifts toward portability

Enterprise AI architecture is now essentially this: don't marry a single provider, build for portability, keep your most sensitive data and logic under your own control.

In practice, this is showing up in three ways.

Firstly, we're seeing a portfolio **approach**. A powerful closed model is maintained for complex reasoning and customer-facing work, while an open-weight model is used for repetitive, high-volume tasks. For businesses handling sensitive information, the appeal is that an open model can be run entirely on their own infrastructure, so data never has to leave the organization.

The second is the rise of the **model-routing layer**, abstraction frameworks that let organizations swap models without rewriting their applications. Palantir's ontology pitch sits here as does the emerging crop of agent orchestration tools that treat the LLM as a pluggable component behind a standardized interface.

The third is the **open-weight movement** itself. Nvidia **shipped Nemotron 3 Ultra** in June under a permissive Linux Foundation license. Meta's Llama continues to expand. Mistral's Forge platform lets enterprises train custom models on their own data. And Mistral is **teasing an upcoming open-weight model** this summer, with early access opening in July.

deployment and governance layer; it benefits directly if enterprises treat models as interchangeable commodities. Mensch's Mistral sells open-weight models and a training platform, and it benefits directly if enterprises distrust closed providers. Zoho's Sridhar Vembu, who endorsed Karp's position publicly last week, has his own reasons for wanting enterprises to own their AI infrastructure rather than rent it from Silicon Valley.

But the fact that multiple executives across different market segments are saying companies need to own their data, maintain deployment flexibility, and not hand their competitive advantage to a provider who might become their competitor suggests the argument is resonating.

What developers should watch

If your most sensitive data is flowing through a third-party API with terms of service that can change, you've made a governance decision that your compliance team may not have fully evaluated.

For the engineering teams actually building on foundation models, the takeaway is that if your most sensitive data is flowing through a third-party API with terms of service that can change, you've made a governance decision that your compliance team may not have fully evaluated.

The engineering choice, to abstract the model layer, to evaluate open-weight options alongside closed APIs, to think about deployment portability the same way you think about cloud portability, is increasingly strategic.



artificial intelligence, developer tools, foundation models, and emerging technologies, with a particular focus...

[Read more from Amanda Caswell →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY



TRENDING STORIES

1. Palantir's Alex Karp and Mistral's Arthur Mensch agree: AI lock-in is coming for enterprises
2. AWS, Microsoft, and Google agree the session is the new unit of compute. They disagree on how to isolate it.
3. Microsoft just admitted its biggest AI mistake — and spent \$2.5 billion fixing it
4. The \$1.3 million theft that exposed AI's blind spot
5. Why cheaper models alone won't save your AI budget

INSIGHTS FROM OUR SPONSORS



Red Hat
OpenShift

[Provide access to Red Hat documentation in environments with limited connectivity](#)

6 July 2026

[BackendTLSPolicy expands Gateway API transport security](#)

July 2026

6 July 2026 **yugabyteDB**

Meko is Open for Self Sign-Up: Here's How to Get Started!

6 July 2026

YugabyteDB: The Data Backbone for Thousands of Agents

22 June 2026

A PostgreSQL Database for Every Agent

18 June 2026

**SUSE**

Reflections from Brussels, Utrecht and Paris

6 July 2026

SUSE AI Factory with NVIDIA is now generally available

6 July 2026

Your Patch Cycle Is Already Behind.

1 July 2026

TNS DAILY NEWSLETTER

**Receive a free roundup of the
most recent TNS articles in
your inbox each day.**

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors



Community created roadmaps, articles,
resources and journeys for developers

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



© The New Stack 2026

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)