

РАЗРАБОТКА ИИ / ИНСТРУМЕНТЫ ДЛЯ РАЗРАБОТЧИКОВ / РАЗРАБОТКА ПРОГРАММНОГО ОБЕСПЕЧЕНИЯ

Однопроходный ИИ-код не умер, но «высокий уровень логического вывода» — это следующий рубеж

Руководители CloudBees и Sparq спорят о том, стоит ли использовать модели ИИ с высоким уровнем логического вывода и оправдывают ли они затраты, или же однопроходные инструменты по-прежнему справляются со своей задачей.

21 июля 2026 г. в 13:13 от [Адриан Бриджуотер](#)



Ардиан Праномо

Спросите у ИИ-модели, что будет следующим после «бекона-двойного», и она почти единогласно ответит: «чизбургер». Это однопроходный результат, соответствующий шаблону, и именно его мы и ожидаем.

мыслительные процессы.

Выход за рамки однопроходного решения больше похож на человеческое мышление, поскольку он разбивает сложные задачи на более управляемые, последовательно выполняемые промежуточные этапы (с возвратом к предыдущим шагам в случае ошибки) для получения ответа.

Компромисс заключается в большем объеме более сложного кода, в котором сложнее обнаружить ошибки при проверке кода, но сама модель, как правило, работает лучше.

what comes after
"bacon double"

Теперь возникает вопрос: когда мы внедрим (или, возможно, даже будем настаивать на том, чтобы)

высокоуровневые рассуждения стали частью набора передовых моделей и

станет ли это следующим препятствием, которое нужно преодолеть?

Cheeseburger. 🍔



OpenAI

Однопроходное кодирование — для простых задач

Генеральный директор CloudBees Мориц Пласниг говорит *The New Stack*, что однопроходное кодирование с помощью ИИ — это как «отвечать первым, что приходит в голову» в споре о рассуждениях. Это значит, что такой подход подходит для простых задач, но не для тех, которые требуют глубокого понимания или неоднозначности.

Однопроходное ИИ-программирование — это как «отвечать первым, что приходит в голову».

«Я не думаю, что однопроходное программирование умерло, но я думаю, что оно будет использоваться только для самых простых задач, — говорит Пласниг. — Гораздо интереснее, когда команды грамотно распределяют работу: низкоуровневые модели — для шаблонных [рутинных и дсказуемых] задач, а высокоуровневые — для всего, что связано с архитектурными сложностями или критическими вопросами безопасности, а

TRENDING STORIES

1. Musk open-sourced Grok Build to fight Anthropic. Anthropic pays him \$1.25 billion a month.
2. Move code review before the code
3. Linux creator Linus Torvalds tells AI haters to walk away from Linux, or go fork it
4. How to kill the code review
5. "Time to clean up human slop": Why AI now reviews code better than your teammate.

Plassnig confirms that he fully expects the “current trajectory of developers demanding more high-reasoning to evolve into the harness itself”, before those same software engineers then being happy to choose the right model for the task at hand.

“What’s interesting is teams routing work intelligently: rather than just treating ‘more reasoning’ as a universal upgrade.”

Sometimes the first response is the right one

General manager of ‘The Shop’ (a practical use case company division) at AI software operations engineering company Sparq, Jackson Stakeman, tells *The New Stack* that he’s somewhat (but not wholly) in consensus with CloudBees CEO Plassnig’s thoughts. Rather than defining single-pass AI coding as blurting the first thing that comes to mind, he defines it as “responding when you think you actually have something of value to add” to a query.

He reminds us that single-pass blurt or not, sometimes the first response is the right one, so pretending otherwise just adds cost and latency to problems that didn’t need it.

“We’re still early in AI adoption at enterprise scale. The simple workloads are the thick part of the bell curve, and the complex multi-agent orchestration tasks are the long tail. For companies that aren’t primarily technology businesses, but that use technology to drive value, what they mostly need from AI is for it to work.”

“We’re still early in AI adoption at enterprise scale. The simple workloads are the thick part of the bell curve, and the complex multi-agent orchestration tasks are the long tail.”

The catch is that the model becomes a black box

In his essentially practical, hands-on role, Stakeman underlines the fact that high-reasoning tools let organizations that never built strong prompt engineering skills skip that step, for a price.

“The catch is that the model becomes a black box,” he warns. “It might solve your problem, but it won’t tell you whether a custom orchestration would have solved it better. In the conversations we’re having with customers, it’s rare to hear about a custom LangChain-style build that’s actually kept pace with how fast the frontier models keep improving.”

He restates his view that frontier model labs seem to be moving faster than most in-house stacks can keep up. It’s a development that’s starting to force a financial conversation that’s only just getting going: real skepticism about whether the premium on models like GPT-5.6 Sol or Fable 5 are worth it for a given task.

As of now, **OpenAI appears to talk** more directly about “medium, max, and scientific” reasoning, rather than high reasoning, even in relation to GPT-5.6 Sol. In a somewhat similar vein, **Anthropic talks about chain-of-thought faithfulness** when discussing why reasoning models don’t always say what they think.

“In our view, the deciding factor isn’t really the model; it’s whether a team has a clear way to measure ROI on AI spend. Teams with that mechanism can experiment and compare a cheap, fast model against a high-reasoning one, task by task. Teams without it tend to default to whatever’s marketed as ‘best’, and those are exactly the teams most likely to get steered away from high-reasoning the moment

High reasoning is not rocket science, or FrontierMath

By way of clarification, high reasoning is not the same as FrontierMath, but there is a degree of overlap between these two occasionally connected cousins.

OpenAI commissioned Epoch AI to produce 300 math questions for the FrontierMath benchmark, which **the nonadvocacy research nonprofit confirmed** at the start of 2025. In terms of use, FrontierMath can act as a benchmark for evaluating advanced mathematical reasoning in AI of the type that features high reasoning.

What do developers think of high-reasoning? It might still be early days; there's precious little ascerbic commentary out there just now.

That may be down to the still-emerging nature of these functions, or it could be because "high-reasoning" is sometimes used to describe **governmental policymaking** and even **Machiavellianism**.

What (arguably) matters most in the immediate months (and yes, this is AI, so it could be weeks or days) ahead is how well enterprise organizations are able to assess any given AI tool for its worth. That measure of worthiness comes down to correctness and effectiveness, multiplied by cost and divided by its ability to integrate and orchestrate before it attempts to accelerate and (dare we even say the word?) innovate.

TNS



Adrian Bridgwater is a technology journalist with three decades of press experience. He has an extensive background in communications, starting in print media, newspapers and also television. Primarily working as an analysis writer dedicated to a software application development 'beat',...

[Read more from Adrian Bridgwater →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY

TRENDING STORIES

1. Musk open-sourced Grok Build to fight Anthropic. Anthropic pays him \$1.25 billion a month.
2. Move code review before the code
3. Linux creator Linus Torvalds tells AI haters to walk away from Linux, or go fork it
4. How to kill the code review
5. "Time to clean up human slop": Why AI now reviews code better than your teammate.

INSIGHTS FROM OUR SPONSORS

[Open Protocols with the Strands Agents SDK](#)

16 July 2026

[Building secure AI agents at scale: Introducing Loom for AWS](#)

9 July 2026

[Introducing MCP server for Registry of Open Data on AWS](#)

7 July 2026

[Your agent wants to search like a 2010 quant](#)

7 July 2026

[Re-autoresearching MSMARCO BM25, on Vespa](#)

ay 2026



AI Will Not End Standard Software. It Will Reveal Why It Exists.

10 July 2026

How any nines Is Making a9s Hub Enterprise-Grade for Kubernetes at Scale

2 July 2026

Deploying Applications to Kubernetes from Source Code: Five Practical Approaches Compared

22 June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the most recent TNS articles in your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

FOLLOW TNS



ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)