


АГЕНТЫ ИИ / ИНФРАСТРУКТУРА ИИ / БЕЗОПАСНОСТЬ

Проблема идентификации агентов ИИ, о которой никто не говорит

Неясная идентификация приводит к тому, что агенты ИИ не проходят проверку на безопасность. Узнайте о четырех решениях, необходимых для обеспечения безопасности агентных систем на уровне платформы.

26 июня 2026 г. в 9:00 [Джексон Коннелл](#)



mana graphic для Unsplash+

[IBM](#) спонсировала этот пост.

Многие агентные проекты успешно проходят этап разработки. Затем наступает проверка безопасности — и вот тут-то все и может застопориться. Неясные модели идентификации и слишком широкие права доступа быстро становятся препятствиями.

проблем выполняет весь рабочий процесс. Затем служба безопасности задает простой вопрос: *Под чьим идентификатором это выполняется?* Ответ ставит процесс на паузу: это общая учетная запись с широкими правами доступа, без четкого владельца, без аудита и без контроля минимальных привилегий.

Коренная проблема не так сложна. Она заключается в неопределенности личности и плохо разграниченных правах доступа. И эта проблема быстро усугубляется. Исследование **2026 Tech Leader Study**, проведенное совместно с Oxford Economics и IBM, показало, что опрошенные компании планируют внедрить в среднем 1661 ИИ-агента, что на 38% больше, чем сегодня. Каждый новый агент — это еще одна личность, которую нужно защитить, и без четких границ проблема быстро усугубляется.

В результате многие агентные системы сосредоточены на том, что могут делать агенты, не определяя, что они должны делать и под чьим руководством. Кроме того, у агентов нет фиксированного набора разрешений. Они запрашивают доступ, вызывают новые инструменты и берут на себя новые роли в процессе работы, поэтому пути доступа комбинируются таким образом, что никто не может их ни предоставить, ни проверить. Без верифицируемой личности невозможна подотчетность, что затрудняет применение принципа наименьших привилегий, отслеживание и реагирование на инциденты.

«Каждый новый агент — это еще одна личность, которую нужно защитить, и без четких границ проблема быстро усугубляется.»

Чтобы устранить эти пробелы, мы подготовили это руководство для разработчиков, архитекторов и DevOps-инженеров, создающих агентные системы, а также для ИТ-руководителей, ответственных за их утверждение.



Cut complexity, fuel growth. IBM automation transforms tangled IT stacks into intelligent, streamlined systems. By infusing AI

[Learn More →](#)

HEAR MORE FROM OUR SPONSOR

svo1975pvn1982@gmail.com

[Submit](#)

The four identity decisions every agentic system must make

Identity decisions can't be treated as an afterthought. Identity shapes how agents authenticate, what they can access, and how their actions are controlled and audited over time. Get it wrong early, and you're building on a shaky foundation.

Here are the four decisions that matter most:

TRENDING STORIES

1. [Cordyceps flaw pattern is more proof CI/CD is part of the attack surface](#)
2. [After Fable 5 ban, Anthropic and 19 organizations launch open source security body](#)
3. [Aikido acquires Root to backport open source fixes without forcing upgrades](#)
4. [How to dramatically improve enterprise security alert tuning to battle cyberattacks](#)
5. [The call is coming from inside your pipeline: the anatomy of a Codecov attack](#)

Workload identity vs. shared service accounts

Shared service accounts are easy, and that's exactly what makes them dangerous. When multiple agents act under a single identity, it becomes hard to tell what

“Shared service accounts are easy, and that’s exactly what makes them dangerous.”

Workload identity assigns each agent **its own identity**. Permissions stay scoped, and actions are attributable. It requires more setup but creates isolation and auditability.

Static API keys vs. short-lived credentials

Static API keys tend to stick around forever. They get hardcoded into apps, passed around between systems, and rarely rotated — which makes them a persistent vulnerability waiting to be exploited.

Short-lived credentials work differently. They’re issued on demand, scoped to a specific task, and expire automatically. In practice, this often relies on identity federation (for example, using OIDC tokens) combined with systems that can issue dynamic credentials at runtime, rather than storing long-lived secrets in code or configuration

Direct credential handoff vs. brokered session access

Handing credentials directly to an agent is simple. It’s also opaque. You don’t have a natural point to evaluate policy or understand what’s happening in real time.

Brokered access introduces a control point into the flow. Requests go through a broker, policies are evaluated in real time, and temporary credentials are issued per session. It adds infrastructure, but restores visibility and policy enforcement.

Fragmented logging vs. full identity lineage

Most systems log what happened. Far fewer capture who initiated it or how an action propagated through a chain of agents and services.

Full identity lineage connects every step. You can trace an operation from triggers to outcomes, which can make debugging faster and enable more credible incident response. The catch is that this requires consistent identity propagation and structured logging from the beginning—it’s hard to retrofit.

vulnerabilities.

Nightfall AI reports that organizations expose nearly 350 secrets per 100 employees each year, with 35% of exposed API keys still active. Combine that with persistent credentials and shared identities, and the potential **blast radius grows fast**.

The pattern is consistent: shared accounts and long-lived keys are faster to build but harder to secure. Workload identity and short-lived credentials require more upfront investment but can deliver more security over time.

Debugging breaches by feel

Think about what happens when an agent running on a shared account with a long-lived key suddenly spikes its data access. Was it a bug? A breach? Routine behavior? Hard to say. Revoking the key might stop the issue, but it could also break a half-dozen unrelated workflows in the process. You're now debugging by feel.

Shortcuts reduce friction at the start and accumulate risk over time.

Standardize identity at the platform layer

The answer isn't to rebuild authentication, **authorization** and auditing from scratch for every agent you ship. That's not scalable.

Instead, standardize identity at the platform layer—centralized identity providers, policy engines and a credential broker to enforce secure defaults and make compliance straightforward rather than a constant negotiation.

“Shortcuts reduce friction at the start and accumulate risk over time.”

Agentic AI works in production when identity is designed up front and enforced at time, rather than assumed from a prior login. When projects are treated as an

IBM supports this through an integrated identity-first approach that spans secrets management, secured access, and identity governance—helping organizations scale agentic systems securely without adding operational complexity.

[Learn how IBM approaches identity-first security for agentic systems](#)

© Copyright IBM Corporation 2026. IBM and the IBM logo are trademarks of IBM Corp., registered in many jurisdictions worldwide. Examples presented are illustrative only. Actual results will vary based on client configurations and conditions; therefore, general expected results cannot be provided.

TNS



Jackson Connell is a Senior Product Marketing Manager at IBM, where he leads security and governance across identity, agentic AI security, proactive risk management, and compliance. He works across IBM's security portfolio, helping enterprises address the governance and trust challenges...

[Read more from Jackson Connell →](#)

SHARE THIS STORY



TRENDING STORIES

1. Cordyceps flaw pattern is more proof CI/CD is part of the attack surface
2. After Fable 5 ban, Anthropic and 19 organizations launch open source security body
3. Aikido acquires Root to backport open source fixes without forcing upgrades
4. How to dramatically improve enterprise security alert tuning to battle cyberattacks

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in
your inbox each day.

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

FOLLOW TNS



ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)