


АГЕНТЫ ИИ / РАЗРАБОТКА ИИ / РАЗРАБОТКА ПЛАТФОРМЫ

Coinbase, Shopify и Ramp создали собственных агентов для написания кода. Все три компании по-прежнему платят Anthropic.

Coinbase, Shopify и Ramp создали собственных агентов для написания кода на основе ИИ, но не отказались от Claude Code. Настоящая борьба идет за оболочку агента, а не за модель.

7 августа 2026 года, 14:16 [Джанакирам MSV](#)



Дж. Р. Корпа для Unsplash+

Инженерные команды предприятий переходят на единую архитектуру для разработки программного обеспечения с использованием искусственного интеллекта. Coinbase, Shopify и Ramp создали внутренних агентов для своих работчиков. Однако ни один из этих агентов не заменил коммерческие ...струменты, такие как **Claude Code**, Codex или Cursor.

Уровень, который эти компании решили взять под свой контроль, — это не большая языковая модель. Это агентская платформа — среда выполнения, которая обеспечивает контекст, разрешения, организацию рабочих процессов, доступ к инструментам и верификацию. Передовые модели от Anthropic, OpenAI и Google по-прежнему обеспечивают механизм логического вывода, в то время как предприятия всё больше берут на себя всё, что с ним связано.

Передовые модели от Anthropic, OpenAI и Google по-прежнему обеспечивают механизм логического вывода, в то время как предприятия все больше берут на себя все сопутствующие задачи.

Компания LangChain недавно признала, что эта архитектура стала воспроизводимой благодаря открытому исходному коду **Open SWE**, и назвала ее публичной реализацией того же шаблона, который уже используется внутри Stripe, Ramp и Coinbase.

Конкурентное преимущество находится за пределами модели

Несмотря на то, что Forge, River и Inspect разрабатывались независимо друг от друга, они решают очень похожие задачи.

Forge от Coinbase интегрирует Slack, GitHub и Linear в рабочий процесс, который позволяет преобразовать сообщение об ошибке в пул-реквест с минимальным вмешательством разработчика. После выявления проблемы Forge получает необходимый контекст репозитория, генерирует код, проверяет изменения и даже создает мобильную сборку для проверки.

River от Shopify использует тот же архитектурный подход. Компания считает, что своим успехом она во многом обязана не самому агенту, а базовой платформе для разработчиков. River работает на основе воспроизводимой «хенерной» среды, поддерживаемой монорепо-репозиторием для всей компании, локальными сеансами выполнения, изолированными средами, управлением

объединенном запросе на слияние.

ТРЕНДОВЫЕ ИСТОРИИ

1. Every software company will become a dev tools company
2. Coinbase, Shopify and Ramp all built their own coding agents. All three still pay Anthropic.
3. "Just rewrite it": What platform teams really think about modernization
4. Develop like you deploy: closing the Kubernetes local-to-cluster gap
5. Platform Engineering vs. DevOps Misses the Point

Shopify reports that River now participates in one out of every eight merged pull requests.

Ramp's **Inspect** extends the pattern further by executing inside sandboxed virtual machines that closely mirror developers' local environments. The platform integrates with Datadog, Sentry, LaunchDarkly, Buildkite, and other engineering systems while remaining model-agnostic. Inspect supports multiple frontier models alongside **MCP** servers, custom tools, and proprietary workflows.

Despite differences in implementation, all three platforms converge on the same architecture: durable execution environments, secure access to enterprise systems, workflow orchestration, and automated verification before code reaches production.

Enterprises are insourcing the harness

This is not unlike what happened with internal developer platforms over the past decade.

Enterprises built their own cloud infrastructure after adopting AWS or Azure. Instead, they built opinionated platforms on top of public cloud services that

The same pattern is now emerging around AI.

The language model becomes another infrastructure dependency, while the enterprise-owned harness determines how that intelligence is applied. Routing requests, enforcing permissions, injecting organizational context, executing tools, and validating results become responsibilities of the platform rather than the model provider.

Model selection consequently becomes a configuration choice instead of a workflow decision made by individual developers.

Cost optimization becomes a platform capability

Coinbase's experience illustrates why enterprises increasingly want to control this layer.

According to **Chintan Turakhia**, Senior Director of Production Engineering for Wallet and AI Engineering at Coinbase, the company's infrastructure now processes more tokens than before while reducing overall spending.

Coinbase co-founder and CEO **Brian Armstrong** attributes the improvement to multiple architectural optimizations rather than a single breakthrough. Coinbase introduced intelligent routing, lower-cost default models, prompt caching, smaller context windows, and better visibility into AI spending. The company has also experimented with open-weight models such as GLM 5.2 and Kimi K2.7.

These optimizations are only possible because Coinbase owns the gateway sitting between developers and foundation models.

Instead of asking thousands of engineers to switch tools, platform teams can update routing policies centrally. Model upgrades, pricing changes, and experimentation become infrastructure operations rather than organization-wide migrations.

Commercial coding assistants still have a role

Owning the harness has not eliminated demand for commercial coding assistants.

Codex in addition to River. Anthropic also ~~lists~~ **lists** Ramp among its enterprise customers, where nearly half of the engineering team uses Claude Code weekly.

The emerging division of responsibilities is becoming clearer.

Internal agents primarily handle asynchronous engineering workflows that begin in Slack or issue trackers, interact with internal systems, and validate changes against production telemetry.

Commercial assistants continue to dominate interactive development sessions in which developers work directly within their editors or terminals.

Rather than replacing each other, the two approaches complement different engineering workflows.

Platform ownership matters more than model ownership

For most enterprises, the question is no longer whether to build or buy an AI coding assistant.

The more important decision is whether to own the orchestration layer that sits between developers and foundation models.

Organizations with mature platform engineering teams can justify investing in proprietary harnesses because they gain centralized governance, model portability, security enforcement and cost optimization. Smaller engineering organizations are likely to continue relying primarily on commercial tools, where vendors absorb the operational complexity.

Requirement	Recommended option	Rationale
Asynchronous work that must query internal databases, feature flags, and telemetry to verify itself	Internal agent	Proprietary integrations, credentials, and verification steps become first-class rather than extensions bolted onto a vendor client
Interactive work on a familiar repository at the editor or terminal	Commercial assistant such as Claude Code or Codex	Engineers already know the ergonomics, and the vendor absorbs the operational complexity

across a large engineering organization	both	move once, without asking engineers to change tools
Teams without a dedicated platform engineering group	Commercial assistant with per-developer quotas	The investment behind Forge, River, and Inspect assumes people whose job is to maintain it

This mirrors the evolution of cloud platforms, where enterprises selectively built higher-level abstractions without attempting to replace the underlying infrastructure providers.

AI economics remain unpredictable

Owning the harness does not eliminate infrastructure costs.

Walmart recently **introduced** usage limits for its internal coding assistant after demand exceeded available budgets. Earlier this year, Uber **exhausted** its annual AI coding budget within four months, while GitHub Copilot has **shifted** to a usage-based credit and developer-quota model.

Recent **research** from Stanford University and Microsoft Research helps explain why these costs remain difficult to predict. Their analysis of agentic software engineering tasks found that autonomous coding workflows can consume up to one thousand times more tokens than interactive code generation. Even identical tasks frequently exhibited order-of-magnitude differences in token consumption, making budgeting significantly harder than in conventional cloud infrastructure.

Cost optimization therefore becomes another responsibility of the platform layer rather than the model itself.

The platform becomes the strategic asset

The most significant architectural shift is that enterprises are beginning to treat AI agents as part of their **internal developer platform**.

A decade ago, organizations differentiated themselves through deployment pipelines, infrastructure automation, and platform engineering. Today, they are extending the same philosophy to AI-assisted software development.

security, execution, verification, and economics.

Model providers will compete to become the preferred reasoning engine inside platforms they do not control.

If common agent runtimes and open protocols continue to mature while enterprises keep their policy and context layers proprietary, competition will increasingly move above the foundation model. Model providers will compete to become the preferred reasoning engine inside platforms they do not control.

That may prove to be the defining architectural shift of enterprise AI. The strategic advantage lies in moving away from the model itself and toward the platform that determines how, when, and where it is used.

TNS



Janakiram MSV (Jani) is a practicing architect, research analyst, and advisor to Silicon Valley startups. He focuses on the convergence of modern infrastructure powered by cloud-native technology and machine intelligence driven by generative AI. Before becoming an entrepreneur, he spent...

[Read more from Janakiram MSV →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic, LaunchDarkly.

SHARE THIS STORY



TRENDING STORIES

1. Every software company will become a dev tools company
2. Coinbase, Shopify and Ramp all built their own coding agents. All three still pay Anthropic.

Just rewrite it": What platform teams really think about modernization

INSIGHTS FROM OUR SPONSORS



VMware Tanzu

August 19: Democratizing AI and Analytics in Financial Services with VMware Tanzu & Strongly.AI

4 August 2026

Using AI for Product Management – The Tanzu Platform product management experience

4 August 2026

Timing is Everything: Bringing Nanosecond Precision to Tanzu Greenplum PXF Parquet Pipelines

28 July 2026



TiDB
Powered by PingCAP

The Complete Agent State Stack: Memory, Files, and Serverless Database Persistence for AI Apps

7 August 2026

Why Attend TiDB SCaiLE 2026: Same Complexity, Different Clock Speeds

6 August 2026

Open Source Data Layers Will Win the AI Future

5 August 2026



DuploCloud

End-to-End DevOps: Evolving Strategies & Automation Trends

January 2024

10 NOVEMBER 2023

How the Terraform License Change Impacts Your DevOps

21 September 2023

TNS DAILY NEWSLETTER

Receive a free roundup of the most recent TNS articles in your inbox each day.

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data

Open Source

Networking

Storage

Security

Software Development

WebAssembly

OPERATIONS

AI Operations

CI/CD

Cloud Services

DevOps

Kubernetes

Observability

Operations

Platform Engineering

CHANNELS

Podcasts

Ebooks

Events

Webinars

Newsletter

TNS RSS Feeds

THE NEW STACK

About / Contact

Sponsors

Advertise With Us

Contributions

 roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap

Backend Developer Roadmap

Devops Roadmap

FOLLOW TNS



© The New Stack 2026

Disclosures

Terms of Use

Advertising Terms & Conditions

Privacy Policy

Cookie Policy

