

glean

Year in Review 2025

Table of Contents

03

Introduction

04

2025 by the numbers

05

The Glean Work AI Platform

06

Market recognition

07

The 10 shifts that defined
Work AI in 2025

15

Glean's platform evolution

18

Our ecosystem

20

Our relentless focus on
customers

21

Guiding principles for building
Work AI

22

Looking ahead to 2026

Introduction

There is no shortage of AI in the enterprise. What's scarce is AI that organizations trust enough to put at the center of how they operate and to run the work their businesses depend on.

In 2025, Glean became that layer of trusted infrastructure for our customers: a Work AI platform that unifies company knowledge across tools, data, models, and agents, and provides a single, governed source of enterprise context.

Today, our customers rely on Glean to understand and coordinate how work happens—bringing together their knowledge, their processes, and how their people work. With more than \$250M in annual revenue (>150% year-over-year growth), Glean now sits inside some of their highest-stakes workflows: account planning and renewals, forecasting and operating reviews, regulatory and customer audits, security incidents and post-mortems, and high-stakes research and diligence.

Our own organization scaled to meet that responsibility in 2025. We grew to more than 1,100 employees, with customers spanning more than 27 different countries, and shipped more than 260 new capabilities. We expanded support for open standards like MCP, deepened our connector and governance stacks, our underlying context graph, and broadened model choice—so customers can connect all their systems, and any frontier model, through a single, trusted control plane for Work AI.

This year-in-review captures what we learned, the key shifts that defined Work AI in 2025, and how we're building toward 2026.

Arvind Jain

Arvind Jain
Founder & CEO, Glean



2025 by the numbers

In 2025, Glean became the trusted Work AI platform for our customers, powering more of their daily work with greater scale, depth, and reach across the enterprise.

>\$250M ARR

>150% YoY revenue growth

40% DAU/MAU

typical SaaS: 10–20%

20T tokens-per-year run rate

1,600% YoY growth

100+ out-of-the-box connectors

6 average queries per user per day

vs 4 Google searches per user per day

27B+ documents indexed since inception

270M Glean Assistant actions

210% YoY growth


80% of customers deploy Glean across 5+ departments



The Glean Work AI Platform

Glean is the Work AI layer where data, tools, models, and agents come together under a single set of security and access controls. Under the hood, the platform is powered by several core components:

- **Glean Assistant.** The primary interface for employees to ask questions, find information, and create content—always grounded in company knowledge and existing permissions.
- **Glean Agents.** A common, governed foundation where department-specific and third-party agents run side by side, so mission-critical workflows follow consistent safety, privacy, and access controls.
- **Enterprise Graph.** The Glean Assistant and agent platform are powered by the Enterprise Graph, a connected map of people, data, processes, and tasks that links documents, tickets, conversations, and systems into a single view of how work flows across the organization.
- **Connectors.** An ecosystem of 100+ connectors and 100+ actions that integrate with tools like Slack, Teams, Google Workspace, Microsoft 365, Salesforce, and Databricks, allowing agents to operate inside existing systems while honoring source-system permissions and governance.
- **Glean Protect:** Glean Protect provides fine-grained control over AI access and actions, with policy enforcement and auditability for security, compliance, and risk teams.



Market recognition

In 2025, Glean's momentum extended beyond product deployment to shape how the market discusses Work AI and context-rich platforms.

- Glean was also named to multiple lists recognizing the most innovative and fastest-growing AI companies, including [Fast Company's Most Innovative Companies](#), [CNBC's Disruptor 50](#), [Bloomberg's AI Startups to Watch](#), and [Forbes' AI 50](#).
- We were recognized as an Emerging Leader in the 2025 AI Knowledge Management Apps / General Productivity Emerging Market Quadrant.
- We continued to earn top customer ratings on Gartner Peer Insights and G2 for relevance of results, governance, and ease of deployment.

“Companies often struggle to mine their own experience and institutional knowledge to stay competitive. When workers don’t know where to look for information, otherwise crucial data is useless. That’s where Glean steps in...”

Fast Company

“The growing enterprise information stack that is dispersed across diverse systems adds complexity and solutions like Glean, enables organizations to deploy AI powered knowledge retrieval solution focused streamlined for customer workflows and systems...”

Fortune

“Glean’s rise through the AI boom has been uniquely tied to the challenges that enterprises face when trying to apply AI.”

Bloomberg

What our customers say

The clearest evidence of impact came from how enterprises described the change in their day-to-day work.

“Bringing together all of the internal knowledge sources through a chat interface has made it much easier for people to do their jobs and find the information they need.”

— John Feldmeier, CIO, eBay

“We looked at frankly all of the major players in the place and Glean just blew them away from a user experience standpoint. Now we're in production across thousands of users at Dell. We're saving hours a week per person. That means thousands of hours per week of productivity back into the business. The results have just been tremendous.”

— Sam Grocott, SVP PMM and AI, Dell

“Glean plays a critical role in uniting a distributed workforce here at Zillow. Best of all, it goes beyond just enterprise search—it’s a solution that fully lets us leverage all our data through the latest AI models. With an 80% adoption rate of Glean, it’s helping everyone work closer together than ever before.”

— Toby Roberts, SVP Engineering, Zillow

The 10 shifts that defined Work AI in 2025

In 2025, we saw several fundamental shifts in how enterprises approached Work AI—from the platforms they chose to the way they governed agents and organized teams. These are ten of the biggest shifts that defined the year.

Context graphs became the competitive moat

1

Despite breathtaking progress in foundation models, most organizations still struggled to turn AI into business results. Benchmarks like MCP-Universe showed that even frontier models such as GPT-5 fail on more than half of complex, real-world orchestration tasks—not because they can't reason, but because they lack enterprise context.

That's why context graphs moved from architectural nice-to-have to competitive moat in 2025. A context graph is a connected map of the organization: it links people, content, systems, and events so AI can see not just what exists—documents, tickets, dashboards—but how work actually happens over time: how incidents unfold, how deals progress, how approvals flow across teams.

Glean's years-long investment in its context platform and enterprise context graph means customers can pair the latest models with real company context, so agents behave more like knowledgeable coworkers than generic copilots. In recent blind tests on ~280 complex enterprise queries, human graders preferred Glean's answers as correct 1.9× more often than ChatGPT's and 1.6× more often than Claude's.

Organizations chose open platforms over model lock-in

2

By 2025, it became clear that the real risk wasn't picking "the wrong" model—it was hard-wiring your architecture to any single model while the frontier kept shifting. Models leapfrogged one another so quickly that leading enterprises began prioritizing open, model-agnostic architectures.

In these setups, deep reasoning could go to one model, high-volume summarization to a cheaper one, and multimodal work to a third—without rewriting workflows, rebuilding agents, or revisiting data governance every time the model landscape changed. The platform became the stable layer; models became interchangeable components. Glean was built for exactly this world: an open Work AI platform with a model-agnostic agentic engine and model hub, so customers can swap and mix providers while keeping a single, governed layer for enterprise context, security, and orchestration.

AI infrastructure became horizontal and composable

3

Early enterprise AI systems were built as vertical stacks: a single vendor bundling models, tools, and agents into one tightly coupled product. That model broke down as AI capabilities and business needs multiplied; each new workflow, system integration, or model upgrade required custom work, slowing progress and locking teams into rigid architectures.

By 2025, leading organizations recognized that this approach wouldn't hold. They invested in a horizontal AI foundation for search, enterprise context, security, and orchestration, rather than rebuilding those capabilities inside every product. Tools, models, and agents plugged into this foundation using open standards like MCP.

At Canva, for example, teams already had a broad ecosystem of tools—but understood that LLMs perform best when they have context. With Glean's easy-to-configure MCP server, they now pull in policies and coding guidelines directly into agent workflows, influencing outputs and unlocking more creative (and compliant) workflows end-to-end.

Agents moved from chatbots to autonomous, multi-agent systems

4

As models became swappable and AI infrastructure more horizontal, the next challenge moved up the stack: how to organize AI at the “work” layer. A single assistant in a chat window couldn’t reliably handle complex, multi-step workflows that spanned teams and systems. Leading enterprises turned instead to multi-agent systems—networks of specialized agents that could own specific workflows end-to-end, coordinate with each other, and reliably drive work to completion.

At Zillow, thousands of Glean agents now support everyday work: one reconstructs incident timelines, another synthesizes UX research (“Ask Z”), another prepares annual reviews. They monitor for appropriate triggers, activate when needed, and hand off to each other—completing most of the heavy lifting before managers or engineers even engage.



“In five to ten years, the way we work will radically change—all our teams will have their own Glean AI agents assisting them.”

Benson Tan

Senior Manager for Automation Solutions

AI only works where work lives

5

In a work world where workers switch between apps and tabs up to 1,200 times per day, the primary constraint on AI impact in 2025 wasn’t capability—it was integration and overcoming the “toggle tax.” When AI lived in a separate application or tab, employees treated it as an occasional detour rather than a core part of their workflow. When it was embedded directly into the systems and processes they already used, it shifted from experiment to infrastructure.

At Booking.com, AI wasn’t launched as “one more destination.” Glean was embedded into existing IT support and content production workflows. Ticket resolution dropped from minutes to seconds, and video script timelines compressed from eight weeks to two—because AI operated inside the tools and rituals teams already trusted, not outside them.

Every major function adopted specialized agents for its own workflows

6

The most consequential AI work in 2025 moved from central innovation teams into the functions themselves. Each major function began building and adopting specialized agents for its core workflows, often starting from a shared agent library and adapting agents to their own metrics and rituals.

Function	Top agents	Business outcomes	Example
Engineering	• <u>Pull Request Review Agent</u> • <u>Engineering Self-evaluation Agent</u> • <u>Engineering Onboarding Agent</u>	• Accelerate delivery timelines • Improve code quality • Reduce time to first commit • Automate documentation creation	At Nextdoor, engineers use Glean agents to answer questions about current code, identify last deploys and code limitations, and understand feature implementations—doubling individual productivity by replacing manual digging with direct answers.

Function	Top agents	Business outcomes	Example
Product	• Sprint Planning Agent • Feature Prioritization Agent • <u>Meeting Recap Agent</u> • Product Usage Metrics Agent	• Ship features faster • Improve feature product market fit • Improve customer adoption rates • Accelerate sprint planning and execution	At Zillow, the AskZ agent indexes and democratizes access to customer research and insights so the whole organization can use data to make product decisions.
Sales	• <u>Account Snapshot Agent</u> • <u>Deal Strategy Agent</u> • <u>Prospect Outreach Agent</u>	• Increase win rates & rep participation • Increase pipeline generation • Accelerate time to first deal • Improve forecast accuracy	Motive’s sales team uses an Account Plan Creation agent that combines CRM data, meeting notes, and call transcripts to cut planning time from three days to two hours—turning static annual plans into living artifacts reps can refresh in a few clicks.
IT	• <u>IT Help Desk Agent</u> • <u>Support Documentation Agent</u>	• Optimize IT support efficiency • Accelerate business innovation • Reduce time to proficiency • Reduce total cost of ownership	A publicly traded cloud security company uses Glean agents as an IT help desk bot and onboarding copilot—cutting L1/L2 ticket volume by 31%, reducing onboarding time by 1.5 weeks, and saving ~\$120K annually plus 3 hours per IT employee per week.
Service & Support	• <u>Support Ticket Timeline Agent</u> • <u>Ticket Next Steps Agent</u> • <u>Customer Sentiment Analysis Agent</u>	• Reduce operational costs • Improve customer satisfaction • Increase first call resolution rate • Reduce inefficient call transfers	T-Mobile’s IT team uses a Glean-powered Internal Service Desk Agent to open, close, and update tickets automatically, escalate based on sentiment, and route complex cases—reducing time to transfer cases to live agents.
Marketing	• <u>SEO Keyword Research Agent</u> • <u>Content Drafting and Publishing Agent</u> • <u>Customer Testimonials Agent</u>	• Enhance customer insights • Increase campaign effectiveness • Automate content generation • Streamline event execution	Booking.com’s marketing team uses Glean to accelerate video script creation—reducing creation time per video by 75% and increasing output from 2 to 5 videos per month.

Function	Top agents	Business outcomes	Example
HR	• HR Policy Q&A Agent • <u>Intelligent Reminders Agent</u> • Employee Onboarding Agent	• Optimize HR operations • Accelerate talent acquisition • Increase employee engagement • Reduce employee attrition	Pure Storage's HR team built Glean agents to answer global benefit and policy questions and support new hires—reducing new-hire-related support tickets by over 50% in the first three months.
Legal	• Contract Review Agent • Security & Privacy Q&A Agent • <u>Delegation Tracker Agent</u>	• Accelerate contract reviews • Improve contract management • Manage regulatory workflows • Reduce external counsel spend	A publicly traded oil & gas company uses Glean to quickly surface matter-related documents and prep for meetings, saving 6 hours per lawyer each week and reducing outside counsel spend by \$84K.
Operations	• <u>IT Help Desk Agent</u> • <u>Meeting Action Tracking Agent</u> • Inventory and Supplier Management Agent	• Enhance operational efficiency • Optimize maintenance operations • Bolster supplier management • Reduce SLA misses	A leading consumer brand uses Glean agents to vet and onboard new suppliers, saving 115+ hours per month, and boosting “ease of onboarding suppliers” scores by 34%.
Finance	• Financial Reporting and Data Analysis Agent • Account Snapshot and Status Agent	• Improve financial planning • Accelerate financial reporting • Simplify financial operations • Streamline investment research	A Fortune 500 finance team uses Glean to bring invoicing purchase-order data from multiple systems into one place, saving ~2 hours per employee each week, and reducing \$100k in redundant IT spent.
Security	• Security Threat Modeling Agent • Incident Response Assistant Agent • Vulnerability Detection Agent	• Accelerate incident response time • Improve incident response quality • Reduce document creation time • Speed post-mortem analysis	A leading fintech company uses Glean agents to streamline incident investigations and documentation, reducing time spent documenting incidents by 37%.

Governance became the prerequisite for scale, not an afterthought

7

As AI agents moved into real operational workflows in 2025, governance became the limiting factor. Without clear roles, permissions, and audit trails, projects stalled in security and legal reviews.

Leading organizations changed their approach: they built governance into the architecture from day one. At DBS bank, for example, teams expanded agent development across functions while using Glean Protect to automatically enforce data permissions. Experimentation scaled rapidly while access remained predictable, auditable, and policy-compliant—turning pilots into durable, production deployments.

New roles emerged to own AI outcomes

8

As AI moved into core workflows, “AI” could no longer be a side responsibility. Organizations began appointing specific individuals to help lead AI transformations, including:

- AI champions were embedded in Sales, Customer Success, HR, and Engineering. They translated between domain experts and technical teams, identified candidate workflows, shaped early AI use cases, gathered feedback, and converted scattered experiments into repeatable playbooks.
- AI outcomes managers took explicit ownership of business results from AI, not just deployments. They tied AI initiatives to metrics like cycle time, NRR, case resolution, or time-to-productivity, and operated like product managers for AI outcomes—deciding where to scale, pivot, or shut down efforts. At Glean, our AI Outcomes Managers work with customers to define baselines and success metrics and codify what works into repeatable playbooks that can be rolled out across regions and business units.

The real bottlenecks were organizational, not technical

9

As headlines piled up about failed AI pilots in 2025, a pattern became clear: even with the most capable technology, impact stalled without deliberate investment in human and organizational change.

- **TIME** paired deployments with classic change management: naming clear owners, investing in training, and treating enablement as an ongoing program rather than a one-time launch.
- **Zillow** made AI part of the culture, running “AI Days,” sharing examples in dedicated channels, and converting early adopters into teachers for their peers.
- **Pure Storage** identified internal champions explicitly responsible for driving adoption.

Leading AI organizations understood that changing how people work was as critical as changing what technology could do.

Impact was measured in workflows, not AI usage


10

In 2025, leading organizations recognized the importance of measuring AI success in terms of clear business outcomes instead: cycle times, resolution rates, participation, NRR, deal velocity, and time-to-productivity.

They stopped asking “How many people used AI this week?” and started asking, for example:

- How much faster did we close incidents or audits?
- How many hours did we remove from annual reviews or account planning?
- How much did participation, coverage, or quality improve in core rituals?

In this lens, agents weren’t generic AI add-ons, but intentional rewires of how work gets done, with named owners and clear measures of success. The organizations that focused on outcomes rather than raw usage are the ones that turned AI into a compounding advantage.



Glean's platform evolution: from enterprise search to leading Work AI platform

From day one, Glean's mission has been to develop the deepest possible understanding of each organization. Between 2021 and 2025, Glean evolved from AI-powered enterprise search to a full Work AI platform with agents at scale.


We began by helping employees find information—documents, messages, tickets, and people—on top of secure, governed enterprise search. By 2023, we layered Glean Assistant, LLMs, and retrieval-augmented generation (RAG) on that foundation so people could get insight by finding answers and summarizing information and communications.

In 2024, we extended this into a broader Work AI platform that combined the assistant, agents, and APIs to automate personal work like drafting content, summarizing communications, research, data analysis, and custom chatbots.

By 2025, that foundation supported agents running at scale, powered by deep enterprise context, a horizontal orchestration layer, and governance built in from the start.

Glean pioneered AI at work


2021 Enterprise Search



Find information:

- Find docs
- Find messages
- Find people


2023 Enterprise Assistant



Get insight:

- Find answers
- Summarize information & comms


2024 Work AI Platform Building Agents



Automate personal work:

- Draft content
- Summarize comms
- Research
- Data analysis
- Custom chatbots

2025 Work AI Platform Agents at scale



Automate business processes:

- Departmental task automation
- Team of agents for every employee
- Developers building with security + data

Platform capabilities introduced in 2025

Structured data comprehension

We introduced structured data comprehension so anyone can ask a business question in plain language and get a trusted answer. Glean now pulls together data from tools like Salesforce, Snowflake, Databricks, and spreadsheets—alongside unstructured content—so people can ask things like “Which Q3 renewals are at risk?” and see a live view across systems, without waiting on a dashboard or writing SQL.

Contextual image understanding, surfacing, and generation

We built contextual image understanding so people can quickly see the chart or diagram they need without digging through decks. When someone asks for “the latest system design diagram for our billing service” or “the Q4 revenue waterfall chart,” Glean can surface or generate that visual directly in the answer—so time goes into reviewing and deciding, not hunting for slides.

Deep Research capabilities

We developed Deep Research so people can hand big questions to an agent and get back a finished brief, not a pile of links. The agent breaks a broad question into smaller research tasks, reviews the right sources, and returns a clear, context-rich summary—ideal for competitive analysis, strategic planning, and market research that used to take days of manual work.

Universal knowledge (web search)

We added universal knowledge so people can see internal context and outside information in a single view. Glean combines data across 100+ enterprise apps with real-time web results, so a rep can prep for a meeting with internal deal notes and the customer’s latest news, or an engineer can debug using both internal runbooks and public best practices.

Glean Canvas

We created Glean Canvas so people can turn a blank page into a solid first draft in one place. Canvas gives a side-by-side workspace to pull in company knowledge, draft, and revise, then send the output straight to Google Docs, Word, email, or Slack—removing copy-paste work and speeding up everything from briefs to reviews.

Unified Agent Builder


We introduced Unified Agent Builder so teams can turn repetitive workflows into agents without writing code. Teams describe what they want the agent to do in plain language, refine it visually, and publish it on shared infrastructure—while Glean automatically applies the right permissions and guardrails so those agents are safe to run in production.

Our ecosystem

Glean's value multiplies through deep integrations with the tools, models, and platforms organizations already rely on. By meeting people and data where they already work, Glean turns a fragmented stack into a single, trusted layer for search, agents, and automation—reducing the toggle tax, shortening time to value, and preventing risky AI sprawl.

Integrations

AI is only as powerful as the context it can access. Glean's connectors and actions span enterprise content systems, collaboration tools, data platforms, and industry-specific applications—including Slack, Microsoft Teams, Google Workspace, Microsoft 365, Salesforce, Databricks, and more than 100 other apps—plus custom and proprietary sources via open APIs. This breadth enables customers to give agents a unified view of their knowledge, rather than stitching together point integrations for each workflow.



“Glean connects every corner of our Microsoft ecosystem—and everything beyond it—with a minimal learning curve, giving our teams the clarity to build smarter and deliver faster.”

Rico Bertucci

VP of Innovation, McCarthy Building Companies

Model partners

OpenAI | Anthropic | Google | Amazon Bedrock | VertexAI

We maintain deliberate model agnosticism, partnering with frontier models like OpenAI, Anthropic (Claude), and Google (Gemini), and others through customers' preferred clouds. This allows customers to select the optimal model for each task—reasoning, speed, multimodal, or cost—while maintaining one consistent layer for context, security, and user experience, rather than rebuilding agents with every model change.

Development & security

Snowflake | Databricks | Palo Alto Networks

Together with Snowflake, Databricks, and Palo Alto Networks, Glean brings Work AI to where data and security already live. Customers can activate Glean on top of their governed data platforms and security controls—so agents reason over trusted Snowflake and Databricks data, while Palo Alto Networks helps enforce enterprise-grade protection. The result: AI that is grounded in your most trusted data, secured by your existing defenses, and ready for production.

Work AI Institute

Stanford | Harvard | University of California at Berkeley | University of California at Santa Barbara | Emory University | University College London | University of North Carolina at Charlotte | University of Notre Dame

In 2025, we announced the Work AI Institute, Glean's research center dedicated to understanding how AI transforms work and equipping leaders with playbooks and frameworks to navigate the evolving world of work. Our first report, the AI Transformation 100, shared 100 concrete moves leaders are making to transform their work with AI and was accompanied by a Harvard Business Review article on the five tensions leaders need to navigate as they bring AI into their organizations.



Our relentless focus on customers

In FY26, being “customer-obsessed” meant letting real deployments and outcomes with customers drive what we built. That mindset shaped how we designed spaces to connect with customers—and to celebrate their successes. We brought that to life through Glean:GO, our first user conference; our Customer Advisory Days; and an ongoing drumbeat of Glean:LIVE events and monthly product Drops.



Guest lectures at Glean Academy

We launched Glean Academy Guest Lectures as an ongoing “classroom for Work AI,” helping customers master not just the technology, but the change management needed to make it stick. Sessions dig into how real companies redesign workflows, govern AI safely, bring skeptical teams along, and prove ROI with business-grade metrics—featuring leaders from organizations like Reddit, T-Mobile, Pure Storage, Nextdoor, and more.



The Gleaniverse

The Gleaniverse is our community hub for Work AI practitioners, a place where customers share playbooks, agents, and real-world lessons from deploying Work AI. From case studies and frameworks to live discussions on governance and organizational change, the Gleaniverse turns individual successes into shared insights.

Guiding principles for building Work AI

Working with hundreds of organizations deploying agents at scale, five principles proved non-negotiable.

“Bringing together all of the internal knowledge sources through a chat interface has made it much easier for people to do their jobs and find the information they need.”

John Feldmeier
CIO, eBay

1. Context before models

Work AI should start from the reality of your organization, not the promise of a model. The durable advantage comes from a rich context layer—connecting systems, cleaning and classifying content, and mapping how people, work, and decisions relate. When that foundation is strong, models become interchangeable components; when it’s weak, even the most advanced models fall short.

2. Openness and choice by design

Design for a world where models continuously leapfrog one another across reasoning, speed, cost, and modality. Work AI platforms should be open by default across clouds, models, tools, and third-party agents so organizations can add or swap capabilities without rebuilding their stack. In practice, that means model-agnostic routing, open standards like MCP, and APIs that let existing and future systems plug into the same context and governance layer.

3. Governance enables scale

Treat governance as core infrastructure, not an afterthought. Agents without clear permissions, policies, and audit trails don’t scale—they get slowed, constrained, or shut down by security and legal. Governance should be first and foremost the platform’s responsibility, not the customer’s.

4. Build for builders, not just end users

Enable every function to build and adapt its own agents, not just central innovation teams. Platforms should make it easy for a spectrum of builders—from low-code business users to senior engineers—to create, compose, and govern agents on shared infrastructure. Agent libraries, natural-language and visual builders, and clear evaluation frameworks turn “AI projects” into a repeatable internal capability rather than one-off experiments.

5. Amplify the best of humanity

Design AI to elevate human work, not replace it. Agents should handle what machines do best—searching, aggregating, summarizing, routing—so people can focus on what only they can do: judgment, creativity, relationships, and leadership. Success should be measured not just in tickets closed or hours saved, but in fewer low-value tasks, deeper customer and team interactions, and more time spent on decisions and ideas that propel the organization forward.

Looking ahead to 2026



In 2026, winning with AI will have less to do with model horsepower and more to do with how well organizations combine three capabilities: context, control, and change.

- Context means letting AI see the whole business, not just a single system.
- Control means making it safe to deploy AI broadly, not just in isolated pilots.
- Change means deliberately redesigning roles, workflows, and incentives so humans and agents work side by side, rather than simply layering AI on top of existing processes.

Glean is building for that world: an open ecosystem for tools and models, safeguards strong enough that risk leaders can move from “no” to “how,” and agents that remove sludge from work so people can focus on what only they can do.

That’s the future we’re committed to at Glean in 2026 and beyond.

[Request a demo](#)