



Вселенная

Предиктивная аналитика с помощью SAS и R

Авторы: Рамчандра С. Мангрулкар, Паллави Виджай Чаван

Тип файла: pdf

Размер: 2,3 МБ

Язык: Английский

Страниц: 176

Предиктивная аналитика с помощью SAS и R: основные концепции, инструменты и реализация

Введение

Предиктивная аналитика преобразует исторические и текущие данные в прогнозы о том, что может произойти в будущем. Она сочетает в себе **статистику, машинное обучение, подготовку данных, теорию вероятности и вычислительные методы** для выявления закономерностей, которые могут помочь в принятии инженерных, бизнес-, научных и операционных решений. 🇮🇳 🤖



Двумя широко используемыми средами для предиктивной аналитики являются **SAS** и **R**. SAS предоставляет интегрированную коммерческую экосистему с отработанными процедурами, корпоративным управлением, отчетностью и возможностями управления моделями. R — это язык статистического программирования с открытым исходным кодом и огромной экосистемой пакетов для статистического анализа, визуализации, машинного обучения и исследований.


Для студентов и специалистов в области инженерии понимание обеих сред имеет большое значение, поскольку предиктивная аналитика — это не просто выбор алгоритма. Для успешного решения необходим полный цикл:


Данные → Очистка → Исследование → Признаки → Модель → Валидация → Прогнозирование → Мониторинг

Цель состоит в том, чтобы разработать модели, которые будут хорошо обобщать новые наблюдения, а не просто запоминать исторические данные. 🎯

PREDICTIVE MODELING WORKFLOW







Теоретическая база

Предиктивная аналитика строится на нескольких математических и статистических основах.

Статистика и теория вероятностей

Предположим, у нас есть набор данных:

$$[D=\{(x_i, y_i)\}_{i=1}^n]$$

где:

- (x_i) — входные переменные,
- (y_i) — целевая переменная,
- (n) — количество наблюдений.

Прогностическая модель пытается оценить:

$$[\hat{y}=f(X)]$$

где (f) определяется на основе исторических данных.

Для задач классификации модель может оценивать вероятность:

$$[P(Y=1|X)]$$

Например, инженерная организация может оценить вероятность того, что оборудование выйдет из строя в течение следующих 30 дней.

Обучение с учителем

В предиктивной аналитике обычно используется обучение с учителем, при котором исторические наблюдения содержат известные результаты.

Примеры:

- Линейная регрессия
- Логистическая регрессия
- Деревья решений
- Случайные леса
- Градиентный бустинг
- Нейронные сети
- Метод опорных векторов

Обобщение модели

Важнейший инженерный принцип — различие между **эффективностью обучения** и **эффективностью в реальных условиях**.

Модель может достичь:

[Точность_{обучения}=98%]

при этом значительно хуже работать с новыми данными.

Такая ситуация обычно связана с **переобучением**.

Поэтому в надёжном процессе прогнозирования данные разделяются на обучающую, проверочную и тестовую выборки.

Определение

Что такое предиктивная аналитика?

Предиктивная аналитика — это систематическое использование исторических данных, статистических методов, алгоритмов машинного обучения и вычислительных инструментов для прогнозирования будущих или неизвестных результатов.

Она отличается от описательной аналитики.

Analytics Type	Main Question	Example
Descriptive	What happened?	Monthly production decreased 8%
Diagnostic	Why did it happen?	Equipment downtime increased
Predictive	What may happen?	Failure probability is 72%
Prescriptive	What should we do?	Schedule preventive maintenance

SAS and R in Predictive Analytics

SAS is particularly strong in structured enterprise analytics, statistical procedures, governance, reporting, and production environments.


R provides extensive flexibility for statistical computing, visualization, experimentation, research, and machine learning.

A simplified comparison is:

[SAS = Enterprise\ Analytics + Governance + Statistical\ Procedures]

[R = Open\ Source + Statistics + Visualization + Machine\ Learning]

Step-by-Step Predictive Analytics Implementation



Step 1: Define the Engineering Problem

Before opening SAS or R, define the problem mathematically.

For example:

Predict whether a manufacturing component will fail during the next 30 days.

The target variable could be:

```
[Y=  
\begin{cases}  
1 & \text{Failure}\  
0 & \text{No Failure}  
\end{cases}]
```

Step 2: Collect Data

Potential variables include:

- Temperature
- Vibration
- Pressure
- Operating hours
- Maintenance history
- Load
- Production rate
- Previous failures

Data can originate from databases, sensors, CSV files, APIs, or enterprise systems.

Step 3: Clean the Dataset

Real-world datasets frequently contain:

- Missing values
- Duplicate observations
- Outliers
- Incorrect data types
- Inconsistent units
- Invalid measurements

For example, temperature values might accidentally contain both °C and °F.

A simple conversion is:

$$[C=(F-32)\times\frac{5}{9}]$$

Ignoring such inconsistencies can severely damage model performance.

Step 4: Explore the Data

Exploratory data analysis helps engineers understand relationships.

Useful statistics include:

$$[\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i]$$

for the mean, and:

$$[s^2 = \frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n-1}]$$

for sample variance.

Visualization can reveal relationships that numerical summaries hide.

Step 5: Engineer Features

Feature engineering transforms raw measurements into useful predictive variables.

For example:

$$[\text{VibrationRate} = \frac{\text{Vibration}_t - \text{Vibration}_{t-1}}{\Delta t}]$$

A rapidly increasing vibration level could provide stronger predictive information than vibration alone.

Step 6: Split the Data

A common structure is:

- 60–70% training
- 15–20% validation
- 15–20% testing

For time-series engineering applications, random splitting may be inappropriate. Earlier observations should generally be used to predict later observations.

Step 7: Select a Model

The model should match the problem.

For continuous outcomes:

$$[y = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p + \epsilon]$$

Linear regression may be appropriate.

For binary outcomes:

$$[P(Y=1)=\frac{1}{1+e^{-z}}]$$

where:

$$[z=\beta_0+\beta_1x_1+\cdots+\beta_px_p]$$

Logistic regression can be useful.

Step 8: Train the Model

In R, engineers can use packages such as `caret`, `tidymodels`, `ranger`, and `xgboost`.

A conceptual R workflow might look like:

```
1 | model <- glm(  
2 |   failure ~ temperature + vibration + pressure,  
3 |   data = training_data,  
4 |   family = binomial  
5 | )  
6 |  
7 | pred <- predict(  
8 |   model,  
9 |   newdata = test_data,  
10 |   type = "response"  
11 | )
```

SAS provides procedures such as PROC LOGISTIC, PROC REG, PROC GLM, PROC HPFOREST, and other analytical procedures for predictive modeling.

Step 9: Evaluate Performance

Accuracy alone is not sufficient.

For classification, engineers can examine:

$$[\text{Precision}=\frac{TP}{TP+FP}]$$

$$[\text{Recall}=\frac{TP}{TP+FN}]$$

and:

$$[F1=2\frac{\text{Precision}\times\text{Recall}}{\text{Precision}+\text{Recall}}]$$

For regression, common measures include:

$$[MAE=\frac{1}{n}\sum|y_i-\hat{y}_i|]$$

and:

$$[RMSE=\sqrt{\frac{1}{n}\sum(y_i-\hat{y}_i)^2}]$$

Step 10: Deploy and Monitor


A predictive model is not finished when its accuracy is measured.


Production systems should monitor:


- Prediction quality
- Data drift
- Feature drift
- Model drift
- Missing values
- Processing failures
- Business or engineering outcomes

⚙️ **Model monitoring is an engineering requirement, not an optional feature.**

Comparison: SAS vs. R







Feature	SAS	R
Licensing	Commercial	Open source
Statistical analysis	Excellent	Excellent
Visualization	Strong	Excellent
Machine learning	Strong	Extensive ecosystem
Enterprise governance	Very strong	Depends on infrastructure
Customization	Strong	Extremely flexible
Community packages	Large	Very large
Learning curve	Moderate	Moderate

Feature	SAS	R
Research flexibility	Strong	Excellent
Enterprise integration	Excellent	Depends on environment
Cost	Licensing required	No software license
Reproducibility	Strong	Strong

When SAS May Be Preferable

SAS can be attractive when an organization requires:


- Enterprise governance
- Centralized analytics
- Established reporting
- Formal workflows
- Regulatory controls
- Large-scale organizational deployment

When R May Be Preferable


R can be especially attractive for:

- Academic research
- Statistical experimentation
- Custom modeling
- Data visualization
- Open-source projects
- Rapid prototyping
- Advanced statistical methods


Diagrams, Tables, and Predictive Modeling Architecture



PREDICTIVE MAINTENANCE ARCHITECTURE



A general predictive analytics architecture can be represented as:



Model Selection Table

Problem	Suitable Models	Common Metrics
Continuous prediction	Linear regression, Random Forest	MAE, RMSE, (R^2)
Binary classification	Logistic regression, Trees	Accuracy, F1, AUC
Multiple classes	Random Forest, Boosting	Macro F1, Accuracy
Time series	ARIMA, ETS, ML	MAE, RMSE, MAPE

Problem	Suitable Models	Common Metrics
Failure prediction	Logistic, Trees, Boosting	Recall, AUC
Customer prediction	Logistic, Boosting	AUC, F1

Examples

Example 1: Predicting Equipment Failure

Suppose an industrial dataset contains:

Variable	Example
Temperature	87 °C
Vibration	8.2 mm/s
Pressure	12.5 bar
Operating Hours	8,400
Maintenance Age	140 days
Failure	1

A classification model can estimate:

$[P(\text{Failure}=1|X)=0.81]$

The resulting **81% probability** does not mean failure is guaranteed. It means the model estimates a relatively high risk based on learned patterns.

Example 2: Predicting Energy Consumption

An energy model could use:

$[\text{Energy}=f(\text{Temperature},\text{Load},\text{Time},\text{Production})]$

The output might be:

$[\hat{E}=14,250\text{ kWh}]$

Engineers could then compare predicted consumption with actual consumption.

Example 3: Regression in R

A simple regression model can estimate an engineering quantity from several explanatory variables:

```
1 | model <- lm(  
2 |   energy ~ temperature + load + production,  
3 |   data = train  
4 | )  
5 |  
6 | summary(model)
```

The coefficients help indicate how the predicted response changes when explanatory variables change, assuming the model assumptions are reasonably satisfied.

Real-World Applications

Predictive Maintenance

Manufacturing companies can analyze sensor information to estimate equipment failure risk.

Benefits may include:

- Reduced unplanned downtime
- Better maintenance scheduling
- Improved asset utilization
- Reduced maintenance waste

Energy Engineering

Predictive models can forecast:

- Electricity consumption
- Peak demand
- Equipment efficiency
- Renewable-energy output

Civil Engineering

Predictive analytics can support:

- Structural health monitoring
- Pavement deterioration prediction
- Construction risk analysis
- Project delay prediction

Automotive Engineering

Models can estimate:

- Component failure
- Battery degradation
- Fuel consumption
- Driver behavior

Healthcare Engineering and Analytics

Predictive methods can support operational forecasting, resource planning, and risk stratification, provided that appropriate privacy, validation, and regulatory requirements are followed.

Common Mistakes

Using Too Many Variables

More variables do not automatically mean a better model.

If irrelevant features are included, the model can become unnecessarily complex.

Ignoring Data Leakage

Data leakage occurs when information unavailable at prediction time accidentally enters the model.

This can produce impressive but misleading performance.

Optimizing Only for Accuracy

Consider a failure-detection problem where only 2% of machines fail.

A model predicting “no failure” every time could achieve:

[Accuracy=98%]

yet be practically useless.

Ignoring Class Imbalance

Rare-event problems require suitable evaluation techniques, sampling strategies, thresholds, or cost-sensitive approaches.

Treating Correlation as Causation

If:

$[Correlation(X,Y)\neq 0]$

this does not automatically prove:

$[X\rightarrow Y]$

Predictive models identify useful relationships; they do not automatically establish causal mechanisms.

Challenges & Solutions

Challenge	Problem	Solution
Missing data	Incomplete observations	Imputation or appropriate exclusion
Data imbalance	Minority class ignored	Resampling, weighting, threshold optimization
Overfitting	Poor generalization	Cross-validation and regularization
Data leakage	Unrealistic performance	Strict feature-time controls
Model drift	Performance declines	Continuous monitoring
Interpretability	Difficult decisions	Explainable models and diagnostics
Large datasets	Computational limitations	Efficient processing and scalable infrastructure

Challenge	Problem	Solution
Poor data quality	Unreliable predictions	Data validation pipelines

Explainability Challenge

Engineers often need to understand **why** a model produces a prediction.

For example:

Why did the model assign a 78% probability of equipment failure?

Methods such as feature importance, partial dependence, permutation importance, and SHAP-based explanations can help investigate model behavior.

Case Study: Predictive Maintenance

Consider a hypothetical manufacturing facility operating 500 industrial pumps.

The engineering team collects:

- Vibration measurements
- Temperature
- Pressure
- Flow rate
- Motor current
- Operating hours
- Maintenance records

The team creates a binary target:

[Y=1]

if the pump experiences a failure within 30 days.

Model Development

The engineers first build a logistic regression baseline.

They then compare it with a tree-based model.

Model	Recall	Precision	F1
Logistic Regression	0.72	0.64	0.68
Decision Tree	0.76	0.61	0.68
Random Forest	0.84	0.69	0.76

The random forest provides the strongest balance in this hypothetical example.

However, the team should not automatically deploy it solely because its F1 score is highest.

They must also consider:

- Interpretability
- Computational cost
- False alarm costs
- Missed failure costs
- Maintenance capacity
- Data drift
- Operational integration

The engineering decision should therefore combine **statistical performance with operational requirements**. 🛠️📈

Essential Tips

Start With a Baseline

Always build a simple model first.

A baseline gives you something against which more complicated models can be evaluated.

Use Domain Knowledge

A machine-learning algorithm cannot replace engineering understanding.

An experienced engineer may recognize that a sudden pressure change is physically meaningful even when a generic algorithm does not immediately reveal the reason.

Validate Properly

Use cross-validation where appropriate, but respect the structure of the data.

Time-dependent data should generally preserve temporal order.

Track Experiments

Record:

- Dataset version
- Features
- Algorithms
- Hyperparameters
- Metrics
- Random seeds
- Training date

Reproducibility is especially important in professional engineering environments.

Keep the Model Simple When Possible

If two models have similar performance:

[Performance_A $\approx$ Performance_B]

prefer the model that is easier to understand, maintain, validate, and deploy—unless there is a strong reason to choose otherwise.

FAQs

What is predictive analytics?

Predictive analytics uses historical and current data, statistics, and machine-learning methods to estimate future or unknown outcomes.

Is SAS better than [R](#) for predictive analytics?

Neither is universally better. SAS is particularly strong for enterprise analytics and governed environments, while R offers extensive flexibility, statistical functionality, visualization, and open-source packages.

Is R difficult for engineering students?

R has a learning curve, but students familiar with basic programming and statistics can learn it progressively. Starting with data frames, visualization, regression, and model evaluation is a practical path.

Can SAS and R be used together?

Yes. Organizations can use both environments within broader analytics workflows. For example, R may be used for experimentation while SAS supports established enterprise analytical processes.

Which models should beginners learn first?

A useful sequence is:

1. Linear regression
2. Logistic regression
3. Decision trees
4. Random forests
5. Gradient boosting
6. Time-series models
7. Neural networks

Understanding the fundamentals is more important than learning dozens of algorithms.

What is the most important part of predictive analytics?

Data quality and problem definition. A sophisticated algorithm cannot reliably compensate for poorly defined targets, biased datasets, data leakage, or incorrect measurements.

Can predictive analytics guarantee future outcomes?

No. Predictive models estimate probabilities or expected values. Uncertainty remains because future conditions may differ from historical conditions.

What should engineers monitor after deployment?

Engineers should monitor prediction accuracy, data quality, feature distributions, model drift, system performance, and changes in the underlying operating environment.

Conclusion

Predictive analytics with **SAS and R** provides engineers and analysts with powerful approaches for transforming data into actionable predictions. The most important concept is that predictive analytics is a complete engineering workflow—not simply the act of running a machine-learning algorithm.


A successful implementation begins with a clearly defined problem, continues through reliable data preparation and feature engineering, and then progresses to model selection, validation, deployment, and monitoring.

SAS offers a mature enterprise-oriented environment with strong statistical procedures, governance, and organizational capabilities. R offers remarkable flexibility, an extensive package ecosystem, powerful visualization, and strong support for statistical research and machine learning.

Ultimately, the best predictive analytics system is not necessarily the most complicated one. It is the system that produces **reliable, validated, interpretable, maintainable, and operationally useful predictions**.

For engineering students, learning the mathematical foundations alongside SAS and R creates a strong analytical foundation. For professionals, combining domain expertise with rigorous predictive modeling can turn raw measurements into early warnings, better forecasts, optimized operations, and smarter engineering decisions. 🚀📊⚙️

Related Posts:



Preview PDF Book

[← Previous Book](#)

[Next Book →](#)

EngiVerse

Explore a vast library of free engineering ebooks curated to help you expand your technical knowledge, conduct groundbreaking research, and advance your career. Our mission is to create a hub of knowledge that fosters innovation, empowers aspiring engineers, and bridges the gap between education and industry.

[Home](#) [Books](#) [About Us](#) [Contact](#)
[Privacy Policy](#) [DMCA](#)

copyright 2025 EngiVerse. All Right Reserved