




Why Attend?

Discover how data enables intelligence during the 7th annual Supermicro Open Storage, which features 12 information-packed sessions. This year's event focuses on implementing enterprise AI, building infrastructure to support AI inference and agentic AI, and managing data for AI, CSPs and enterprises. All sessions provide actionable recommendations from industry veterans.

These technically focused panels bring together industry experts from across Supermicro's ecosystem, including the leading data management, SSD and HDD storage media and silicon manufacturers. Thirty-eight industry leaders from 21 companies are featured, including AMD, DDN, Hammerspace, Intel, MinIO, Nutanix, Scality, Solidigm, VAST Data, WD and many more. Sessions with neocloud provider Crusoe and CSP Iron Mountain add practitioner perspective.

All sessions air at 10:00 a.m. PDT and are available on demand following the broadcasts. Register for free to view any of the sessions.

Agenda



AUGUST 11



Optimizing AI Inference Performance and Cost with Multi-Tier Storage

AI inference storage requires a tradeoff between performance and cost. Achieving the targeted performance at the budgeted cost requires careful engineering and selecting the best storage technology for each task. Combining a high-performance parallel file system using all-flash media with an object storage tier based primarily on HDDs provides both high performance and optimized TCO. This session brings together leaders in file and object storage, flash media, HDDs and CPUs from WEKA, Scality, Samsung Semiconductor, WD and Supermicro to describe how it all comes together in an engineered solution.

AUGUST 12



Vertical AI Solutions

While they typically use common AI hardware and software, vertical industries such as financial services, life sciences/pharma and manufacturing have unique requirements, including response times, data and workflows. In this panel with DDN, Kioxia, WD and Supermicro speakers examine financial market exchanges and genomics workflows as examples of how the software, SSD and HDD storage media, as well as system requirements, are addressed and where commonalities and differences emerge.

AUGUST 13



Moving AI from POC to Production

Like all IT initiatives, AI projects require user acceptance testing, integration testing and deployment. However, they also face additional hurdles before reaching production, including cost and token economics, the selection of scalable infrastructure, data readiness, access and governance, and user onboarding. This panel with Nutanix, MinIO, PEAK:AIO and Supermicro discusses these challenges and offers best practices based on past deployments.

AUGUST 18



Infrastructure for Neocloud AI

Neoclouds, which specialize in large-scale AI workloads, offer the scale of hyperscalers while maintaining an AI-specific focus. This session focuses on how neoclouds provide differentiated services, scale and storage to support customers' AI workloads. Representatives from Crusoe discuss how the neocloud differentiates its features and functionality alongside partners Supermicro, VAST Data and Kioxia.

AUGUST 19



Architecting AI Inference and Data Management for Maximum Scale



Supermicro, and Kioxia can provide a highly scalable solution for AI inference. The session also explains why KV Cache and AI Data Platform (AIDP) are critical for improving inference response time and optimizing the design of an AI Factory. And the benefits that the combine solution brings for improving time-to-first-token responses and optimizing the data and compute design of AIDP on Supermicro all-flash storage systems.

AUGUST 20



HAMMERSPACE

SANDISK™

Enterprise Storage Modernization for AI

Enterprises often have a variety of siloed legacy storage systems. Adding new capacity involves extending these legacy systems or deploying greenfield infrastructure. AI introduces new requirements for incorporating existing storage systems and data while adding purpose-built capacity. This session features Hammerspace, Sandisk and Supermicro discussing how to integrate existing and new capacity with emerging AI workflows.

AUGUST 25



VAST



SOLIDIGM.

Breaking the Context Wall: Storage for Scalable Agentic AI

Agentic AI at scale will require high-capacity, low-latency context memory to store and retrieve key value, or KV, cache data used in the inference workflow. New storage tiers and architectures have been proposed to solve this problem by balancing the reprocessing of inference queries against the retrieval of previously computed tokens. This new tier, positioned between the local GPU system SSDs and network storage, has been called “context memory” and allows long context tokens to persist in a large-scale storage array. This session with VAST Data, Solidigm and Supermicro describes the implementation, uses and trade-offs of the context memory tier.

AUGUST 26



MINIO

Enabling AI with Data Lakes and Lakehouses

Data lakes and lakehouses have found a new purpose in AI beyond traditional analytics and data warehousing. Modern data lakes are built on object storage infrastructure and use industry-standard data formats to ensure interoperability. The transactional infrastructure capabilities of data lakehouses also support AI workflows natively, including data pipeline processing, inference workloads and context memory serving. Speakers from MinIO, AMD and Supermicro discuss the changing roles of data lakes and lakehouses.

AUGUST 27



SOLIDIGM.

AI Data Platforms for Enterprise AI

AI data platforms represent a new category of AI infrastructure that handles the pre-processing of enterprise data, including data ingestion, normalization, vectorization and the hand-off of data to the enterprise AI factory. These scalable systems, which incorporate both data management and GPU infrastructure, can start small for workgroups and scale as demand grows. Preintegrated software and hardware enable faster on-site deployment and operational readiness. Built-in workflows for RAG, visual search and document processing, among other capabilities, provide an accessible starting point for AI. Experts from DDN, Solidigm and Supermicro discuss AI data platform use cases.



SEPTEMBER 1



Edge-Cloud Storage-as-a-Service for CSPs and Enterprise

The next generation of cloud storage incorporates on-premises edge services that provide a low-latency local storage node that can be centrally managed by the CSP. In this session, CSP Iron Mountain describes the next phase of Iron Cloud, including an on-prem solution for local backup and retrieval that also moves data to the Iron Cloud backend for long-term retention. This solution incorporates Scality object storage and uses Supermicro backend and edge systems.

SEPTEMBER 2



Managing Unstructured Data for AI

Unstructured data accounts for 80%-90% of an enterprise's unique data. Most enterprises retain only 5% of the data they generate due to cost and the rapid expansion of AI workflows, IoT sensors and other machine-generated data. Much of this data is valuable for future training, as well as compliance audits, debugging and understanding AI outputs. Aggregating, normalizing and processing this data form the foundation of enterprise AI implementation. Leaders from Cloudian, Hammerspace, Seagate and Supermicro discuss best practices for managing and retaining this data.

SEPTEMBER 3



Enterprise AI Implementation Challenges and Solutions

Enterprise AI, especially agentic AI, introduces new IT challenges, including data readiness and governance, hardware and software infrastructure, deployment models, scaling and economics, and operations. This session with Nutanix, AMD and Supermicro brings together experts in AI infrastructure from solution, silicon and systems perspectives to discuss on-prem and hybrid deployment models, solution integration and more.

Diamond



Gold





Speakers

