


ИНФРАСТРУКТУРА ИИ / ЭКСПЛУАТАЦИЯ ИИ / СТРАТЕГИЯ ИИ

## «Аварийный выключатель ИИ» предполагает, что вы знаете, что именно пытаетесь отключить

Ввести «аварийный выключатель» для ИИ проще, чем реализовать. Для реального сдерживания требуется полный контроль над сложной корпоративной инфраструктурой.

7 августа 2026 года, 9:00 [Робин Татам](#)



Осаруге Игбиноба для Unsplash+

[Perforce](#) спонсировала этот пост.

Выражение «аварийный выключатель ИИ» вошло в обиход, потому что оно позволяет людям просто говорить о сложном страхе.

По мере того как системы искусственного интеллекта становятся все более автономными и их все сложнее оценивать с точки зрения привычных операционных принципов, четко прописанная возможность вмешательства выглядит обнадеживающе. Если что-то начинает вести себя так, что создает

выключатель».

Недавний случай, когда модели OpenAI вышли за пределы изолированной тестовой среды и попали в Hugging Face, стал конкретным примером, вызывающим опасения. CNBC сообщает, что этот инцидент помог продвинуть двухпартийный законопроект, согласно которому некоторые компании, работающие в сфере ИИ, должны иметь возможность отключать, замедлять или приостанавливать работу своих моделей, а Министерство внутренней безопасности должно получить право отдавать распоряжения о замедлении или отключении в случаях, когда речь идет о потенциальном катастрофическом ущербе.

Политическая реакция вполне объяснима. Когда появляется новая категория рисков, особенно та, которую общественность еще не знает, как оценивать, лидеры ищут способ превратить абстрактную проблему в нечто, требующее конкретных действий. В данном случае таким языком стал термин «отключение».

Люди, работающие в крупных компаниях, как правило, за формулировками политики видят другой вопрос. Если поступит приказ о закрытии, что именно будет закрыто?



Perforce — это надежный партнер, который помогает управлять разработкой программного обеспечения в эпоху внедрения искусственного интеллекта. Его решения обеспечивают контроль кода, качества, инфраструктуры и данных, позволяя внедрять инновации без рисков. Создано для критически важных приложений, от которых зависит прибыль, где провал недопустим.

Узнать больше →

УЗНАЙТЕ БОЛЬШЕ ОТ НАШЕГО СПОНСОРА

svo1975pvn1982@gmail.com

В современной производственной среде ответ на этот вопрос обычно подразумевает отслеживание более чем одной системы. Сервис с поддержкой искусственного интеллекта может зависеть от конечных точек, API, облачных ресурсов, систем идентификации, реестров пакетов, конвейеров данных, автоматизации рабочих процессов, инструментов логирования и последующих приложений, которые обрабатывают выходные данные модели.

Некоторые из этих зависимостей могут принадлежать разным командам. Другие могут находиться за пределами компании. Некоторые из них могли начинаться как эксперименты, а затем стать частью производственного процесса, не подвергаясь такой же тщательной проверке, как исходная архитектура. К тому времени, когда сервис становится достаточно важным, чтобы вызвать вопросы в сфере управления, он уже может не напоминать изолированное приложение с одним владельцем и чистой операционной поверхностью.

#### TRENDING STORIES

1. [Why your company should \(try to\) build its own AI SRE](#)
2. [The blank-check AI coding era is dead. Here's what comes next.](#)
3. [Temporal 5x'd AI spend and doubled revenue. CEO Samar Abbas says he can't prove they're connected.](#)
4. [AI Agents Are a Security Ticking Time Bomb](#)
5. [Why OpenAI and Anthropic are hiring forward deployed engineer teams](#)

Writing shutdown authority into legislation is far simpler than carrying that decision through a production estate shaped by years of migrations, exceptions, integrations, acquisitions, temporary fixes, and team-level decisions. That implementation gap is where the issue becomes most relevant to infrastructure teams.

---

topics still deserve attention, especially as organizations formalize where AI may be used, which data can be shared, and how employees should evaluate generated output.

---

“Writing shutdown authority into legislation is far simpler than carrying that decision through a production estate shaped by years of migrations, exceptions, integrations, acquisitions, temporary fixes, and team-level decisions.”

---

As AI moves deeper into production workflows, the discussion also needs to involve a more practical concern: when an AI-enabled system **creates unacceptable risk**, can the organization understand the affected environment well enough to constrain it quickly, consistently, and with evidence?

The phrase “kill switch” may drive the public discussion, but the practical answer lives in the systems surrounding the AI capability.

## The limits of a single control

Emergency stops are the kind of control most people picture when they hear the phrase “kill switch.” They make sense in physical systems. Manufacturing equipment, industrial machinery, and certain safety-critical devices can be designed with direct shutdown mechanisms. Software estates already stretch that metaphor, and enterprise AI stretches it further.

The model may be the most visible part of the discussion, although it is rarely the full surface area. An AI assistant used in software delivery might have access to repositories, CI/CD tools, artifact stores, ticketing systems, secrets, test environments, and deployment workflows. An AI agent used in IT operations might read telemetry, recommend remediation, open change requests, call automation scripts, or modify infrastructure through approved orchestration paths.

Stopping one part of that chain can leave other paths untouched. Turning off a service may not revoke the credentials it uses. Suspending inference may leave downstream systems acting on outdated outputs. Interrupting the wrong dependency can create a separate service incident while the original risk remains

---

into a sequence of technical tradeoffs.

A credible response plan must account for the system as it exists now, not as it looked during an architecture review months earlier. Infrastructure teams bring useful skepticism to that exercise because they are used to tracing scope, access, ownership, dependencies, and verification paths under pressure. They also know that many environments contain a gap between documented intent and production behavior.

Once AI is embedded in business workflows, those operational details become part of the governance conversation. They expose the places where policy language has moved faster than the infrastructure knowledge needed to make policy executable.

## Before containment comes discovery

Much of the public conversation assumes organizations know where AI is running, which is a generous assumption in many enterprise environments.

AI can enter an enterprise through obvious channels, such as internally approved model providers or purpose-built applications. It also arrives through less visible paths. A SaaS product adds an AI feature. A development team experiments with an open source model. A hosted API gets attached to an internal tool. A vendor introduces an AI capability inside software the company already approved. Over time, the line between an “AI system” and a system that happens to use AI becomes harder to define.

This is where the kill switch metaphor starts to show its limits. If the relevant systems are discovered during the response, the team is already behind. Dependency questions, business impact, access paths, and evidence collection all become harder when the basic inventory is still being assembled.

---

**“If the relevant systems are discovered during the response, the team is already behind.”**

---

Infrastructure teams have seen versions of this problem before. During incident response, a service thought to be isolated turns out to have undocumented



---

data, and actual operating conditions refuse to line up cleanly. AI adds a new category of concern, but the underlying visibility problem is familiar.

The challenge quickly expands beyond the model itself. Teams need to understand which systems call external models, where generated content influences workflows, which accounts and automation paths sit between a recommendation and an action, and how third-party AI capabilities have found their way into the environment.

Most discussions start with how to stop risky AI behavior. In many environments, the more revealing question comes earlier: can the organization produce a reliable picture of where AI touches the estate, which systems depend on it, and which workflows would keep moving if access changed?

## Containment depends on the state of the system

Most infrastructure teams know that containment is less a single action than a set of operating conditions. The difficult work happens before the incident, when teams decide what should change if risk reaches a level that requires intervention.

Under ordinary circumstances, a service operates with a defined set of permissions, connections, dependencies, and logging requirements. Under restricted operation, selected assumptions change while investigators preserve evidence and determine whether the risk has been contained. A team might turn off endpoints, suspend integrations, limit external network access, revoke or rotate credentials, increase logging, or isolate workloads while the situation is being investigated.

The details vary because environments vary. That is exactly why generic answers tend to fall apart. A useful containment plan must match the systems it is intended to govern, including the dependencies that surround them and the business processes that rely on them.

Inventory sounds mundane until a response effort depends on it. In my work with **infrastructure and compliance** teams, I've repeatedly seen organizations struggle to maintain an accurate picture of their environments as cloud resources, Kubernetes clusters, SaaS services, and AI projects multiply faster than governance processes can track them during normal operations, which creates audit and support headaches. During a containment event, those same blind spots slow

## Governance eventually reaches production

Governance efforts often begin with documentation. Committees are formed to define responsibilities, agree on escalation paths, and establish a common language for discussing risk before an incident forces the issue.

The conversation shifts once someone asks whether the control can be demonstrated. A document can describe who has authority to suspend an AI-enabled service. Still, it cannot turn off an integration, revoke a credential, increase logging, or prove that a set of systems entered a restricted condition. A risk register can identify a containment scenario, although it cannot document which nodes changed, when they changed, and whether they remain aligned with the required configuration.

Security, compliance, and infrastructure teams know this gap well. It appears in patching programs, configuration baselines, incident response exercises, supply chain reviews, and disaster recovery planning. Written controls tend to be cleaner than the real-world environments they describe. Production systems reflect years of accumulated decisions, exceptions, migrations, temporary fixes, acquired assets, and workarounds that may outlive the original reasons they existed.

---

**“Written controls tend to be cleaner than the real-world environments they describe.”**

---

AI increases urgency because some systems are **becoming more autonomous** and more connected to business workflows. It also adds outside pressure. When an incident becomes visible enough to prompt legislative action, boards and customers start asking sharper questions. A company may be able to point to an AI policy, but boards, customers, and regulators eventually want to understand exactly how that policy translates into action.

If leadership declares an AI-enabled workflow should be restricted, the discussion moves quickly from oversight to execution. Teams need to know where to intervene, which systems are affected, who owns the required changes, how completion will be verified, and what evidence remains once the response is over.

---

customer-facing systems, or environments connected to critical business operations.

## Why a mandate will not solve the estate problem

A federal shutdown authority, if enacted, would place legal pressure on a narrow class of powerful AI providers. It would not remove the implementation burden for organizations that adopt, integrate, fine-tune, host, or embed AI systems inside their own environments.

Even if a major AI provider can throttle or suspend a model, each enterprise still must understand its own exposure in the context of its applications, workflows, dependencies, and operating assumptions. Which applications depend on that model? Which workflows fail open or fail closed if access is restricted? Which internal systems contain cached outputs, calculated decisions, or agent-created changes? Which business processes need manual fallback when an AI service is unavailable?

Policy debates often treat AI control as if the decisive action happens at the model layer. Sometimes it will; however, in many business environments, the risk will live in the connections that surround the model. A hosted AI service may be suspended while local workflows, scripts, integrations, and access tokens continue following their last known configuration.

---

**“A serious AI containment strategy has more in common with mature infrastructure management than with an emergency stop button.”**

---

A serious AI containment strategy has more in common with mature infrastructure management than with an emergency stop button. It requires an up-to-date inventory of AI-adjacent systems, a map of dependencies and access paths, predefined restricted conditions for high-risk services, tested procedures for applying those conditions, and evidence that changes were enforced. Ownership also needs to be clear, since fast action becomes difficult when authority is scattered across teams.



and business impact.

---

## Infrastructure teams belong earlier in the conversation

Spend enough time running infrastructure, and you develop a complicated relationship with documentation. Most organizations have diagrams, inventories, and governance processes, and all of them serve a purpose. The challenge is that production systems keep evolving long after those artifacts are created. Acquisitions introduce systems that do not fit cleanly into existing models. Applications gain integrations nobody anticipated during the original design process. Temporary exceptions become permanent. Cloud resources intended to live for a week are still running after a year.

Most of this happens for defensible reasons, usually in support of uptime mandates, delivery pressure, customer needs, or business continuity SLAs. The result is that operational knowledge becomes dispersed across people (some of whom will inevitably have moved on), tickets, runbooks, monitoring systems, and memory rather than living neatly in one place.

Infrastructure teams spend their days tracing dependencies, untangling ownership questions, and figuring out how systems behave outside a design review. Bringing that perspective into AI governance conversations early can prevent containment plans from depending on assumptions that did not translate into production. It also helps organizations understand the difference between disabling a model, restricting access to a service, isolating a workload, and preserving evidence during an investigation.

Scale complicates things further. A manual action that works effectively for ten systems may fail across hundreds or thousands. A change one expert can perform during business hours may become fragile if that person is unavailable when an event occurs. A runbook that looks adequate in a tabletop exercise may not survive a live environment where dependencies have changed, and the current ownership is unclear.

The phrase “kill switch” will probably remain part of the public debate because it is simple, memorable, and familiar. Practitioners do not have to accept the metaphor literally to leverage the attention it creates. They can redirect the conversation toward more useful questions: what restricted operation would mean for a given service, which dependencies would have to change, which controls can be applied

These questions are less dramatic than a big red button, but their answers are also most likely to improve readiness.

## Control starts before the incident

The Hugging Face incident gave the industry a vivid story, and Washington responded with the language of shutdown authority. That reaction is understandable. Leaders want mechanisms that sound equal to the risk, especially when the public conversation moves faster than the technical details can be explained.

By the time an organization begins thinking about containment, much of the hard work should already be done. Teams should already understand what is running, who owns it, what depends on it, and how changes will ripple through the environment.

AI can reduce certain workflow bottlenecks, but it also exposes weak inventory, unclear ownership, and brittle operating assumptions faster than many teams are prepared to handle. A future incident will not pause while teams locate assets, clarify ownership, identify credentials, or discover that a service dependency was never documented.

The current debate may be framed around new kill switches. For most organizations, the more useful work starts with building and maintaining an accurate picture of the systems, dependencies, and workflows that already exist across the estate.

TNS



Robin Tatam is Senior Technical Marketer and Evangelist for Perforce Puppet and a Certified Information Security Manager (CISM). He writes and speaks about the practical challenges of managing infrastructure at scale, particularly where automation, governance, and operational reality intersect. His...

[Read more from Robin Tatam →](#)

---

## TRENDING STORIES

1. Why your company should (try to) build its own AI SRE
2. The blank-check AI coding era is dead. Here's what comes next.
3. Temporal 5x'd AI spend and doubled revenue. CEO Samar Abbas says he can't prove they're connected.
4. AI Agents Are a Security Ticking Time Bomb
5. Why OpenAI and Anthropic are hiring forward deployed engineer teams

## INSIGHTS FROM OUR SPONSOR



Perforce is the trusted partner to govern software delivery in the age of AI execution. Its solutions enforce guardrails across code, quality, infrastructure, and data—enabling innovation without introducing risk. Built for high-stakes, revenue-critical applications where failure isn't an option.

[Learn More →](#)

---

TNS DAILY NEWSLETTER

# most recent TNS articles in your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

## ARCHITECTURE

Cloud Native Ecosystem  
Containers  
Databases  
Edge Computing  
Infrastructure as Code  
Linux  
Microservices  
Open Source  
Networking  
Storage

## ENGINEERING

AI  
AI Engineering  
API Management  
Backend development  
Data  
Frontend Development  
Large Language Models  
Security  
Software Development  
WebAssembly

## OPERATIONS

AI Operations  
CI/CD  
Cloud Services  
DevOps  
Kubernetes

## CHANNELS

Podcasts  
Ebooks  
Events  
Webinars  
Newsletter

## THE NEW STACK

[About / Contact](#)  
[Sponsors](#)  
[Advertise With Us](#)  
[Contributions](#)



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

[Frontend Developer Roadmap](#)  
[Backend Developer Roadmap](#)  
[Devops Roadmap](#)

## FOLLOW TNS



© The New Stack 2026

[Disclosures](#)  
[Terms of Use](#)  
[Advertising Terms & Conditions](#)  
[Privacy Policy](#)  
[Cookie Policy](#)