📖
Вселенная

Data Science для бизнеса (2 книги в 1) Машинное обучение и аналитика данных

Авторы: Райли Адамс, Мэтт Хендерсон

Тип файла: pdf

Размер: 3,6 МБ

Язык: Английский

Страниц: 315

Data Science для бизнеса (2 книги в 1) Машинное обучение и аналитика данных

Введение

Данные стали одним из самых ценных ресурсов в современной инженерии и бизнесе. Организации ежедневно собирают информацию с веб-сайтов, датчиков, машин, из данных о взаимодействии с клиентами, финансовых систем, мобильных приложений и операционных платформ. Настоящая проблема заключается не в том, чтобы просто собрать эту информацию, а в том, чтобы использовать ее для принятия полезных решений. 📊⚙️

Наука о данных для бизнеса объединяет аналитику данных, статистику, программирование, машинное обучение, визуализацию и экспертные знания в предметной области для решения практических бизнес-задач. Вместо того чтобы рассматривать науку о данных как чисто теоретическую дисциплину, организации могут использовать ее для поиска ответов на такие вопросы, как:

- Какие продукты, скорее всего, будут пользоваться спросом?
- Почему уходят клиенты?
- Какое оборудование может потребовать ремонта?
- Как оптимизировать складские запасы?
- Какие маркетинговые кампании привлекают ценных клиентов?
- Как выявить операционные риски на более ранних этапах?
- Какие бизнес-результаты могут возникнуть в будущем?

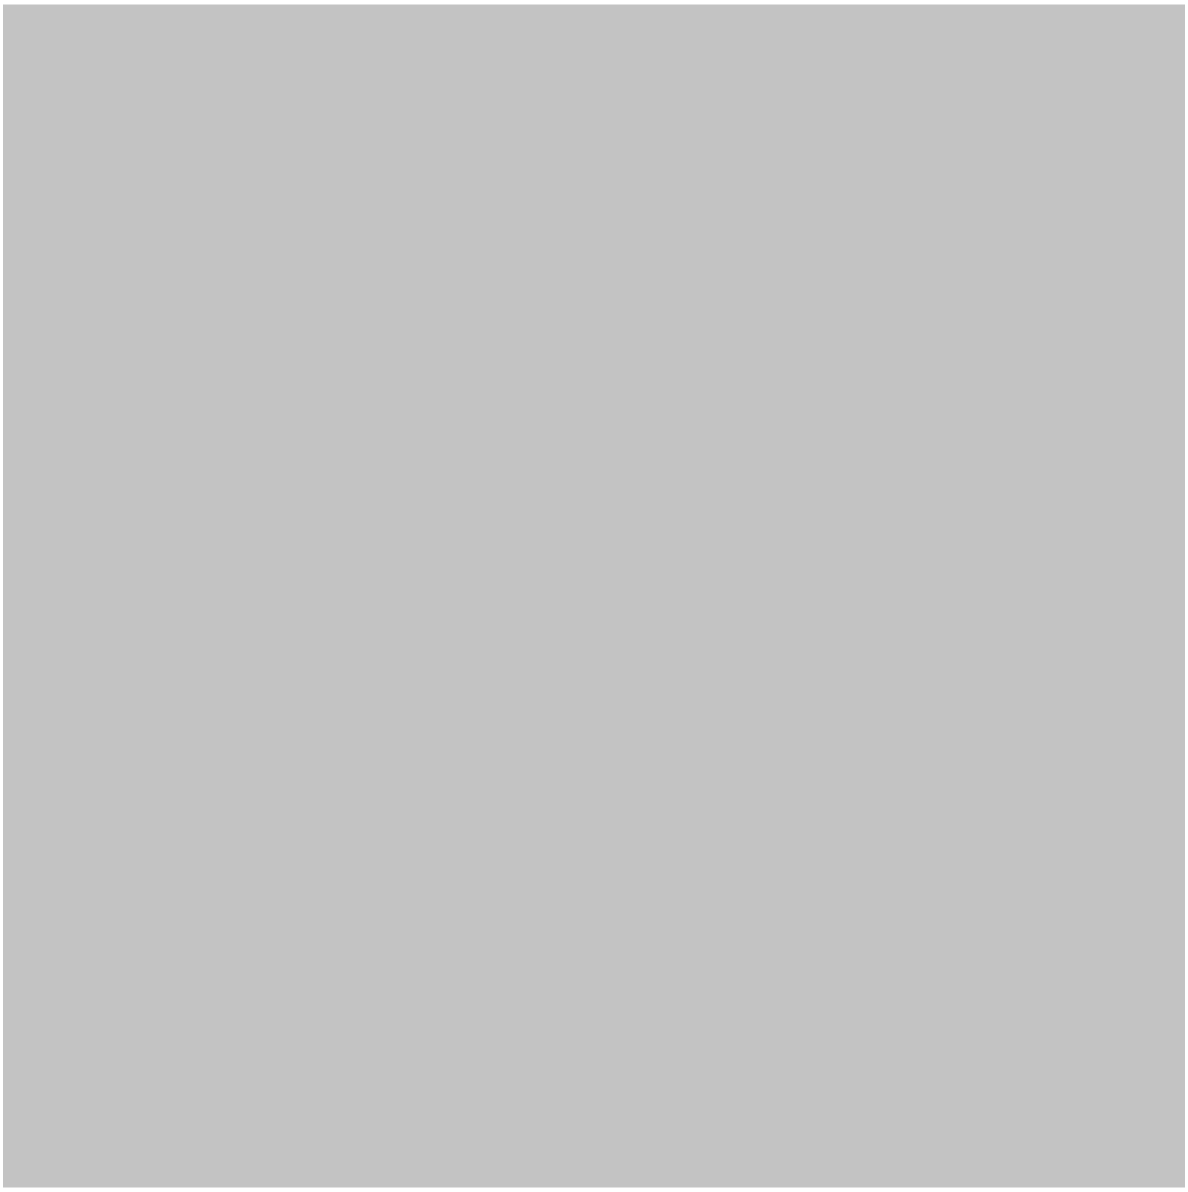

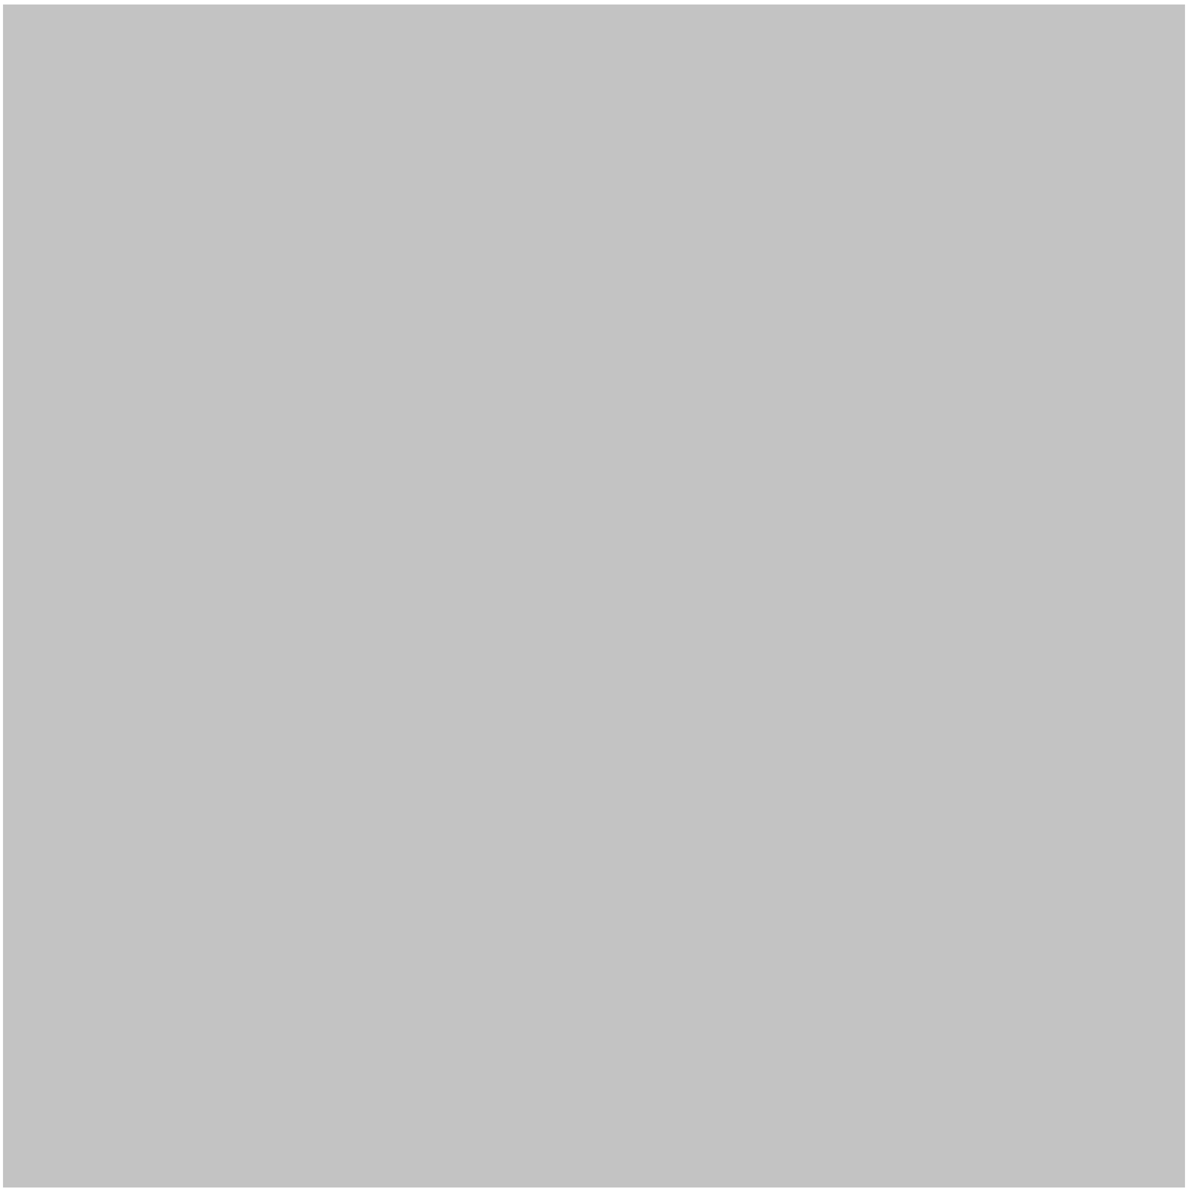

Практический подход сочетает в себе **машинное обучение + анализ данных + знание бизнеса**. Аналитика помогает организациям понять, что произошло и почему, а машинное обучение — выявить закономерности и сделать прогнозы.

В этой статье тема рассматривается в практическом ключе и состоит из двух частей: **«Машинное обучение» и «Аналитика данных»**. Она предназначена для студентов, инженеров, аналитиков, менеджеров, разработчиков и специалистов, которые хотят понять, как наука о данных может помочь в принятии реальных бизнес-решений.

Теоретическая основа

От бизнес-данных к бизнес-аналитике

Традиционные компании часто полагались на отчеты, подготовленные вручную. Аналитики собирали информацию, создавали электронные таблицы, строили диаграммы и представляли исторические результаты лицам, принимающим решения.

Современная наука о данных расширяет этот процесс.

Типичная организация, работающая с данными, может собирать информацию из:

- Системы взаимодействия с клиентами
- Платформы планирования ресурсов предприятия
- Веб-сайты и приложения
- Промышленные датчики
- Финансовые операции
- Взаимодействие в социальных сетях
- Системы управления цепочками поставок
- Маркетинговые платформы
- Облачные приложения
- Устройства Интернета вещей

Затем информация может пройти несколько этапов обработки: сбор, хранение, очистка, анализ, моделирование, визуализация и принятие решений.

Analytics and Machine Learning

Data analytics generally focuses on understanding available information.

Machine learning goes a step further by allowing computer systems to identify patterns from data and use those patterns for tasks such as classification, prediction, recommendation, and anomaly detection.

The two disciplines are complementary rather than competing.

For example, an engineering company could analyze historical equipment failures to understand which operating conditions are associated with breakdowns. A machine-learning system could then learn from those patterns and help identify equipment that may require inspection.

Definition

What Is Data Science for Business?

Data Science for Business is the systematic use of data, analytical methods, computational techniques, and machine-learning approaches to support business decisions, improve operations, manage risk, understand customers, and create measurable value.

It combines several disciplines:

Discipline	Contribution
Data Analytics	Understands historical and current performance
Statistics	Measures uncertainty and relationships
Machine Learning	Finds patterns and produces predictions
Programming	Automates data processing and analysis
Data Engineering	Builds reliable data pipelines
Visualization	Communicates insights clearly
Business Knowledge	Connects technical results to practical decisions

The Two-Part Perspective

A useful way to understand the topic is through two connected components:

Part 1 — Data Analytics 📊

Focuses on questions such as:

> *What happened?*

> *Why did it happen?*

> *What is happening now?*

Part 2 — Machine Learning 🤖

Focuses on questions such as:

> *What is likely to happen?*

> *Which patterns can be detected automatically?*

> *Which customers, machines, products, or transactions require attention?*

Together, these capabilities create a stronger decision-support system.

Step-by-Step Data Science for Business Process

Step 1: Define the Business Problem

The first step is not selecting an algorithm.

It is defining the business problem.

For example, "We need artificial intelligence" is not a sufficiently precise objective. A better objective might be:

> *"We want to identify customers who are likely to discontinue their subscriptions so that the retention team can contact them earlier."*

This gives the data-science team a clear purpose.

Step 2: Identify Relevant Data

Next, determine which information can help solve the problem.

For customer retention, useful data might include:

- Purchase history
- Subscription activity
- Customer-support interactions
- Product usage
- Account age
- Previous cancellations
- Marketing engagement

Step 3: Clean the Data

Real-world data is rarely perfect.

It may contain:

- Missing values
- Duplicate records
- Incorrect formats
- Outliers
- Inconsistent categories
- Invalid timestamps
- Data-entry errors

Cleaning is therefore one of the most important stages in the workflow.

Step 4: Explore the Data

Exploratory data analysis helps analysts discover patterns.

They may examine:

- Distributions
- Trends
- Correlations
- Customer segments
- Seasonal behavior
- Product performance
- Geographic differences

Visualization can make complicated datasets easier to understand.

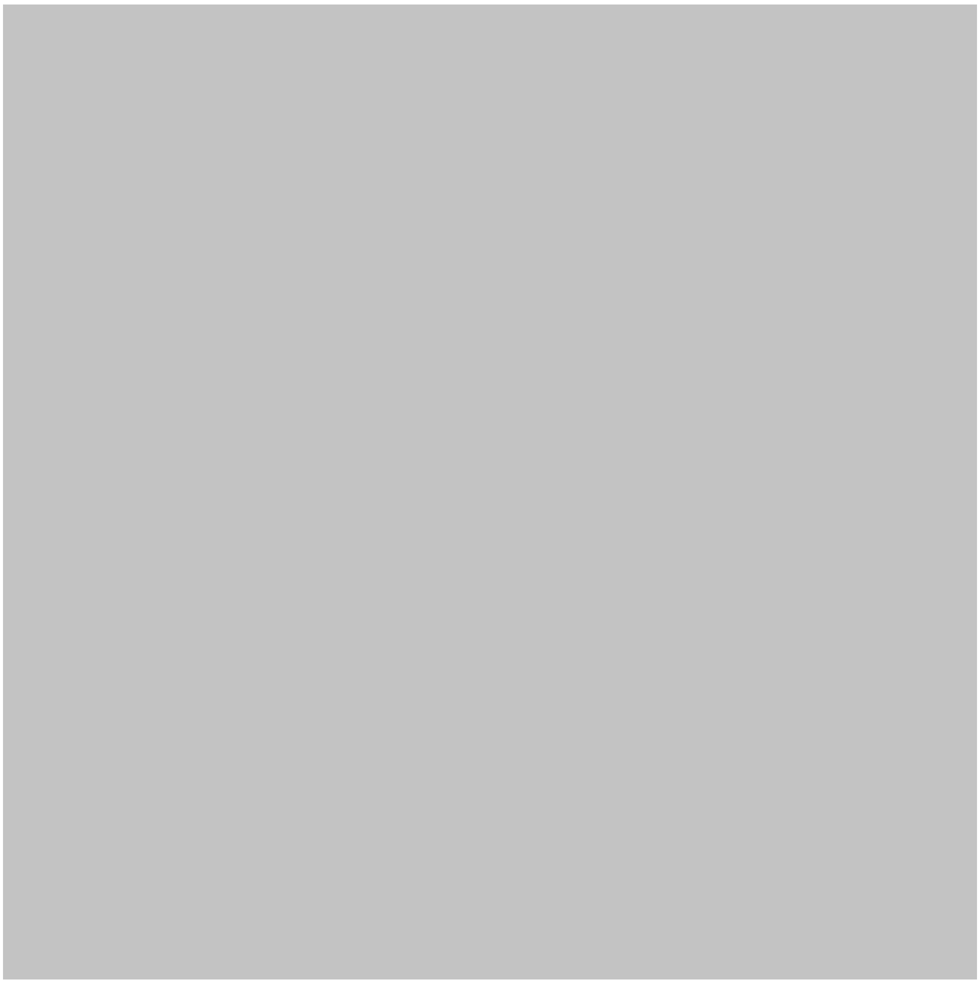

Step 5: Build an Analytical or Machine-Learning Model

Depending on the problem, a team may use:

- Regression
- Classification
- Clustering
- Decision trees
- Random forests
- Gradient boosting
- Neural networks
- Time-series methods
- Recommendation systems
- Anomaly detection

The most sophisticated algorithm is not necessarily the best solution.

Step 6: Evaluate the Result

A model should be evaluated using appropriate performance measures and business criteria.

For example, a fraud-detection system should not be judged only by how many transactions it classifies correctly. False alarms can also create significant operational costs.

Step 7: Deploy the Solution

A successful model needs to become part of a business process.

It might appear as:

- A dashboard
- An automated alert
- A recommendation engine
- A customer-service tool
- A forecasting system
- A maintenance platform
- A financial-risk application

Step 8: Monitor and Improve

Business environments change.

Customer behavior, market conditions, product offerings, regulations, and economic conditions can all influence data.

Therefore, machine-learning systems should be monitored continuously.

Comparison

Data Analytics vs Machine Learning

Feature	Data Analytics	Machine Learning
Primary purpose	Understand data	Learn patterns and make predictions
Typical question	What happened?	What may happen?

Feature	Data Analytics	Machine Learning
Human involvement	Often high	Can be partially automated
Output	Reports, insights, dashboards	Predictions, classifications, recommendations
Common users	Analysts and managers	Data scientists and engineers
Example	Analyze monthly sales	Predict future customer demand

Traditional Reporting vs Data Science

Traditional reporting usually describes historical performance. Data science can combine historical analysis with predictive and automated capabilities.

However, traditional reporting remains valuable. A business should not replace reliable reporting simply because machine learning is available.

Diagrams and Tables

Data Science Business Architecture

```
                 BUSINESS PROBLEM
                        │
                        ▼
                 DATA COLLECTION
                        │
                        ▼
                  DATA STORAGE
                        │
                        ▼
                DATA PREPARATION
                        │
              ┌─────────┴─────────┐
              ▼                   ▼
       DATA ANALYTICS      MACHINE LEARNING
              │                   │
              └─────────┬─────────┘
                        ▼
                BUSINESS INSIGHT
```

```
                     |
                     ▼
              DECISION / ACTION
                     |
                     ▼
               MONITORING
                     |
                     └──────► CONTINUOUS IMPROVEMENT
```

Business Analytics Maturity

Level	Capability	Example
1	Descriptive	Monthly sales report
2	Diagnostic	Identifying reasons for declining sales
3	Predictive	Forecasting demand
4	Prescriptive	Recommending an operational action
5	Intelligent automation	Automatically responding to selected conditions

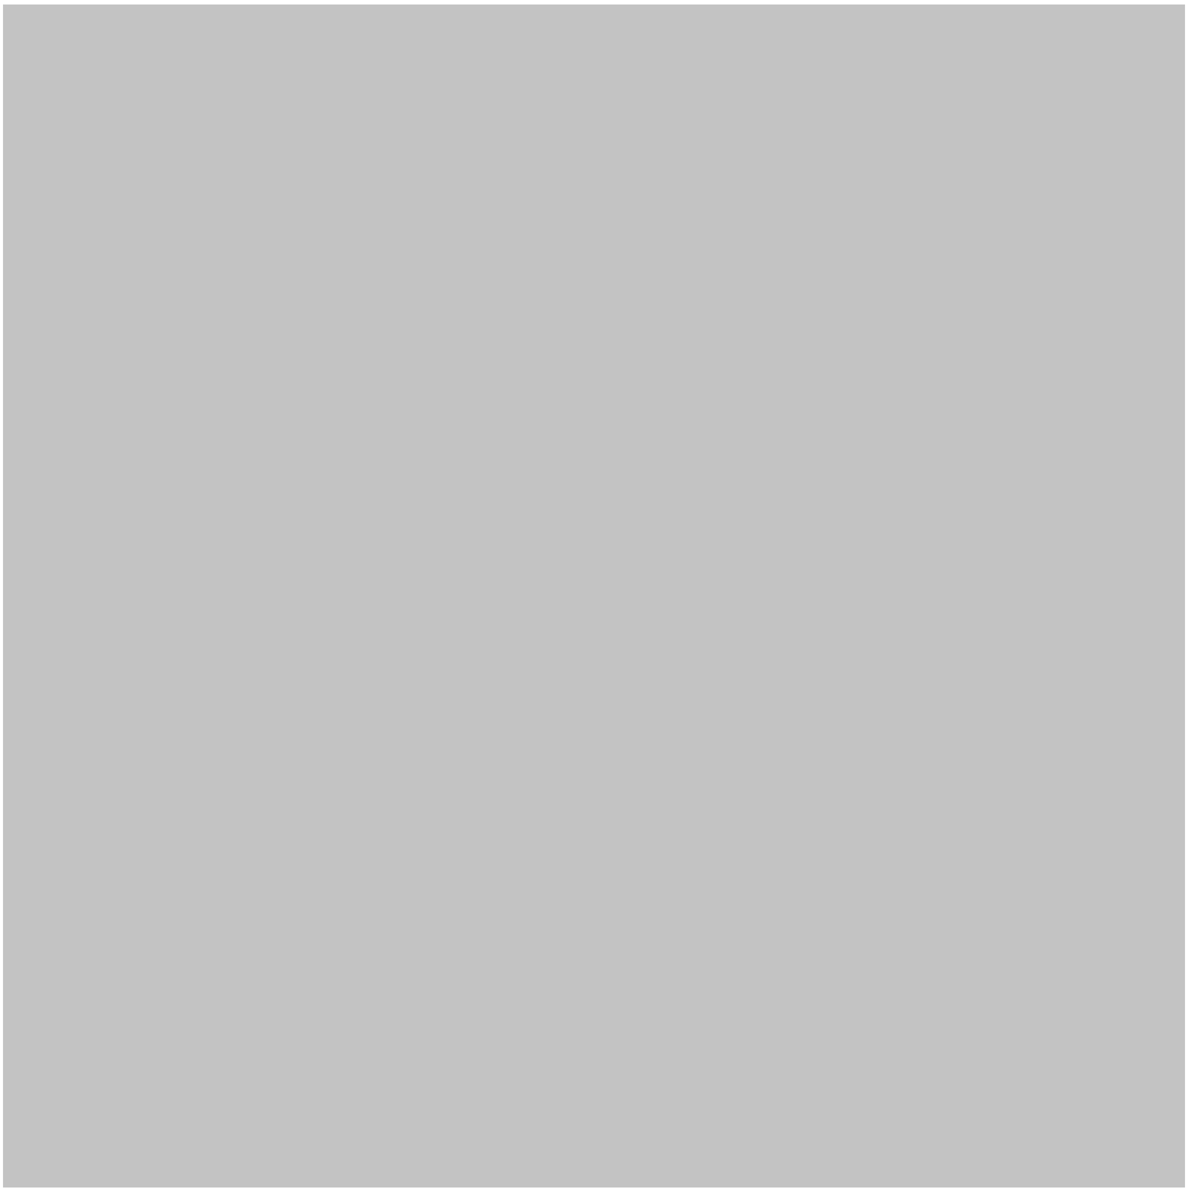

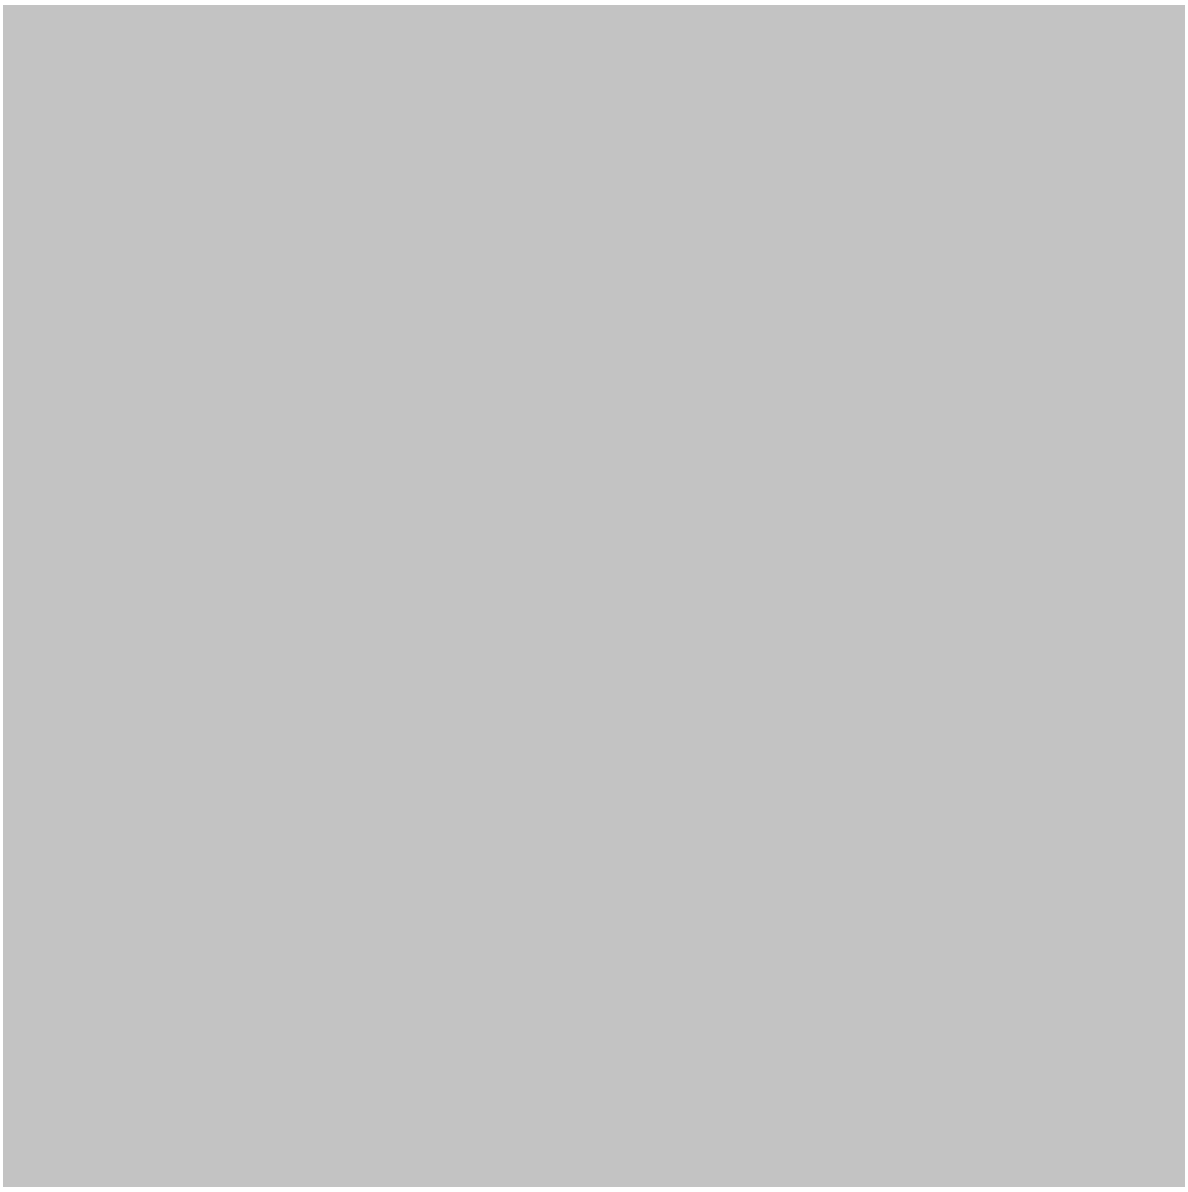

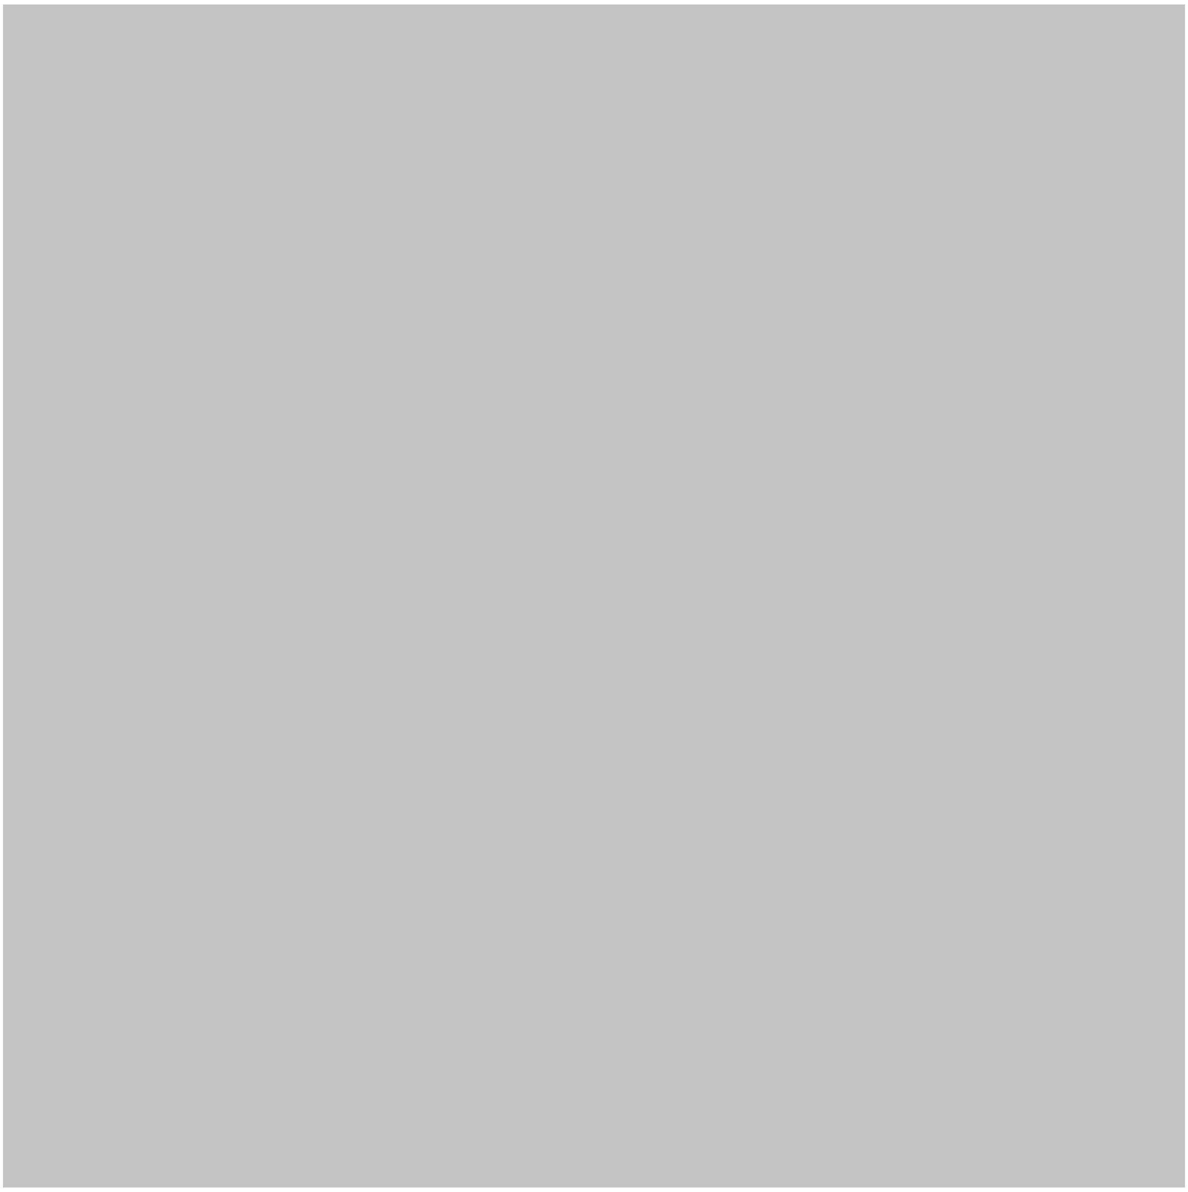

Examples

Customer Churn

A telecommunications company can analyze customer behavior and identify patterns associated with cancellations.

A machine-learning model can then prioritize customers who appear more likely to leave.

The retention team can focus its resources on those customers.

Predictive Maintenance

An industrial facility can collect information from pumps, motors, compressors, and other equipment.

Analytics can reveal unusual operating patterns.

Machine learning can help identify equipment behavior associated with previous failures.

Engineers can then investigate before a major breakdown occurs.

Retail Demand Forecasting

A retailer can analyze historical sales, seasonal behavior, product categories, promotions, and regional trends.

Machine learning can help estimate future demand.

The organization can use those predictions to improve purchasing and inventory decisions.

Fraud Detection

Financial institutions process large numbers of transactions.

Data science can identify unusual transaction patterns and flag potentially suspicious activity for further review.

Real-World Application

Manufacturing 🏭

Manufacturing companies can combine machine data, production records, quality measurements, and maintenance information.

Applications include:

- Predictive maintenance
- Quality inspection
- Production optimization
- Energy monitoring
- Supply-chain forecasting
- Defect detection

Healthcare

Data science can support operational and analytical activities such as resource planning, patient-flow analysis, and research.

Healthcare applications require particularly careful attention to privacy, security, validation, and regulatory requirements.

Finance

Financial organizations use analytics and machine learning for:

- Fraud detection
- Risk analysis
- Customer segmentation
- Credit-related decision support
- Market analysis
- Operational monitoring

Transportation

Transportation companies can analyze:

- Vehicle utilization
- Route performance
- Fuel consumption
- Maintenance requirements
- Passenger demand
- Delivery patterns

Engineering and Construction

Engineering organizations can use data science to analyze project schedules, material usage, equipment performance, safety indicators, costs, and resource allocation.

This creates opportunities for better project planning and early identification of operational problems.

Common Mistakes

Starting With the Algorithm

One of the most common mistakes is choosing a machine-learning algorithm before understanding the business problem.

Better approach: define the decision first and select the technology afterward.

Ignoring Data Quality

Poor-quality data can produce unreliable conclusions even when advanced algorithms are used.

Solution: establish data-quality checks before modeling.

Measuring Technical Performance Only

A model can perform well statistically while producing little business value.

For example, a prediction system may be accurate but too slow to influence real-time decisions.

Solution: connect model evaluation with business outcomes.

Overcomplicating the Solution

A complex neural network may be unnecessary when a simpler analytical method can solve the problem.

Solution: begin with an understandable baseline.

Forgetting Human Oversight

Automated predictions should not automatically become business decisions in every situation.

Human review can remain important for high-impact decisions.

Challenges and Solutions

Challenge	Practical Solution
Missing data	Establish validation and cleaning procedures
Fragmented systems	Build standardized data pipelines
Lack of expertise	Combine technical and domain teams

Challenge	Practical Solution
Poor model adoption	Involve users early
Privacy concerns	Apply governance and access controls
Changing data	Monitor model performance
Unclear ROI	Define measurable business objectives
Model complexity	Prefer appropriate, explainable methods

Data Governance

Data governance should cover:

- Ownership
- Access
- Quality
- Security
- Privacy
- Retention
- Documentation
- Regulatory requirements

A technically impressive system can still fail if its data governance is weak.

Case Study

Predictive Maintenance in a Manufacturing Plant

Consider a manufacturing company operating hundreds of industrial machines.

The company historically maintained equipment according to fixed schedules. This approach had two problems.

First, some machines were serviced even when they were operating normally. Second, unexpected failures could interrupt production.

The engineering team decided to develop a data-science solution.

Data Collection

The company collected information from machine sensors and maintenance records.

The dataset included information about:

- Operating conditions
- Temperature
- Vibration
- Maintenance events
- Production cycles
- Previous failures
- Equipment age

Analytics Stage

Engineers first examined historical data.

They discovered that certain combinations of operating conditions frequently appeared before maintenance events.

Instead of immediately deploying an advanced model, the team created dashboards showing machine behavior and historical patterns.

Machine-Learning Stage

The next stage used historical maintenance information to train a predictive system.

The system generated risk indicators for machines requiring additional inspection.

Operational Integration

The predictions were integrated into the maintenance workflow.

Maintenance engineers received prioritized alerts rather than reviewing every machine with equal priority.

Business Impact

The value of the project was not simply the machine-learning model.

The complete solution connected:

Sensors → Data → Analytics → Prediction → Engineering Action

This illustrates an important principle:

> *Data science creates value when insights are converted into useful actions.*

Essential Tips

For Students 🎓

Start with fundamentals:

- Statistics
- Python
- SQL
- Data visualization
- Data cleaning
- Machine learning
- Business fundamentals

Do not focus exclusively on memorizing algorithms.

Build projects using realistic datasets and explain the business purpose of every project.

For Engineers ⚙️

Connect data science with engineering knowledge.

Your domain expertise can help determine whether a discovered pattern is physically meaningful or merely a statistical coincidence.

For Professionals 💼

Focus on business outcomes.

Before launching a project, ask:

1. What decision will this system improve?
2. Who will use the result?
3. What data is available?
4. How will success be measured?
5. What happens after the prediction is generated?

For Organizations

Create cross-functional teams.

A strong data-science project may involve:

Domain experts + Data analysts + Data engineers + Data scientists + Software engineers + Business stakeholders

No single role needs to solve every part of the problem.

FAQs

What is Data Science for Business?

Data Science for Business is the application of analytics, statistics, programming, machine learning, and domain knowledge to improve business decisions and operational outcomes.

Is machine learning the same as data analytics?

No. Data analytics primarily focuses on understanding and interpreting data, while machine learning focuses on learning patterns that can support prediction, classification, recommendation, or automation.

Do I need advanced mathematics to learn business data science?

A strong mathematical foundation is useful, especially for advanced machine learning. However, beginners can start with practical analytics, visualization, SQL, Python, and fundamental statistics before progressing to more advanced mathematical concepts.

Which programming language is commonly used?

Python is widely used because of its extensive ecosystem for data analysis, visualization, machine learning, and automation. SQL is also extremely important for working with structured business data.

Can small businesses use data science?

Yes. Small organizations can begin with relatively simple solutions such as sales dashboards, customer segmentation, inventory analysis, website analytics, and demand forecasting.

What makes a machine-learning project successful?

Technical accuracy is only one factor. A successful project should solve a meaningful problem, use reliable data, integrate into an actual workflow, provide understandable results, and create measurable value.

Is data science useful for engineering?

Absolutely. Engineering applications include predictive maintenance, quality control, energy optimization, process monitoring, forecasting, simulation support, and equipment analytics.

What should I learn first?

Start with **data analysis, SQL, Python, statistics, visualization, and data cleaning**. Then progress toward supervised learning, unsupervised learning, model evaluation, deployment, and machine-learning operations.

Conclusion

Data Science for Business is not simply about creating sophisticated algorithms. It is about transforming data into reliable information, information into insight, and insight into better decisions. 🚀

The combination of **Data Analytics + Machine Learning** provides organizations with a powerful framework for understanding current performance while anticipating future possibilities.

Analytics helps answer:

“What happened, and why?”

Machine learning helps explore:

“What is likely to happen next?”

Business expertise then addresses the most important question:

“What should we do about it?”

For students, this field offers an opportunity to combine engineering, programming, statistics, and problem-solving. For professionals, it provides practical tools for improving operations, forecasting demand, managing risk, understanding customers, and developing intelligent systems.

The strongest data-science solutions are rarely the ones with the most complicated algorithms. They are the solutions that use **high-quality data, appropriate methods, clear communication, responsible governance, and strong business understanding** to produce useful results.

In a world increasingly driven by connected systems and digital information, learning how to turn data into actionable engineering and business intelligence is becoming an essential professional capability. 📊🤖⚙️

Related Posts:

Data Science: Techniques and Intelligent Applications

Time Series for Data Science: Analysis and Forecasting

Statistical Theory 2nd Edition

Power Pivot and Power BI

An Introduction to Statistics: An Active Learning...

Preview PDF Book

← Previous Book

Next Book →