



Вселенная

[Главная](#)

[Книги](#)

[О нас](#)

[Контакты](#)

[Политика конфиденциальности](#)

[DMCA](#)

Элементы науки о данных

Автор: Аллен Дауни

Тип файла: pdf

Размер: 10,8 МБ

Язык: Английский

Страниц: 290

Основы науки о данных: начало работы с наукой о данных и Python

Введение

 **Наука о данных** стала одной из важнейших технических дисциплин в современной инженерии, бизнесе, научных исследованиях, здравоохранении, финансах и технологиях. Организации ежедневно генерируют огромные объемы информации, но необработанные данные имеют ограниченную ценность до тех пор, пока их не преобразуют в полезные знания.

Наука о данных — это системный подход к преобразованию необработанных наблюдений в **инсайты, прогнозы и решения**. Python стал одним из самых

популярных языков программирования для этого процесса, поскольку он сочетает в себе относительно простой синтаксис и мощную экосистему научных библиотек и библиотек для машинного обучения.

Новичкам наука о данных может показаться сложной, поскольку она объединяет в себе несколько областей:

- 🐍 Программирование на Python
- 📐 Математика и статистика
- 🗄️ Управление данными
- 📊 Визуализация данных
- 🤖 Машинное обучение
- 🔬 Научные вычисления
- 🧠 Решение задач и знание предметной области


Однако эти элементы не обязательно изучать одновременно. Структурированный подход позволяет студентам и инженерам переходить от базовых операций в Python к сложным аналитическим системам.







DATA SCIENCE PROJECT WORKFLOW

From Raw Data to Actionable Insights




DATA SCIENCE WORKFLOW





Python and the Basics of Data Science



Основная идея проста:

Исходные данные → Обработка → Анализ → Выводы → Решение

В этой статье объясняются основные элементы науки о данных и дается практическое руководство по использованию Python для решения инженерных и аналитических задач.

Теоретическая база

Что такое наука о данных?

Наука о данных — это междисциплинарная область, которая объединяет программирование, статистику, математику, вычислительные методы и экспертные знания в предметной области для извлечения значимой информации из данных.

Упрощенное математическое представление выглядит так:

Данные → Модель → Прогнозирование/анализ

Предположим, инженерная компания ведет учет:

- Температура
- Давление
- Вибрация
- Скорость машины
- Энергопотребление
- Неисправности

Специалист по анализу данных может исследовать взаимосвязи между этими переменными и разработать модель, способную оценить вероятность выхода оборудования из строя.

Жизненный цикл анализа данных

Типичный проект по анализу данных состоит из нескольких этапов:

1. Определение проблемы
2. Сбор данных
3. Очистка данных
4. Исследовательский анализ
5. Разработка признаков
6. Разработка модели
7. Оценка модели
8. Развертывание
9. Мониторинг

⚙️ Важно отметить, что это не всегда прямая линия. Инженеры часто возвращаются к предыдущим этапам, когда обнаруживают недостающие данные, неверные предположения или низкую эффективность модели.

Почему Python?

Python особенно полезен тем, что его экосистема включает специализированные инструменты практически для каждого этапа рабочего процесса.

Цель	Распространенная технология Python
Численное вычисление	NumPy

Цель	Распространенная технология Python
Манипулирование данными	панды
Visualization	Matplotlib
Statistical analysis	SciPy
Machine learning	scikit-learn
Deep learning	PyTorch / TensorFlow
Interactive analysis	Jupyter
Data storage	SQL connectors / databases

The strength of Python comes less from the language alone and more from the combination of its libraries.

Definition

Essential Elements of Data Science

The major elements can be organized into six interconnected areas.

1. Data

Data is the raw material of data science.

It may be:

- Numerical
- Categorical
- Textual
- Image-based
- Audio
- Sensor measurements
- Time-series information
- Geospatial information

2. Statistics

Statistics helps answer questions such as:

What happened?

How frequently did it happen?

Is the observed relationship meaningful?

Important concepts include:

- Mean
- Median
- Variance
- Standard deviation
- Probability
- Correlation
- Regression
- Hypothesis testing

3. Programming

Programming allows analysts to automate repetitive operations.

Instead of manually processing 1,000,000 measurements, Python can perform the same operation programmatically.

4. Visualization

Charts help humans understand patterns that may be difficult to detect in tables.

Examples include:

 Line charts

 Bar charts

 Scatter plots

 Heat maps

 Box plots

5. Machine Learning

Machine learning allows computers to learn patterns from historical data and use those patterns to make predictions or classifications.

6. Domain Knowledge

A technically accurate model can still produce poor decisions if the engineer does not understand the application.

For example, an engineer analyzing vibration data must understand machinery behavior, not simply the mathematics behind a machine-learning algorithm.

Step-by-Step Explanation

Step 1: Define the Problem

Before writing Python code, define the engineering or business question.

For example:

Can historical temperature and vibration measurements be used to predict machine failure?

This question determines what data is required and what type of model might be appropriate.

Step 2: Collect Data

Data can originate from:

- Databases
- Sensors
- CSV files
- APIs
- Laboratory experiments
- Enterprise systems
- Public datasets

A simple Python example for loading a CSV file is:

```
1 | import pandas as pd
2 |
3 | data = pd.read_csv("machine_data.csv")
4 |
5 | print(data.head())
```

Step 3: Inspect the Dataset

Before analysis, inspect the structure.


```
1 | print(data.shape)
2 | print(data.columns)
3 | print(data.info())
```

These commands can reveal the number of observations, available variables, and potential data-quality problems.

Step 4: Clean the Data

Real datasets are rarely perfect.

They can contain:

- Missing values
- Duplicate records
- Incorrect measurements
- Outliers
- Inconsistent units
- Incorrect data types

For example:

```
1 | data = data.drop_duplicates()
2 | data = data.dropna()
```

However, blindly deleting missing observations is not always appropriate. The correct method depends on why the data is missing.

Step 5: Explore the Data

Exploratory Data Analysis, or **EDA**, helps identify patterns.

```
1 | print(data.describe())
```

Visualization can provide additional insight:

```
1 | import matplotlib.pyplot as plt
2 |
3 | data["temperature"].plot()
4 | plt.title("Temperature Measurements")
5 | plt.xlabel("Observation")
6 |
7 |
```

```
plt.ylabel("Temperature")
plt.show()
```

Step 6: Identify Relationships

Suppose machine temperature increases before failure.

A scatter plot can help investigate this relationship.

```
1 | plt.scatter(data["temperature"], data["failure"])
2 | plt.xlabel("Temperature")
3 | plt.ylabel("Failure")
4 | plt.show()
```

Step 7: Build a Model

A classification model might predict whether a failure will occur.

```
1 | from sklearn.model_selection import train_test_split
2 | from sklearn.ensemble import RandomForestClassifier
3 |
4 | X = data[["temperature", "vibration", "pressure"]]
5 | y = data["failure"]
6 |
7 | X_train, X_test, y_train, y_test = train_test_split(
8 |     X, y, test_size=0.2, random_state=42
9 | )
10 |
11 | model = RandomForestClassifier(random_state=42)
12 | model.fit(X_train, y_train)
```

Step 8: Evaluate the Model

A model should never be considered successful simply because it produces predictions.




Engineers should evaluate appropriate metrics such as:



- Accuracy
- Precision
- Recall
- F1-score
- Mean absolute error
- Root mean squared error

- ROC-AUC

The metric depends on the problem.

The Power of Python in Data Science





Comparison

Traditional Data Analysis vs Data Science

Feature	Traditional Analysis	Data Science
Primary goal	Understand existing data	Understand and predict
Programming	Sometimes limited	Usually important
Statistics	Important	Highly important

Feature	Traditional Analysis	Data Science
Machine learning	Limited	Frequently used
Automation	Moderate	High
Data volume	Small to medium	Small to massive
Visualization	Common	Essential
Deployment	Less common	Often required

Python vs Other Languages


Characteristic	Python	R	MATLAB
Beginner friendliness	★★★★★	★★★★	★★★★
Data analysis	Excellent	Excellent	Excellent
Machine learning	Excellent	Excellent	Very good
Engineering computing	Excellent	Good	Excellent
General programming	Excellent	Moderate	Good
AI ecosystem	Excellent	Good	Good

Python is therefore particularly attractive to students who want a combination of **engineering, programming, analytics, and artificial intelligence**.


Diagrams and Tables

Data Science Pipeline





Typical Python Data Science Stack



Data Types

Data Type	Example	Possible Analysis
Numerical	Pressure = 25.4	Regression
Categorical	Material = Steel	Classification
Time series	Temperature/hour	Forecasting
Text	Customer review	NLP
Image	X-ray / inspection image	Computer vision
Sensor	Vibration signal	Predictive maintenance

Examples

Example 1: Calculating an Average

Suppose an engineer records five temperatures:

```

1 | temperatures = [21, 23, 25, 24, 27]
2 |
3 | average = sum(temperatures) / len(temperatures)
4 |
5 | print(average)

```

The result provides a basic statistical summary.

Example 2: Using pandas

```

1 | import pandas as pd
2 |
3 | df = pd.DataFrame({
4 |     "speed": [1000, 1200, 1400, 1600],
5 |     "temperature": [40, 45, 51, 59]
6 | })
7 |
8 | print(df.describe())

```

This creates a structured dataset and generates descriptive statistics.

Example 3: Correlation

```
1 | print(df.corr())
```

Correlation can help identify variables that move together.

⚠ But correlation does **not automatically prove causation**.

Real-World Application

Predictive Maintenance

Industrial equipment produces enormous amounts of sensor information.

Data science can combine:

Temperature + vibration + pressure + operating time → Failure prediction

A predictive-maintenance system might warn engineers before a bearing, pump, turbine, or motor reaches a critical condition.

Energy Optimization

Data scientists can analyze energy consumption and determine when machines or buildings consume excessive electricity.

Machine-learning models can estimate future demand and help optimize operating schedules.

Structural Engineering

Engineers can combine sensor measurements with historical inspection data to identify unusual structural behavior.

Potential inputs include:

- Strain
- Displacement
- Acceleration
- Temperature
- Wind speed

Healthcare Engineering

Data science can support medical-image analysis, hospital resource planning, and predictive analytics while requiring careful attention to privacy, validation, and regulatory requirements.

Financial Engineering

Financial institutions use statistical models and machine learning for forecasting, risk analysis, fraud detection, and portfolio-related analytics.

Common Mistakes

Starting With Machine Learning

❌ One of the biggest beginner mistakes is immediately trying to train sophisticated models.

A better progression is:

Python → Statistics → Data Cleaning → Visualization → Machine Learning

Ignoring Data Quality

A highly sophisticated algorithm cannot reliably compensate for severely flawed input data.

The principle is simple:

Poor data can produce poor decisions.

Using Too Many Libraries

Beginners sometimes install dozens of libraries before understanding the fundamentals.

Start with a focused toolkit:

- Python
- Jupyter
- NumPy

- pandas
- Matplotlib
- scikit-learn

Confusing Correlation With Causation

Two variables can be correlated without one causing the other.

Always investigate the underlying mechanism.

Evaluating Only Accuracy

A model with 99% accuracy may still be useless when the important event occurs only 1% of the time.

For imbalanced engineering problems, precision, recall, F1-score, and other metrics may be more informative.

Challenges & Solutions

Challenge: Missing Data

Solution: Determine why values are missing and choose an appropriate strategy such as imputation, removal, or model-based treatment.

Challenge: Large Datasets

Solution: Use efficient data structures, optimized queries, chunk processing, databases, and distributed computing when necessary.

Challenge: Overfitting

Overfitting occurs when a model learns training data too specifically and performs poorly on unseen observations.

Solution: Use:

- Train/test separation
- Cross-validation
- Regularization
- Feature selection
- Appropriate model complexity

Challenge: Difficult-to-Interpret Models

Complex models can be difficult to explain.

Solution: Use interpretable models when appropriate and apply explainability techniques where necessary.

Challenge: Changing Real-World Data

A model trained today may perform differently next year because operating conditions change.

Solution: Monitor model performance continuously and retrain when justified.

Case Study

Predicting Industrial Pump Failure

Imagine a water-treatment facility operating 100 industrial pumps.

Each pump produces:

- Temperature readings
- Vibration measurements
- Pressure readings
- Flow rate
- Operating hours
- Maintenance history

The engineering team wants to reduce unexpected failures.

Stage 1: Data Collection

Historical sensor readings are combined with maintenance records.

Stage 2: Data Cleaning

Incorrect sensor values are identified, duplicated records are removed, and missing observations are investigated.

Stage 3: Feature Engineering

The team creates additional variables such as:

Average vibration over 24 hours

Temperature increase over time

Operating hours since maintenance

Stage 4: Model Development

A classification model is trained using historical examples labeled as:

0 = No failure

1 = Failure

Stage 5: Evaluation

The team evaluates the model using precision, recall, F1-score, and a confusion matrix.

Stage 6: Deployment

When the model detects a high-risk condition, maintenance personnel receive an alert.

The objective is not merely to produce a high machine-learning score. The real objective is to **reduce downtime, improve safety, and optimize maintenance resources**.

Essential Tips

Build Projects

📌 Theory becomes much easier when connected to practical problems.

Good beginner projects include:

- House-price analysis
- Energy consumption prediction
- Weather-data analysis
- Student-performance analysis
- Industrial sensor analysis
- Customer segmentation

Learn Statistics

You do not need to become a theoretical mathematician, but understanding probability, distributions, correlation, regression, and uncertainty is extremely valuable.

Practice Data Cleaning

Real-world data science spends significant effort on data preparation.

Learning how to handle messy datasets is therefore more useful than memorizing dozens of algorithms.

Learn SQL

Python is powerful, but much professional data lives in relational databases.

SQL + Python is an extremely useful combination.

Understand the Engineering Problem

For engineering students and professionals, domain expertise can become a major advantage.

A data scientist who understands the physical system behind the data can often ask better questions and recognize unrealistic results.

Document Your Work

Keep track of:

- Dataset versions
- Assumptions
- Feature definitions
- Model parameters
- Evaluation results
- Software versions

This improves reproducibility and professional quality.

FAQs

What are the main elements of [data science](#)?

The major elements include data collection, data cleaning, statistics, exploratory analysis, visualization, feature engineering, machine learning, model evaluation, and deployment.

Is Python difficult for beginners?

Python is generally considered beginner-friendly because its syntax is relatively readable. However, becoming proficient in professional data science requires consistent practice.

Do I need advanced mathematics to learn data science?

Not at the beginning. Basic algebra, statistics, probability, and logical reasoning provide a strong foundation. More advanced machine learning and research applications may require deeper mathematics.

Which Python libraries should I learn first?

A practical starting set is **NumPy, pandas, Matplotlib, and scikit-learn**. Jupyter is also useful for interactive experimentation.

Is data science useful for engineers?

Yes. Data science can support predictive maintenance, process optimization, quality control, energy analysis, simulation, monitoring, forecasting, and intelligent automation.

Should I learn Python or machine learning first?

Learn basic Python first. Once you can work comfortably with variables, functions, loops, data structures, files, and libraries, move into data analysis and then machine learning.

What is the difference between data science and machine learning?

Data science is the broader discipline. Machine learning is one component of data science that focuses on algorithms capable of learning patterns from data.

Can data science guarantee accurate predictions?

No. Predictions always contain uncertainty. Model quality depends on data quality, assumptions, feature selection, model choice, validation, and changes in the real-

world system.

Conclusion

🚀 **Data science is not simply about writing Python code or training machine-learning models.** It is a complete problem-solving discipline that combines data, statistics, programming, visualization, computational methods, and domain expertise.

For beginners, the best route is gradual:




Python → Data Structures → NumPy/pandas → Statistics → Visualization → Machine Learning → Real Projects

For engineering professionals, the next level is learning how to integrate these techniques into actual systems involving sensors, databases, simulations, optimization, forecasting, and decision-making.

The most important principle is to start with the **problem**, not the algorithm. Once the engineering question is clearly defined, Python provides an exceptionally flexible environment for collecting, transforming, analyzing, visualizing, and modeling data.

Whether your goal is predictive maintenance, energy optimization, financial analytics, scientific computing, artificial intelligence, or engineering automation, understanding the fundamental elements of data science gives you a powerful foundation for building practical and intelligent solutions. 🔄📊⚙️🤖

Related Posts:



Preview PDF Book

[← Previous Book](#)

[Next Book →](#)

EngiVerse

Explore a vast library of free engineering ebooks curated to help you expand your technical knowledge, conduct groundbreaking research, and advance your career. Our mission is to create a hub of knowledge that fosters innovation, empowers aspiring engineers, and bridges the gap between education and industry.

[Home](#) [Books](#) [About Us](#) [Contact](#)
[Privacy Policy](#) [DMCA](#)

copyright 2025 EngiVerse. All Right Reserved