

Introducing Gemini 3.7 Flash

Aug 13, 2026 | 11 min read



Our most intelligent workhorse model yet for coding and agents.

Tulsee Doshi

Senior Director, Product Management, on behalf of the Gemini team



 [Listen to article](#)

6:23 minutes 

Today, we're building on the progress of our widely used Flash series by introducing Gemini 3.7 Flash, our most intelligent workhorse model yet for coding and agents.

This release comes just three weeks after [Gemini 3.6 Flash](#), and is a direct result of developer feedback and algorithmic innovations that we look forward to bringing to future models. 3.7 Flash delivers substantial improvements across software engineering, knowledge work, and web development workflows — with an introductory price of half the original 3.6 Flash cost per million tokens.

Better intelligence for complex workflows




< >

3.7 Flash shows strong gains over 3.6 Flash in coding tasks like debugging and issue resolution. It also achieves higher first-pass code accuracy and has improved performance in generating production-ready code as seen in [FrontierCode 1.1 Main](#) (43.6% vs 34.4%) and [DeepSWE v1.1](#) (65.3% vs 49.0%).


In web development, 3.7 Flash generates more functional layouts and feature-complete apps in fewer prompts. For UI generation, the model shows high design adherence and parity based on a reference input, whether it's a screenshot, an image, or a full design system. It outperforms 3.6 Flash on Arena.ai's [WebDev Arena](#) with an Elo score of 1588 vs 1538.

For knowledge-dense fields like finance, law, and biosciences, 3.7 Flash delivers improved reasoning and accuracy. It significantly outperforms 3.6 Flash on the GDP.pdf benchmark (34.0% vs 22.0%), an eval for testing a model's ability to process complex documents. It also surpasses 3.6 Flash in [AutomationBench](#), demonstrating it can more effectively complete real-world business workflows (30.4% vs 17.0%).



From a simple text prompt to a fully playable 3D game. We used Gemini 3.7 Flash combined with Nano Banana to dynamically generate characters, items, and...

+




Stunning, interactive landing pages generated i
Flash to orchestrate sub-agents, using Gemini C

< >

Better developer experience and price

3.7 Flash is available through the end of the year at an introductory price¹ of \$0.75/1M input tokens and \$3.75/1M output tokens. This price combined with the enhanced model performance enables developers and customers to scale production-ready agents cost effectively.



Gemini Spark, available to Google AI Pro and Ultra subscribers in over [160 countries](#), will be using Gemini 3.7 Flash starting today. We launched Spark at I/O as your personal AI agent that runs 24/7, taking action on your behalf while under your direction. This model update makes Spark more efficient for knowledge work with improved tool use for Google Workspace apps, delivering improved accuracy and output quality for complex, multi-skill workflows.

Gemini Spark helps you turn your ideas into action

With 3.7 Flash, Gemini Spark can turn ideas into action more efficiently by consolidating files, drafting emails, and updating status documents.

+

Built with safety in mind

We continually work to improve the coverage and robustness of [Frontier Safety safeguards](#). Gemini 3.7 Flash is shipping with updated safeguards against misuse in the domains of Chemical, Biological, Radiological, and Nuclear (CBRN) and cyber offense, while enabling beneficial use cases, in accordance with [our approach to bioresilience](#) and our [cyber program](#).

For more information, see the 3.7 Flash [model card](#).

Try it today

- **Developers:** Explore agent-first workflows in [Google Antigravity](#) or start building today in the Gemini API via [Google AI Studio](#) and [Android Studio](#). Get started with our [developer guide](#).
- **Enterprises:** Access 3.7 Flash in [Gemini Enterprise Agent Platform](#) and the [Gemini Enterprise](#) app.
- **Individuals:** Available via Spark, your 24/7 personal agent in the Gemini app for Google AI Pro and Ultra subscribers in [supported countries](#).

Detailed benchmarks

Benchmark		Gemini 3.7 Flash	Gemini 3.6 Flash	Claude Sonnet 5	GPT-5.6 Terra	Muse Spark 1.2
Input price	\$/M tokens	\$0.75*	\$0.75*	\$2.00	\$2.00	\$1.25
Output price	\$/M tokens	\$3.75*	\$3.75*	\$10.00	\$12.00	\$4.25
Artificial Analysis Intelligence Index	Composite model intelligence	56	52	55	57	57
FrontierCode 1.1 Main	Production code quality	Score 43.6%	34.4%	42.7%	41.3%	-
DeepSWE v1.1	Long-horizon software engineering	65.3%	48.6%	53.8%	69.6%	54.9%
Code Arena	Web development	Elo 1588	1538	1541	1523	1535
Terminal-bench 2.1	Agentic terminal coding	85.8%	78.0%	80.4%	87.4%	82.9%
Terminal-bench 3.0	General agent capabilities	14.9%	5.4%	14.6%	20.8%	-
AutomationBench	Enterprise workflow automation	Private set 30.4%	17.0%	10.7%	23.6%	-
GPVval-AA v2	Knowledge work	Elo 1525	1422	1598	1578	1628
Harvey LAB-AA	Complex legal workflows	90.7%	85.1%	90.1%	85.2%	-
GDRpdf	Expert PDF document comprehension	34.0%	22.0%	28.0%	24.7%	16.0%
CharXiv Reasoning	Information synthesis from complex charts	No tools 84.5%	85.2%	77.0%	85.9%	-
		With tools 88.7%	89.4%	88.3%	-	-
LVBench	Long video understanding	85.4%	84.2%	68.5%	78.9%	-
GDM-MRCR v2 (8-needle)	Long context performance	128k (average) 97.0%	91.8%	81.5%	93.5%	-
OSWorld-2.0	Agentic computer use	47.9%	33.8%	-	50.2%	-
Agent's Last Exam	Multi-modal desktop and GUI agent tasks	Pass rate 26.3%	24.2%	33.3%	28.0%	-
HLE-Verified	Multidisciplinary expert reasoning	53.6%	51.2%	31.0%	51.1%	-
BioMysteryBench	Bioinformatics research reasoning	Human solvable 87.1%	80.6%	87.5%	83.8%	-
		Human difficult 43.5%	41.2%	34.1%	49.4%	-
LABBench2	Biology real-world research tasks	82.1%	76.1%	80.1%	81.2%	-

Methodology: deepmind.google/models/evals-methodology/gemini-3-7-flash

* For 3.6 and 3.7 Flash, introductory price expires on December 31, 2026. Starting January 1, 2027, \$1.50/1M input tokens and \$7.50/1M output tokens will apply.

POSTED IN:

Gemini models

More Information



Related stories



Gemini models

Omni experts share what excites them most about the model.

By Lindsey Lanquist

Gemini models

See what 5 builders are making with Gemini Omni

By Lindsey Lanquist

AI

The latest AI news we announced in July 2026

By News from Google Team

< >

Get the latest news from Google in your inbox

Sign up for our newsletters with product updates, event information, special offers, and more.

Email address

Subscribe

Your information will be used in accordance with [Google's privacy policy](#). You may opt out at any time.