



Вселенная

Think Stats: исследовательский анализ данных. 2-е издание

Автор: Аллен Б. Дауни

Тип файла: pdf

Размер: 15,8 МБ

Язык: Английский

Страниц: 226

Think Stats: Exploratory Data Analysis 2nd Edition — практическое руководство по статистике для инженеров на Python  

Введение

Современная инженерия все больше опирается на данные. Независимо от того, отслеживает ли инженер работу оборудования, анализирует ли энергопотребление, оценивает ли качество производства, изучает ли сетевой трафик или создает интеллектуальную систему, одних только необработанных измерений редко бывает достаточно. Инженерам нужны методы преобразования наблюдений в полезную информацию.





Именно здесь на первый план выходит **исследовательский анализ данных (exploratory data analysis, EDA)**. Вместо того чтобы сразу применять сложные математические модели, EDA побуждает инженеров сначала понять, как на самом деле выглядят их данные: их распределение, вариативность, необычные наблюдения, взаимосвязи и возможные источники искажений.


«**Статистика: исследовательский анализ данных**», **2-е издание**, написанная Алленом Б. Дауни и опубликованная издательством O'Reilly в 2014 году, рассматривает статистику с точки зрения программирования. Второе издание представляет собой 226-страничное введение, предназначенное для читателей с базовыми навыками работы с Python, которые хотят изучить теорию вероятностей, статистику и практический анализ данных.

Книга особенно интересна студентам инженерных и технических специальностей, поскольку в ней программирование рассматривается как инструмент для понимания статистических идей, а не как чисто теоретическая математика.


Exploratory Data Analysis (EDA) with Python

Validate and understand your dataset with sanity checks and statistical exploration





Data Visualization Techniques in EDA



Основная идея проста:

**Данные → Исследование → Доказательства → Статистические рассуждения →
Инженерное решение**

Этот рабочий процесс может помочь новичкам развить статистическую интуицию, а опытным инженерам — получить практическую основу для изучения наборов данных.

Теоретическая база

Что такое исследовательский анализ данных?

Исследовательский анализ данных — это систематическое изучение набора данных перед тем, как делать однозначные выводы или строить сложные прогностические модели.

Инженер может начать с простого вопроса:

Влияет ли повышение рабочей температуры на выход компонентов из строя?

Вместо того чтобы сразу подгонять данные под модель машинного обучения, EDA задает вопросы:

- Какие значения температуры были зафиксированы?
- Сколько наблюдений доступно?
- Не хватает ли каких-либо измерений?
- Какова типичная температура?
- Насколько сильно она варьируется?
- Есть ли экстремальные значения?
- Изменяется ли частота отказов в зависимости от температуры?
- Есть ли другие переменные, влияющие на эту взаимосвязь?

Этот процесс снижает риск формирования выводов на основе неверно интерпретированных данных.

Статистическое мышление

Традиционная статистика иногда может показаться слишком математизированной. Think Stats предлагает вычислительный подход, используя эксперименты на Python и реальные наборы данных, чтобы сделать статистические концепции более осязаемыми. Во втором издании книги рассматриваются распределения, функции распределения вероятностей, функции распределения, моделирование распределений, взаимосвязи между переменными, оценка, регрессия, анализ временных рядов, анализ выживаемости и аналитические методы.

Для инженеров это ценно тем, что статистические концепции становятся связанными с измеримыми величинами.

Например:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

представляет собой выборочное среднее.

В то же время дисперсия показывает, насколько сильно значения отклоняются от этого среднего:

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

Это не просто формулы. Они отвечают на практические инженерные вопросы о **средней тенденции и вариативности**.

Definition

Think Stats: Exploratory Data Analysis 2nd Edition

Think Stats, 2nd Edition is a practical introduction to statistics and probability for Python programmers, with an emphasis on computational exploration of real datasets.

Rather than requiring advanced mathematical statistics as a starting point, the book uses programming, experiments, simulations, visualizations, and case studies to develop understanding. The official description specifically emphasizes exploring real datasets, answering interesting questions, and using short programs to investigate statistical behavior.

Core concepts

The second edition introduces ideas such as:

Concept	Engineering purpose
Distribution	Understand how measurements are spread
Mean & median	Describe central tendency
Variance	Quantify variability
Percentiles	Identify relative positions
PMF	Study discrete outcomes
CDF	Understand cumulative probabilities
Correlation	Examine relationships
Regression	Model relationships
Resampling	Investigate sampling uncertainty
Hypothesis testing	Evaluate statistical evidence
Time-series analysis	Study measurements over time
Survival analysis	Examine time-to-event behavior

The objective is not simply to calculate statistics, but to understand **what those statistics mean**.

Step-by-Step Explanation: Applying Think Stats Principles

Step 1: Define the engineering question

Begin with a specific question.

For example:

“Does increasing machine temperature correspond to increased vibration?”

Avoid starting with:

“Let’s analyse everything in this dataset.”

A focused question determines which variables and statistical techniques matter.

Step 2: Import the data

A typical Python workflow may begin with a CSV dataset:

```
import pandas as pd
data = pd.read_csv("machine_data.csv")
print(data.head())
print(data.info())
```

At this stage, the goal is not prediction.

The goal is **inspection**.

Step 3: Validate the measurements

Check:

- Missing values
- Incorrect units
- Duplicate observations
- Impossible measurements
- Incorrect data types
- Sensor errors
- Unexpected categories

For example, if temperature should be between -20°C and 150°C but the dataset contains 8500°C , something needs investigation.

Step 4: Explore individual variables

A histogram can reveal whether measurements are approximately symmetric, skewed, multimodal, or dominated by unusual observations.

```
import matplotlib.pyplot as plt
data["temperature"].hist(bins=20)
plt.xlabel("Temperature (°C)")
plt.ylabel("Frequency")
plt.show()
```

This is much more informative than looking at the mean alone.

Step 5: Calculate summary statistics

Useful quantities include:

Mean, Median, Variance, Standard Deviation

In Python:

```
print(data["temperature"].describe())
```

The median can be particularly useful when extreme measurements distort the mean.

Step 6: Explore relationships

If two variables may be related, a scatter plot provides a quick visual diagnostic.

```
plt.scatter(data["temperature"], data["vibration"])
plt.xlabel("Temperature (°C)")
plt.ylabel("Vibration")
plt.show()
```

A pattern may suggest a relationship, but **correlation does not automatically establish causation**.

Step 7: Test the engineering hypothesis


Once a pattern has been identified, statistical methods can be used to determine whether the evidence is sufficiently strong.

This may involve:

- Resampling
- Confidence intervals
- Hypothesis tests

- Correlation analysis
- Regression

The result should be interpreted in engineering context rather than treated as an isolated mathematical answer.



Reading data

```
In [4]: train = pd.read_csv("train.csv.zip")
        test = pd.read_csv("test.csv.zip")
```


Exploratory Data Analysis

```
In [5]: print("Target column has", train['target'].isna().sum(), "missing values")
        print("comment_text column has", train['comment_text'].isna().sum(), "missing values")
```


```
Target column has 0 missing values
comment_text column has 0 missing values
```

```
In [7]: # see distribution of the comment_text length
        train['comment_text'].apply(lambda x : len(str(x).split())).hist()
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1d78266e438>
```



Steps for Performing Exploratory Data Analysis



Comparison: Think Stats 2nd Edition vs Traditional Statistics Learning

Conceptual comparison

Approach	Traditional statistics course	Think Stats approach
Starting point	Mathematical concepts	Data and practical questions
Main tool	Equations	Python + statistics
Learning method	Lectures and formulas	Experiments and datasets
Probability	Often theoretical	Computationally explored
Visualization	Supporting technique	Major exploratory tool
Programming	Sometimes optional	Central to workflow
Real datasets	Variable	Strong emphasis

Approach	Traditional statistics course	Think Stats approach
Simulations	May be limited	Important learning method
Engineering relevance	Depends on course	Strong practical connection


Why this matters to engineers



An engineer frequently receives imperfect datasets rather than perfectly structured textbook examples.

The ability to write a short Python program, inspect a distribution, simulate a process, and test an assumption can therefore be more useful than memorizing dozens of formulas.

Diagrams and Tables: The EDA Engineering Pipeline

A practical EDA pipeline can be represented as:





Choosing the right visualization

Engineering question	Useful visualization
How are values distributed?	Histogram
Where is the median?	Box plot

Engineering question	Useful visualization
Are two variables related?	Scatter plot
How does a variable change over time?	Line plot
How do categories compare?	Bar chart
Are there unusual observations?	Box plot / scatter plot
How do cumulative probabilities behave?	CDF

Different plots reveal different structures. For example, histograms are useful for distributions, scatter plots for relationships, line plots for ordered trends, and bar plots for category comparisons.

Examples

Example 1: Manufacturing

Imagine a factory measuring the diameter of metal components.

The target diameter is:

$D=25.00\text{ mm}$

Suppose 10,000 components are measured.

The engineer can calculate:

$\bar{D}$

and

s

but these values alone do not tell the entire story.

A histogram could reveal two distinct peaks, suggesting that two production machines may be operating differently.

That discovery might be invisible if the engineer only reports the overall average.

Example 2: Energy engineering

Suppose an engineer records building electricity consumption every 15 minutes.

EDA can identify:

- Daily peaks
- Weekend behavior
- Seasonal patterns
- Unusual consumption
- Relationship between temperature and electricity demand

A time-series visualization can expose patterns that a single annual average would completely hide.

Example 3: Civil engineering

Consider concrete compressive-strength measurements.

If the dataset contains:

$fc' = [31, 32, 30, 33, 31, 32, 58]$

the value 58 MPa deserves investigation.

It may represent:

- A genuinely strong sample
- A different concrete mix
- A measurement problem
- A transcription error

An engineer should investigate the observation rather than automatically delete it.

Real-World Applications

Mechanical engineering

EDA can support:

- Predictive maintenance
- Vibration monitoring
- Failure analysis

- Fatigue testing
- Manufacturing quality control

Engineers can investigate relationships among temperature, vibration, load, speed, and component lifetime.

Electrical engineering ⚡

Possible applications include:

- Power-quality analysis
- Voltage monitoring
- Load forecasting
- Battery testing
- Sensor-data analysis

Civil engineering 🏗️

EDA can be applied to:

- Material testing
- Structural monitoring
- Traffic measurements
- Soil properties
- Construction quality control

Software and data engineering 💻

The same principles apply to:

- Server response times
- Network latency
- Error rates
- User activity
- Database performance

A histogram of response times, for example, may reveal a long tail that is hidden by the average response time.

Common Mistakes

Mistake 1: Trusting the mean blindly

The mean can be strongly affected by extreme values.

Use the median and distribution visualization when appropriate.

Mistake 2: Confusing correlation with causation

If:

$r=0.85$

between two variables, that indicates a strong linear association, not necessarily a causal relationship.

A third variable may influence both.

Mistake 3: Removing outliers automatically

An outlier is not automatically an error.

It may represent an important engineering event.

Mistake 4: Ignoring data quality

Garbage data can produce sophisticated but meaningless conclusions.

Mistake 5: Building models too early

A complicated machine-learning model cannot compensate for a misunderstood dataset.

Mistake 6: Using inappropriate visualizations

A chart should answer a question rather than simply decorate a report.

Challenges & Solutions

Challenge	Practical solution
Missing measurements	Investigate the collection process
Extreme values	Validate before removing
Small sample size	Use uncertainty-aware interpretation

Challenge	Practical solution
Non-normal distribution	Examine empirical distributions
Confounding variables	Use stratification or multivariable models
Measurement noise	Examine repeated observations
Time dependence	Use time-series methods
Data-entry errors	Apply validation rules
Misleading charts	Match visualization to the variable type

One of the strongest lessons of computational statistics is that **experimentation can make abstract concepts tangible**. Think Stats explicitly uses short programs and experiments to help readers develop statistical understanding.

Case Study: Investigating Birth-Weight Data 🧠 📊

One of the important case-study contexts associated with Think Stats is the **National Survey of Family Growth (NSFG)**.

Instead of presenting statistics as disconnected exercises, the book uses real-world survey data to demonstrate the process of importing data, examining variables, transforming information, validating observations, and interpreting results. The second-edition chapter structure begins with exploratory data analysis and then progresses through distributions, probability, relationships, estimation, regression, time series, survival analysis, and analytical methods.

Consider an engineering-style interpretation of the workflow.

Suppose the dataset contains a measurement X and we want to determine whether two groups differ.

First:

μ_1 =mean of Group 1

and

μ_2 =mean of Group 2

Then examine:

$$\Delta = \mu_1 - \mu_2$$

But the important question is not simply whether $\Delta \approx 0$.

The analyst should ask:

 **How large is the difference?**

How variable are the measurements?

How much uncertainty exists?

Could sampling variation explain the observed difference?

This way of thinking is transferable to engineering experiments, where two manufacturing processes, materials, sensors, or operating conditions may need to be compared.

Essential Tips for Students and Professionals

Start with questions, not algorithms

Before opening Python, write down the engineering question.

Plot before modelling

A simple graph can reveal:

- Skewness
- Clusters
- Outliers
- Trends
- Nonlinear relationships

Understand your variables

Know the units, measurement method, physical meaning, and expected range.

Use simulation to build intuition

Random sampling and resampling can make probability concepts easier to understand.

Keep a reproducible workflow

Save:

- Original data
- Cleaning steps
- Python scripts
- Analysis outputs
- Figures
- Assumptions

Combine statistical and engineering knowledge

A statistically unusual value may be completely reasonable from a physical perspective.

The strongest analysis combines:

Statistical evidence+Domain knowledge

Use Python as a thinking tool

The major benefit is not simply learning syntax. It is learning how to use computation to ask better questions about data.

The official second-edition resources provide the book online along with supporting code, while the author also provides the current third edition separately.

FAQs

What is Think Stats 2nd Edition?

Think Stats: Exploratory Data Analysis, 2nd Edition is a practical introduction to probability, statistics, and exploratory data analysis using Python. It is written by Allen B. Downey and published by O'Reilly.

Is Think Stats 2nd Edition suitable for beginners?

Yes. The book is classified by O'Reilly as beginner-level and assumes basic Python knowledge rather than advanced statistical mathematics.

Do I need advanced mathematics?

No. Basic mathematics is sufficient to begin. The computational approach helps readers understand statistical ideas through programming and experiments.

Is the book useful for engineering students?

Yes. Its emphasis on real datasets, distributions, visualization, statistical inference, regression, and time-series analysis makes it applicable to many engineering disciplines.

What Python skills should I know first?

You should be comfortable with basic Python concepts such as variables, functions, loops, lists, dictionaries, and importing modules. Familiarity with NumPy and pandas is useful.

What is the difference between EDA and [machine learning](#)?

EDA focuses on **understanding the data**. Machine learning focuses on learning patterns that can be used for prediction or other computational objectives.

EDA generally comes before sophisticated modelling.

Is the second edition still useful?

Yes. The second edition remains a useful introduction to computational statistics and EDA. However, Allen Downey now provides a third edition, so readers should be aware that the second edition is no longer the author's current edition.

Where can I find the official Think Stats resources?

The author's Think Stats 2nd Edition website provides online reading and supporting code.

Conclusion

Think Stats: Exploratory Data Analysis, 2nd Edition provides an accessible bridge between programming and statistical thinking. Its greatest strength is the idea that

statistics should not be treated only as a collection of equations. Instead, engineers can use code to investigate data, run experiments, visualize distributions, test assumptions, and develop evidence-based conclusions.

For an engineering student, this approach creates a practical foundation for data science. For a professional engineer, it provides a repeatable workflow for examining measurements before committing to complex models.

The essential process can be summarized as:



Question→Data→Validation→EDA→Statistical Evidence→Engineering Decision

 **The key lesson is simple:** don't let a model speak before you understand the data.

That principle remains valuable across mechanical engineering, electrical engineering, civil engineering, manufacturing, energy systems, software engineering, and modern data-driven engineering.

Note: This article is an original educational overview and does not reproduce the book's text. The factual book details above are based on the official author/O'Reilly information and bibliographic sources.

Related Posts:



Preview PDF Book

[← Previous Book](#)

[Next Book →](#)

EngiVerse

Explore a vast library of free engineering ebooks curated to help you expand your technical knowledge, conduct groundbreaking research, and advance your career. Our mission is to create a hub of knowledge that fosters innovation, empowers aspiring engineers, and bridges the gap between education and industry.

[Home](#) [Books](#) [About Us](#) [Contact](#)
[Privacy Policy](#) [DMCA](#)

copyright 2025 EngiVerse. All Right Reserved