

10 сентября 2026 года Продукт Релиз

Создавайте более естественные голосовые интерфейсы с помощью GPT-Live-1 в API

GPT-Live-1 добавляет в API естественные полнодуплексные диалоги ChatGPT с возможностью более гибкого управления тем, как голосовые агенты говорят и ведут себя.

OpenAI

Впервые представленная в ChatGPT, GPT-Live-1 способна одновременно слушать и говорить, и, как в случае с Codex и ChatGPT Work, может делегировать более сложные задачи и действия моделям и инструментам, с которыми она работает.

При выпуске API для GPT-Live-1 мы сосредоточились на новых возможностях, которые позволяют разработчикам управлять голосовым взаимодействием и настраивать его в соответствии с потребностями пользователей, рабочими процессами и целями. Одна из ключевых особенностей GPT-Live-1 — плавная обработка прерываний — уже приносит пользу бизнесу: по предварительным оценкам Speak, GPT-Live-1 дает учащимся больше времени на обдумывание ответа, сокращая количество прерываний почти на 80 % по сравнению с предыдущими системами с пошаговым управлением.

Ключевые преимущества GPT-Live-1 в API:

- **Обработка прерываний:** Улучшает обработку прерываний с помощью единой модели, которая анализирует входящий и исходящий аудиосигнал, избегая задержек и ненадежных переключений между последовательностями STT–LLM–TTS.
- **Делегирование рассуждений и вызовов инструментов:** GPT-Live-1 может делегировать рассуждения и вызовы инструментов серверной текстовой модели, такой как GPT-6 Astra, или сторонней модели.
- **Тон, темп и стиль:** Позволяет разработчикам задавать тон, темп и стиль общения агента с помощью системных подсказок.
- **Бесшумное управление контекстом и фоновым шумом:** Лучше справляется с фоновым шумом и тишиной, не прерывая разговор и не озвучивая каждый шаг.
- **Надежность при длительных сеансах:** Улучшает сохранение контекста и качество общения при длительных взаимодействиях.
- **Поддержка телефонии:** Позволяет использовать полнодуплексных голосовых агентов для телефонных звонков — от бронирования столика в ресторане до поддержки клиентов.

Попробуйте GPT-Live-1

Начните разговор и ведите себя естественно. Прерывайте собеседника, смейтесь, меняйте свое мнение — попробуйте сделать это дома или в шумном месте, например в кафе или на городской улице.

Начало сеанса

Посмотрите, на что это способно ▾

Эта демонстрация ограничена по времени. Используя ее, вы соглашаетесь с [Условиями](#) OpenAI и подтверждаете, что ознакомились с нашей [Политикой конфиденциальности](#).

Yelp Cognition Picsart Heygen OpenAI Presence



Hi, I'd like to make a reservation for Saturday at 7 p.m. for four people.

Hi, thanks for calling Stella Lago. How can I help you?





Yelp Host seamlessly secures a reservation with GPT-Live-1 handling background noise, side conversations, and interruptions.

Simplify your voice-agent architecture and reduce voice latency

Traditional voice agents stitch together speech-to-text, a reasoning model, and text-to-speech. Each handoff adds latency and creates more opportunities to lose timing, context, or the natural rhythm of a conversation. Developers are often the ones left coordinating those stages, including what happens when someone interrupts, pauses, or changes direction.

GPT-Live-1 handles listening and speaking in a single model, simplifying the voice layer. It can respond to interruptions and acknowledgements as they happen, while delegating deeper reasoning to the back end. This lets the conversation continue while work happens in the background.

“Compared to our cascaded build, GPT-Live-1 simplified our code base by 80% and removed 23K lines of code. This enabled natural, real-time patient conversations & freed our team to improve the experience from booking an appointment to navigating care.”

—Tony Stoyanov, Co-Founder & CTO

Developers choose the models, tools, and agent harness behind the conversation. For example, they might pair GPT-Live-1 with a model like Luna for high-volume tasks like scheduling or order updates, and use a model like Astra for complex customer issues that require reasoning. That flexibility lets developers match reasoning depth, speed, and cost to each task.

GPT-Live-1 natively provides ASR transcripts and response text. It also offers strong alphanumeric understanding and supports keyword biasing. Although GPT-Live-1 is not a turn-based model, it natively supports turn detection, so developers can continue to build around explicit turn boundaries.

Measuring the full-duplex advantage

Across our evaluations, GPT-Live-1 improves Full Duplex Bench performance by 30 percentage points over GPT-Realtime-2.1, with large gains in turn-taking latency and interactive behavior. Paired with GPT-6 Astra at medium reasoning effort, it also ranks #1 on Tau3, which measures frontier voice-agent intelligence on end-to-end tasks.

OpenAI

Evaluates spoken customer-service tasks in airline, retail, and telecom domains. Pass@1 measures task success; the headline gives each domain equal weight.

* GPT Live backend: Astra (medium).

Evaluates spoken banking support with knowledge retrieval and account tools. Pass@1 is the fraction of 97 banking_knowledge tasks completed successfully.

* GPT Live backend: Astra (medium).

Evaluates pause handling, conversational turn taking, interruptions, and backchannels.

Measures how quickly the agent starts its reply after the user finishes a turn.

OpenAI

Tests tool use from spoken requests containing natural pauses, hesitations, and self-corrections. Pass@1 scores the tool-call sequence.

* GPT Live backend: Terra (low).

Evaluates the spoken answer to tool-using requests containing pauses, hesitations, and self-corrections. Scores how well the answer matches the reference intent.

* GPT Live backend: Terra (low).

What customers are saying

OpenAI

our traditional voice architecture. When Yelp Host uses GPT-Live-1 to answer calls, like reservations and food orders, we're seeing meaningful improvements in call handling rates. Callers are also speaking fuller, more natural sentences, which tells us the experience on the other end of the phone feels genuinely different.”

—Alex Levy, Chief Technology Officer

New voice options

Developers need voices that fit their product and sound natural to the people using it. With GPT-Live-1, we’re expanding from a small set of real-time voices to a broader selection across accents, dialects, and languages giving developers more choice in how their assistants sound.

Listen to the new voices

- Quartz
- Ripple
- Vesper
- Willow
- Stone
- Gleam
- Meridian
- Bossa
- T



Listen

00:07

Australian English influenced

We’ll continue to expand voice options and language availability over the coming months.

Pricing & Availability

GPT-Live-1 is available in the API today at \$0.05 per minute for the front-end voice layer. Pair it with the backend model and agent harness that fit your product, then build a voice experience that can scale with the work it needs to do.

Connecting GPT-Live-1 To Codex



```
1 import { Codex } from "@openai/codex-sdk";
2
3 const thread = new Codex().startThread({
4   workingDirectory: "./repo",
5   sandboxMode: "read-only",
6   approvalPolicy: "never",
7 });
8
9 async function answer(live, delegationId, context) {
10   const { finalResponse } = await thread.run(
11     `Answer the latest question using this repo.
12     Reply in two short spoken sentences.\n${context}`
13   );
14
15   live.send({
16     type: "session.commentary.append",
17     delegation_id: delegationId,
18     content: finalResponse,
19   });
20 }
```

Connecting GPT-Live-1 to Codex. This excerpt shows how an application passes conversation context to Codex and returns its answer to GPT-Live-1. Connection setup and delegation handling are omitted.

For custom voice access, [contact sales](#) to learn more about eligibility and the request process.

Another way to build voice workflows on top of GPT-Live-1 is with [OpenAI Presence](#), which uses the model to power real-time voice interactions. Presence helps enterprises deploy trusted AI agents that can answer questions, resolve issues, use company

API Platform 2026

Author

OpenAI

Keep reading

View all



Now everyone can put data to work

Product Sep 10, 2026



Introducing ChatGPT for Financial Services

Product Sep 10, 2026



Introducing the Agents API

Product Sep 10, 2026

Research

Research Index

Research Overview

Products

ChatGPT ↗

ChatGPT Business

Business

Overview

Solutions ↗

Company

About Us

Our Charter

More

Stories

Academy

Latest Advancements	ChatGPT for Education ↗	Partner Network	Support	Podcast
GPT-6	Codex	Contact Sales	Help Center ↗	RSS
GPT-5.6	Release Notes	Developers		Terms & Policies
GPT-5.5	API Platform	Apps SDK ↗		Terms of Use
GPT-5.4	Overview	Open Models		Privacy Policy
	API Log In ↗	Docs ↗		Other Policies
Safety	Docs ↗	Resources ↗		
Safety Approach		Developer Forum ↗		
Deployment Safety ↗				
Security & Privacy				
Trust & Transparency				