

DeepSeek Plans 160,000 Huawei Ascend Chips for 1GW Ulanqab Site

By [Alexis Dufresne](#) · Published September 4, 2026 at 10:57 UTC · Updated September 4, 2026 at 11:53 UTC

DeepSeek

China AI

Chips

NVIDIA

ai-business

Free AI alerts in your inbox

Get →

TL;DR

- DeepSeek plans to deploy at least 160,000 Huawei Ascend 950DT accelerators at a roughly 1 GW data center in Ulanqab, Inner Mongolia.
- The Hangzhou startup wants at least part of the capacity operating by the end of 2027 or early 2028, subject to Huawei's production.
- DeepSeek intends to use the 950DT for operating its models, not training, even though Huawei designed and marketed the chip for training.

EDITOR'S NOTE

A 160,000-accelerator commitment on a single site is the clearest signal yet that a top Chinese lab will run user-facing inference on Huawei rather than Nvidia; any China-inference revenue model built on the incumbent should be reweighted now.

DeepSeek plans to deploy at least 160,000 of Huawei's top AI accelerators at a data center it is building in Ulanqab, Inner Mongolia, [Bloomberg reported](#). If it lands as described, it would be 'one of the largest known clusters of Huawei AI chips' anywhere.

your@email.com

Get them

from other operators, according to the report. It wants at least part of the capacity 'operating by the end of 2027 or early 2028.' The silicon in question is Huawei's next-generation Ascend 950DT, which Bloomberg says DeepSeek intends to use for 'operating its models, though the timeline for installation will depend on Huawei's production capabilities.'

The unusual choice is what the 950DT is being reserved for. 'The startup doesn't currently plan to use the 950DT chips for training,' Bloomberg writes, 'even though Huawei designed and marketed these processors to serve that more demanding function.' Inference, not training, is the workload DeepSeek wants hosted domestically at this scale.

Ulanqab already hosts Alibaba, VNET Group and Kuaishou, and more than 90% of the area's computing capacity is dedicated to AI, [Yahoo Finance's writeup](#) of the same reporting notes. Neither Bloomberg nor the follow-on coverage names a dollar figure, a financing structure or the power-procurement partner. It is our latest entry in a long run of [China AI stories](#) documenting how quickly the domestic accelerator stack is being wired into production, more than 340 of them in the last 90 days alone.

EDITOR'S TAKE

Watch Huawei's late-2026 950DT shipment and yield disclosures: if the ramp slips, DeepSeek's late-2027 partial-capacity target moves with it and the 160,000-chip figure becomes intent rather than delivery.

- Does reserving the 950DT for inference imply DeepSeek is still training on Nvidia today, and what is its plan when its current training silicon ages out?

Get the next one before it breaks

AI Weekly tracks stories like this 3x/week. Breaking alerts only when they matter. 50,000+ AI professionals.

Originally reported by [bloomberg.com](#)

[Read the original article →](#)

Original headline: DeepSeek Orders 160,000 Huawei Ascend 950DT Chips for 1GW Inner Mongolia AI Data Center

Don't miss the next one

3x/week. Breaking AI alerts only when they matter. Free, no spam, opt out anytime.

Subscribe

☒ Also get breaking AI alerts