

What Claude Doesn't Know About Your Business (and How to Fix That)

EBOOK

Introduction

If you use Claude daily, you already know what it's good at: making sense of the information you give it. However, its blind spot is everything you don't provide: your contracts, tickets, wikis or specs that exist across multiple systems.

Claude was built to reason, not to store and govern your business knowledge. That job belongs to agentic retrieval-augmented generation (RAG) technology provided within the Progress® Agentic RAG platform, connected via Model Context Protocol (MCP).


The Progress Agentic RAG solution retrieves the correct, source-cited context automatically, so there's no need to copy and paste it into Claude.

Agentic RAG is the foundation for every AI experience your team creates, with Anthropic models handling the reasoning and generation on top. It provides a single knowledge base to support any number of Claude-powered experiences—for example, retrieval feeds Claude only what's relevant instead of stuffing context windows, so your tokens go toward reasoning rather than re-explaining your business every time you use AI.

In this ebook, you'll learn how to combine Claude's reasoning with the Progress Agentic RAG solution to deliver trusted, context-rich AI experiences that scale from personal productivity to enterprise workflows.

What's Ahead

- What Claude is built for, and where a knowledge layer must take over
- How the Progress Agentic RAG solution turns scattered content into usable context
- How the MCP connects between Claude's reasoning and that context
- How to build custom AI experiences on the solution, with Claude generating the answers
- How every answer is scored for trust, automatically with RAG Evaluation Metrics Intelligence (REMi)
- Why smarter retrieval allows you to get more out of your Claude tokens
- How to scale AI from personal use to a production workflow



What Claude Is Built For, and Where a Knowledge Layer Has to Take Over

Claude's core skill is reasoning: reading, synthesizing, comparing, drafting and explaining based on whatever context it digests. You can provide Claude with a contract, a codebase or a messy set of notes, and it can summarize it and provide clear insights. That's the job it's trained for and the reason people use it.

What Claude doesn't do is treat your business documents as a source of truth. Its memory can recall your preferences and working style across chats, but that's derived from personal context, not your company's contracts, tickets or product specs. Therefore, answers lack citations to original sources.

If you ask Claude a question that is outside the context you've provided, it may come up empty-handed. Worse, it might hallucinate and generate a plausible-sounding answer that isn't grounded in truth.

And that gap can't be patched with a bigger context window or a smarter prompt; it requires an entirely different approach that:

1. Ingests and indexes content as it changes
2. Retrieves the right pieces for a given question
3. Attaches a verifiable source to every answer


This is the key responsibility of an enterprise knowledge layer: continuously

maintaining, governing and auditing context. And it's not something a reasoning model can do as a side function.

Here's the handoff point:

- Claude reasons over context; it doesn't source, store or govern it
- A knowledge layer sources, stores and governs content; it doesn't reason over it

Put together, you get an answer that's both intelligent and traceable back to where it came from. Skip the knowledge layer, and Claude is only ever as good as what you happened to paste into your assets that day.



How Progress Agentic RAG Turns Scattered Content into Usable Context

Claude's reasoning is only as good as the context it receives. This is where the Progress Agentic RAG solution comes in: it creates that context by transforming scattered enterprise content into a unified, retrievable knowledge layer.

This starts with ingestion. Pre-built connectors pull in content from SharePoint, Google Drive, Dropbox, Progress® ShareFile® software or CMS platforms, and Sync Agents keep that content current automatically. No one has to re-enter a spec sheet manually every time the information changes. Coverage runs across 60+ file formats and any language, as well as multi-modal content such as audio and video, so a recorded customer call or a scanned PDF is just as searchable as a Word doc.

Ingestion alone isn't enough, though; it just moves the pile of data into one place. The harder problem is making that data answerable. Agentic RAG technology indexes content across three layers at once—semantic vectors, full-text and knowledge graphs—so retrieval captures meaning and relationships between concepts, not just keyword matches.

On top of that sit 30+ retrieval strategies, including hybrid search, structured queries, reranking and more, each of which is tunable per use case, without the need to reindex everything from scratch.

By the time a question reaches Claude, the heavy lifting is already done. Rather than having to search through tons of raw documents, Claude can work with a curated set of relevant, source-backed passages where:

- Ingestion keeps content current across every system and format
- Multi-layer indexing captures meaning, not just keywords
- Retrieval strategies surface exactly what a given question needs

Claude still does the reasoning. Agentic RAG just makes sure it has the right and most relevant content to work with.

MCP: The Connection Between Claude's Reasoning and Your Context

Ingestion and indexing solve half the problem: your business' content is now organized and retrievable. The other half is getting it in front of Claude at the moment a question is asked, without a custom integration needed for every new system or use case.

MCP is Anthropic's open standard for connecting AI systems to the places data lives. Instead of building a one-off pipeline between Claude and every tool you use, MCP gives you a single connection pattern.

Agentic RAG exposes its retrieval layer as an MCP server, and Claude, as an MCP client, calls into it mid-conversation. When you ask a question, Claude doesn't just answer from what's in the chat window; it queries Agentic RAG, gets back the relevant passages and reasons over those alongside everything else in context.

As a result, the following happens:

- **Answer quality improves:** Claude is reasoning over retrieved, governed content instead of guessing at what it doesn't know.
- **Auditability is guaranteed:** Agentic RAG technology logs every retrieval and preserves a clear audit trail so you can trace every answer back to its source documents.
- **Tokens stop paying for re-explaining yourself:** You're not re-uploading the same spec sheet or pasting the same background into every new conversation—MCP fetches it fresh each time, so your tokens go toward the actual question, not restating context Claude should already be able to reach.

Because MCP is a standard rather than a custom build, the same connection works whether you're asking Claude a one-off question or wiring retrieval into an agent workflow. Instead of rebuilding the foundation for every new Claude workflow, you build it once and reuse it.

Building Custom AI Experiences on Agentic RAG, with Claude Generating the Answers

So far, we've described what happens inside a single Claude conversation. But the process doesn't have to be confined to a chat window. When your content is ingested, indexed and exposed through MCP, it becomes a shared knowledge foundation that powers every AI experience you build—not just your conversations with Claude.


For example, your support team builds a customer-facing assistant, while a research group is building a copilot over internal reports. Someone else stands up a generative search experience on the company site. Each of these is a different application with a different interface and audience, but none of them re-does the ingestion, indexing or retrieval work. They all call into the same governed layer, and Claude does the reasoning and generation for whichever experience is asking.

That reuse makes all the difference. The alternative, standing up a separate retrieval pipeline for every new AI idea, is exactly the “too-many-RAG-projects problem” that stalls roadmaps after the first use case ships. A modular architecture with separate ingestion, retrieval, tuning and evaluation enables you to configure a new use case against the existing knowledge layer instead of building one from scratch.

What this means in practice:

- **One layer, many experiences:** A sales assistant today, a support agent next quarter, a research copilot after that—all built on content that's ingested and indexed once
- **Claude stays the generation engine everywhere:** Whatever the experience, the reasoning and writing is still Claude's job; retrieval just changes what it's reasoning over
- **No rebuild per use case:** Every new AI application reuses the same connectors, indexing, retrieval and tuning instead of starting from scratch

The result is an AI roadmap that compounds instead of restarting so every new experience you build makes the next one faster and easier to ship.



REMi: The Answer-Quality Layer Claude Doesn't Have

Claude is confident by design—it's built to give you a clear, well-reasoned answer, not to hedge or flag its own uncertainty by default. In a single conversation, that's useful; across an entire organization running hundreds or thousands of queries a day, it's a blind spot.

Claude has no built-in way to tell you whether last Tuesday's answers were grounded in your documents, if quality is drifting on a particular topic or which responses a compliance auditor should review first. Each conversation stands alone, with no shared knowledge accumulated across sessions.

That's not something a reasoning model can fix from the inside. Checking your own homework isn't the same skill as doing the homework. It requires a separate system, watching every answer, scoring it against the same standard and rolling that up into something a team can act on.


REMi is a proprietary, fine-tuned evaluation model that scores every single answer against three metrics:

1. **Groundedness:** Is the answer supported by retrieved content?
2. **Context relevance:** Was the right content retrieved?
3. **Answer relevance:** Does the response address the question asked?

REMi runs automatically on every query, turning the question, "Was this answer good?" from a one-time judgment into a continuously measured, organization-wide signal. And Claude by itself can't do that.

What this results in:

- **No more sampling and hoping:** Instead of a reviewer spot-checking 20 transcripts a week, every answer is scored so quality problems don't hide in the 99% nobody reviewed.
- **Drift is visible:** Health dashboards surface when groundedness or relevance dips on a topic, source or team before it becomes a pattern of bad answers users just stop trusting.
- **Governance is based on evidence:** Audit logs tie every query and score back to a source, giving compliance and risk teams something they can review and rely on. Claude provides an answer to your question; REMi verifies it.





Getting More from Your Claude Tokens Through Smarter Retrieval

Every token you send to Claude costs something, and when you want a better answer you may be tempted to feed it more information, such as an entire document, email thread or every version of a product spec. While that works, it's an expensive way to provide Claude context, and it doesn't scale as agentic workflows multiply the number of calls per query.

Retrieval changes the economics because it's selective by design. Instead of stuffing a maximum-size context window with everything that might be relevant, Agentic RAG uses 30+ retrieval strategies to surface only the passages that actually answer the question.

In contrast, Claude spends its tokens reasoning over ten relevant paragraphs instead of parsing through ten thousand words of surrounding noise. It's cheaper per query and often more accurate, because feeding Claude less irrelevant content reduces the chance of Claude latching onto the wrong passage.

Retrieval also creates a second lever that has nothing to do with prompt size. Because retrieval and orchestration sit above any single model, not every query needs Claude's most capable (and most expensive) model. Routine lookups and straightforward questions with clear, well-retrieved context can route queries to smaller or cheaper models that perform just as well with targeted context. Claude can then reserve its full reasoning power for the queries that actually need it. That means:

- **Retrieval trims what you pay for.** You're sending Claude the needle, not the whole haystack, on every query.
- **Model routing trims what you pay for it with.** Not every question needs your most expensive model if retrieval already did the hard work of finding the right context.
- **Pipeline reuse compounds the savings.** The same ingested, indexed content serves every use case, so you're not paying to re-process the same documents for each new AI experience you build.

The result? Higher-quality answers and lower, more predictable costs. This makes it possible to expand AI from a proof of concept into a production capability that can be used across the enterprise.

Scaling from Personal Use to a Production Workflow

Everything up to this point has been about improving your Claude conversations with more relevant context, lower token cost and answers you can trust. But there's a big difference between using Claude well yourself and building a Claude-powered workflow an entire team can rely on without learning the setup behind it.

That's where most AI efforts stall: a workflow that succeeds for one person asking thoughtful questions doesn't automatically scale to serving customers, support regulated processes or run autonomously.

The good news is that everything described in this ebook is designed to scale. The knowledge layer, the MCP connection and REMi's scoring aren't personal workarounds; they're shared infrastructure. So whether one person or an entire support team asks a question, the same ingested content, the retrieval and trust scoring remain in place—only the users and the volume change.

Here's what needs to be in place:

- **Move from ad hoc to governed.** Access control, audit logs and compliance certifications (e.g., SOC 2, ISO 27001, GDPR) need to be in place before an answer reaches a customer or regulator, not added after something goes wrong.
- **Move from occasional checks to continuous ones.** REMi's automated scoring replaces the manual spot-checks that work fine for personal use, but can't scale to production volume.
- **Move from one workflow to a reusable foundation:** The connectors, indexing and retrieval tuning built for your first use case are the starting point for the next one, instead of a one-off that gets rebuilt from scratch.

None of this requires starting over. It requires recognizing that the layer beneath your daily Claude usage is already there to support your next use case.

Ready to see what your Claude conversations have been missing?



Start a free trial and connect Progress Agentic RAG to your own content, with no sales call required.

About Progress Software

Progress Software (Nasdaq: PRGS) empowers organizations to achieve transformational success in the face of disruptive change. Our software enables our customers to develop, deploy and manage responsible AI-powered applications and personalized digital experiences with agility and ease. Businesses of all sizes get a trusted provider in Progress, with the products, expertise and vision they need to turn AI disruption into a competitive advantage. Millions of developers and technologists at hundreds of thousands of organizations depend on Progress every day. Learn more at www.progress.com

© 2026 Progress Software Corporation and/or its subsidiaries or affiliates.
All rights reserved. Rev 2026/08 | 1216435018901201

 /progresssw
 /progresssw
 /progresssw
 /progress-software
 /progress_sw_