

FineBooks scores 14 open OCR models on 2,165 old book pages

By [Alexis Dufresne](#) · Published September 3, 2026 at 18:12 UTC · Updated September 3, 2026 at 18:15 UTC

Free AI alerts in your inbox

Get →

TL;DR

- rednote-hilab's dots.mocr (3B) leads the FineBooks BHL leaderboard at 97.6% reading accuracy across 2,165 expert-transcribed pages.
- Sub-2B models OvisOCR2 (0.9B, 96.9%) and PaddleOCR-VL-1.6 (1B, 96.1%) come next at under \$0.50 per thousand pages.
- Ground truth covers only Antiqua typefaces in English, French, German and Latin; Fraktur, handwriting and multi-column layouts are excluded.

EDITOR'S NOTE

Small open OCR models are now cheap and accurate enough to reprocess large public-domain corpora for LLM training, but silent modernization of archaic characters means libraries still cannot drop them into scholarly transcription workflows.

FineBooks, a collaboration between Hugging Face and EleutherAI, has published [a public leaderboard](#) that scores 14 open OCR models against 2,165 expert-transcribed pages from the Biodiversity Heritage Library, drawn from 18th- and 19th-century printed books in English, French, German and Latin. The top model, rednote-hilab's dots.mocr at 3B

your@email.com

Get them

50.170.

The write-up frames this as a training-data problem, not a library problem. "Public domain books are one of the largest sources of long-form text that can be freely used for training open AI models — but only if the text quality is good enough," Sebastian Majstorovic and [Daniel van Strien](#) write in the [accompanying post](#). Full runs across the ground truth cost between \$0.50 and \$3.00 on Hugging Face Jobs; the cheapest model, PaddleOCR-VL-1.6, comes in at \$0.34 per thousand pages. Van Strien is one of two researchers we track who circulated the leaderboard on launch.

The leaderboard reports two CER variants: a "diplomatic" score that counts silent modernization of archaic letterforms like the long s (ſ) as an error, and a "reading" score that tolerates it. That split is why the same page can look like a solved problem for AI training and an open one for scholarly transcription. The ground truth itself covers only Antiqua typefaces in four languages. Fraktur, non-Latin scripts, handwriting and multi-column layouts sit outside scope.

"In most cases, however, these new OCR models have been primarily trained on and optimized for modern documents," the authors note.

Get the next one before it breaks

AI Weekly tracks stories like this 3x/week. Breaking alerts only when they matter. 50,000+ AI professionals.

your@email.com

Subscribe

✓ Also get breaking AI alerts

A 15.9M-parameter specialist OCR model outperforms many VLMs hundreds of times larger. On a 2,165-page historical eval dataset, Kraken PP-OCRv6 ranks #4 on reading CER and #1 when long-s, ligatures and case are preserve...

[View on Bluesky →](#)



Daniel van Strien @danielvanstrien.bsky.social :

A 15.9M-parameter specialist OCR model outperforms many VLMs hundreds of times larger. On a 2,165-page historical eval dataset, Kraken PP-O... →

Originally reported by [huggingface.co](#)

[Read the original article →](#)

Original headline: BHL OCR Leaderboard - a Hugging Face Space by finebooks

Don't miss the next one

3x/week. Breaking AI alerts only when they matter. Free, no spam, opt out anytime.

your@email.com

Subscribe

✓ Also get breaking AI alerts