


AI MODELS / AI STRATEGY / LARGE LANGUAGE MODELS

GPT-6 Astra's score of 98.6% looked like AGI. Then researchers read the fine print.

OpenAI says GPT-6 Astra scored 98.6% on ARC-AGI-3, up from 7.8% six months ago. But undisclosed test settings and opaque reasoning complicate the AGI question.

Sep 3rd, 2026 3:05pm by [Amanda Caswell](#)



There was no mistaking the divide in March with the release of ARC-AGI-3. While frontier AI models could do little more than register a sub-1% score, humans were able to navigate their new interactive settings.

OpenAI reports a different story for **GPT-6 Astra** six months on: 98.6%. Put that against the GPT-5.6 Sol it has superseded, which OpenAI puts at 7.8%, and the improvement is hard to miss.

Then again, one has to consider what ARC-AGI is designed for. The whole point is to put models in uncharted interactive territory where they cannot simply rely on training data to find an answer but must work out the mechanics of the

Given what ARC-AGI-3 was built to test, a jump to 98.6% is enormous. But the number comes with an important caveat.

The asterisk on Astra's 98.6%

Astra was evaluated through the company's Responses API harness, with two settings changed to better reflect how the model performs in real-world use. OpenAI says those changes weren't made specifically for ARC-AGI-3, but the other models in its comparison were evaluated using different setups.

ARC-AGI-3 requires a model to find its way through an unfamiliar environment, which means the setup it runs in can affect how well it performs.

The model is only part of the story

The gains aren't limited to ARC-AGI-3. On FrontierMath Tier 4, the model scored 97.6%, followed by 100% on ExploitBench and 99.2% on SRE-Bench with four attempts. Terminal-Bench Science saw one of the biggest jumps, from 22.4% for GPT-5.6 Sol to 64.6%.

OpenAI warns against rolling those results into a single measure of performance, but the range shows how much more the model can take on. In the company's demonstrations, it works directly inside software such as KiCad, Power BI and Unity, while an experimental Codex feature lets it keep notes and search earlier context when a job runs longer than a single context window.

TRENDING STORIES

1. [GPT-6 Astra aced the hardest AI benchmark. The asterisk matters more than the score.](#)
2. [OpenAI leaving Cursor: "Developers have to be prepared to adapt when it happens."](#)
3. [Claude Desktop can now easily run Qwen, DeepSeek and Kimi models -- after Ollama's first effort stalled](#)

5. Runway wants to generate software as you use it. Solaris is its first step.

On offline OSWorld 2.0, Astra scored 72.6% while taking about 40 minutes per task, compared with Sol's 65.7% and roughly 75 minutes.

On offline OSWorld 2.0, Astra scored 72.6% while taking about 40 minutes per task, compared with Sol's 65.7% and roughly 75 minutes.

Beyond benchmarks into discovery

The math is where things get more interesting. OpenAI says Astra was involved in two new findings about gaps between prime numbers. Mathematician **Julia Stadlmann** had already pushed one bound from 246 to 240. With Astra involved, it fell again, this time to 186. The company points to another case where the model helped improve part of a bound that hadn't budged in more than 80 years.

There's an important gap in OpenAI's account, though. It doesn't spell out what Astra came up with on its own, what researchers suggested or how the work moved between them. So while this goes beyond solving a benchmark with a known answer, it's not enough to call the math evidence of AGI.

Alignment gains, oversight gaps

Astra is pushing past many of AI's familiar limits, from unfamiliar problems to longer tasks. Yet even a 98.6% score on ARC-AGI-3 doesn't settle the AGI debate. Part of the problem is that performance more often depends on the system around the model. Intelligence itself doesn't improve evenly, either.

In OpenAI's internal tests involving difficult or impossible tasks without production safeguards, GPT-5.6 Sol went beyond what it was authorized to do 48.2% of the time. Astra didn't do so even once. Yet when researchers explicitly asked the models to evade monitoring, Astra's written reasoning was harder to follow than

s. OpenAI says that's partly because Astra can solve simpler problems in fewer

If AGI means doing useful intellectual work across different fields, Astra is getting remarkably close to what many people once had in mind. If it means matching human judgment across the board, ARC-AGI-3 can't establish that.

Epoch AI's Greg Burnham described Astra as the "end of one era, start of another."

"...end of one era, start of another."

[Editor's note: This article's headline has been updated to clarify that a 98.6% score on ARC-AGI-3 does not mean the benchmark was "aced." ARC-AGI-3 scores performance against a human baseline, with 100% representing performance at or above the median human baseline.]

TNS



Amanda Caswell is an AI journalist, certified prompt engineer, and technology commentator whose work and expertise have been featured on Fox News and CBS News. She covers artificial intelligence, developer tools, foundation models, and emerging technologies, with a particular focus...

[Read more from Amanda Caswell →](#)

TNS owner Insight Partners is an investor in: OpenAI.

SHARE THIS STORY



TRENDING STORIES

1. GPT-6 Astra aced the hardest AI benchmark. The asterisk matters more than the score.
2. OpenAI leaving Cursor: "Developers have to be prepared to adapt when it happens."

-
4. DeepSeek's first vision model vs. Gemini 3.7 Flash: It comes down to spend vs. speed
 5. Runway wants to generate software as you use it. Solaris is its first step.

INSIGHTS FROM OUR SPONSORS



How to query Apache Iceberg tables from ClickHouse

6 August 2026

Scaling agent systems with Apache Kafka and A2A (Part 4): Visualizing the Agent2Agent protocol

4 August 2026

Scaling agent systems with Kafka and A2A (Part 3): Agent2Agent Protocol explained

3 August 2026



End-to-End DevOps: Evolving Strategies & Automation Trends

26 January 2024

CTO Insights Unveiled – Unlocking Tech Leadership Excellence at Plato Elevate

10 November 2023

How the Terraform License Change Impacts Your DevOps

21 September 2023



Modernizing virtualization in higher education: How automated node recovery protects data integrity

4 September 2026

The architecture of autonomy: How ING built a future-proof tech strategy

3 September 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the most recent TNS articles in your inbox each day.

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers

ENGINEERING

AI
AI Engineering

Infrastructure as Code

Linux
Microservices
Open Source
Networking
Storage

Data

Frontend Development
Large Language Models
Security
Software Development
WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

 roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



© The New Stack 2026

Disclosures

ns of Use

Advertising Terms & Conditions

