

Представляем Agentic Document Extraction 2-го поколения



Команда LandingAI
8 сентября 2026 года

Поделиться На : [in](#) [X](#)




Agentic Document Extraction *Second Generation*

Page

Block

Line

Word



A New Foundation for the Agentic Era

Команда LandingAI рада представить обзор всех инноваций, реализованных в **Agentic Document Extraction (2-е поколение)**. Этот пост поможет вам разобраться во всех доступных функциях и узнать, как начать их использовать.



дите за запуском



Переосмысление обработки документов в эпоху агентов

Извлечение документов с помощью агентов теперь **быстрее, точнее и эффективнее**, чем когда-либо прежде. С точки зрения клиента, это обеспечивает более высокую производительность при оптимизации затрат на каждый документ. Клиенты получают платформу с более удобными для агентов результатами и более подробными цитатами, которые доступны в больших масштабах. Это не просто обновление. Это фундаментальное переосмысление и перестройка того, как ИИ понимает документы и работает с ними в эпоху агентов.

В будущем Agentic Document Extraction будет работать на основе моделей **семейства DPT-3**, в которое входят DPT-3 Pro и DPT-3 Verity. Эти новые базовые модели созданы с нуля для того, чтобы читать документы так, как это делает внимательный человек, начиная с общего макета и структуры и заканчивая отдельными словами и ячейками таблиц. Модели DPT-3 Parse имеют полностью переработанный ответ API, который так нравится ИИ-агентам. Мы тщательно протестировали новый ответ API с помощью новых интерфейса командной строки ADE и навыка ADE Claude и обнаружили, что ИИ-агенты способны создавать очень сложные рабочие процессы для анализа документов и интерфейсные приложения, используя только подсказки на естественном языке.

Традиционная обработка документов была создана для эпохи до появления агентов. Технология Agentic Document Extraction от LandingAI создана для эпохи агентов.

Три темы и многое другое

За последний год вы рассказали нам, что именно вам нужно от платформы для обработки документов, чтобы справляться с самыми сложными задачами. Вы спросили, и мы прислушались, внедряя инновации для снижения затрат и создавая возможности для точной и масштабной прослеживаемости.

В этом выпуске есть **3 тематических сообщения**, которые помогут вам разобраться в новых возможностях коллекции. Кроме того, в **разделе Overflow** представлен длинный список точечных улучшений, которые не вписываются в общую тему. Ищите ссылки на более подробные ресурсы, где они доступны. Все, о чем здесь говорится, также описано в официальном [журнале изменений](#) и [руководстве по миграции](#).

Темы:

1. **Доступность при масштабировании:** значительное снижение цен означает, что теперь вы можете точно и экономично автоматизировать весь процесс обработки документов.
2. **Выходные данные, готовые для использования агентами:** мы рассматриваем ИИ-агента как основного пользователя технологии Agentic Document Extraction. Выходные данные API специально созданы для использования ИИ-агентами.
3. **Атомарное заземление:** модели DPT-3 обеспечивают гораздо более точное заземление выходных данных. DPT-3 Pro заземляет данные до уровня строк, а DPT-3 Verity — до уровня слов, с оценкой достоверности для каждого слова. Это открывает множество новых возможностей для использования и дальнейшего применения.

1. Доступность в больших масштабах при снижении цен на 25–80 %

Высокая точность анализа документов — это обязательное условие. После достижения точности на первый план обычно выходят два других показателя — стоимость и скорость. Agentic Document Extraction Gen2 создан для того, чтобы автоматизировать весь ваш документооборот. Мы добились этого с помощью нескольких рычагов ценообразования.

Благодаря всем этим инновациям мы прогнозируем снижение затрат на **25–80%** для клиентов, использующих смешанные рабочие нагрузки. По сравнению с DPT-2, парсинг с



помощью DPT-3 Verity на стандартном уровне требует на 80–90% меньше кредитов для обработки типичного делового документа, например счета.

1A. Переменное ценообразование в зависимости от количества возвращенных символов

В Parse реализована новая концепция ценообразования, основанная на сложности документа, измеряемой количеством выходных символов. При использовании моделей DPT-3 вы платите за количество возвращенных символов, а не за количество страниц в документе.

Страница с небольшим количеством контента стоит меньше, а страница с большим количеством контента — дороже. В каждом ответе Parse указывается количество символов в выводе и итоговая стоимость в кредитах, так что вы можете точно знать, сколько стоит каждая страница и почему. Пользователи DPT-2 помнят, что каждая страница стоила ровно 3 кредита.

1B. Сэкономьте деньги, подождав немного

Новые уровни обслуживания позволяют сэкономить 50 % на задачах Parse и Extract за счет ожидания. Уровень обслуживания определяет, как быстро вы получите результат выполнения каждой задачи. **Приоритет** предназначен для случаев, когда результат ожидает человек или агент. **Стандартный** уровень используется по умолчанию; он работает асинхронно и подходит для производственных конвейеров, где ответ может прийти в течение от нескольких минут до нескольких часов.

Относительная стоимость составляет 1,0x для уровня «Приоритет» и 0,5x для уровня «Стандартный».

Note that the [ADE Visual Playground](#) uses the Priority service tier since a person (you) is waiting on the result. Cut credit usage in half by switching to API-based usage and using the default, Standard tier.

1C. Select the least expensive, capable model

With this launch we are introducing DPT-3 Pro and DPT-3 Verity (currently in public preview). Both are parsing models, both use complexity-based pricing, and both offer the Priority and Standard service tiers. DPT-3 Verity charges roughly 40% of the credits that DPT-3 Pro charges.

- **DPT-3 Verity** (verity: truth, faithfulness to the source) transcribes digitally created documents exactly as written, deterministically, and returns a confidence score and a bounding box for every word. Use it for high-volume digitally generated files containing primarily text, tables, and simple form fields.
- **DPT-3 Pro** reads the layout of a page before it reads the words, detecting multiple block types from tables and figures to marginalia and signatures and returning them in reading order. Charts and figures are described in detail with generative capabilities. DPT-3 Pro handles scanned and handwritten pages, non-Latin scripts, and math formulas.

You can select which model handles your parsing work based on the complexity of the document. By the Fall of 2026, you'll also have a choice to use automated intelligent routing between models. This is cost optimization designed in, so you don't have to build it yourself.

Recap of what's new

Models

- **DPT-3 Pro**, a high-accuracy model capable of handling extreme document complexity.
 - **DPT-3 Verity**, a high-accuracy, lower-cost model for transcription of digital text and tables.
- [Public Preview]**

Service Tiers

- **Parse** offers two service tiers with Standard priced at 0.5x compared to Priority.
- **Extract** offers two service tiers with Standard priced at 0.5x compared to Priority.

Pricing philosophy

- Complexity-based pricing for Parse. Parse charges **variable credits** based on the number of characters returned.
- Users have control over which pages and which block types they want included in the output.

Why it matters

The economics often decide which workflows organizations choose to automate at all and how they distribute workloads. When we first launched Agentic Document Extraction we quickly won the difficult workloads, but many customers sent their simpler pages to other document AI services due to price. We simply didn't have compelling pricing for basic pages - which are the majority of pages in most industries.

With the introduction of two parsing models, service tiers, and complexity-based pricing, **price per page has come down by as much as 80%** and parsing is now possible at **less than one cent per page** with the Verity model on the Standard tier.

Go deeper into the details

- Understanding [pricing core concepts](#)
- Compare v1 and v2 [pricing on real documents and use cases](#)
- Understanding [service tiers](#)
- Complete [credit consumption documentation](#)

2. Agent-ready and agent-friendly outputs

An agent is only as reliable as the context it receives. Agentic Document Extraction Gen2 produces the kind of output agents need. The output is built for machines to consume, meaning that it is predictable and highly structured. The DPT-3 API response is genuinely new, rebuilt from the ground up with AI agents in mind.

For existing customers this means **faster delivery** of use cases across the enterprise with the **help of coding agents**. But it also means setting aside the time to make the switch to the new API response structure.

Your code from ADE Gen1 will not work with Gen2. Hopefully the pricing improvements above have already motivated you to [migrate](#). If you are not convinced yet, this section will introduce you to how we think about agent-ready outputs and why they matter.

2A. Logical, hierarchical structure

With this release we are **introducing blocks and retiring chunks**. A block is a logical grouping of related content on a page. It could be a table, a paragraph of text, or a bar chart together with its title, legend, and caption. Pages contain one or more blocks. The structure is **pages → blocks → tables → cells** or **pages → blocks → lines → words**.

Also new: every block carries a **stable ID** based on its type, and **Markdown span pointers** locate it on the page. The redesigned API response includes a new top-level object named `structure`. Within the `structure` you'll find `page`, `range { start, end }`, `box { xmin ... ymax }`, and `atomic_grounding`. AI agents are able to traverse this structure very easily.

2B. Highly standardized markdown outputs

The Markdown output from both Parse models is now much more highly structured and standardized. This means that customers won't have to plan for multiple edge cases and can expect a conforming response each time.

The most notable changes are for Figures, Attestations, Checkboxes, and Tables.

- Figures use standard image syntax with subtype labels ([Chart], [Diagram], [Photograph],
Навык АДЕ Клод учит вашего агента строить [Подробнее](#)
- Within Figures, the transcribed text is clearly generated from generated text, with all
- Attestations are structured as Signed or Not Signed. The verbose descriptions of the past are gone.
- Checkbox outputs are structured and standardized as [X] or [] , including checkboxes inside tables.
- Tables can be delivered as HTML, which preserves structure that plain Markdown flattens.

Recap of what's new

Response Structure

- Hierarchical parse layout model
- Stable block IDs and Markdown span pointers that link back to the source
- New v2 Parse API and response structure for DPT-3. Available via the [v2 Parse endpoints](#)
- New v2 Extract API and response structure. Available via the [v2 Extract endpoints](#)
- More compact output. Cleaner for agents to consume and, because Parse pricing tracks output characters, cheaper too.

Standardization

- Consistent **key-value pairs**, intact and predictable.
- Standardized Markdown for **figures**: standard image syntax with subtype labels ([Chart], [Diagram], [Photograph], [Illustration]).
- Standardized Markdown for **checkboxes**: one consistent representation, including checkboxes inside tables.
- Standardized Markdown for **attestations**: structured cleanly as Signed or Not Signed.

- Tables are delivered as HTML by default with option to select Markdown instead.

Generative vs Literal Outputs

- For figures, the literal transcription of the text on the figure is **kept separate** from the model's generated description of it. Customers can always tell which text came directly from the document.

Why it matters

Structured, grounded, predictable output is what lets an agent act on a document without a person re-reading it first. When every value has a stable ID and a pointer back to its place on the page, downstream systems can retrieve, reason, and cross-reference without guessing. And they behave consistently across runs.

Standardized Markdown reduces the chance an agent makes an error due to formatting differences. As a bonus, more signal in fewer characters also lowers costs.

Go deeper into the details

- Understanding the [Parse API response shape](#)
- Understanding the [structure of pages and blocks](#)
- Find [blocks by ID](#)
- Deep dive on [detecting signatures, stamps, and seals](#)

3. Atomic grounding, and what it unlocks

Trust and traceability are what make scale safe. A system can read a million pages, but if you can't check any single answer against its source, you can never use that system in a regulated workflow.

In Gen2, we are introducing the term **atomic grounding**, which is grounding at the smallest structural unit of a document. The Gen2 Parse models both return atomic grounding as part of their response.

The Gen2 Extract model returns atomic citations drawn from the atomic grounding. Atomic citations allow you to trace any extracted value back to a specific word on a specific page. As before, Extract V2 offers you an infinite, fully custom schema with output formatting options. Any schemas you built with Extract V1 will carry over.

All DPT-3 family Parse models offer fine-grained grounding:

- **Table-cell grounding** works exactly as it did in Gen1 but has become more detailed. Every cell in a table has an identifier that can be referenced. Now each table cell also has a bounding box location associated with it.
- **Line-level grounding** is returned by DPT-3 Pro. Every text line within a block has its own Markdown span pointers and bounding box coordinates. This applies to all block types.
- **Word-level grounding** with **word-level confidence scores** is returned by DPT-3 Verity. Every word has its own Markdown span pointers and bounding box coordinates. Every word has an associated confidence score.

Recap of what's new

- Atomic grounding output from both Parse models.
 - DPT-3 Pro at the line level. Page → block → line
 - DPT-3 Verity at the word level. Page → block → line → word
- Table-cell grounding now provides bounding box coordinates per table cell.
- Extract V2 returns fine-grained citations drawn from the fine-grained grounding of Parse V2.

Why it matters

Atomic citations are the foundation for a class of capabilities that couldn't be built on DPT-2. Any use case that requires traceability down to the word level now becomes possible. This includes:

- PII and PHI can be localized to the exact line, word, or table cell, then intentionally obscured or removed based on the location coordinates.
- The parse outputs of two documents can be compared and version differences can be localized and highlighted.
- Users can build UIs to rapidly do human reviews of extracted values and/or edit Markdown in place.

Atomic grounding from Parse is what makes each of these buildable. With this level of traceability, organizations can automate more use cases and operate them at higher volumes.

Go deeper into the details

- Understanding [atomic grounding](#)

- Understanding the [Parse Outputs](#)

Overflow: More in this release

As mentioned before, this is not an incremental update. There are many, many new capabilities to make you aware of that didn't fit neatly into the 3 themes.

Major changes and enhancements not previously covered:

- [ADE CLI](#) to parse documents and extract fields from the command line.
- Updated [ADE Claude Skill](#) available via the plugin Marketplace.
- New [ade-python Python client library](#) carrying the v2 APIs.
- New [ade-typescript TypeScript client library](#) carrying the v2 APIs.
- Improved [Parse Jobs API](#) for async parsing at scale. Submit large documents (up to 6,000 pages or 1 GB per PDF) and retrieve results when they're ready.
- New [Extract Jobs API](#) (to mirror the Parse Jobs API) for asynchronous extraction jobs. Available on the Priority and Standard tiers.
- Separate [v2 \(DPT-3\) documentation](#) not co-mingled with the previous documentation for v1 (DPT-2).
- More **Enterprise deployment options**. In addition to US and EU cloud, choose from your own VPC on AWS, Azure, or Google Cloud, Snowflake, or on-premises, including air-gapped environments.

Minor changes and enhancements not previously covered:

- **New Playground URL** at ade.landing.ai. Update your bookmarks.
- Ability to **configure Parse options in the Playground** to mirror the `API options` parameter.
- **Redesigned Usage dashboard**. Daily and monthly consumption charts plus per-job logging.
- Ability to set **credit usage alerts**.
- The Playground user interface offers **low-confidence highlighting**.
- **Improved form handling**, especially checkbox detection, including checkboxes inside tables.
- **Unified async job model**. Parse Jobs and Extract Jobs now share a single response envelope and job-id format, with the option to store job output for later retrieval.
- **Standardized error handling**. Consistent error responses in a `{"code", "message"}` shape.

- **Correct handling of rotated images.** EXIF orientation is now respected, so bounding boxes line up with the document as you see it - important for scans and phone photos.
- **Notifications of finished jobs** via registered webhooks.
- Ability to **restrict API access by IP Address** on the Enterprise plan.

For the full list of changes, see the [changelog](#).

A new foundation for document intelligence in the agentic era

ADE Gen2 is a new foundation for the agentic era: autonomous enough to do the reading at scale, auditable enough to trust with the results. Users are getting a new model, a new API response, pricing that makes the whole document mix economical, and grounding that ties every answer back to the page.

This is bigger than automating paperwork. As AI moves from answering questions to executing work, documents become the foundation, and agents can only make good decisions when the information they act on is accurate, grounded, and trustworthy.

ADE Gen2 is available today. We are excited to see what you build with it.

Get started free at ade.landing.ai.



Company

[About Us](#)

[Leadership](#)

[Careers](#)

[Partnerships](#)

Agentic APIs

[Parse](#)

[Extract](#)

[Build Schema](#)

[Classify](#)

Contact Us

Enterprise Sales

Section

Split

Resources

Docs

Blogs

Events

Onboarding

Support

Trust center

Pricing

Security & compliance

Github

Industries

Financial services

Insurance

Healthcare

Energy & Utilities

Legal

Logistics

2026 ©LandingAI. All rights reserved. Privacy Policy | Legal

LandingAI

