

Last Translation Benchmark ships MT hard cases with fail rules

By [Alexis Dufresne](#) · Published September 4, 2026 at 09:12 UTC · Updated September 4, 2026 at 09:45 UTC

Free AI alerts in your inbox

Get →

TL;DR

- Last Translation Benchmark ships peer-reviewed hard examples, each paired with handcrafted verification rules that flag concrete failure cases instead of assigning a metric score.
- LTBv1 covers submissions accepted before September 1st 2026, spans texts, images, audio and videos, and is designed to accept ongoing contributions.
- The authors argue MT is stuck: standard benchmarks are saturating, automatic metrics are gameable, and human evaluation lacks reproducibility and scalability.

The [Last Translation Benchmark](#), posted on arxiv, is a collection of examples "that break leading machine translation models," each paired with "handcrafted verification rules describing concrete failure cases." The dataset is designed to stay live and accept ongoing contributions. LTBv1, the current release, covers submissions accepted before September 1st 2026 and spans texts, images, audio and videos.

The paper is blunt about why it exists. The authors write that "standard benchmarks for machine translation are approaching saturation," and that automatic metrics are "unreliable, vulnerable to reward-hacking, and

your@email.com

Get them

Instead of assigning a score, each item defines what a failing translation would look like on that specific example, so a model either trips a rule or it does not.

Led by Vilém Zouhar and Niyati Bafna, the collaboration lists hundreds of contributing authors. Two researchers on our radar had shared the arxiv link within a day of its posting.

EDITOR'S NOTE

If per-example verification rules gain traction as an evaluation standard, translation vendors lose the cover that saturated automatic metrics offer and get graded on named failure modes instead.

Get the next one before it breaks

AI Weekly tracks stories like this 3x/week. Breaking alerts only when they matter. 50,000+ AI professionals.

your@email.com

Subscribe☒ Also get breaking AI alerts**SHARED ON BLUESKY BY 2 AI EXPERTS**

V **Vilém Zouhar** @zouharvi.bsky.social :

Machine translation is not solved and it will take a while for it to be done arxiv.org/abs/2609.04173

→

M **Marco** @mcognetta.bsky.social **AMPLIFIED**

@zouhar.bsky.social

Machine translation is not solved and it will take a while for it to be done
arxiv.org/abs/2609.04173

Originally reported by **arxiv.org**

[Read the original article →](#)

Original headline: *Last Translation Benchmark*

Don't miss the next one

3x/week. Breaking AI alerts only when they matter. Free, no spam, opt out anytime.

Subscribe

☒ Also get breaking AI alerts