


AI INFRASTRUCTURE / HARDWARE / OPEN SOURCE

Nvidia's \$12.9B Hugging Face deal has an open-source problem

Nvidia's reported \$12.9B Hugging Face acquisition could reshape how developers deploy open AI models — and which hardware gets the inside track.

Aug 27th, 2026 11:35am by [Amanda Caswell](#)



Nvidia has reportedly agreed to buy Hugging Face for \$12.9 billion, putting one of the biggest names in AI hardware in charge of a platform developers rely on to find and run open models.

The Information first reported the deal Wednesday, citing a person familiar with the agreement. Nvidia and Hugging Face had not publicly confirmed it as of publication.

Hugging Face doesn't push developers toward one chipmaker, which is what makes the acquisition interesting. Its Optimum libraries work with Nvidia's TensorRT-LLM also support hardware from AMD, Intel, and AWS. Projects such as Optimum

The company also plays a role in what happens after a developer chooses a model, including how easily they can get it running on the hardware they want to use. The problem for Nvidia might be this: It is buying a platform whose value depends on openness and hardware neutrality, but if the purchase means Nvidia hardware is favored, that value might diminish.

The company also plays a role in what happens after a developer chooses a model, including how easily they can get it running on the hardware they want to use.

Hugging Face already sits between the model and the chip

Hugging Face has expanded well beyond file hosting. With Inference Endpoints, developers can deploy a model from the Hub while Hugging Face handles the underlying infrastructure.

Those hosted deployments can run on AWS, Microsoft Azure or Google Cloud, but most of the GPU options Hugging Face lists are Nvidia chips, including the T4, L4 and A100. That gives developers a wider choice of hardware through Hugging Face's open-source libraries than through its hosted services.

If the deal goes through, Nvidia would own both sides of that experience.

TRENDING STORIES

1. Nvidia's \$12.9B Hugging Face deal has an open-source problem
2. When agents build, deploy, and maintain, persistence becomes the hard problem
3. Agentic AI has a latency problem that more compute won't solve
4. Cut GPU inference cold start from 8 minutes to less than a minute

Deployment defaults favor Nvidia

NIM (Nvidia Inference Microservices) already works with models hosted on Hugging Face. Developers can point NIM to an hf:// repository path and pull the model directly from the Hub.

Owning Hugging Face would give Nvidia more room to bring NIM and CUDA-optimized containers directly into the deployment experience. Nvidia hasn't announced plans to make NIM the default, and support for AMD and Intel could remain exactly where it is.

Owning Hugging Face would give Nvidia more room to bring NIM and CUDA-optimized containers directly into the deployment experience.

The bigger question is what happens over time. Nvidia could provide earlier support for new models on its own hardware or make deployment easier. At the same time, AMD, Intel, and AWS may have to reconsider how much engineering work they want to contribute to integrations maintained within a competitor-owned platform. Some of that work could eventually move elsewhere.

For developers, the difference may come down to which path requires less work. A competing chip doesn't have to disappear from Hugging Face to become less appealing if an Nvidia model deployment takes fewer steps. We've seen a similar fight over the layers between AI models and the developers using them, with **Cloudflare building more of that infrastructure itself**.

Open models counter custom chips

That tension also helps explain why Hugging Face could be worth considerably more to Nvidia than its revenue alone would suggest.

Nvidia has been expanding its own Nemotron family of open models while investing heavily across the AI ecosystem. At the same time, some of its biggest customers are working to reduce their dependence on Nvidia hardware.

the hyperscalers. Earlier this month, five European companies committed to purchasing AI compute built around non-Nvidia hardware that hasn't been manufactured yet — a sign that the appetite for alternative accelerators is strong enough to attract forward contracts.

OpenAI this week published results from its new Jalapeño accelerator, which showed 1.5 to 1.9 times more work per watt while cutting end-to-end latency by up to 3.6 times on large open-weight models — although the chip has not yet been deployed at anything approaching Nvidia's scale. A strong open-model ecosystem gives Nvidia a counterweight to that trend.

Open models are often expected to run in very different environments, and Hugging Face helps developers make that possible. A model found on the Hub might end up running on Nvidia hardware, an AMD GPU, or a cloud accelerator.

That flexibility is part of what Nvidia would be buying. Pushing Hugging Face too heavily toward its own hardware could make the platform less useful to developers who rely on it to work across different systems.

Interest in those models is also growing. Models from companies including DeepSeek, Moonshot AI and Z.ai have narrowed the gap with proprietary systems. At the same time, Hugging Face CEO Clément Delangue told *The Information* in June that the company had doubled its number of paying subscribers during the first six months of 2026. Delangue later said the company was "close to profitability."

The Information puts Hugging Face's annualized revenue at about \$150 million. Against a \$12.9 billion price tag, that's a multiple of roughly 86.

So Nvidia would be paying for much more than Hugging Face's current business. It would be buying a place developers already turn to when they want to work with open models, including models that don't have to run on Nvidia hardware.

It would be buying a place developers already turn to when they want to work with open models, including models that don't have to run on Nvidia hardware.

there. Libraries such as transformers and diffusers are open source, and models on the Hub remain subject to their own licenses. Openly licensed models can still be hosted elsewhere, while the underlying libraries can be forked.

Much harder to recreate is the community Hugging Face has built around them. Developers already know where to look for models and have built workflows around the Hub and its integrations.

TNS



Amanda Caswell is an AI journalist, certified prompt engineer, and technology commentator whose work and expertise have been featured on Fox News and CBS News. She covers artificial intelligence, developer tools, foundation models, and emerging technologies, with a particular focus...

[Read more from Amanda Caswell →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY



TRENDING STORIES

1. Nvidia's \$12.9B Hugging Face deal has an open-source problem
2. When agents build, deploy, and maintain, persistence becomes the hard problem
3. Agentic AI has a latency problem that more compute won't solve
4. Cut GPU inference cold start from 8 minutes to less than a minute
5. Developers are coding to a moving target, and nobody knows where AI lands next



Octopus Easy Mode - Runbook Webhooks

1 September 2026

Debunking the hotfix pipeline myth

27 August 2026

Inside Platform Engineering with Nigel Douglas

21 August 2026



Introducing YugabyteDB Resource Governance

2 September 2026

Why Yugabyte Joined the Agentic AI Foundation

20 August 2026

How PostgreSQL CDC Tools Work on YugabyteDB

17 August 2026



Raw WebSockets for AI streaming: when patching stops paying off

3 September 2026

Add resumable streaming and reliable tool calling to your OpenAI agent

2 September 2026

Abyly Pub/Sub now supports React Native push notifications

28 July 2026

Receive a free roundup of the most recent TNS articles in your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

OPERATIONS

DevOps
CI/CD

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

CHANNELS

Podcasts
Ebooks

Kubernetes

Observability

Operations

Platform Engineering

Newsletter

TNS RSS Feeds

THE NEW STACK

About / Contact

Sponsors

Advertise With Us

Contributions

 roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap

Backend Developer Roadmap

Devops Roadmap

FOLLOW TNS



© The New Stack 2026

Disclosures

Terms of Use

Advertising Terms & Conditions

Privacy Policy

Cookie Policy