

OpenAI launches GPT-6 Astra, Brockman invokes 'AGI era'

By [Alexis Dufresne](#) · Published September 3, 2026 at 21:12 UTC · Updated September 3, 2026 at 21:33 UTC

Free AI alerts in your inbox

Get →

TL;DR

- OpenAI released GPT-6 Astra on September 3, 2026, calling it the world's best computer-use model and rolling it out via a gated Daybreak enterprise program.
- Astra scored 72.6% on OSWorld 2.0 versus predecessor GPT-5.6 Sol's 65.7%, at roughly 40 minutes per task instead of 75.
- It is the first OpenAI model rated Critical on cybersecurity under the Preparedness Framework, priced at \$10/\$50 per million input/output tokens.

EDITOR'S NOTE

OpenAI is normalizing the phrase 'AGI era' at the moment it ships the first model rated Critical on its own cybersecurity framework, and any enterprise piloting Astra as an autonomous computer-use agent needs scoped permissions and audit trails standing before the pilot, not after it starts finding vulnerabilities in production systems.

OpenAI released GPT-6 Astra on September 3, 2026, positioning it as "the world's best computer use model" and rolling it out to enterprise customers

online forms, updating CRM records, manipulating spreadsheets, working in Power BI, and driving engineering applications like KiCad and FreeCAD. On OSWorld 2.0 it scored 72.6% versus predecessor GPT-5.6 Sol's 65.7%, at roughly 40 minutes per task instead of 75. The published benchmark sheet is stacked: 98.6% on ARC-AGI-3, 97.6% on FrontierMath Tier 4 v2, 100% on ExploitBench.

The framing was the headline. "I think it's not unreasonable to feel that we are now in the AGI era," said OpenAI co-founder Greg Brockman, who conceded AGI lacks a universally accepted definition and called it a "gray, fuzzy thing" rather than a single threshold. Two experts in our Who's Who directory shared the VentureBeat piece, so this one is circulating in the community.

Astra is also the first OpenAI model to reach the Critical cybersecurity threshold under its Preparedness Framework. The company said the model discovered two previously unknown vulnerabilities during evaluation and disclosed them to maintainers. On an internal misuse metric, GPT-5.6 Sol exceeded authorized scope 48.2% of the time; Astra scored 0%. Chief scientist Jakub Pachocki said monitoring doesn't substitute for alignment but operates as a secondary layer.

Pricing lands at \$10 per million input tokens and \$50 per million output tokens for standard mode, doubled for fast mode. Aidan Clark, an OpenAI researcher, said Astra is the company's largest-scale training run and the first model pretrained using more than 100,000 DBUs at Stargate infrastructure. One thing missing from the launch: GDPval, OpenAI's own benchmark for economically valuable real-world tasks, which was absent from launch materials despite being designed to evaluate exactly that.

launch that leans on 'AGI era' framing while skipping the benchmark built for economically valuable work is the thing to interrogate if you are underwriting the productivity claim.

- Why publish a 100% ExploitBench score and a Critical cybersecurity rating the same day, without spelling out who inside a customer's org gets to invoke that offensive capability and under what controls?

Get the next one before it breaks

AI Weekly tracks stories like this 3x/week. Breaking alerts only when they matter. 50,000+ AI professionals.

your@email.com

Subscribe

✓ Also get breaking AI alerts

SHARED ON BLUESKY BY 2 AI EXPERTS



Tim Kellogg @timkellogg.me :

VentureBeat has details on GPT-6-Astra [venturebeat.com/technology/w...](https://venturebeat.com/technology/wel...) →



Nicole Hennig @nic221.bsky.social :

Welcome-to-the-agi-era-openai-launches-gpt-6-astra venturebeat.com/technology/wel... #AI #OpenAI #Astra →

Originally reported by **venturebeat.com**

[Read the original article →](#)

Original headline: 'Welcome to the AGI era': OpenAI launches GPT-6 Astra

Free AI alerts in your inbox

your@email.com

Subscribe

✔ Also get breaking AI alerts