

АГЕНТЫ ИИ / ИНФРАСТРУКТУРА ИИ / ПРОТОКОЛ КОНТЕКСТА МОДЕЛИ (MODEL CONTEXT PROTOCOL, MCP)

«Принесите это в наш магазин»: как Workday предлагает размещать ИИ-агентов рядом с вашими самыми ценными данными

Технический директор Workday Гейб Монрой утверждает, что ограничения корпоративного ИИ должны быть встроены в механизм логического вывода, где данные отдела кадров и бухгалтерии не могут довольствоваться «достаточно хорошим» результатом.

28 июня 2026 года, 11:00 [Фредерик Лардинуа](#)




Фото Патрика Хендри на Unsplash

Workday, платформа для расчета заработной платы и управления кадровыми данными, уже некоторое время использует искусственный интеллект и чат-ботов, но если в других компаниях допускается небольшая погрешность, то

Гейб Монрой, технический директор Workday, говорит *The New Stack*, что корпоративный ИИ должен пройти проверку, прежде чем кто-то подпустит его к своим HR- и финансовым данным.

«Существует не так много систем, которые были бы более критичными — или менее снисходительными — к ошибкам, чем те, что работают с людьми и деньгами», — говорит он. Здесь неприемлемо, когда система «работает большую часть времени», — говорит Монрой в интервью.

На своей **конференции разработчиков DevCon** в начале июня компания Workday рассказала о своих планах по преодолению этого барьера. Компания представила **Agent-Ready Tools** — набор коннекторов, которые позволяют агентам взаимодействовать с платформой через протокол Model Context Protocol (MCP), а также агент для разработчиков, который позволяет создавать приложения и агентов для Workday на простом языке, и **Agent Passport**, который тестирует и проверяет агентов перед их запуском в эксплуатацию и продолжает отслеживать их работу. Первым партнером по аттестации стала компания Cisco.

Ограничения должны быть в механизме логического вывода

Монрой большую часть своей карьеры проработал в сфере инфраструктуры и разработки: в компаниях Deis, Microsoft, DigitalOcean, а в последнее время — в Google. В Google он занимался созданием инфраструктуры для крупных лабораторий искусственного интеллекта, чтобы выполнять логический вывод в больших масштабах. Для человека, который так много внимания уделял инфраструктуре, переход в Workday может показаться странным шагом, но Монрой утверждает, что на данном этапе безопасность больших языковых моделей является — или, по крайней мере, должна являться — частью базовой инфраструктуры для предприятий.

«Ставки выше... в мире людей и денег»

Мире Workday, в мире людей и денег ставки выше, и именно этим я был по-настоящему увлечен в Workday — и я действительно рассматриваю это как

искусственного интеллекта, чтобы они могли делать масштабные логические выводы.

“[Inferencing] involves prefill and decode, and a whole bunch of really technical machinery in place to stream tokens out to end users, but what is nowhere in that stack today is the concept of native LLM-level enforced guardrails — guardrails that are part of the core inference.”

TRENDING STORIES

1. [The Anthropic leader who built Claude Code says he ditched prompting — now he just writes loops.](#)
2. [Greptile, Cursor, and Devin agree that agents should run their code. What they run it against matters.](#)
3. [Anthropic gives @Claude a permanent seat in your Slack channels](#)
4. [Claude Code vs. Cursor vs. Codex vs. Antigravity — six months in](#)
5. ["Bring it to our shop": Workday's pitch for keeping AI agents close to your most valuable data](#)

In his view, making it safe for enterprises to operate inference at scale has to be done at the inference engine layer. Agent gateways and similar add-ons that wrap the model from the outside sit at the wrong layer, he argues.

For a system of record like Workday, the system must strictly enforce guardrails on who the user is, what their budget authority is, and where they sit in the org chart. Those constraints, Monroy says, can be wired deep into the inference rather than checked after the fact. “These are things that we have the opportunity to wire deeply into the inferences at a very, very low level, in a way that produces much safer outcomes,” he says.

Workday’s recent [acquisition of Pipedream](#) partially plays into this. With Pipedream, an agent can reach out to third-party systems outside Workday to, for example, pull a policy document from Google Drive, and the platform can then verify that this specific agent has all necessary access rights to do so.

course, and while many SaaS companies are currently **building out their own agent platforms and agent orchestration services**, seemingly all offering the same services, Monroy argues that orchestration should happen close to the data sources.



Workday CTO Gabe Monroy

“If you’re trying to run an agent interaction that is integrating with people and money, that orchestration loop should probably happen closer to Workday, ideally on the Workday,” he says. “I do think there is something to be said for inference engines having proximity to the underlying system, because when it comes to a low-level inference engine and runtime perspective, there are some differentiated things you can do at the inference engine level that are only possible due to proximity to.”

“I do think there is something to be said for inference engines having proximity to the underlying system...”

He likened it to car repair. If somebody hands you a toolbox, you may or may not be able to fix your car’s problem. “Maybe you’ll get the outcome, but I tell you what, if you really want your car repair done right, bring it to our shop. We got the hydraulic lift, we got the trained engineers who know how to do this stuff, and we got the tools — but our tools are the best in the industry.”

For the part of an agent’s orchestration loop that has to run next to HR and finance data, “you really should be running that in our shop.”

For more general-purpose workflows, though, he believes there is room for general-purpose platforms. Not everything needs to run on Workday, after all, and Workday does offer MCP servers that third-party tools can tap into to access its data and tools as well.

Some degree, every system-of-record vendor is making some version of the **proximity-and-context case** right now. Owning the context for AI agents to work

Meanwhile, there is also a group of other companies positioning themselves as neutral parties benefiting from the fact that these other vendors are now opening their platforms to third-party agents. It seems unlikely that, in the long run, enterprises will want to manage multiple agent platforms, but at this point, it is still far from clear in which direction this pendulum will swing.

Happy to lose the tooling fight

If proximity is where Workday thinks it can win, there is one area where the company doesn't want to compete: developer tools.

"We're not going to differentiate Workday on building better tools for developers," Monroy says. "We're going to differentiate on safety, on trust, on the inference engine." It's telling that the demos at DevCon ran on Claude Code, Cursor, OpenAI's Codex, and Google's Antigravity rather than on anything Workday-branded. "On the tooling front, I'm happy to have Claude Code and Codex and others," he says.

It's no secret that **developer loyalty to AI tools doesn't really exist** at this point. Developers definitely don't want a one-off tool from a single vendor, and Monroy notes that he isn't going to ask them to adopt one. Instead, Workday wants to own skills, which he calls "the underlying lingua franca that we're using to converge across the system."

Workday is happy to be the system of record for its vertical and to let others fight over the tools.

TNS



Before joining The New Stack as its senior editor for AI, Frederic was the enterprise editor at TechCrunch, where he covered everything from the rise of the cloud and the earliest days of Kubernetes to the advent of quantum computing....

[Read more from Frederic Lardinois →](#)

TNS owner Insight Partners is an investor in: OpenAI.

TRENDING STORIES

1. The Anthropic leader who built Claude Code says he ditched prompting — now he just writes loops.
2. Greptile, Cursor, and Devin agree that agents should run their code. What they run it against matters.
3. Anthropic gives @Claude a permanent seat in your Slack channels
4. Claude Code vs. Cursor vs. Codex vs. Antigravity — six months in
5. "Bring it to our shop": Workday's pitch for keeping AI agents close to your most valuable data

INSIGHTS FROM OUR SPONSORS



Model-agnostic AI is key to business continuity as models keep changing

26 June 2026

Announcing Advanced Security for SonarQube Cloud Team plan

24 June 2026

Now available: SonarQube plugin for Codex

24 June 2026



Agents hate friction: early thoughts on building for agents

29 June 2026

ClickHouse achieves AWS Cloud Operations Competency

26 June 2026



Connect Sakana AI (Fugu) with Web Search API

30 June 2026

SerpApi Weekly Changelog: June 22-28, 2026

29 June 2026

Trending Travel Destinations using Python & SerpApi

26 June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in
your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



© The New Stack 2026

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)