


ИИ / ОТКРЫТЫЙ ИСХОДНЫЙ КОД / БЕЗОПАСНОСТЬ

После запрета Fable 5 компания Anthropic и 19 организаций запускают проект по обеспечению безопасности с открытым исходным кодом

Проект Akrites, размещенный на платформе Linux Foundation, направлен на устранение уязвимостей в системах безопасности с открытым исходным кодом до того, как ими воспользуются злоумышленники. Среди учредителей проекта — Anthropic, Google, OpenAI и Microsoft.

26 июня 2026 года, 14:36 [Пол Соерс](#)



Появление передовых моделей искусственного интеллекта, способных сканировать крупные проекты с открытым исходным кодом и выявлять множество уязвимостей за один проход, дало защитникам беспрецедентный инструмент, но в то же время открыло доступ к нему и для злоумышленников.

работать сообща.

Результатом стал проект **Akrites**, запущенный в **четверг** Linux Foundation, который выступает в качестве координационного органа по обнаружению, устранению и раскрытию уязвимостей в критически важном программном обеспечении с открытым исходным кодом.

В число учредителей входят около 20 организаций, в том числе AWS, Anthropic, Google, Microsoft и ее дочерняя компания GitHub, OpenAI, Cisco, Red Hat, NVIDIA, Chainguard, Sonatype, Ericsson, Vodafone, Citi и JPMorgan Chase. Инициатива названа в честь **акритов** — солдат, охранявших самые отдаленные границы Византийской империи — места, которые были наиболее уязвимы, чаще всего подвергались нападениям и зависели от того, кто придет их защищать.

Запуск пришелся на нестабильный период в сфере безопасности ИИ. Еще в апреле **Anthropic выпустила Claude Mythos** в рамках проекта Glasswing, сделав свою самую мощную модель доступной для небольшой группы доверенных партнеров специально для защиты от киберугроз. Затем, в начале июня, **Anthropic выпустила Fable 5** и **Mythos 5** — первые общедоступные модели класса Mythos со встроенными средствами защиты от неправомерного использования. Три дня спустя правительство США приостановило использование обеих моделей после того, как исследователи нашли способ использовать их для помощи в кибератаках.

В частности, Anthropic является одним из основателей Akrites.

Разрозненность в сфере безопасности: проблема самостоятельности

Модель безопасности с открытым исходным кодом долгое время опиралась на свободную, децентрализованную сеть разработчиков, исследователей и организаций, которые выявляют проблемы и сообщают о них. Когда на поиск серьезной уязвимости уходили недели работы экспертов, у защитников было время, чтобы опередить злоумышленников, но искусственный интеллект устранил этот пробел.

-
2. A public Sentry key is all it takes to hijack Claude Code, Cursor, and Codex
 3. Demo: HashiCorp's Consul Service Intentions Improve Security
 4. Checkmk offers a permanent scanner for Log4j
 5. First Line of Defense: Developer Security Tools in the IDE

When multiple organizations independently scan the same widely used library and each files their own report, maintainers face a wall of duplicates, and the real, exploitable findings get buried in the noise. Worse, every additional party sitting on knowledge of an unpatched vulnerability increases the chances it leaks before a fix exists.

Varun Badhwar, CEO of software supply chain security company **Endor Labs** and a founding member of Akrites, says that AI tools have already surfaced thousands of validated open source vulnerabilities in recent months, with fewer than 5% patched — a figure from his company's own data that hasn't been independently verified. The hard part, he says, was never the discovery itself.

"For years we have believed finding vulnerabilities was never the hard part. Fixing them was. AI has made that gap impossible to ignore."

"For years, we have believed finding vulnerabilities was never the hard part. Fixing them was," Badhwar says in a statement. "AI has made that gap impossible to ignore."

The existing model — each organization working separately, filing its own reports — is itself the problem Akrites is designed to fix. **Jason Clinton**, deputy chief information security officer at Anthropic, argues the model has simply been left behind.

by how quickly AI can now find vulnerabilities.”

“Open source projects collectively underpin much of the internet, and the existing model for coordinated disclosure has been outpaced by how quickly AI can now find vulnerabilities,” Clinton says. “Getting ahead of that requires the industry to coordinate on findings and get fixes upstream before they’re disclosed and exploited.”

Patch first, publish second

The core of the Akrites initiative is a shared Security Incident Response Team (SIRT) that acts as a single point of coordination for the industry. Rather than maintainers receiving a dozen separate reports about the same flaw from a dozen different organizations, the SIRT consolidates findings, validates which are genuine and exploitable, and manages a single coordinated fix and disclosure process. It uses established industry standards — CVE, CVSS, among others — and operates under strict confidentiality rules from the moment a finding comes in.

When a patch is ready, it goes back into the original project on the maintainer’s terms. For projects with no active maintainer, Akrites will step in as a fallback so a fix can still reach everyone who depends on the code.

JPMorgan Chase CISO **Pat Opet** explains the underlying logic: success should be measured by fixes reaching live systems, rather than by patches being published.

“We owe maintainers a single, reliable signal: confirmed vulnerabilities, well-tested proposed fixes, and a predictable partner they can trust.”

“AI has massively compressed the time between vulnerability discovery and exploitation to near real-time, which means we have to compress the time from fix to deployment,” Opet says. “We owe maintainers a single, reliable signal: confirmed vulnerabilities, well-tested proposed fixes, and a predictable partner they can trust, rather than a flood of duplicative, conflicting reports.”

infrastructure operators and the vendors they depend on; General, for organizations that want to contribute without committing large engineering resources; and Associate, for open-source foundations and projects at no cost.

Seed funding comes from **Alpha-Omega**, an Open Source Security Foundation (**OpenSSF**) project **under the Linux Foundation**, backed by Anthropic, AWS, Google, Microsoft, OpenAI, and others, with an annual budget of over \$7 million. Microsoft's Azure CTO **Mark Russinovich** pointed to Alpha-Omega as proof of what coordinated industry action can achieve.

"OpenSSF and Alpha-Omega demonstrated what is possible when industry comes together to strengthen open source security," Russinovich says. "Building on our experience co-founding these organizations, Akrites was created to address the emerging inflection point of AI-powered vulnerability discovery and defense."

TNS



Paul is an experienced technology journalist covering some of the biggest stories from Europe and beyond, most recently at TechCrunch where he covered startups, enterprise, Big Tech, infrastructure, open source, AI, regulation, and more. Based in London, these days Paul...

[Read more from Paul Sawers →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY



TRENDING STORIES

1. After Fable 5 ban, Anthropic and 19 organizations launch open source security body
2. A public Sentry key is all it takes to hijack Claude Code, Cursor, and Codex
- ↻ Demo: HashiCorp's Consul Service Intentions Improve Security
- .. Checkmk offers a permanent scanner for Log4j

INSIGHTS FROM OUR SPONSORS



When Amazon Aurora Stops Scaling: The Ceiling You Can't Buy Your Way Past

25 June 2026

The Fintech AI Platform: Why Agents Need One Database, Not Five

23 June 2026

Database Branching for AI Agents: How TINE Solves the Schema Drift Problem

15 June 2026



Signadot

Comparing Kubernetes Development Tools in 2026: Telepresence, mirrord, Okteto, and Signadot

18 June 2026

Watch Claude Code Validate a Microservices Change in a Closed Loop

4 June 2026

Signadot Plans

20 May 2026



Building a Stateful IT Service Desk Agent with LangGraph on Amazon EKS

26 June 2026

Open Governance for MySQL: A Step Forward for the Community

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in
your inbox each day.

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Kubernetes
Microservices

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models

Storage

WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

 roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



© The New Stack 2026

Disclosures
Terms of Use
Advertising Terms & Conditions
Privacy Policy
Cookie Policy

