

АГЕНТЫ ИИ / РАЗРАБОТКА ИИ / МОДЕЛИ ИИ

Искусственный интеллект Alibaba писал код 16 дней подряд, и каждый коммит публиковался на GitHub

Qwen3.8-Max от Alibaba потратил 16 дней на самостоятельную разработку инструмента командной строки — 265 коммитов, 127 пул-реквестов, без помощи человека. Журнал аудита GitHub находится в открытом доступе.

3 августа 2026 года, 15:23 [Аманда Касвелл](#)




Фото: Kvistholt Photography, Unsplash

Alibaba запустила Qwen3.8-Max, мультимодальную модель с 2,4 триллионами параметров, предназначенную для решения сложных задач, на выполнение которых может уйти несколько дней. Модель доступна через QwenCloud и Alibaba Cloud Model Studio по цене 2 доллара за миллион входных токенов и 6 долларов за миллион выходных токенов.

модель содержит 2,4 триллиона параметров, для каждого токена активируется примерно 95 миллиардов.

Архитектура «смесь экспертов» задействует те части модели, которые необходимы для решения каждой конкретной задачи, вместо того чтобы каждый раз использовать все 2,4 триллиона параметров. По словам представителей компании, на следующей неделе веса модели будут опубликованы на Hugging Face и ModelScope, что сделает ее первой моделью Qwen-Max с возможностью загрузки весов.

Архитектура модели, основанная на сочетании экспертных знаний, задействует те части модели, которые необходимы для выполнения каждой задачи, но разработчикам все равно приходится иметь дело с огромной моделью.

Как работает Qwen3.8-Max

Даже при том, что одновременно активны только 95 миллиардов параметров, все веса должны храниться и распределяться между несколькими узлами с графическими процессорами и большим объемом памяти. На практике это делает самостоятельный хостинг недоступным для большинства разработчиков и делает Qwen3.8-Max более реалистичным вариантом для крупных организаций и провайдеров логического вывода, у которых есть необходимая инфраструктура.

Qwen3.8-Max — реалистичный вариант для крупных организаций и поставщиков логического вывода, у которых есть необходимая инфраструктура.

Таким образом, Alibaba приблизилась к новому Kimi K3 от Moonshot AI, у которого **2,8 триллиона параметров**. Alibaba владеет **36% акций Moonshot AI**, так что это не совсем независимые конкуренты. Kimi K3 возглавил рейтинг интент-разработчиков на Arena уже через несколько часов после своего

подписок. При выпуске чипа с открытым исходным кодом возникла именно такая проблема: спрос превысил возможности графических процессоров в течение 48 часов, а параметры модели в 2,8 терафлопс сделали самостоятельный хостинг непрактичным для всех, кроме самых крупных магазинов. Qwen3.8-Max теперь участвует в той же гонке.

Другие долгоиграющие демонстрации

Согласно тестам Alibaba, Qwen3.8-Max потратил **16 дней на создание приложения с интерфейсом командной строки**, пять дней на воспроизведение результатов научной работы и улучшил дизайн чипа после примерно **500 итераций**. Alibaba не предоставила достаточно информации, чтобы сторонние исследователи могли проверить эти результаты.

Alibaba's most ambitious software demo had Qwen3.8-Max spend 16 days building "oh-my-cli" on its own. During that time, it made **265 commits and opened 127 pull requests** as it worked through **151 GitHub issues**. The repository and its audit trail are public under **qwen-code-dev-bot/oh-my-cli**, so developers do not have to rely solely on Alibaba's account of what happened.

TRENDING STORIES

1. [Opus 5 vs. Fable 5: What does half the price buy?](#)
2. [Alibaba's AI coded for 16 days straight and every commit is on GitHub](#)
3. [DeepSeek's smaller model just outperformed its own flagship](#)
4. [OpenAI fixed GPT-5.6 Sol's most frustrating flaw: Burning limits while it waits](#)
5. [Grok 4.5 vs. Claude Opus 4.8: Costs and what works, not the spec sheet](#)

The difference between the model and its harness is key for long-running jobs.

These parts can decide if an agent runs for 10 days or runs after just 10 minutes.

That infrastructure layer is what **Anthropic acquired the CI/CD startup Mendral to build out**, and it's the same lesson **GoDaddy learned when it opened its registrar to AI agents**; the guardrails had to come first.

Alibaba says it trained Qwen3.8-Max on complicated projects that an agent might actually encounter. The model performed similarly when tested with QwenWork, Claude Code, Codex, OpenClaw and Hermes.

The company has **published guides** showing developers how to connect it to those tools. Developers should soon be able to put that claim to the test. Model Studio works with the OpenAI and Anthropic API formats, so teams can bring Qwen3.8-Max into coding tools they already use without rebuilding their setup from scratch.

Context window limits

In one experiment, the model was given the paper “Unified Data Selection for LLM Reasoning” without any code. Alibaba says the agent wrote about 7,600 lines of code, started 33 GPU training jobs, and reproduced the paper’s six main results. It also tested 18 more ideas and improved the reported AIME24 result by 2.7 points.

In another test, Qwen3.8-Max optimized a cryptographic circuit over about 500 iterations. Alibaba says the model cut the design from **8,298 logic gates to 678**. After running it through OpenROAD, an open-source chip design tool, the physical area dropped by 81%.

This was not a one-shot coding test. The agent had to inspect each result and use it to decide what to try next. But even a million-token context window can only take an agent so far. Developers will want to see whether Qwen3.8-Max can remember earlier decisions and instructions as a project grows.

Sending nearly one million tokens at every step would cause other problems — and as **the economics of agentic AI have shown**, cheaper per-token pricing does not automatically translate to cheaper projects when autonomous agents are consuming tokens at rates that dwarf human-supervised sessions.

far. Developers will want to see whether Qwen3.8-Max can remember earlier decisions and instructions as a project grows.

Benchmarks need outside verification

Alibaba's internal benchmarks show Qwen3.8-Max performing close to models from Anthropic and OpenAI. According to Alibaba, it scored 86.6 on Terminal-Bench 2.1, compared to 88.8 for OpenAI's GPT-5.6 Sol (max). Independent runs of the same benchmark produce somewhat different scores for both models, so these figures may shift once outside labs repeat the tests. It also placed **fifth in Text Arena and second in Vision Arena**. In **Alibaba's own comparison table covering 31 tests**, the model leads in six; Claude Fable 5 or GPT-5.6 Sol comes out ahead in most of the rest.

These results make Qwen3.8-Max a strong contender, but they do not yet prove the model can maintain production software on its own for weeks. Once Alibaba releases the weights, developers can find out how well Qwen3.8-Max works outside the company's own setup. If its performance drops when it is moved to different infrastructure or connected to other coding tools, that will tell us far more than its place on a leaderboard.

TNS



Amanda Caswell is an AI journalist, certified prompt engineer, and technology commentator whose work and expertise have been featured on Fox News and CBS News. She covers artificial intelligence, developer tools, foundation models, and emerging technologies, with a particular focus...

[Read more from Amanda Caswell →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY



3. DeepSeek's smaller model just outperformed its own flagship
4. OpenAI fixed GPT-5.6 Sol's most frustrating flaw: Burning limits while it waits
5. Grok 4.5 vs. Claude Opus 4.8: Costs and what works, not the spec sheet

INSIGHTS FROM OUR SPONSORS



Octopus Deploy

**Connecting Octopus Cloud to your internal systems —
without opening a single port**

3 August 2026

Octopus Easy Mode - Claude Agent

31 July 2026

**How to promote a release from Development to
Production With Argo CD and Octopus Deploy**

30 July 2026



supermetal

Iceberg CDC Merge on Read for Every Query Engine

28 June 2026

Announcing MotherDuck and DuckDB Targets

17 June 2026

Announcing Apache Doris Target

15 June 2026

PagerDuty

**'Celebrating a Baton Pass at the NYSE's Closing Bell by
Dan DiLullo**

Operations: Key Takeaways From PagerDuty on Tour 2026 by PagerDuty

23 July 2026

The 13 Questions CEOs Ask After an Incident (And What IT Leaders Must Be Ready to Answer) by PagerDuty

22 July 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the most recent TNS articles in your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem

ENGINEERING

AI

Edge Computing

Infrastructure as Code

Linux

Microservices

Open Source

Networking

Storage

Backend development

Data

Frontend Development

Large Language Models

Security

Software Development

WebAssembly

OPERATIONS

AI Operations

CI/CD

Cloud Services

DevOps

Kubernetes

Observability

Operations

Platform Engineering

CHANNELS

Podcasts

Ebooks

Events

Webinars

Newsletter

TNS RSS Feeds

THE NEW STACK

About / Contact

Sponsors

Advertise With Us

Contributions



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap

Backend Developer Roadmap

Devops Roadmap

FOLLOW TNS



