📖
Вселенная

Прикладная наука о данных

Авторы: Мартин Брашлер, Тило Штадельманн, Курт Штокингер

Тип файла: pdf

Размер: 16,3 МБ

Язык: Английский

Страниц: 465

Прикладная наука о данных: извлеченные уроки для бизнеса, основанного на данных

Введение: от сбора данных к бизнес-ценности 📊🚀

Современные компании ежедневно генерируют огромные объемы информации. Транзакции клиентов, активность на веб-сайтах, инженерные измерения, данные о цепочке поставок, финансовые данные, взаимодействие в социальных сетях, датчики Интернета вещей и операционные журналы — все это потенциальные источники ценной информации.

Но **наличие данных не делает бизнес автоматически ориентированным на данные**.

Настоящая инженерная задача состоит в том, чтобы преобразовать необработанную информацию в достоверные данные, а затем на их основе принимать решения и добиваться измеримых бизнес-результатов. Принятие решений на основе данных означает, что при принятии решений используются факты, показатели и данные, а не только интуиция.


Dashboard
Analytics
20%
35 %
3%
25 %
2%
25 %
3%
45 %
4%
Decision Model
Decision
Option A
Result 1
Result 4
Result 4


Analytics
Home
Dashboard
Reporting
Customization
Help
Log Out
Audience
Country
City
Status
Sex
Age
Interests
Collaboration
Behavior
Technology
Mobile
Custom 1
Custom 2
Custom 3
Income
Education
September
Checklist
Monthly Sales
$132,953
Estimate
$242,451
Yearly Sales
$1,842M
Estimate
$2,742M
Debt-to-Equity
Profit Margin
21,53%
Retention Rate
83,73%
Sales By Product
New Customers
New Sales
KPI
65%
37%
22%
75%
32%

Прикладная наука о данных находится на стыке **статистики, информатики, математики, инженерии и бизнес-стратегии**. Ее цель — не просто создание сложных моделей машинного обучения. Успешная наука о данных связывает бизнес-проблему с данными, анализом, экспериментами, прогнозированием и действиями.

Полезный подход к рассмотрению процесса:

Бизнес-задача → Данные → Анализ → Модель → Инсайт → Решение → Результат → Обратная связь 🔄

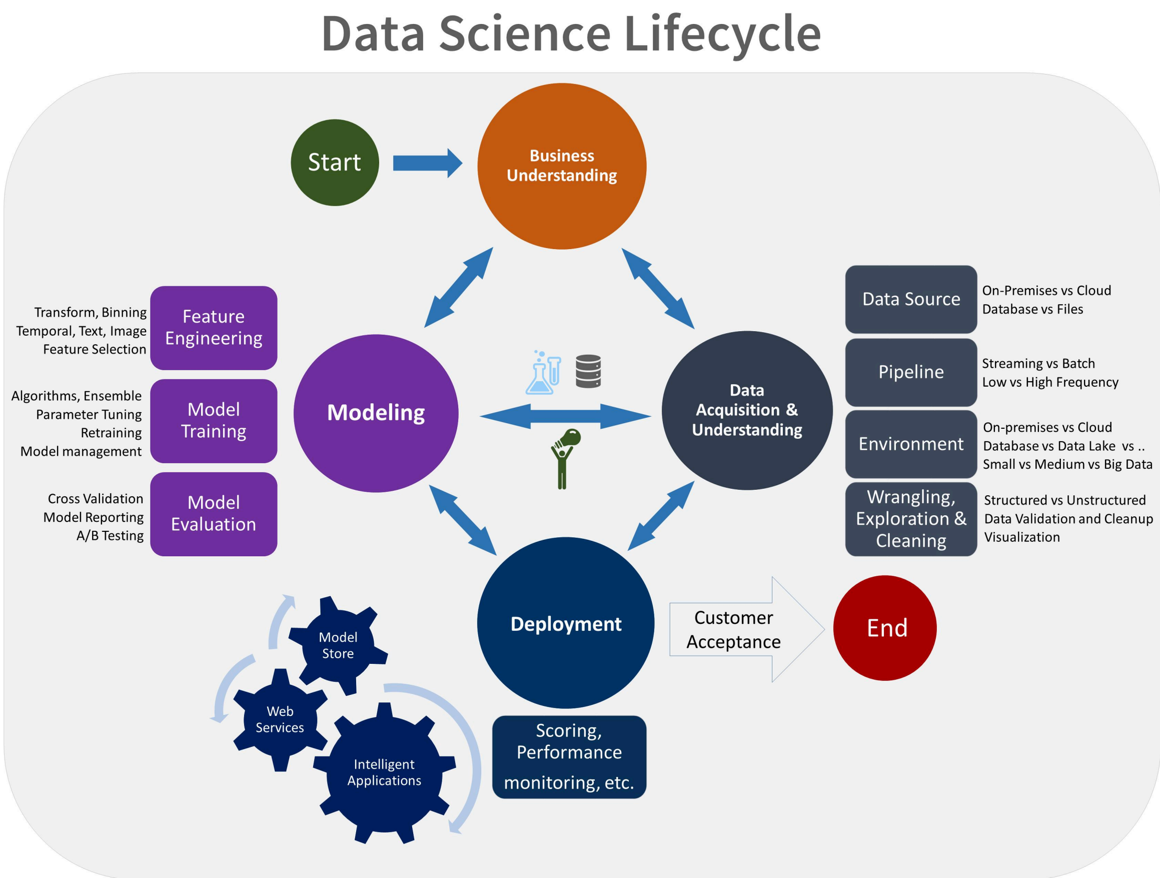

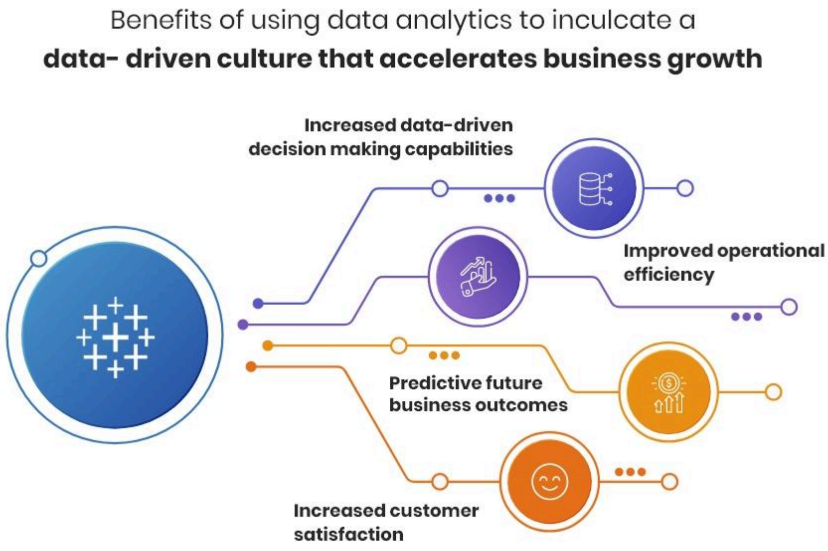
Benefits of using data analytics to inculcate a
data- driven culture that accelerates business growth
Increased data-driven decision making capabilities
Improved operational efficiency
Predictive future business outcomes
Increased customer satisfaction


Using Data To Drive Better Business Outcomes
Connect analytics to real, measurable results across your organization.
Make informed decisions
Measure real business results
Showcase tangible proof

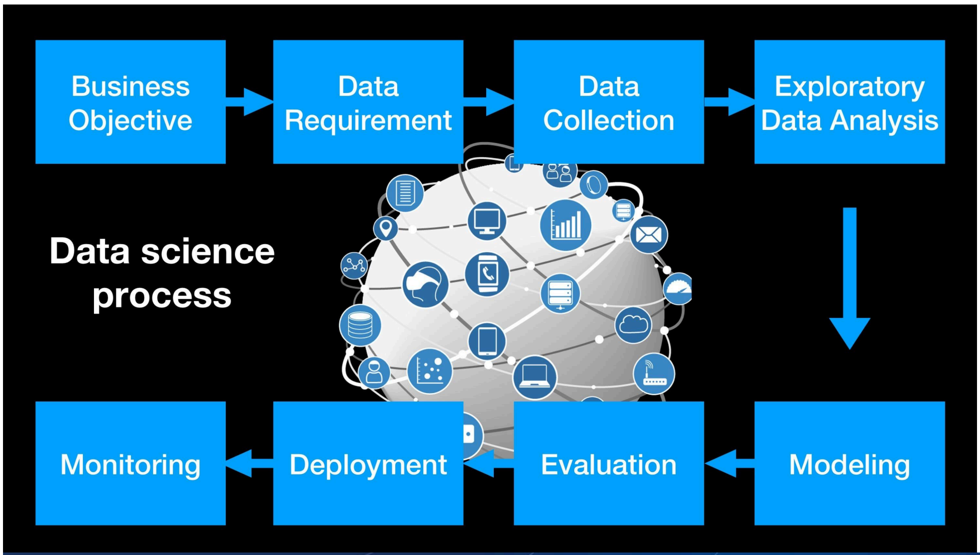

Таким образом, самые важные уроки можно извлечь не только из успешных моделей, но и из провальных проектов, некачественных данных, неправильно подобранных ключевых показателей эффективности, проблем с развертыванием и ситуаций, когда простой статистический метод работает лучше, чем продвинутая система искусственного интеллекта.

Теоретическая база

Что такое прикладная наука о данных?

Наука о данных сочетает в себе аналитические и вычислительные методы для извлечения полезной информации из структурированных и неструктурированных данных.

Упрощенное математическое представление:

[D \rightarrow P(D) \rightarrow M(D) \rightarrow I \rightarrow A]

где:

- (D) = raw data
- (P(D)) = processed data
- (M(D)) = analytical or machine-learning model
- (I) = insight
- (A) = business action

The final objective is not necessarily the model itself. The objective is **better action**.

For example, a retailer might use historical sales data to estimate:

[\hat{Y}_{t+1}=f(X_1,X_2,\ldots,X_n)]

where (\hat{Y}_{t+1}) represents predicted future demand and (X_1,\ldots,X_n) represent variables such as price, season, promotions, inventory, and location.

The prediction becomes valuable only when the business uses it to improve purchasing, inventory allocation, pricing, or logistics.

Descriptive, Predictive, and Prescriptive Analytics

Applied data science commonly operates across three analytical levels.

Analytics Type	Main Question	Example
Descriptive	What happened?	Sales fell 8% last month
Predictive	What may happen?	Demand may increase next month
Prescriptive	What should we do?	Increase inventory by 12%
Diagnostic	Why did it happen?	Sales declined because of stock-outs

A mature data-driven organization gradually moves from simply **reporting history** toward predicting future conditions and recommending appropriate actions.

The Data Science Lifecycle 🔄

A practical lifecycle is iterative rather than strictly linear:

Business Understanding → Data Collection → Cleaning → Exploration → Feature Engineering → Modeling → Evaluation → Deployment → Monitoring → Improvement

Research and industry practice increasingly emphasize the importance of connecting analytics with actual decision processes rather than treating dashboards or models as isolated technical products.

Definition

Applied Data Science in Business

Applied data science is the practical use of statistical analysis, programming, machine learning, data engineering, visualization, experimentation, and domain knowledge to solve measurable real-world problems.

A successful applied data-science project should answer five questions:

1. 🚀 **What business problem are we solving?**
2. 🚀 **What evidence do we have?**
3. **What prediction or insight can we produce?**
4. **What decision will change because of it?**
5. **How will we measure whether the decision worked?**

This leads to an important engineering principle:

> ***A technically accurate model can still be a business failure.***

For example, a classification model with 95% accuracy sounds impressive. However, if the minority class represents only 2% of customers, predicting “normal” for everyone could achieve approximately 98% accuracy while being completely useless for fraud detection.

Therefore:

[\text{Model Quality} \neq \text{Business Value}]

Instead:

[\text{Business Value}=f(\text{Model Quality},\text{Adoption},\text{Action},\text{Economic Impact})]

Step-by-Step Explanation: Building a Data-Driven Solution 🛠️

Step 1: Define the Business Question

Do not begin with:

"We have AI. What can we do with it?"

Begin with:

"What decision is difficult, expensive, slow, or uncertain?"

Examples include:

- Which customers are likely to leave?
- How much inventory should be ordered?
- 🚀 Which leads are most valuable?
- Which machine is likely to fail?
- Which marketing campaign generates profitable customers?

A precise business question creates a measurable analytical target.

Step 2: Identify the Data

Potential data sources include:

- SQL databases
- CRM systems
- ERP platforms
- Web analytics
- IoT sensors
- Financial systems
- Customer-support records
- APIs
- CSV/Excel files
- Cloud data platforms

The important lesson is that **more data is not necessarily better data**.

Step 3: Clean and Validate the Data 🧹

Typical problems include:

- Missing values
- Duplicate records
- Incorrect units
- Outliers

- Inconsistent dates
- Incorrect categories
- Measurement errors
- Data leakage

For numerical data, basic statistics such as mean and standard deviation can reveal unexpected behavior:

$$z=\frac{x-\mu}{\sigma}$$

A large absolute (z)-score may indicate an unusual observation, although it does not automatically mean the observation is wrong.

Step 4: Explore the Data

Exploratory Data Analysis (EDA) helps engineers understand relationships before modeling.

Useful techniques include:

- Histograms
- Box plots
- Scatter plots
- Correlation matrices
- Time-series analysis
- Grouped statistics
- Distribution analysis

For example, if customer churn appears higher among customers with low engagement, the relationship can be investigated before constructing a predictive model.

Step 5: Engineer Useful Features

Raw variables are often insufficient.

Suppose a company has:

- purchase date
- customer registration date
- order amount
- number of orders

Feature engineering can sometimes produce more practical improvements than changing algorithms.

Step 6: Build a Baseline

Before selecting an advanced algorithm, create a simple baseline.

For regression:

- Mean prediction
- Linear regression
- Decision tree

For classification:

- Majority-class prediction
- Logistic regression
- Simple decision tree

The baseline establishes a reference point.

If a complex model improves performance only marginally, the additional complexity may not be justified.

Step 7: Evaluate the Model

Model selection should match the business problem.

For regression:

$$MAE=\frac{1}{n}\sum_{i=1}^{n}|y_i-\hat{y}_i|$$

For classification, useful metrics include:

- Precision
- Recall
- F1-score
- ROC-AUC
- PR-AUC
- Accuracy

The correct metric depends on the consequences of errors.

For example, missing a fraudulent transaction may be much more expensive than incorrectly flagging a legitimate transaction.

Step 8: Deploy and Monitor

A model that exists only in a notebook is not yet a production solution.

Deployment may involve:

Database → Data Pipeline → Model → API/Application → Business User → Decision

After deployment, performance must be monitored because business conditions change.

A model trained on 2024 customer behavior may behave differently in 2026.

This phenomenon is often associated with **data drift** or **concept drift**.

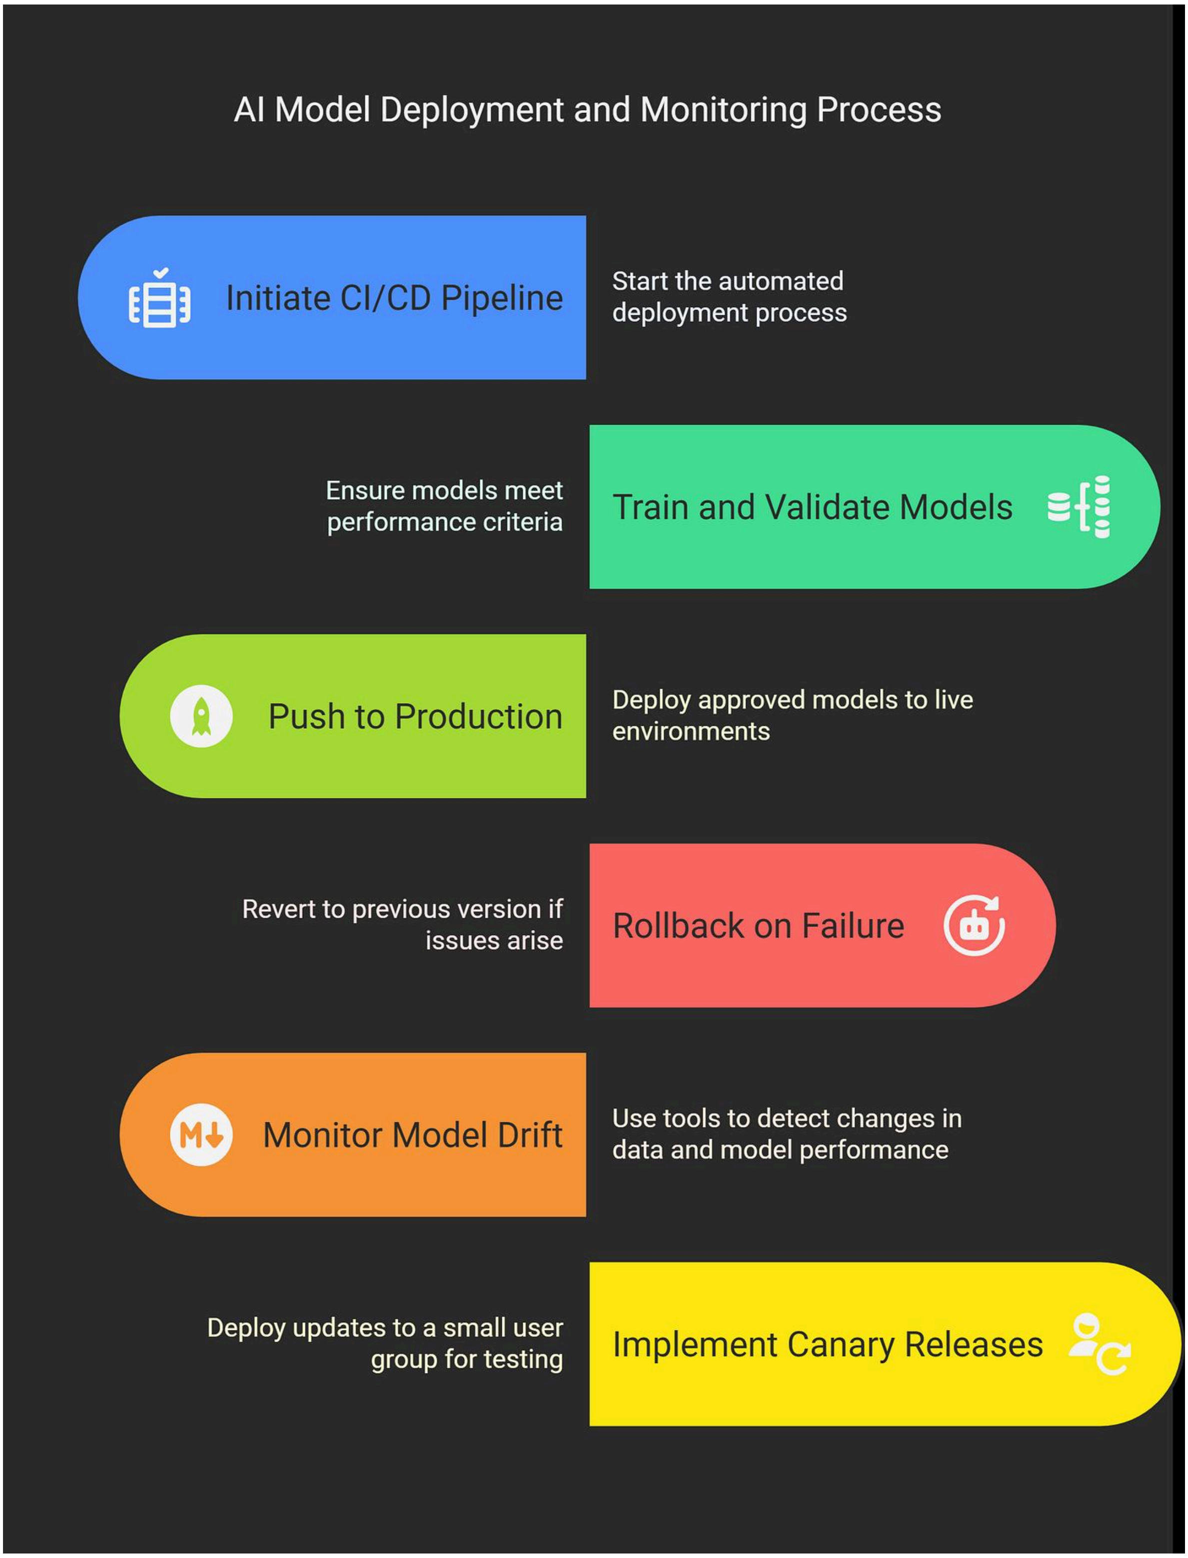
AI Model Deployment and Monitoring Process
Initiate CI/CD Pipeline
Start the automated deployment process
Train and Validate Models
Ensure models meet performance criteria
Push to Production
Deploy approved models to live environments
Rollback on Failure
Revert to previous version if issues arise
Monitor Model Drift
Use tools to detect changes in data and model performance
Implement Canary Releases
Deploy updates to a small user group for testing

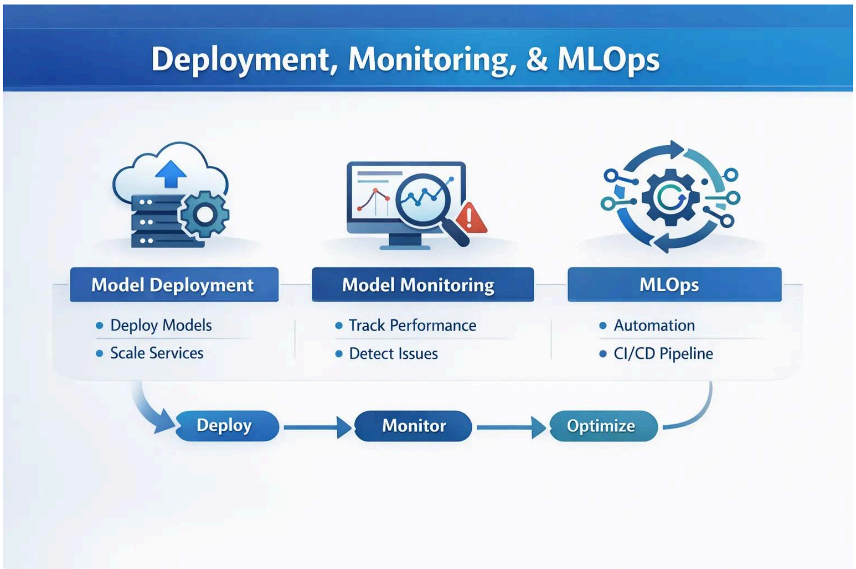

Deployment, Monitoring, & MLOps
Model Deployment
Deploy Models
Scale Services
Model Monitoring
Track Performance
Detect Issues
MLOps
Automation
CI/CD Pipeline
Deploy
Monitor
Optimize

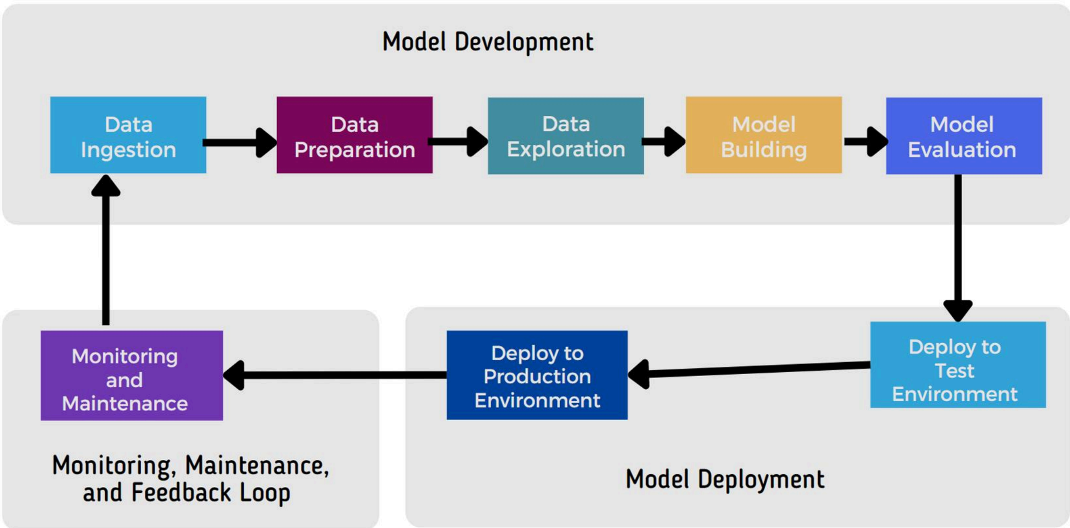

Model Development
Data Ingestion
Data Preparation
Data Exploration
Model Building
Model Evaluation
Deploy to Test Environment
Deploy to Production Environment
Monitoring and Maintenance
Monitoring, Maintenance, and Feedback Loop
Model Deployment

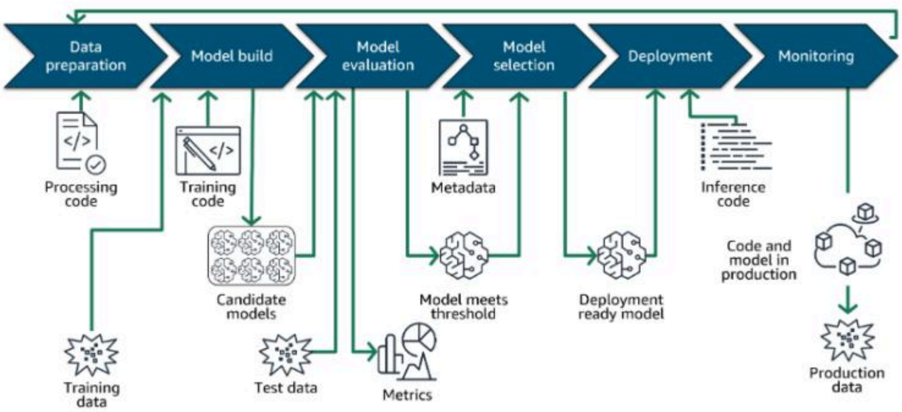

Data preparation
Model build
Model evaluation
Model selection
Deployment
Monitoring
Processing code
Training code
Metadata
Inference code
Candidate models
Model meets threshold
Deployment ready model
Code and model in production
Training data
Test data
Metrics
Production data

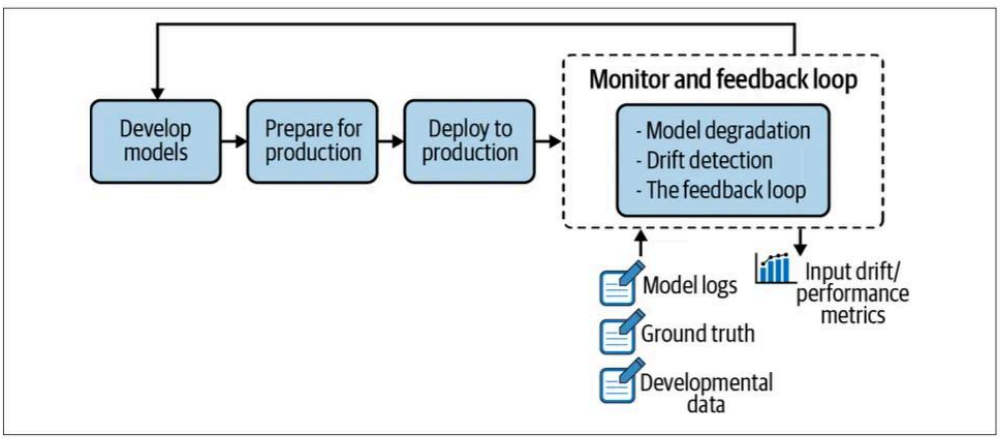

Develop models
Prepare for production
Deploy to production
Monitor and feedback loop
- Model degradation
- Drift detection
- The feedback loop
Model logs
Ground truth
Developmental data
Input drift/ performance metrics

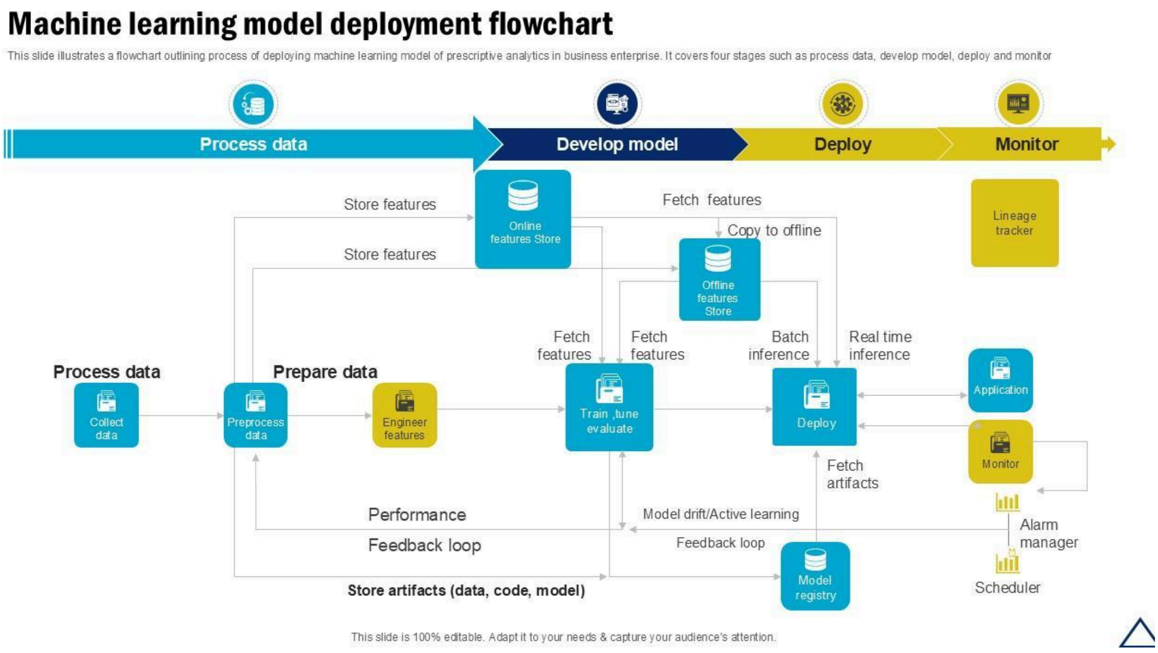

Comparison: Traditional Business Analysis vs Applied Data Science

Factor	Traditional Analysis	Applied Data Science
Primary focus	Historical performance	Historical + future behavior
Main tools	Reports, spreadsheets	Python, SQL, ML, BI
Typical question	What happened?	What will happen?
Automation	Moderate	High
Prediction	Limited	Core capability
Experimentation	Sometimes	Frequently
Data volume	Small–medium	Medium–massive
Decision support	Human interpretation	Human + analytical systems

However, applied data science should **not replace traditional analysis**.

The strongest organizations combine both.

A financial analyst may understand business constraints that a machine-learning model cannot see. Likewise, a data scientist can identify statistical patterns that would be difficult to discover manually.

Diagrams & Tables: The Data-to-Decision Pipeline 📈

Conceptual Architecture

```
┌────────────────────────┐
│ Business Problem     │
└───────────┬────────────┘
          ↓
┌────────────────────────┐
│ Data Sources         │
│ SQL / APIs / IoT     │
└───────────┬────────────┘
          ↓
┌────────────────────────┐
│ Data Engineering     │
│ ETL / ELT / QA       │
└───────────┬────────────┘
          ↓
┌────────────────────────┐
│ Analytics / ML       │
└───────────┬────────────┘
          ↓
┌────────────────────────┐
│ Insight / Forecast   │
└───────────┬────────────┘
          ↓
┌────────────────────────┐
│ Business Decision    │
└───────────┬────────────┘
          ↓
┌────────────────────────┐
│ Measured Outcome     │
└───────────┬────────────┘
          │
          └────────────→ Feedback → New Data
```

KPI Design

A good KPI should connect directly to a business objective.

Business Goal	KPI	Possible Data Science Use
Increase revenue	Revenue per customer	Customer segmentation
Reduce churn	Churn rate	Churn prediction
Reduce costs	Cost per transaction	Optimization
Improve operations	Downtime	Predictive maintenance
Improve marketing	Conversion rate	Propensity modeling
Improve inventory	Stock-out rate	Demand forecasting

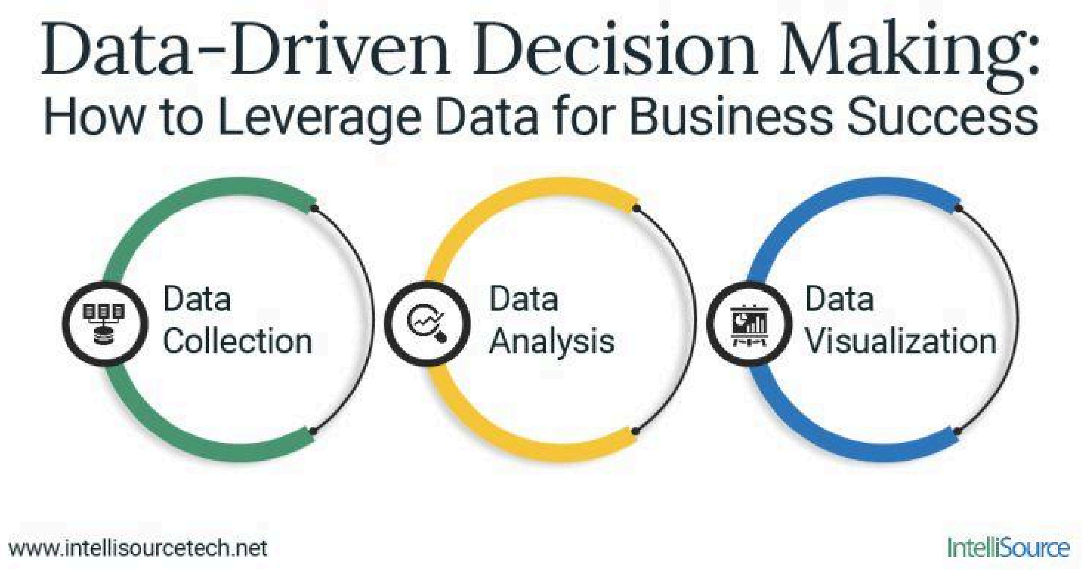
Data-Driven Decision Making:
How to Leverage Data for Business Success
Data Collection
Data Analysis
Data Visualization
www.intellisourcetech.net
IntelliSource


Monthly Sales
Retention
81%
69%
85%
65%
37%
Profit and Loss summary

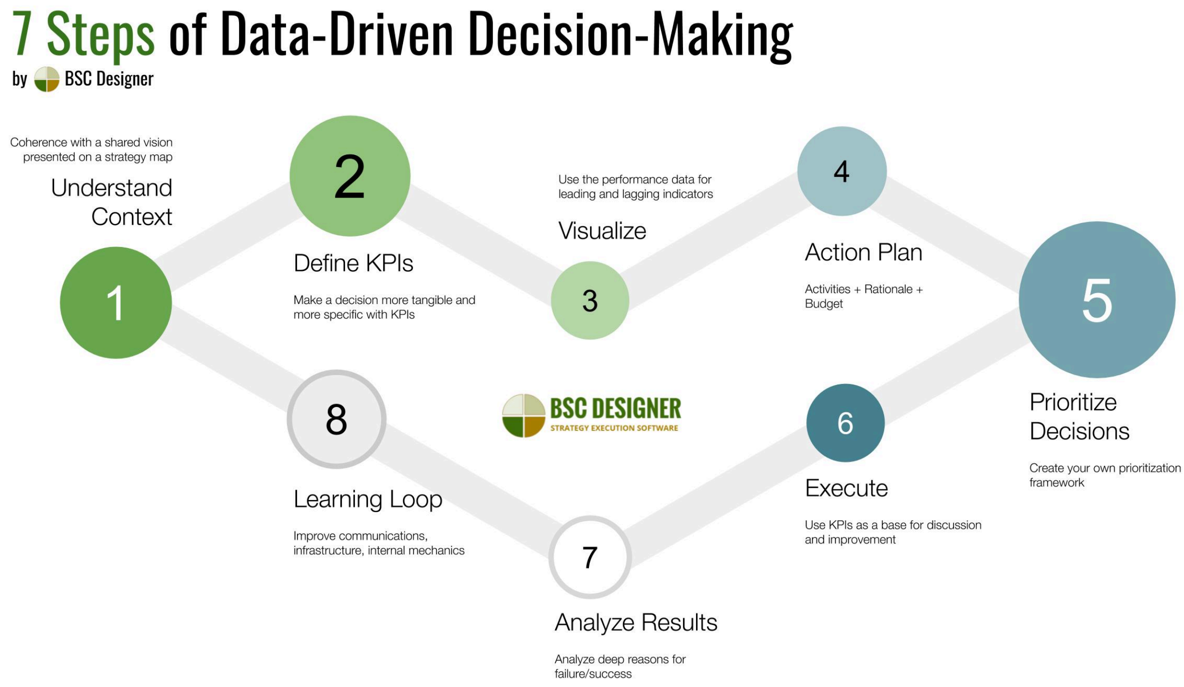

Examples

Example 1: Customer Churn

A telecommunications company wants to identify customers at high risk of cancellation.

Possible variables include:

- Contract duration
- Monthly expenditure
- Number of complaints
- Service usage
- Payment history
- Customer-support interactions

A model generates:

[P(\text{Churn}|X)]

A customer with:

[P(\text{Churn}|X)=0.82]

could be placed in a high-risk segment.

The business could then test retention strategies.

The important lesson is that the model should support an **intervention**, not merely produce a probability.

Example 2: Predictive Maintenance ⚙️

An engineering company monitors vibration, temperature, pressure, and operating hours.

Suppose vibration increases progressively:

$$V(t)=V_0+kt$$

If the system detects a pattern associated with previous failures, maintenance can be scheduled before catastrophic downtime occurs.

This can potentially reduce:

- Unplanned downtime
- Emergency maintenance
- Production losses
- Safety risks
- Spare-parts delays

Example 3: Demand Forecasting

A retailer can estimate future demand using historical sales, seasonality, promotions, prices, holidays, and regional behavior.

A simple model might represent demand as:

$$D_t=\beta_0+\beta_1P_t+\beta_2S_t+\beta_3A_t+\epsilon_t$$

where:

- P_t = price
- S_t = seasonal effect
- A_t = advertising activity
- ϵ_t = unexplained variation

The resulting forecast can support inventory planning.

Real-World Applications 🌍

Financial Services

Data science can support:

- Fraud detection
- Credit-risk analysis
- Customer segmentation
- Financial forecasting
- Transaction monitoring

Manufacturing

Applications include:

- Predictive maintenance
- Quality control
- Production optimization
- Energy monitoring
- Supply-chain forecasting

Healthcare

Potential applications include:

- Resource forecasting
- Medical-image analysis
- Patient-risk modeling
- Operational optimization

Retail and E-Commerce

Businesses can apply data science to:

- Recommendation systems
- Dynamic pricing
- Customer segmentation
- Demand forecasting
- Marketing attribution

Engineering and Infrastructure

Engineers can use data-driven systems for:

- Structural monitoring
- Energy optimization
- Traffic prediction
- Equipment reliability

- Smart-building systems

Common Mistakes ⚠️

Starting With the Algorithm

Choosing XGBoost, neural networks, or another sophisticated algorithm before understanding the business problem is a common mistake.

Solution: define the decision first.

Ignoring Data Quality

A powerful model cannot compensate for fundamentally unreliable measurements.

Solution: create data-quality checks before modeling.

Optimizing Only Accuracy

Accuracy can be misleading, especially with imbalanced datasets.

Solution: select metrics based on business consequences.

Building Dashboards Nobody Uses

A dashboard containing 40 charts does not necessarily create value.

Solution: design around specific decisions and users.

Confusing Correlation With Causation

If:

[corr(X,Y)\approx0.8]

that does not prove (X) causes (Y).

Solution: use experiments, causal methods, domain knowledge, and controlled comparisons where appropriate.

Challenges & Solutions

Challenge	Why It Matters	Solution
Poor data quality	Produces unreliable insights	Automated validation
Data silos	Prevents complete analysis	Data integration
Skill shortages	Slows implementation	Training + cross-functional teams
Model drift	Reduces production accuracy	Continuous monitoring
Privacy concerns	Creates legal and ethical risk	Governance + minimization
Lack of adoption	Insights are ignored	User-centered design
Excessive complexity	Raises cost and maintenance	Start simple
Weak KPIs	Success becomes unclear	Define measurable outcomes

Data-driven transformation also requires organizational change. Research on managerial decision making has found that intuition can still override analytics because of uncertainty, trust, knowledge, and organizational factors.

Case Study: Predictive Maintenance in a Manufacturing Plant 🏭

Consider a hypothetical manufacturing facility operating 200 industrial machines.

The engineering team experiences frequent unexpected failures. Each failure causes approximately:

[C_f = $8,000]

in downtime, emergency labor, and replacement components.

Assume the plant experiences 30 major failures annually:

$$Annual\ Loss=30\times \$8,000=\$240,000$$

Data Collection

Sensors record:

- Temperature
- Vibration
- Pressure
- Motor current
- Operating hours
- Maintenance history

Model Development

Engineers construct a failure-risk model:

$$P(Failure|X_1,X_2,\ldots,X_n)$$

The model identifies machines with elevated risk.

Operational Integration

Instead of simply displaying:

“Machine #47 has a 78% failure probability.”

the system generates an actionable recommendation:

“Inspect Machine #47 within the next maintenance window.”

That distinction is critical.

The data-science system has moved from **prediction to decision support**.

Measuring Success

Suppose annual failures decline from 30 to 18.

The avoided failures are:

$$30-18=12$$

Potential avoided cost:

$$12\times \$8,000=\$96,000$$

If the complete analytics system costs $40,000 annually:

$$Net\ Benefit=\$96,000-\$40,000=\$56,000$$

This is the type of measurement that matters to management.

The model is valuable because it produces an economic outcome—not because it has an impressive algorithm.

Essential Tips for Students and Professionals 💡

For Beginners

Start with:

SQL → Statistics → Python → Data Visualization → Machine Learning

Do not rush directly into deep learning.

Understand:

- Mean and variance
- Probability
- Correlation
- Regression
- Sampling
- Hypothesis testing
- Data cleaning
- Visualization

For Advanced Engineers

Focus on the complete production system:

Data Engineering + ML + MLOps + Governance + Business Strategy

Learn how to handle:

- Feature pipelines
- Model versioning

- Data drift
- Model monitoring
- APIs
- Cloud infrastructure
- Experiment tracking
- Explainability
- Security

The 70/20/10 Principle

A practical mindset is:

70% → Understand the problem and data

20% → Build and evaluate the analytical solution

10% → Communicate and operationalize the result

The exact percentages are not universal, but the principle is powerful: **data science is much more than model training.**

FAQs

What is applied data science?

Applied data science is the practical use of statistics, programming, machine learning, data engineering, and domain knowledge to solve real-world problems and improve decisions.

Is data science only about artificial intelligence?

No. Data science includes statistics, SQL, visualization, experimentation, forecasting, optimization, data engineering, and machine learning. AI is one part of the broader discipline.

What is the most important skill in business data science?

Problem formulation is one of the most important skills. A data scientist must understand what decision the analysis is intended to improve.

Why is data quality so important?

Poor-quality data can produce biased, unstable, or misleading results. A sophisticated algorithm cannot automatically correct incorrect measurements or systematically missing information.

Should businesses always use machine learning?

No. Sometimes a simple regression model, SQL query, statistical test, or well-designed dashboard provides a better solution than machine learning.

What is the difference between a dashboard and a data-science model?

A dashboard generally communicates and monitors information, while a data-science model can estimate, classify, predict, or optimize outcomes. Modern business systems can combine both.

How can companies measure the success of data science?

Measure business outcomes such as increased revenue, reduced cost, lower downtime, improved conversion, reduced churn, faster operations, or improved customer experience.

Will data science replace engineers and analysts?

Data science is more likely to change how professionals work than eliminate the need for domain expertise. Engineers, analysts, managers, and data scientists each contribute different forms of knowledge.

Conclusion

Applied data science is fundamentally an **engineering discipline for turning uncertainty into better decisions**. 📊⚙️

The biggest lesson is simple:

[\boxed{\text{Data} \rightarrow \text{Insight} \rightarrow \text{Decision} \rightarrow \text{Measured Value}}]

Businesses should not measure success by the number of dashboards created, the complexity of a machine-learning algorithm, or the size of a dataset.

They should ask:

Did the analysis improve the decision?

Did the decision improve the outcome?

Can the improvement be measured?

The strongest data-driven organizations therefore build a continuous feedback loop. Data reveals what is happening, statistics explains patterns, machine learning helps predict possibilities, engineers build reliable systems, and business leaders turn those insights into action.

The future of applied data science is consequently not simply about creating smarter models. It is about creating **smarter organizations**—where reliable data, engineering discipline, human expertise, experimentation, and measurable business objectives work together. 🚀

That is the central lesson of the data-driven business: **the ultimate product of data science is not a model—it is a better decision.**

Related Posts:

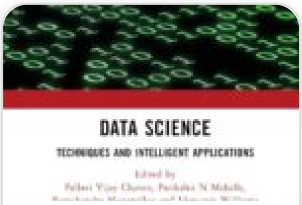

Data Science: Techniques and Intelligent Applications

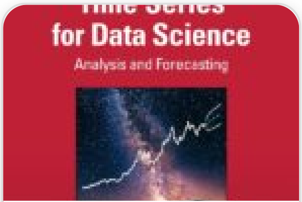

Time Series for Data Science: Analysis and Forecasting

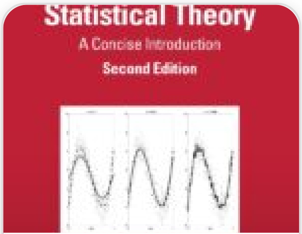

Statistical Theory 2nd Edition

Power Pivot and Power BI

An Introduction to Statistics: An Active Learning...

Preview PDF Book

← Previous Book

Next Book →

EngiVerse

Explore a vast library of free engineering ebooks curated to help you expand your technical knowledge, conduct groundbreaking research, and advance your career. Our mission is to create a hub of knowledge that fosters innovation, empowers aspiring engineers, and bridges the gap between education and industry.

Home Books About Us Contact Privacy Policy DMCA

copyright 2025 EngiVerse. All Right Reserved