

АГЕНТЫ ИИ / ЭКОСИСТЕМА CLOUD NATIVE / ТЕСТИРОВАНИЕ ПРОГРАММНОГО ОБЕСПЕЧЕНИЯ

Грептил, Курсор и Девин согласны с тем, что агенты должны запускать свой код. Важно, против чего они его запускают.

Агенты по написанию кода теперь запускают его, чтобы убедиться, что он работает. Но в облачных командах песочницы с макетами не позволяют выявить ошибки, возникающие при взаимодействии сервисов.

27 июня 2026 года, 11:00 [Арджун Айер](#)



Ирина Тригубова для Unsplash+

[Signadot](#) спонсировала этот пост.

В отрасли признали, что масштабный код для доставки грузов требует проверки во время выполнения, и этот процесс идет быстрыми темпами. Однако для большинства команд возможности такой проверки остаются ограниченными.

В агентной разработке формируется консенсус: чтобы массово выпускать **код агентов**, нужно дать им возможность проверять свой код во время

Ограничение заключается в том, что мы не знаем, действительно ли работает то, что создал агент, и ответ на этот вопрос мы получим только после запуска кода. Поэтому отрасль переносит проверку из очереди на рассмотрение человеком в собственный цикл работы агента, предоставляя агентам возможность выполнять свою работу и проверять ее до того, как человек увидит изменения.

«В отрасли признали, что масштабный код для доставки грузов требует проверки во время выполнения, и этот процесс идет быстрыми темпами. Однако для большинства команд возможности такой проверки остаются ограниченными».

Greptile, Cursor, Codex и Devin теперь запускают код, написанный агентом, а не просто читают его.

Это правильное направление. Как только **агенты открывают пул-реквесты** быстрее, чем их могут проверить люди, проверка становится узким местом. Внутренние агенты Stripe **отправляют более 1000 проверенных пул-реквестов в неделю**. Агент, который может запускать собственный код, выявлять ошибки и исправлять их до передачи работы, — это то, что не дает этой проблеме полностью лечь на плечи ваших старших инженеров.



Signadot — это платформа на основе Kubernetes, которая позволяет агентам искусственного интеллекта для написания кода выполнять масштабную проверку. Сочетая быстрые, масштабируемые эфемерные среды с системой проверки, разработанной для сложных распределенных систем, Signadot обеспечивает высокую скорость генерации кода и безопасное слияние пул-реквестов.

[Learn More →](#)

Representation, Mirror, Okteto, and Signadot

18 June 2026

Watch Claude Code Validate a Microservices Change in a Closed Loop

4 June 2026

Signadot Plans

20 May 2026

HEAR MORE FROM OUR SPONSOR

svo1975pvn1982@gmail.com

Submit

For cloud-native teams, this is not a nice-to-have. And it is where the approach most companies are adopting hits a ceiling. The common pattern gives each agent its own sandboxed environment, a snapshot that reproduces what a developer can run locally. For a self-contained application, that is enough.

For an agent working in a distributed system, it verifies one service against mocks and misses behavior that only appears when the change interacts with the rest of the system.

TRENDING STORIES

1. The Anthropic leader who built Claude Code says he ditched prompting — now he just writes loops.
2. Greptile, Cursor, and Devin agree that agents should run their code. What they run it against matters.
3. Anthropic gives @Claude a permanent seat in your Slack channels
4. Claude Code vs. Cursor vs. Codex vs. Antigravity — six months in
5. "Bring it to our shop": Workday's pitch for keeping AI agents close to your most valuable data

against is the real system or a stand-in for it.

The industry is converging on runtime

The clearest signal is that agent tools have stopped just reading the diff and started running it. Static analysis has a ceiling: It can reason about what code says, not what it does. A race condition that needs a real request, a regression that only appears after a page renders — none of it shows up in the text.

The trend is visible everywhere. Greptile **recently announced TREX**, a new feature that executes each change in a disposable sandboxed environment and returns logs, traces, and screenshots of what happened. **Cursor's cloud agents** clone the repo into their own VM to build and test. **OpenAI's Codex Cloud** does the same. **Devin** runs in a full environment with its own shell and test runner. Different products, one move: Give the agent a place to run its work and prove it before handoff.

Runtime verification is exactly the layer most agentic stacks are missing, and pushing it into the agent's own loop is how review stays fast as volume climbs. The open question is the level of runtime fidelity it takes to actually call a change verified.

Where it has to go further for cloud-native teams

The approach all these tools have taken is to give an agent roughly what a developer has at their desk: its own service, the dependencies it can start locally, and mocks standing in for everything else. That is a real and useful capability, and it is the correct direction.

For a cloud-native application built from many services, it covers less than half of it. The rest of the surface it misses is the part that matters most: How the change behaves relative to the rest of the system. A mock only returns the answer someone expected in advance. It can confirm an assumption but never catch a wrong one. And the expensive bugs in these systems are exactly the ones that break an assumption where two services meet:

- Contract drift between services
- A serialization mismatch at a boundary
- A retry or timeout policy that misbehaves two hops away


something when the change runs against the real services, real data, and real traffic around it.

The same goes for the non-functional behavior teams care about most: Performance and load regressions, resource contention, and runtime security issues surface only when the change runs inside a realistic system. That is the layer a sandboxed environment cannot reach.

The intuitive fix, cloning the whole system into every agent's environment, does not hold up. Standing up dozens of stateful services with their data and config, fresh for each agent on every iteration, is impractical at the volume and speed with which agents generate code. And a copy is still a snapshot, not the live system.

Where the sandbox model stops

One service runs for real. Everything else is mocked.



Cross-service behavior is never tested


An architecture that reaches the whole system

The way through is to stop giving each agent its own copy of the system and instead let every agent run verification against one shared, production-like system with isolation.

The pattern is straightforward. A shared cluster runs all the real dependencies, the full set of services as they behave in production. To verify a change, you deploy the modified service into that baseline and use request-level isolation to keep each agent's traffic on its own path. The agent's requests hit its version of the

This gives the agent a realistic runtime in which its change runs alongside the live services, data, and policies it has to work with. Integration and system behavior become observable because the interactions are real rather than mocked.

It also scales to the way agents work. Because each change is just one service layered onto the running system rather than a private copy of everything, many agents can verify at once, cheaply, without colliding. The environment is lightweight and short-lived, and it runs on infrastructure the team already operates. This is the model we build at **Signadot**. Today, that means Kubernetes, though the same idea will carry over as the tooling matures.



It fits the loop the agent tools already use. The agent writes, runs, reads the failure, and tries again, all before the pull request. What changes is the bar for a passing run. Not green against mocks the agent wrote to match its own assumptions, but correct against the real system those assumptions describe.

The trend is right. The reach is the question.

time verification is becoming a first-class part of how software ships, and that is the correct evolution. The teams furthest ahead already treat the agent's job as

“Not green against mocks the agent wrote to match its own assumptions, but correct against the real system those assumptions describe.”

The question for anyone running a **cloud-native system** is how far that proof reaches. An isolated per-agent environment verifies a change on its own, and for many changes, that is enough.

Proving a change is right when it has to work alongside everything around it is a different problem. It takes integration and system-level verification against a runtime that holds the rest of the system, not a stand-in for it.

The teams that get the most out of background agents will be the ones whose verification loop reaches the whole system, not just the service in front of them.

TNS



Arjun Iyer, CEO of Signadot, is a seasoned expert in the cloud native realm with a deep passion for enhancing the developer experience. Boasting over 25 years of industry experience, Arjun has a rich history of developing internet-scale software and...

[Read more from Arjun Iyer →](#)

TNS owner Insight Partners is an investor in: OpenAI.

SHARE THIS STORY



TRENDING STORIES

1. The Anthropic leader who built Claude Code says he ditched prompting — now he just writes loops.

Greptile, Cursor, and Devin agree that agents should run their code. What they run it against matters.

5. "Bring it to our shop": Workday's pitch for keeping AI agents close to your most valuable data

INSIGHTS FROM OUR SPONSOR



Signadot is a Kubernetes-native platform that empowers AI coding agents to verify code at scale. Combining fast, scalable ephemeral environments with a validation framework built for complex distributed systems, Signadot ensures high-velocity code generation results in safely merged pull requests.

[Learn More →](#)

Watch Claude Code Validate a Microservices Change in a Closed Loop

4 June 2026

Signadot Plans

20 May 2026

AI Code Testing for Cloud-Native Systems: Why Unit Tests Aren't Enough

8 May 2026

Introducing the signadot-validate Skill: Close the Loop for Coding Agents Building Microservices

8 May 2026

Closing the Loop for Coding Agents in Cloud-Native Systems

6 May 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in
your inbox each day.

svo1975pnv1982@gmail.com

SUBSCRIBE

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds



roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



© The New Stack 2026

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)