


ИНФРАСТРУКТУРА ИИ / МОДЕЛИ ИИ / СТРАТЕГИЯ ИИ / УПРАВЛЕНИЕ API

Palantir и Nvidia хотят изменить правила владения государственным ИИ

Новый движок Palantir и Nvidia запускает открытые модели Nemotron в изолированных сетях, смещая акцент с вопроса о том, какой ИИ использовать, на вопрос о том, какой ИИ принадлежит компании.

29 июня 2026 года, 13:55 [Аманда Касвелл](#)



До недавнего времени использование искусственного интеллекта означало обращение к чужой модели. Но подключение приложения к API от OpenAI, Anthropic или Google не подходит для государственных структур и операторов критически важной инфраструктуры. Если данные по закону или в силу особенностей работы не могут покидать защищенную сеть, то размещение конечной точки в чужом облаке не подходит, независимо от того, идет ли речь об анализе разведывательных данных, управлении энергосистемой или ведении медицинских карт.

модель, а устройство для ее развертывания и владения ею».

Именно через эту призму следует рассматривать **заявление Palantir** в понедельник. Компания представила «интеллектуальный движок», созданный на основе искусственного интеллекта Nvidia и **открытых моделей Nemotron**, для запуска, настройки и постоянного совершенствования ИИ в изолированных и других защищенных средах, при этом оставляя данные и весовые коэффициенты моделей в распоряжении заказчика. Самое показательное здесь то, что Palantir поставляет не модель, а инструмент для ее развертывания и владения ею.

От вызова ИИ до его эксплуатации

Большинство организаций в течение многих лет будут использовать гибридные модели, размещенные как на собственных, так и на внешних серверах, но наряду с этим подходом формируется и другой. Вместо того чтобы рассматривать модель как внешнюю службу, приложение взаимодействует с внутренней платформой ИИ, которая направляет запросы к одной или нескольким моделям, работающим на графических процессорах организации. Данные никогда не покидают периметр, что дает службам безопасности и контроля прямой контроль над управлением, аудитом и хранением.

Таким образом, мы видим, что вопрос смещается с *какой модели нам следует позвонить* на *в какие модели нам следует инвестировать и управлять ими самостоятельно*. И, что не менее важно, *во сколько нам обходится их эксплуатация?*

Что на самом деле находится в коробке

Здесь анонс становится более конкретным, чем в пресс-релизе, и здесь стоит разобраться, что такое Nemotron.

Nemotron is Nvidia's family of open-weight models, released in 2026 in three sizes — Nano (~31.6B parameters), Super (120B), and Ultra (550B). All three use a hybrid Mamba-Transformer mixture-of-experts design that activates only about a tenth of their parameters per token, so they run far cheaper than their headline sizes suggest, with context windows up to a million tokens.

name.cpp, or as Nvidia NIM microservices, the containerized, TensorRT-LLM-optimized path that ships inside the Nvidia AI Enterprise suite. Palantir, notably, was already on Nvidia's published list of early Nemotron adopters, so this is less a cold start than a formalization.

TRENDING STORIES

1. [AWS, Microsoft, and Google agree the session is the new unit of compute. They disagree on how to isolate it.](#)
2. [GitLab just surveyed 1,500 developers. Here's why it matters for your codebase.](#)
3. [Neoclouds, sovereign AI and Postgres: The new operating model for regulated enterprises](#)
4. [Palantir and Nvidia want to change who owns government AI](#)
5. [OpenAI wants to claim more of the AI stack with Jalapeño, its first custom chip](#)

An important note for builders weighing the move: Nemotron doesn't top the raw-capability leaderboards; open families like DeepSeek, Qwen, and Kimi K2 generally score higher on absolute benchmarks. Nemotron's pitch is efficiency on Nvidia silicon and genuine openness, which is exactly the axis that matters when the deployment target is your own hardware behind an air gap rather than a hosted endpoint.

Palantir's contribution is the layer that turns "download the weights" into "run this in a classified environment and keep improving it." The company describes three engineering surfaces: *deployment* (getting base and customized models into air-gapped and classified networks), *context* (prompts, workflow structure, and model behavior in production), and *model* (changing the weights themselves on proprietary data and mission outcomes). Underneath sits Palantir's existing stack — AIP, Ontology, Foundry, and Apollo — handling data authorization, enforced isolation, and auditability.

The "self-improving" claim is a telemetry loop: The engine captures usage and the data, then uses it to post-train and align the model toward the tasks where it

The catch: Owning the stack means owning the stack

The sovereignty pitch is genuinely attractive, and it has a price tag the announcement doesn't dwell on. Owning your models means owning everything around them. That's GPU capital expenditure and the power and cooling to match; an inference stack you keep patched and performant; a model lifecycle — fine-tuning, evaluation, rollback — that you staff and run yourself; and a security burden that doesn't go away just because the box is disconnected.

“Owning your models means owning everything around them.”

Ultra-class models, in particular, are not modest: running a 550B model, even at ~10% activation, requires multi-GPU server nodes, not a spare rack.

For an agency that legally cannot use a hosted API, that cost is simply the cost of doing the work, and the calculus is easy. For a commercial enterprise that *could* use a hosted endpoint, the math is a real trade — control and data residency on one side, capex and operational headcount on the other. The right answer is workload-specific, and anyone selling it as obvious is selling something.

Read the announcement for what it is

It's also worth being clear about what this news is not. It's a packaging-and-positioning announcement between two companies with aligned incentives: NVIDIA sells more GPUs and more AI Enterprise licenses, Palantir sells more platform, and “sovereign AI” is the banner both have been marching under all year. There's no named agency customer, no contract value, and no benchmarks. Palantir CEO Alex Karp says many U.S. clients already use these models and frames the appeal as avoiding proprietary insight being baked into the weights of **closed models**; NVIDIA's Jensen Huang calls open source foundational to national security. Both quotes are doing strategic work. The substance here is the integration and the go-to-market, not a technical breakthrough.

Why it still matters

...p away the government framing, and the same pattern is spreading through finance, healthcare, manufacturing — anywhere data residency and compliance are

compute unit, and an inference tooling layer (NIM, Tensort-LLM, VLLM) mature enough that air-gapped serving is an engineering project rather than a research one.

“The model becomes a component. The engine becomes the product.”

The takeaway here for developers is that as open models converge on “good enough,” differentiation shifts to everything around them — deployment, routing, governance, evaluation, security, and lifecycle management. The model becomes a component. The engine becomes the product.

Most organizations will keep calling hosted APIs for plenty of work. But for a growing set of them, the interesting question is how much of the stack they’re prepared to own — and whether they’ve counted the cost of owning it.

TNS



Amanda Caswell is an AI journalist, certified prompt engineer, and technology commentator whose work and expertise have been featured on Fox News and CBS News. She covers artificial intelligence, developer tools, foundation models, and emerging technologies, with a particular focus...

[Read more from Amanda Caswell →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY



TRENDING STORIES

1. AWS, Microsoft, and Google agree the session is the new unit of compute. They disagree on how to isolate it.

3. Neoclouds, sovereign AI and Postgres: The new operating model for regulated enterprises
4. Palantir and Nvidia want to change who owns government AI
5. OpenAI wants to claim more of the AI stack with Jalapeño, its first custom chip

INSIGHTS FROM OUR SPONSORS



How durable sessions unify human-to-human and human-to-agent messages

26 June 2026

Is AI making your teams better, or just busier?

26 June 2026

One person, one vote: building live voting with message annotations

26 June 2026



Deploying Applications to Kubernetes from Source Code: Five Practical Approaches Compared

22 June 2026

Why Crossplane Alone Is Not Database Management, and What Platform Teams Need Instead

19 June 2026

Multi-Account AWS Is a Control-Plane Problem — And Most Teams Learn That the Hard Way

17 May 2026

When Amazon Aurora Stops Scaling, The Ceiling You Can't Buy Your Way Past

25 June 2026

The Fintech AI Platform: Why Agents Need One Database, Not Five

23 June 2026

Database Branching for AI Agents: How TINE Solves the Schema Drift Problem

15 June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the most recent TNS articles in your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

Databases

Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

API Management

Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



[Privacy Policy](#)

[Cookie Policy](#)