

Мы только что появились на Product Hunt — на день раньше. Зарегистрируйтесь и поставьте лайк; мы будем рады любой поддержке.

Upvote Prefactor →

ОЦЕНКА АГЕНТОВ В РЕЖИМЕ РЕАЛЬНОГО ВРЕМЕНИ

Оценивайте своих ИИ агентов

Prefactor оценивает каждый производственный запуск в момент его выполнения (качество, дрейф и риск), а затем использует эти оценки для принятия мер, так что сбой в работе агента **выявляется в процессе**, а не фиксируется постфактум.

[Получите 1 млн бесплатных спенсов →](#)

[Прочитать документацию](#)

Добавляется в ваш стек за считанные минуты: **TypeScript** и **Python SDK**, нативные для LangChain, Claude, Vercel AI, OpenClaw и LiveKit.

AGENT.TS

claims-processor

v2.0

В ПРОДАКШЕНЕ

\$ npm i @prefactor/sdk

const pf = префактор({ агент: "обработчик претензий" })

\$ предварительная оценка - live · запуск # 4821



магистр права в качестве судьи

технический

качественный

0.92



312



ВЫСОКИЙ

1.4с

2.8с

\$0.04

PII · PCI

ПОЛЬЗОВАТЕЛЬСКИЕ ОБЛАСТИ



GitHub · 12 оценок ✓



Sentry · 2 ошибки

■ issue_refund удерживается · принудительно применяется через 42 мс

Уничтожить

Утвердить

P УЧАСТВУЙТЕ В PRODUCT HUNT УЖЕ СЕГОДНЯ

1 000 000 бесплатных переходов для первых 50 регистраций

стоимость использования составляет 2500 долларов США за наш счет. Зарегистрируйтесь, отправьте код **ProductHunt** через чат в приложении или на адрес hello@prefactor.tech, а затем в течение 7 дней зарегистрируйте агента с помощью SDK. Кредитная карта не требуется.



FEATURED ON

Product Hunt



435

НАМ ДОВЕРЯЮТ КОМАНДЫ ИЗ



COMPANY

COMPANY

COMPANY

COMPANY

> \$01 / РАЗРЫВ

проблема: наблюдать ≠ вмешиваться

! РАЗРЫВ

Видеть все. Не останавливать ничего.

В этом и заключается разрыв. Наблюдаемость и мониторинг предоставляют вам информационные панели (отслеживание, оценка, оповещения), а затем возвращают проблему. К тому времени, как вы прочитаете график, агент уже примет меры.



Дашборды не вмешиваются

Они сообщают, что агент слил персональные данные, но это уже произошло. Ничто не мешает следующему сливу.

02



Оценки без конкретики

Неудовлетворительная оценка — это просто еще один показатель. Она ничего не меняет в том, что может сделать агент.

03



Вы — петля

Люди, которые вручную переключают тумблеры, не могут заменить собой больше нескольких агентов.

> §02 / CLOSED LOOP

spine: observe → evaluate → act

THE RELIABILITY TRIFECTA

Observability stops at the dashboard.

Other tools observe and score, then hand you the problem. Prefactor wires evaluations and risk straight into action: pause a risky run for approval, or enforce a policy at runtime, through the SDK or API.



Enforce



Evaluate with the evals you define: LLM-as-judge, technical & qualitative metrics, plus data-risk



Act: hold, approve or block, human-in-the-loop, via SDK or API

A risky agent is **caught, not just charted.**

prefactor · live run

ENFORCING

① OBSERVE customer-support-v3 · run #4821

② EVALUATE evals & risk, on every run

generate_reply llm-as-judge · custom eval

0.96 ✓

issue_refund action · sensitive data

CRITICAL

↳ high-risk action: **held for human approval**

③ ACT **Human-in-the-loop · run paused**

SDK · API

⏸ Paused **before** issue_refund executed

→ Awaiting approval, **enforced at runtime**

> §03 / WORKFLOW

path: install → enforcement

THE WORKFLOW

From **to** .

Five steps, each a feature that's live today. Step through them:

Prefactor

One command connects your workspace and discovers agents across your runtimes: no migration, no rip-and-replace.

2

SDK INTEGRATIONS

Instrument your agents

3

RUNTIME VISIBILITY

See every run in real time

4

EVALS & SCORING

Evaluate what matters

5

ENFORCEMENT

Enforce, automatically or with a human

\$ prefactor init

✓ CLI connected to workspace

scanning agent runtimes...

✓ 4 agents discovered

> §04 / CUSTOM SPANS

enrich: any datasource → run

CUSTOM SPANS

Bring context into the agent run.

Custom spans aren't just markers. Pull context from any datasource (GitHub, Linear, Jira, your database, internal APIs) into the run, so every evaluation is grounded in what actually happened.

```
const span = pf.customSpan('enrich_review_context');
span.attach(await github.getPR(482));
span.attach(await db.customerTier(userId));
span.score({ grounded: true }); // → grounds every eval
```

[View the SDK reference](#) → [Custom spans guide](#) →

custom span · **enrich_review_context**

PULLS CONTEXT FROM



GitHub

PR #482 · diff & checks



Linear

AGT-231 · In Review



Jira

OPS-88 · priority P1



Postgres

customer tier: Enterprise



+ any REST / GraphQL API or internal datasource

→ Attached to the run, available to **every eval** as ground truth

> **\$05 / LIFECYCLE**

promote: dev → staging → prod

THE AGENT DEVELOPMENT LIFECYCLE

Ship agents like software:

Prefactor versions every agent, tracks it against a schema, and promotes it through dev, staging and prod only when its evals pass, so you can compare scores version-to-version and prove each one is better than the last.

Rolling

claims-processor v2.0

schema v3

evals 94% ✓

Ready to promote

PROMOTE
WHEN EVALS PASS →

STAGING

claims-processor v2.0

schema v3

evals 96% ✓

Validated against prod traffic

PROMOTE
WHEN EVALS PASS →

PRODUCTION

claims-processor v2.0

live · scored

healthy

↶ instant rollback to v1.4

Versioning

Immutable versions by semver, commit or tag, with lifecycle states from pending to retired, so you always know exactly which agent ran.

Schema management

Two-tier schema-and-deployment versioning, tracked independently, with conformance drift surfaced when a run breaks its declared shape.

Environments

First-class dev / staging / prod. Every deployment binds a version to an environment, with eval-gated promotion and instant rollback.

[See versioning & promotion in the docs →](#)

WORKS WITH YOUR STACK

One layer over every agent, editor and workflow.

Native SDK integrations for the agent frameworks you build on, plus the coding tools and workflow platforms your team already uses.

AGENT ORCHESTRATION

NATIVE SDK



LangChain



Claude



Vercel AI



OpenClaw



LiveKit

OR ANY STACK:
CORE SDK

TypeScript



Python

CODING TOOLS



VS Code



GitHub Copilot



Cursor



Claude Code

WORKFLOW TOOLS



n8n



+ more coming

Connected through native SDKs, **OpenTelemetry**, and a **TypeScript & Python core SDK** that instruments anything else.

"We had forty agents in production and no honest way to say which ones were still doing their job. Prefactor gave

Head of AI Platform

Global financial services

0

sensitive-data categories detected in agent traffic

0

of runs traced: every call, tool & decision

<5 min

from CLI install to first traced run

HITL

hold & enforce at runtime, via the SDK or API

> §08 / SECURITY

posture: least privilege · auditable

Built for enterprise security.

Agents act with real access, so Prefactor is built around least privilege, full auditability and your existing identity stack: the enforcement layer beneath the reliability story.

✓ Scoped, least-privilege access

✓ Complete auditable record

✓ Sensitive-data detection: 17 categories

✓ SSO for enterprise accounts

⊕ SOC 2 Type II **IN PROGRESS**

⊕ RBAC **ON THE ROADMAP**

> §09 / FAQ

reference: owned terms first

What is agent observability?

Capturing every agent run as structured trace data (each LLM call, tool invocation and decision) so you can see exactly what an agent did, how well it performed, and what it cost.

What is agent evaluation?

How do I instrument my agent with Prefactor?

What are custom spans, and how do they measure quality?

How does human-in-the-loop work?

How does Prefactor enforce at runtime?

Why do AI agents need identity and scoped access?

GET STARTED

See Prefactor on your own agents.

Free for your first 25,000 spans a month. Drop in the SDK and watch your own runs get scored, live. Or take the 30-minute walkthrough with an engineer.

Knowledge

Knowledge

Know your AI agents are doing their job

PRODUCT

Platform

Integrations

Open-source evals

Pricing

Compare

Changelog

Docs

RESOURCES

Agent Risk Audit

Blog

Agentic Contract Whitepaper

MCP Whitepaper

Glossary

Free Tools

Security

COMPANY

About Us

News & Events

Contact Us

Careers

SOLUTIONS

Healthcare

Insurance

For Developers

For Heads of AI

For AI Product

Security & Governance

For Founders

LEGAL

Privacy Policy

Terms of Service

Cookie Policy

Acceptable Use

603060

Email for our newsletter

Subscribe →

 Slack Community

Backed by Antler, Black Nova VC & Func Ventures