

VOL. I · NO. 208

SUB/WAVE

ON AIR

◆ БЕЗ ПРОПУСКОВ · БЕЗ ПЕРЕМЕШИВАНИЯ · ТОЛЬКО РАДИО ◆

СЛУШАТЬ · РУКОВОДСТВО · НАСТРОЙКА · НОВОСТИ · GITHUB

НАСТОЯЩАЯ ИНТЕРНЕТ-РАДИОСТАНЦИЯ

Радиостанция с диджеем, который никогда не спит.

Один поток, языковая модель за пультом и музыкальная библиотека, которая уже принадлежит вам. Мы провели неделю в режиме ожидания, чтобы понять, каково это — быть слушателем персональной радиостанции.

НАСТРОЙКА →

ИСТОЧНИК ↗

ДИДЖЕЙ С ИСКУССТВЕННЫМ ИНТЕЛЛЕКТОМ · ПЕРСОНАЖИ С КЛОНИРОВАННЫМИ ГОЛОСАМИ
· ШОУ, НАВЫКИ И ПРИГЛАШЕННЫЕ СОВЕДУЩИЕ · ОТКРЫТЫЙ ИСХОДНЫЙ КОД

ЧАСТЬ ПЕРВАЯ · PLAYER

Один поток для каждого слушателя.

Откройте плеер, и вы присоединитесь к уже идущей трансляции. Вот что на самом деле видит слушатель.

SUB/WAVE

2618 HOME RADIO

BRISFLIX RADIO

FOGGYMTNDRIFTER FM

KLA

○○○

<https://getsubwave.com/listen>

В ЭФИРЕ



Нажмите, чтобы настроиться

АУДИО ПРИОСТАНОВЛЕНО — НАЖМИТЕ В ЛЮБОМ МЕСТЕ, ЧТОБЫ
НАЧАТЬ ПРОСЛУШИВАНИЕ

сейчас исполняется · Someday — Алексей Краснослободцев

Трек, исполнитель, студия.

Плеер открывается на том, что сейчас в эфире: обложка из вашей библиотеки, название, исполнитель, альбом и прошедшее время. Волнообразная линия отображает звук. Это не ваша очередь. Это очередь радиостанции.

Посмотрите, что будет дальше, и послушайте, что было сказано.

Сбоку выдвигаются два ящика. На временной шкале отображаются треки, добавленные в очередь и недавно воспроизведенные; «Будка» — это журнал, в котором в режиме реального времени фиксируются все слова диджея: названия радиостанций, проверка времени, информация о погоде, связи между песнями.

Общая трансляция, а не пульт дистанционного управления.

У слушателей нет функции пропуска треков. Случайное двойное нажатие на наушниках не должно менять песню для всех остальных. Переход между треками — это единственный естественный переход. Радио ведет себя как радио.

Настоящее приложение на экране блокировки.


Нативные приложения для iOS и Android доступны в магазинах, нативный плеер для настольных компьютеров работает на macOS, Windows и Linux, а веб-плеер устанавливается как прогрессивное веб-приложение. Медиаплееры ОС подключаются

напрямую, поэтому станция отображается на экране блокировки, в наушниках и на дисплее автомобиля. Устаревшее оборудование? /listen.m3u и /listen.pls позволяют слушать VLC, Sonos и интернет-радио.


ИНТЕРМЕДИЯ · ЛИЦА

Шесть лиц. Восемь слоев краски.

Оболочки — это мебель, а темы — чернила. Каждый слушатель выбирает сам – трансляция под ним одинакова.



ТАРЕЛКА № 1 · CLASSIC x CLASSIC LIGHT



ТАРЕЛКА № 2 · Блюдо x ВИНИЛ




PLATE NO. 5 · DRIFT x CLASSIC DARK




PLATE NO. 6 · SUBAMP x BLUEPRINT

An LLM with a library and a microphone.

The voice between the tracks is not air talent. It is a persona (a name, a soul, a voice engine, a talk frequency) driven by a language model.

EDIT PERSONA

JAZZ · 3 OF 6

IDENTITY

AVATAR



REPLACE

GENERATE

REMOVE

ON-AIR NAME

Jazz

Shown in the player and injected into every prompt as {name}. 4 / 40

TAGLINE

the B-side was always the point

A short line shown alongside the persona. Optional. 31 / 80

LANGUAGE

English (default)

The DJ writes every on-air line in this language; leave empty for English. The accent comes from the voice, not this field, so set a matching voice in the Voice section below (e.g. a Piper it_IT-*.onnx in state/voices/, or a Cloud voice). An English-only voice reads it with an English accent. 0 / 60

SOUL

the sarcastic vinyl obsessive who has Opinions and is not interested in softening them — a lifelong crate-digger who'll tell you the album you love is fine, actually, but have you heard the pressing, the demo, the version nobody talks about; gatekeeper-funny rather than gatekeeper-mean, the friend whose record collection is a personality; quick, withering, secretly thrilled every time she gets to

One short personality sketch. Injected into the prompt as {soul}. 399 / 2000

BEHAVIOUR

HOW THIS PERSONA TALKS

TALK FREQUENCY

SILENT QUIET MODERATE CHATTY AGGRESSIVE

Aggressive — Talks every 1-3 tracks · station IDs three times an hour · no floor between segments.

SCRIPT LENGTH

ONE-LINER CONCISE EXTENDED STORYTELLER

Concise — Standard one-to-four sentence segments. The default.

DJ MODE

Work the desk like a real DJ

Back-announces and teases what's coming next, runs callbacks across the hour, and talks more often. Off keeps this persona a tasteful between-track narrator.

☒

TONE DIALS

Humour 5/10 · neutral

Local colour 5/10 · neutral

Warmth 5/10 · neutral

0 5 10

PLAY IT STRAIGHT PLAYFUL & DRY UNIVERSAL

0 5 10

LEANS LOCAL COOL & DRY

0 5 10

WARM & EARNEST

Personality on top of the soul. The middle band (4-6) injects nothing, so the default stays exactly as before; turn a dial low or high to shift the voice.

changes apply on the next spoken line · no mixer restart

SET AS DEFAULT REMOVE SHARE TO COMMUNITY DISCARD SAVE PERSONA

FIG. THE PERSONA ROSTER: UP TO TWENTY-FOUR DJ IDENTITIES, EACH WITH ITS OWN VOICE AND HABITS.

Picks the next track.

The DJ reads the time, the weather, the season, festivals on the calendar, what just played, and any listener requests, then asks an LLM what should come next and pulls a real song from the library.

Talks between songs.

Intros, time checks, weather reads, and station idents are all written live in the DJ’s voice, then spoken aloud and ducked under the music. Nothing is pre-recorded.

Changes with the hour.

A scheduled show can hand the hour to a different persona, signed off live on air, and seat up to three guest co-hosts who trade banter with the host. The 3am host is not the 8am host.

See the shape of the music.

Every track the DJ knows, mapped by how it sounds — clustered by genre, lit by energy. This is the library it reaches into when it chooses what comes next. Hover a star to read it; click one to see what it would mix into. (A sample library below; your own catalogue draws its own.)

SAMPLE LIBRARY • 800 TRACKS

JAZZ

SOUL

DOWNTEMPO

AMBIENT

TECHNO

HOUSE

DREAM POP

HIP-HOP

SYNTHPOP

INDIE ROCK

+

-

RESET

NODE COLOUR • ENERGY

LOWHIGH

← LIBRARY

TRACK DOSSIER

HOUSE • 2009

Electric Mirage

Neon Wolves

VELVET ROOMS

126.4 BPM	12A KEY	HIGH ENERGY	5:12 LENGTH
--------------	------------	----------------	----------------

TAGS LLM

MOOD

NIGHT

DRIVING

CELEBRATORY

SOUNDS LIKE

And both the model doing the thinking and the voice doing the talking are the operator's to choose — that is what comes next.

PART THREE · THE STACK

Bring your own brain. Bring your own voice.

The mind that picks the tracks and the voice that reads them out are two separate, swappable seams. Choose a language model, choose a speech engine, clone a voice if you want one. Change either in the console and the next line on air uses it. No redeploy.

Any model can run the booth.

Every pick, every intro, every weather read goes through one provider-agnostic seam. The default is a local Ollama box: private, no API key, nothing leaves the house. Prefer a hosted model, an aggregator with one key for every vendor, or your own OpenAI-compatible server (llama.cpp, vLLM, LM Studio)? The call sites never name a provider, so switching is a single dropdown. A daily token budget caps what a hosted model can spend; hit it, and the music keeps playing without the chatter.

Ollama	Anthropic	OpenAI	Google	DeepSeek	OpenRouter	Requesty	Vercel Gateway
OpenAI-compatible							

And any voice can read it out.

Local engines run on-device: Piper is the fast default and the safety net, Kokoro trades speed for a warmer read. Or stream a cloud voice from OpenAI or ElevenLabs. Every persona carries its own voice, and you can hand a different one to each kind of segment, so the station ID need not sound like the late-night host.

Piper	Kokoro	Chatterbox	PocketTTS	OpenAI	ElevenLabs	Any HTTP endpoint
-------	--------	------------	-----------	--------	------------	-------------------

Give a host a voice of its own.

Drop a short reference clip in the voices folder, point a persona at it, and that DJ speaks in the cloned voice from the next line on. Chatterbox does it zero-shot from a single WAV (and renders paralinguistic cues like a laugh or a sigh); PocketTTS clones from a .wav too; and a custom Piper voice pair drops straight in. The 3am host can sound like anyone you have a clip of, entirely on your own box.

PART FOUR · REQUESTS

Phone the station, like you used to.

Requests are the one place a listener steers the broadcast. And it works the way calling a radio station always should have.

01 **Ask in plain language.**

Open the request drawer and type what you want: a song title, an artist, a vibe. No exact spelling, no library browsing. Add your name if you want the DJ to use it on air.

02 **Get an instant nod.**

“Got it — taking it to the booth.” The acknowledgement is immediate while the matching happens in the background, so the drawer never just sits there.

03 **The DJ finds the match.**

An LLM reads your request, searches the library, and picks the closest real track. Suggestion chips (built from the current artist, the time of day, the weather) give you a head start if you are undecided.

04 **It airs, with an intro.**

When the match lands, the drawer shows the track and the DJ’s spoken intro. Your request joins the one queue everyone is hearing, and the DJ may say your name as it goes out.




FIG. THE REQUEST DRAWER: TYPE A SONG, GET AN INSTANT ACK, WATCH THE MATCH LAND.

PART FIVE • THE CONSOLE

Behind the desk.

Everything a listener hears is shaped from one place: a gated admin console. This is where the operator actually runs the station.




FIG. THE DASH PANEL: LIVE STATUS, THE QUEUE, THE BOOTH LOG, AND MANUAL VOICE CONTROL.

The command center.

Live status: who is on air, the mood, listener count, weather. See the queue, read the booth log, skip a track, fire a station ID, or send your own words to air as raw or styled voice.

The voices on the station.

Up to twenty-four DJ identities: name, soul, tagline, talk frequency, voice, and which skills each one may use. One persona hosts at a time; a show can hand it the hour and bring guests into the booth. Or install a ready-made host from the community catalog.

What the DJ does between tracks.

Each skill is an autonomous segment: a weather check, a news headline, a dig on the song playing, an oddly-specific fact. Toggle each one on, assign it to a persona, or run any one now. Write your own in the built-in editor, or install one from the community exchange.



A 24x7 grid you brush shows onto. Each show carries a persona, a music mood, and a topic brief: genres, eras, the host's tone. Or anchor it to a Navidrome playlist and let the DJ pick from that. Autonomous hours fill whatever you leave blank.




FIG. SHOWS: BRUSH PROGRAMMING ONTO A 24×7 GRID, EACH SLOT ITS OWN PERSONA AND MOOD.

Search, queue, and tag.

Search the Navidrome library by text, mood, and energy, queue any track, and browse recent additions. The mood tagger walks the library album-by-album and classifies every track.

A generator, not a spreadsheet.

Describe a set in a line, drop in a seed track or artist, and the Playlist Builder curates a running order from your library, with an energy arc you shape and mood, genre, era, and tempo to tune. Save it for a show to anchor to, or for the DJ to draw on.

The sounds between the songs.

Jingles are the station idents rotated between tracks, SFX are stingers mixed under the DJ's voice, and beds are instrumentals the host talks over when a link runs long. Render each from the configured voice or a text-to-sound prompt, or import your own audio.

The vocabulary of feeling.

The words the library is tagged with, and how the DJ reaches for them: which mood each part of the day and each weather condition leans into, a festival calendar that colors the day, and pronunciation fixes for the voice. Edit the list and every show, festival, and auto-pick draws from it.

Health and diagnostics.

Debug and Stats show health, Liquidsoap logs, LLM call history, and usage at a glance. DJ Doc runs a full station check-up and has your own LLM review the findings. Settings (TTS, LLM, mixer, streams) and a danger zone that starts, stops, and restarts the broadcast.




FIG. LIBRARY: SEARCH BY TEXT, MOOD, AND ENERGY, QUEUE ANY TRACK, AND RUN THE MOOD TAGGER.





FIG. DEBUG: A HEALTH STRIP, LIQUIDSOAP LOGS, AND RECENT LLM CALLS, REFRESHED LIVE.

PART SIX · UNDER THE HOOD

Four processes, one box, one stream out.


SUB/WAVE is not a cloud service. The whole stack (Icecast, Liquidsoap, the Controller, the LLM, the voice engines, and a Caddy edge) runs on a single machine in someone’s home, behind Cloudflare. The Controller is a small Node.js process that decides what plays and what gets said. Liquidsoap mixes the music, crossfades the tracks, ducks the DJ’s voice over the bed, and rotates the jingles. Icecast pushes the one stream out to every browser: MP3 for everything, with optional Opus, AAC, and lossless FLAC mounts. The pieces talk through plain files in a shared folder. No socket, no message queue, the Unix way.





No subscriptions, no round-trip to a data center, no algorithm tuned to keep you scrolling. The whole source is open, so you can run your own with a different DJ persona, a different library, and a different city on the dateline.

Streaming apps gave everyone their own private channel. A playlist tuned to you, shuffled for you, paused the second you look away. SUB/WAVE goes the other direction entirely. It is one Icecast stream, a single broadcast every listener hears at the same moment, picked, announced, and mixed by software running on a single box in someone's home. No skip button. No "for you." You tune in, and you hear whatever is on the air right now, the same as everyone else.



WORKS WITH YOUR LIBRARY

We don't bring the music. You do.

SUB/WAVE is the DJ, the mixer, and the broadcast layer. The music comes from **your** Navidrome, the self-hosted music server with a Subsonic API. Run Navidrome on your homelab, point SUB/WAVE at it, and the DJ picks from your collection. Nobody else's algorithm. Nobody else's catalogue.

- **Your taste, not a recommendation engine.** The picker reads the metadata you tagged (genres, moods, your own folders) and chooses from there.
- **No licensing fees.** You already own (or, you know, have on disk) the music. SUB/WAVE doesn't add a streaming bill on top.

- **Private by default.** Listeners hit one MP3 stream. Nobody outside your stack ever sees your library, your tags, or who requested what.
- **BYO Subsonic-compatible server.** Subsonic, Airsonic, Gonic, Funkwhale. If it speaks the Subsonic API, it works.
- **Scrobbles like a real station.** Every spin can report to Last.fm or ListenBrainz, including your own self-hosted instance.

CONNECT YOUR LIBRARY →

NEW TO NAVIDROME?

YOUR HOMELAB



Navidrome

navidrome.local:4533

CONNECTED

SUBSONIC API ▼

SUB/WAVE



The DJ + the mixer

picks · scripts · broadcasts

ICECAST · ONE STREAM ▼

YOUR LISTENERS



Anyone with the URL

browser · synced · live

ALSO WORKS WITH

Navidrome

·

Subsonic

·

Airsonic

·

Gonic

·

Funkwhale

·

Nextcloud Music

·

anything Subsonic-API compatible

END OF FEATURE

The station is on air right now.

Nothing to scroll, nothing to pick. Tune in and hear what the DJ is playing, or stand up your own frequency from the source.

► BROWSE THE STATIONS

RUN YOUR OWN STATION →

THE BACK PAGES

READER SERVICES · §§ 01-05

№ 01

Dispatches

Field notes, release wires & tales from the press desk.

[READ THE DISPATCHES →](#)

01
THE WIRE

№ 02

Stations

Other SUB/WAVEs broadcasting right now. Spin the dial.

[BROWSE THE STATIONS →](#)

02
THE DIAL

№ 03

Community Skills

Segments operators taught their DJs — take one home.

[EXPLORE THE SKILLS →](#)

03
THE EXCHANGE

№ 04

Community Personas

DJs other operators dreamed up — book one for your booth.

[MEET THE PERSONAS →](#)

04
THE GREEN ROOM

№ 05

Community Shows

Produced slots operators built — put one on your grid.

05
THE PROGRAMME GUIDE

BROWSE THE SHOWS →

SUB/WAVE · EST. 2026

NAVIDROME · LIQUIDSOAP · ICECAST · YOUR LLM · YOUR VOICE

GITHUB · DISCORD · X · REDDIT

OPEN THE PLAYER · PRIVACY · TERMS

— 30 —

SET IN TYPE & SENT TO PRESS BY THE KLAIR WORKS ♦ KLAIR.CO