

Документация DigitalOcean

Узнайте, как создавать, развертывать и масштабировать свои приложения с помощью DigitalOcean. Ознакомьтесь с нашими продуктами с помощью технических описаний в документации, примеров кода, справочной информации о наших API, интерфейсе командной строки, клиентских библиотеках и многом другом.

Начало работы



Обзор платформы DigitalOcean

Информация о платформе DigitalOcean, такая как платежные реквизиты, примечания к выпуску, доступность продукта в зависимости от центра обработки данных, планы поддержки и данные учетной записи.



Главная Страница продукта

Информация о функциях, ценах, доступности и ограничениях продуктов DigitalOcean; о том, как использовать продукты через панель управления; об управлении аккаунтом, командами и выставлением счетов; а также сведения о платформе, примечания к выпуску и правила использования продуктов.



Эталонный Дом

Программное управление ресурсами и интеграция с экосистемой разработчиков с помощью интерфейсов командной строки, API и SDK.



Поддержка Дома

Ищете техническую поддержку для своего аккаунта или инфраструктуры DigitalOcean? Просмотрите разделы, посвященные продуктам, или создайте заявку на портале поддержки.

Просмотр по

Жанр продукта



Облако DigitalOcean с поддержкой искусственного интеллекта

Создавайте, обучайте и внедряйте ИИ-агентов с помощью DigitalOcean AI-Native Cloud.



Вычислить

Создавайте приложения так, как вам удобно, с помощью нашего набора вычислительных продуктов, включая виртуальные машины, управляемые контейнеры, PaaS и бессерверные функции.



Хранение

Надежно храните любые объемы данных в облаке и получайте к ним доступ с помощью объектного хранилища Spaces, совместимого с S3, блочного хранилища Volumes или сетевого файлового хранилища на основе NFS.



Контейнеры и изображения

Создавайте резервные копии, загружайте пользовательские образы, используйте предварительно настроенные образы для создания ресурсов и храните образы Docker в частном реестре.



Управляемые базы данных

Запускайте полностью управляемые кластеры баз данных с выбранной вами СУБД и избавьтесь от необходимости выполнять настройку и обслуживание вручную.



Создание сетей

Защитите трафик своих приложений и контролируйте его с помощью VPC, фильтрации трафика и балансировки нагрузки.



Руководство

Отслеживайте работоспособность своей инфраструктуры, URL-адресов и т. д., настраивайте оповещения, чтобы быть в курсе событий, и распределяйте ресурсы по проектам.



Службы передачи данных

Запускайте базы данных приложений, поддерживайте поиск по сходству и рабочие процессы с использованием искусственного интеллекта, а также создавайте возможности для расширенной генерации (RAG) на основе собственного контента.

Средства разработки



doctl Интерфейс командной строки (CLI)

Управляйте ресурсами DigitalOcean из командной строки с помощью doctl — нашего интерфейса командной строки с открытым исходным кодом (CLI).



Обзор API

Программное управление виртуальными машинами, пространствами и другими ресурсами DigitalOcean с помощью обычных HTTP-запросов. Используйте RESTful API для программного управления виртуальными машинами, пространствами и другими ресурсами DigitalOcean.



Paperspace APIs, CLIs, and SDKs Overview

Interact with Paperspace resources programmatically using the Paperspace API or CLI, and find documentation for legacy tools.



Ansible

Automate DigitalOcean infrastructure and configuration management using the open source Ansible framework.



Terraform

Deploy and change many resources simultaneously using the open source Terraform tool.



PyDo

Official Python client for the DigitalOcean API (OpenAPIv3). Install with pip, authenticate with a personal access token, and call API operations via `pydo.Client`.



SDK and Client Libraries

Official and community client libraries for the DigitalOcean API, with installation instructions and quickstart examples.



Open-Source Software

We use and contribute to open source software.



DigitalOcean MCP Servers

Use MCP servers to manage DigitalOcean services from any MCP-compatible client.

Latest Updates

Upcoming Changes 3



- Managed MySQL 8.0 database clusters are scheduled for forced upgrades to MySQL 8.4 beginning 30 October 2026 during [each cluster's maintenance window](#).
- App Platform's XL build resources (8 CPUs and 20 GiB of memory during builds) are now enabled for all apps by default. The `xl-build` flag is now deprecated and will be removed in a future release. Remove `xl-build` from your [app spec](#) to avoid potential errors once the flag is fully retired.
- [DigitalOcean Managed Caching](#) is being discontinued on 30 June 2025. To replace Managed Caching, we are offering [Managed Valkey](#), a Redis-compatible alternative with RDMA and higher throughput. All existing Managed Caching clusters automatically convert to Valkey clusters by 30 June 2025 during your [upgrade window](#), retaining all data.

June 2026 29



30 JUNE

- [Private Droplets](#) are now generally available in all regions. Private Droplets have no public network interface and no public IP address, using VPC-only networking with automatic integration with VPC NAT gateway, VPC peering, and VPC private DNS. See the [Private Droplets documentation](#) for setup instructions and limitations.
- The following agent evaluation metrics are deprecated and should no longer be used:
 - Tone
 - Retrieved Chunk Usage
 - Prompt Perplexity

Use the currently supported metrics listed in [Agent Evaluation Metrics](#) instead. To monitor deployed agent behavior outside of evaluations, use [Agent Metrics and Runtime Logs](#).

- The following Anthropic model is now available on DigitalOcean Inference for [serverless inference](#), [Agent Development Kit](#), and [agents](#):
 - [Claude Sonnet 5](#)

For more information, see the [Available Models page](#).

- Agent evaluations support for the Agent Development Kit (ADK), previously in preview, is now removed. To evaluate agents, use [agent evaluations via the DigitalOcean Control Panel](#) for supported agent types. To monitor ADK agent behavior, use [Agent Metrics and Runtime Logs](#).
- Insights, agent tracing, and conversation logs are deprecated for all agents, including agents created through the Control Panel, CLI, API, and Agent Development Kit (ADK). To monitor deployed agent behavior, use [Agent Metrics and Runtime Logs](#) instead.
- The Agent Evaluations MCP server tool has been renamed to [Evaluations](#).
- Custom metrics are now available for [DigitalOcean Evaluations](#). You can define your own metrics to evaluate model behavior against criteria specific to your use case.
- [DigitalOcean Evaluations](#) is now generally available. Use Evaluations to create test cases, run evaluation datasets, and measure model performance against selected metrics.
- Presets are now available for [DigitalOcean Evaluations](#). You can save and reuse evaluation configurations, including the candidate model, system prompt, hyperparameters, judge model, and metrics.
- Model Evaluations is now renamed to [DigitalOcean Evaluations](#).

29 JUNE

- Serverless Inference now requires a positive prepaid account balance before you can send inference requests. Usage charges are deducted from this balance, and access is suspended if it reaches \$0. You can add a prepayment manually or enable auto-reload to replenish your balance automatically. For more information, see [Manage Serverless Inference Prepayment](#).
- The DigitalOcean Control Panel now supports light and dark themes. From the profile menu in the top right corner of the control panel, you can [set your theme](#) to a light or dark appearance, or match your operating system's appearance setting.

25 JUNE

- You can now attach [Network File Storage](#) shares to multiple VPC networks, and expose specific subdirectories to single VPC networks by creating access points. A share can connect to up to 10 VPC networks in the same region. Access points restrict clients to a specific path within the share, and are isolated from one another so that clients on one cannot see the directories governed by another. This lets you host multiple tenants on a single share, each scoped to their own directory. For details, see [How to Create and Delete Network File Storage Access Points](#).

24 JUNE

- [Single sign-on \(SSO\) with OIDC](#) for DigitalOcean Kubernetes (DOKS) clusters is now in [general availability](#). You can authenticate users to your Kubernetes clusters through an identity provider like Auth0, authentik, JumpCloud, Keycloak, or Okta, instead of using token-based authentication. SSO is configured per cluster with an issuer URL and client ID from your identity provider. You can enable it using [doctl](#), the [DigitalOcean API](#), or [Terraform](#).
- The following Z.ai model is now available on DigitalOcean Inference for [serverless inference](#), [dedicated inference](#), [Agent Development Kit](#), and [agents](#):
 - [GLM-5.1](#)

For more information, see the [Available Models page](#).

- A [Fedora 44](#) (fedora-44-x64) Droplet base image is now available in the Control Panel and through the API.

22 JUNE

- The following Z.ai model is now available on DigitalOcean Inference for [serverless inference](#), [dedicated inference](#), [Agent Development Kit](#), and [agents](#):
 - [GLM-5.2](#)

For more information, see the [Available Models page](#).

18 JUNE

- To comply with [Peru's tax regulation for digital services](#), DigitalOcean charges Value Added Tax (VAT) for customers with a tax location in Peru, beginning on 1 June 2026. These charges appear on invoices issued on and after 1 July 2026. [Learn more about taxes in Peru](#).

17 JUNE

- The following Xiaomi model is now available on DigitalOcean Inference for [serverless inference](#), [dedicated inference](#), [Agent Development Kit](#), and [agents](#):
 - [MiMo-V2.5-Pro](#)

For more information, see the [Available Models page](#).

- DigitalOcean Inference supports server-side tools on serverless inference, dedicated inference, and inference routers. You can add the following tools:
 - Web search, web fetch, knowledge base retrieval, and remote MCP server tools to your requests in the [Chat Completions](#) and [Responses](#) APIs.
 - Provider-native tools such as bash, text editor, computer use, and web fetch for Anthropic models with the [Messages API](#).
 - Function calling and tool search for OpenAI models on the Responses API, and Anthropic models on the Messages API.

Web search and web fetch are in [public preview](#). For more information, see [Use Server-Side Tools](#).

- Updated CentOS Stream 9 and CentOS Stream 10 (`centos-stream-9-x64` , `centos-stream-10-x64`) Droplet base images require a minimum Droplet size of `s-1vcpu-1gb` . You cannot create Droplets with these images on the `s-1vcpu-512mb-10gb` plan because the image exceeds the available disk size for that plan. For more information, see [Linux Images for Droplets](#).

15 JUNE

- The following Anthropic models are deprecated from DigitalOcean Inference as of 15 June 2026:
 - [Claude Opus 4](#)
 - [Claude Sonnet 4](#)

Migrate to [Claude Opus 4.8](#) (`anthropic-claude-opus-4.8`) and [Claude Sonnet 4.6](#) (`anthropic-claude-4.6-sonnet`), respectively, to avoid service disruption. For information on our model deprecation policy and recommended replacement models, see [Model Support Policy](#).

12 JUNE

- The following Xiaomi model is now available on DigitalOcean Inference for [serverless inference](#), [Agent Development Kit](#), and [agents](#):
 - [MiMo-V2.5](#)

For more information, see the [Available Models page](#).

- Claude Fable 5 is no longer available on DigitalOcean Inference. Access to all other Anthropic models remains available. For more information, see [the Anthropic statement on Claude Fable 5](#).

10 JUNE

- We support passthrough tool search on the [Messages API](#) for Anthropic models and the [Responses API](#) for OpenAI models, enabling deferred loading of tools in agentic workflows. There is no additional cost to using tool search. For more information, see [Use Server-Side Tools](#).
- Debian 12 [reached end of life](#) on 10 June 2026. Per our [image deprecation policy](#), this image is available exclusively via the API for the next 30 days before we remove it from our platform.

9 JUNE

- The following Anthropic model is now available on DigitalOcean Inference for [serverless inference](#), [dedicated inference](#), [Agent Development Kit](#), and [agents](#):
 - [Claude Fable 5](#)

For more information, see the [Available Models page](#).

4 JUNE

- The following NVIDIA model is now available on DigitalOcean Inference for [serverless inference](#), [dedicated inference](#), [Agent Development Kit](#), and [agents](#):
 - [Nemotron 3 Ultra](#)

For more information, see the [Available Models page](#).

3 JUNE

- The public endpoint when creating a dedicated inference deployment is now disabled by default. For more information on how to enable access, see [Use Dedicated Inference Endpoint](#).

May 2026 11 ▼

April 2026 63 ▼

March 2026 24 ▼

February 2026 25 ▼

January 2026 12 ▼

[View All Release Notes](#) →

COMPANY

[About](#)
[Careers](#)
[Blog](#)

DOCS

[Docs Home](#)
[API Reference](#)
[CLI Reference](#)
[Release Notes](#)
[llms.txt](#)
[Trust Platform](#)

COMMUNITY

[Tutorials](#)
[Q&A](#)
[Write for DOnations](#)
[Currents Research](#)
[Legal](#)
[Code of Conduct](#)

SUPPORT

[Support Center](#)
[Report Abuse](#)



© 2026 DigitalOcean, LLC. All rights reserved [Cookie Preferences](#)

