

ИИ / ОПЕРАЦИИ С ИСПОЛЬЗОВАНИЕМ ИИ / FINOPS / БОЛЬШИЕ ЯЗЫКОВЫЕ МОДЕЛИ

«Токенамаксинг — это реально, дорого и это распространяется»: бюджеты на ИИ стремительно растут

Lanai Token Tuner помогает предприятиям сократить расходы на искусственный интеллект, сопоставляя затраты на токены с рабочими процессами и определяя, где более дешевые модели могут заменить премиальные.

27 мая 2026 года, 13:27 [Адриан Бриджуотер](#)




Фото Compagnons на Unsplash

В борьбе с токенамаксингом появилось новое оружие.

Tokenmaxxing, конечно же, происходит, когда компания решает, что использование токенов ИИ равнозначно продуктивности. Но использование токенов может быстро превратиться в показатель тщеславия, и компания, которая считает чрезмерное использование токенов прямым показателем

Какое-то время токенмаксинг был на пике популярности, но, похоже, здравый смысл берет верх, и фокус смещается на результат, а не на использование ИИ ради самого процесса. Возьмем недавний случай токенмаксинга в Uber:

Технический директор Uber Неппалли Нага в прошлом месяце заявил *The Information*, что ему «пришлось вернуться к чертежной доске, потому что бюджет, который, как [он] думал, [ему] понадобится, уже исчерпан». Этот бюджет был выделен на использование Uber кода Anthropic Claude.

Со своей стороны, главный операционный директор Uber Эндрю Макдональд ответил через несколько недель, заявив в *интервью Rapid Response*, о котором впервые сообщило издание *Business Insider*, что комментарии Наги о превышении бюджета на Claude стали «последней каплей» для операционной команды.

«Все подумали: “О, момент, когда у меня взорвалась голова”», — сказал Макдональд. «Нам придется начать говорить о потреблении токенов и связанных с этим расходах в сравнении с численностью персонала, а также принимать решения как инженерная организация.

«Если вы не можете напрямую связать это с тем, сколько полезных функций и возможностей вы предоставляете своим пользователям, такую сделку может быть сложнее оправдать».

ТРЕНДОВЫЕ ИСТОРИИ

1. [The reason enterprise outages almost never start where ops teams think](#)
2. [Why OpenAI and Anthropic are hiring forward deployed engineer teams](#)
3. [Google finally builds the AI and agent platform it's been describing for years](#)
4. [Why agentic AI makes the ops platform the most important layer in the enterprise](#)

Лекси Риз, соучредитель и генеральный директор компании Lanai, подчеркивает, что эта проблема существует повсеместно. Uber — лишь одна из последних крупных компаний, столкнувшихся с ней.

«Tokenmaxxing — это реальная проблема, она дорого обходится и затрагивает не только нескольких инженеров или компании», — говорит Риз в интервью *The New Stack*.

Избыточное использование токенов может привести к появлению раздутого кода, разрастанию агентных систем и областей, в которых программные приложения со временем могут стать хрупкими или даже уязвимыми. Кроме того, это дорого и снижает прозрачность общего состояния системы.

Lanai, an AI accountability company, aims to help enterprises understand where AI spend occurs, which workflows AI is applied to, and at what cost.

The company recently debuted **Token Tuner** to identify where lower-cost models can reduce unnecessary token costs. It's the latest tool that developers and leaders can use to control token usage by engineers and end users. The internet is full of top-ten lists for how to reduce token usage. Companies and organizations like **Kong**, **Braintrust**, **LiteLLM**, and **Dynatrace**, among others, offer tools to ensure token usage is being budgeted.

“Tokenmaxxing is real, it's expensive and it's spreading beyond just a few engineers or companies.”

Reese and team have positioned Token Tuner as a service that fills the missing context gap for enterprises by **mapping token spend to workflows**, model choices, efficiency, and value created. The software ties each AI interaction to a measurable outcome and generates a productivity score based on how well each user matched token usage and model choice to the task at hand.

For example, an employee using **Opus 4.7** for email responses is likely to receive a lower efficiency score than if they used a smaller model for the task.

outcomemaxxing to analyze which workflows are actually improving productivity.

Currently in beta, one Lanai Token Tuner user delegated 4.2% of all AI leverage hours across the organization while using only 0.7% of tokens. Their efficiency score was 6.0, indicating they were matching tasks to the right models, while others were burning 10x as many tokens for half the efficiency.

Lanai Chief Product Officer **Mohit Mehta** tells *The New Stack* that Token Tuner is an all-terrain vehicle, i.e., its scoring engine can calculate productivity scores when a single workflow spans multiple models simultaneously.

“Productivity is estimated by the complexity of work delegated to AI as observed through prompt and tool activity by Lanai’s proprietary models,” says Mehta. “The model operates at the level of prompts and tool invocations independent of models and applications.”

Tracking AI usage for business tasks

As we start to place greater emphasis on business results from applied technology deployments (even politicians have begun using the term “**measurable outcomes**” in recent times), we need to question which instrumentation is required at the API layer for Token Tuner to attribute tokens to specific business outcomes.

“Lanai aggregates prompt interactions and associated tool activity for a given session and then runs proprietary models to calculate the task type and associated productivity gain and complexity,” explains Mehta. “This enables customers to go from contextless vendor invoice to connecting intent to value to cost at the interaction level. No custom instrumentation is required for this functionality.”

“Rather than relying on synthetic evaluations, we utilize observed outcome data. Our recommendations are grounded in how actual users within an organization achieve comparable results across different models.”

place to assess output quality equivalence before surfacing the recommendation:

“Rather than relying on synthetic evaluations, we utilize observed outcome data,” clarifies Mehta. “Our recommendations are grounded in how actual users within an organization achieve comparable results across different models.

“Rather than a recommendation like ‘this will work for you,’ we provide empirical evidence that ‘teams in your company performed this exact workflow on Haiku with equal success,’ for example. This represents real-world preference at scale over synthetic benchmarks.”

Key features include workflow-level value visibility, a service that shows which teams, workflows, and use cases are driving AI spend and whether that usage is tied to measurable business value.

Productivity and efficiency measurement compares token spend with the leverage gained by users, teams, and workflows to show where AI creates the most value per dollar. A spend optimization recommendation function identifies runaway workflows, mismatched tasks, and premium model usage for work that lower-cost models could handle.

AI's next killer service: efficiency?

First, the Earth cooled, and we just wanted AI... and the plain old predictive version was fine. Then, the dinosaurs died off, and we wanted domain-specific RAG-based intelligence with what then became agentic AI services that could work for us with human-in-the-loop oversight to ward off the rise of the robots. Now, perhaps, we want AI that is fit for purpose in the most applied sense of the term, so that we don't use it where we don't need to, and we use high-octane services only when we can really justify the turbocharge.

In truth, AI's next killer app factor will come down to a whole lot more than just business efficiency, but this could become a more prevalent part of the mix.

TNS



Adrian Bridgwater is a technology journalist with three decades of press experience. He has an extensive background in communications, starting in print media, newspapers and also

TNS owner Insight Partners is an investor in: Anthropic.

SHARE THIS STORY



TRENDING STORIES

1. The reason enterprise outages almost never start where ops teams think
2. Why OpenAI and Anthropic are hiring forward deployed engineer teams
3. Google finally builds the AI and agent platform it's been describing for years
4. Why agentic AI makes the ops platform the most important layer in the enterprise
5. "Tokenmaxxing is real, expensive & it's spreading": AI budgets are exploding

INSIGHTS FROM OUR SPONSORS



chainguard

Everything we announced during AI Readiness Innovation Week

25 June 2026

Fewer CVEs, more accurate findings: Wiz now scans Chainguard Libraries for Python and Java

25 June 2026

Chainguard Repository adds new policies, Chainguard Libraries for JavaScript is GA

June 2026

Moving from Probabilistic Reasoning to Deterministic Execution

24 June 2026

Kong Mesh 2.14: Finer Zone Proxy Control and Tighter Security

23 June 2026

Why We Need to Stop Prompt Hacking

23 June 2026



Kubeshop™

AI Test Automation: Fix the 3 New Testing Bottlenecks

17 June 2026

Test Observability: Fix Flaky Kubernetes Tests (2026)

11 June 2026

Testkube AI Launch: New Use Cases, OAuth, and Open Source Analysis

9 June 2026

TNS DAILY NEWSLETTER

**Receive a free roundup of the
most recent TNS articles in
your inbox each day.**

svo1975pnv1982@gmail.com

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

FOLLOW TNS



© The New Stack 2026

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)