

АГЕНТЫ ИИ / РАЗРАБОТКА ИИ / РАЗРАБОТКА ПРОГРАММНОГО ОБЕСПЕЧЕНИЯ

ИИ может написать код. Но ответственность по-прежнему лежит на вашей команде.

Почему непрерывная многоуровневая верификация становится тем фактором, который определяет, принесет ли ИИ-кодирование выгоду или приведет к увеличению расходов на обслуживание.

23 июня 2026 г. в 12:00 [Роберт Керли](#)



Компания Puzzle Creative

[Sonar](#) выступила спонсором этого поста. Insight Partners инвестирует в Sonar и TNS.

R индустрии программного обеспечения об искусственном интеллекте до сих пор говорят в основном с точки зрения скорости. Модели и агенты теперь

обеспечения. Но скорость — это не то же самое, что контроль. Сейчас в разговорах об искусственном интеллекте в программировании все чаще поднимается вопрос о том, могут ли команды проверять код с той же тщательностью и последовательностью, с какой они его генерируют.

Именно здесь меняется экономика технического долга. ИИ удешевляет создание долга и удорожает его последующее обнаружение, резко увеличивая количество проблем в коде. Результат работы ИИ может выглядеть корректно и проходить модульные тесты, но при этом не соответствовать архитектурному контексту, стандартам кодирования и требованиям к ремонтпригодности, которые обеспечивают устойчивость программного обеспечения в долгосрочной перспективе.

В результате происходит перераспределение затрат: на создание кода тратится меньше усилий, но больше внимания уделяется проверке, анализу и исправлению ошибок. По сути, искусственный интеллект превратил разработчиков из кодировщиков в верификаторов кода. Но мы обрекаем наши команды на провал, потому что объем кода, который нужно проверить, огромен.

Куда смещаются затраты

Этот сдвиг имеет значение, потому что технический долг был дорогостоящим и до того, как его усугубил искусственный интеллект. По некоторым оценкам, ежегодный ущерб от технического долга в США составляет 1,5 триллиона долларов. В целом проблема охватывает две взаимосвязанные категории: долг на уровне кода, то есть ошибки, уязвимости и «запахи кода», и архитектурный долг, из-за которого системы становятся хрупкими, запутанными и трудно поддающимися развитию.

Вторая категория заслуживает большего внимания, чем ей обычно уделяют. По прогнозам Gartner, к 2027 году архитектурный технический долг составит 80% от всего технического долга. В этом контексте Sonar был признан лидером в магическом квадранте Gartner® 2026™ для инструментов управления техническим долгом, заняв первое место по способности к реализации. Кроме того, некоторые отраслевые исследования показывают, что архитектурный долг накапливается гораздо быстрее, чем долг на уровне кода, поскольку ущерб носит системный, а не локальный характер.

Продуманная функция может замедлить работу разработчика, а дрейф архитектуры — работу всей организации.

Sonar is the industry standard for code verification and automated code review, trusted by 75% of the Fortune 100. Its SonarQube platform analyzes over 750 billion lines of code daily, helping to prevent outages, reduce risk, lower technical debt, and ensure compliance.

[Learn More →](#)

THE LATEST FROM SONAR

[Introducing Sonar Vortex and the SonarQube Remediation Agent](#)

30 June 2026

[Cut your coding agent's cost with Sonar Vortex](#)

30 June 2026

[Model-agnostic AI is key to business continuity as models keep changing](#)

26 June 2026

HEAR MORE FROM OUR SPONSOR

svo1975pvn1982@gmail.com

Submit

“A messy function can slow a developer down; architectural drift can slow an entire organization down.”

This is the real risk with AI-assisted development. AI can produce plausible code at a volume and velocity that overwhelms the practices teams have traditionally relied on to maintain trust. Human review, periodic audits, and retrospectives were built for a world where humans wrote code in smaller increments and at human pace. Those approaches are no longer sufficient in an agentic development environment, where governance must be applied on an ongoing basis as the volume of code increases.

State of Code Developer Survey report, 96% of developers said they distrust AI-generated code, yet only 48% said they consistently verify it. The same research found that 88% reported at least one negative impact of AI on technical debt, and 38% said reviewing AI-generated code takes more effort than reviewing human-written code. That combination should concern engineering leaders. It suggests organizations are increasing software throughput faster than they are increasing confidence.

TRENDING STORIES

1. "Code should be regenerated, not maintained": Codeplain makes the case for spec-driven development
2. Your engineering org needs an AI slop registry
3. Loops are replacing prompts. Verification is about to be your biggest problem.
4. Tutorial: Building a Real-time Flight Tracker with Gemini Function Calling Capability
5. Small Language Models Are Driving a New Wave of Edge Native AI

"96% of developers said they distrust AI-generated code, yet only 48% said they consistently verify it."

The way teams develop code has changed in the agentic era. Verifying code at the pull request (PR) means going back to adjust prompts, re-prompting, and incurring additional token spend to regenerate code. Teams need to stop treating quality and maintainability as something assessed only after a PR is opened. Standards need to shape agents before code is generated, during the code generation loop, and again before it's merged. In short, **verification has to become** continuous and multilayered.

Guide, verify, solve

1. **Guide:** AI systems are only as good as the context they receive. Agents need more than a prompt; they need architectural context, coding standards, and project-specific details. Guidance shifts quality enforcement from post-generation to upfront direction, helping prevent and reduce debt from the initial prompt.
2. **Verify:** Verification is where trust is earned. Teams need deterministic, in-workflow verification that catches issues generated by agents before the code reaches the branch or PR. This is the key shift from the older model of technical debt management. Multi-layered verification ensures that the generated code meets **coding and compliance standards** before it proceeds through the CI/CD pipeline.
3. **Solve:** Detection without automated remediation means the onus is still on developers to resolve problems that are piling up faster than they can manage. Issues found during verification need to be automatically fixed, rechecked, and fed back into a new PR. Otherwise, verification becomes a reporting mechanism rather than an operational discipline.

Technical debt must increasingly be treated as a business liability, not a cleanup task developers occasionally squeeze in between roadmap items. Without proactive management of technical debt at the pace of AI, refactoring projects will begin to outpace the building of value-driven features. AI code slop is real, and the newer, improved LLMs are **often more verbose** than older, less accurate models.

In the AI era, the winners will not be the teams that move fastest at generation. They will be the teams that pair that speed with continuous multilayer verification, so today's output does not become tomorrow's technical debt nightmare.

TNS



Robert Curlee is a software product management leader with 18 years of experience working in the Enterprise IT Services Management, IT Operations, Storage, and Security markets. Having started his career in software engineering, he currently serves as Product Marketing Manager...

[Read more from Robert Curlee →](#)

TRENDING STORIES

1. "Code should be regenerated, not maintained": Codeplain makes the case for spec-driven development
2. Your engineering org needs an AI slop registry
3. Loops are replacing prompts. Verification is about to be your biggest problem.
4. Tutorial: Building a Real-time Flight Tracker with Gemini Function Calling Capability
5. Small Language Models Are Driving a New Wave of Edge Native AI

INSIGHTS FROM OUR SPONSOR



Sonar is the industry standard for code verification and automated code review, trusted by 75% of the Fortune 100. Its SonarQube platform analyzes over 750 billion lines of code daily, helping to prevent outages, reduce risk, lower technical debt, and ensure compliance.

[Learn More →](#)

Introducing Sonar Vortex and the SonarQube Remediation Agent

30 June 2026

Cut your coding agent's cost with Sonar Vortex

30 June 2026

Model-agnostic AI is key to business continuity as models keep changing

30 June 2026

Now available: SonarQube plugin for Codex

24 June 2026

Now available: SonarQube plugin for Cursor

24 June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the most recent TNS articles in your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing

ENGINEERING

AI
AI Engineering
API Management
Backend development

MICROSERVICES

Open Source
Networking
Storage

Large Language Models

Security
Software Development
WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

 roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



© The New Stack 2026

Disclosures
Terms of Use
Advertising Terms & Conditions
Privacy Policy

