


ИИ / АГЕНТЫ ИСКУССТВЕННОГО ИНТЕЛЛЕКТА / МОДЕЛИ С ИСКУССТВЕННЫМ ИНТЕЛЛЕКТОМ

Anthropic Sonnet 5: он сокращает отставание от Opus 4.8 и будет дешевым до августа

Anthropic's Claude Sonnet 5 приближается к Opus 4.8 в агентских бенчмарках и поставляется с вводной ценой API в размере \$ 2 / \$ 10 за миллион токенов до августа.

30 июня 2026 г. в 14:00 [Фредерик Лардинуа](#)



Во вторник компания Anthropic **выпустила Sonnet 5**, последнюю модель в своей основной серии Sonnet.

Компания называет Sonnet 5 «самой агентной моделью Sonnet на сегодняшний день». По сравнению с эталонными показателями ее производительность близка к **Opus 4.8** и значительно превосходит **Sonnet 4.6**. В Anthropic особо отмечают, что модель гораздо лучше справляется с задачами, требующими логического мышления, использования инструментов, написания программного кода и работы со знаниями.

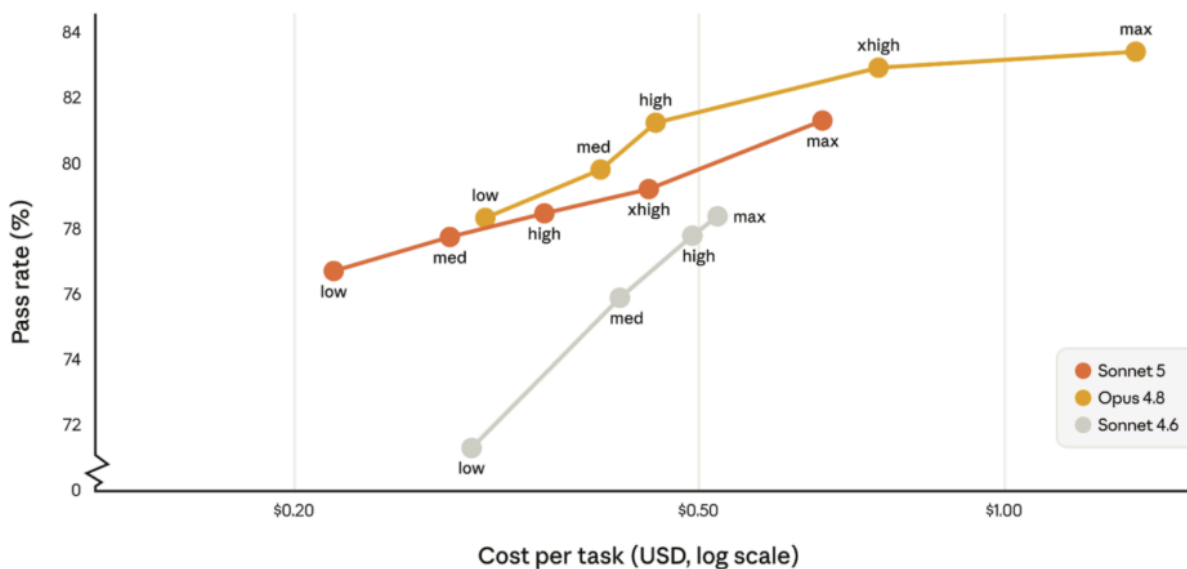
но она достаточно близка к ней, чтобы на данный момент стать более доступной альтернативой Opus 4.8 — ведь Opus 5 не может быть далеко позади (при условии, что его не **ограничивают так же, как Fable 5**).

Anthropic подчеркивает, что Opus 4.8, особенно на более высоких уровнях логического вывода, все же обеспечивает более высокую точность. Но также отмечает, что «Sonnet 5 предоставляет разработчикам более доступные варианты, которые по качеству значительно превосходят то, что было доступно ранее».

	Sonnet 5	Sonnet 4.6	Opus 4.8 For reference
Agentic coding SWE-bench Pro	63.2%	58.1%	69.2%
Agentic coding Terminal-Bench 2.1	80.4%	67.0%	82.7%
Multidisciplinary reasoning Humanity's Last Exam	43.2% no tools	34.6% no tools	49.8% no tools
	57.4% with tools	46.8% with tools	57.9% with tools
Computer use OSWorld-Verified	81.2%	78.5%	83.4%
Knowledge work GDPval-AA v2	1618	1395	1615

Тесты Sonnet 5. Фото: Anthropic

При максимальном уровне логического вывода Extra High Sonnet 5 показывает примерно такие же результаты, как Opus 4.8 на средних и высоких настройках в тестах BrowseComp, разработанных OSWorld и агентном поиске. Но поскольку этот уровень логического вывода требует больше ресурсов, чем Opus 4.8 на сопоставимом уровне, для некоторых задач лучше подойдет Opus 4.8.



Источник: Anthropic

Sonnet 5, по крайней мере в рамках бенчмарка, доступного на данный момент, всегда превосходит Sonnet 4.6.

Однако бенчмарки — это лишь часть общей картины. Поведение модели также влияет на то, как ее воспринимают пользователи. В Anthropic говорят, что тестировщики заметили, что модель теперь часто справляется со сложными задачами, «например, там, где предыдущие версии Sonnet пасовали».

TRENDING STORIES

1. OpenClaw and Hermes agree on what an agent is. They disagree on what controls it.
2. AWS launches a desktop for agents
3. Anthropic Sonnet 5: It closes the gap with Opus 4.8, and is cheap until August
4. Developers are now validating code they didn't write — and may not understand
5. The US government just told OpenAI who's allowed to use the next GPT 5.6 model

capacity from running Opus 4.8), Anthropic is offering an introductory price for API users at \$2 per million input tokens and \$10 per million output tokens until August 31. It will go up to \$3/\$15 per million input/output tokens, as Anthropic has previously charged for Sonnet models.

When asked if this is the first time Anthropic has offered introductory pricing, an spokesperson told *The New Stack*: “We want our customers to test Sonnet 5 against their real workloads at the lowest possible cost during the migration window.”

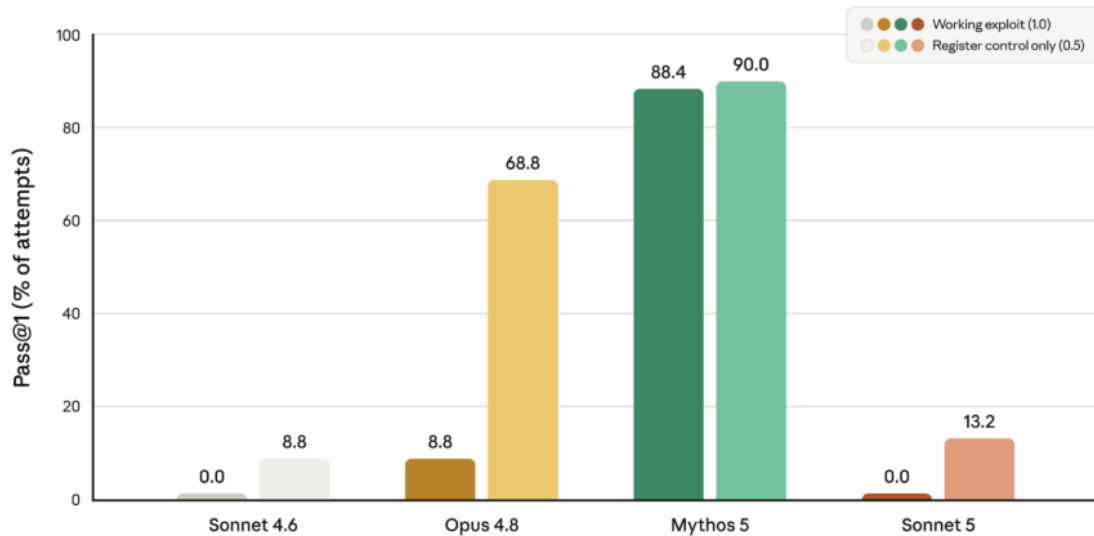
For the time being, Anthropic has also increased rate limits for Chat, Cowork, and Claude Code users to, as the company puts it, “to accommodate the higher token usage of higher effort levels.”

Sonnet 5 safety

In terms of AI safety, which is obviously a focus for Anthropic, especially after Fable 5’s unceremonious departure, Anthropic notes that it didn’t “deliberately train Sonnet 5 on cybersecurity tasks” and that while it can handle some routing cyber tasks, its performance on these tests is far behind Opus 4.8 and, of course, Mythos. Still, Anthropic kept its cyber safeguards on for this model, but since the risk is low, those safeguards are not as strict as those for models like Fable 5.

The company explicitly notes that when trying to find exploits in **Firefox 147**, for example, “Sonnet 5 was never able to develop a full working exploit, but it does show a slightly higher rate of partial success than its predecessor, Sonnet 4.6.”

The risk of the U.S. government pulling Sonnet 5 out of circulation is quite low, then, it seems.



Credit: Anthropic

TNS



Before joining The New Stack as its senior editor for AI, Frederic was the enterprise editor at TechCrunch, where he covered everything from the rise of the cloud and the earliest days of Kubernetes to the advent of quantum computing....

[Read more from Frederic Lardinois →](#)

TNS owner Insight Partners is an investor in: Anthropic.

SHARE THIS STORY



TRENDING STORIES

1. OpenClaw and Hermes agree on what an agent is. They disagree on what controls it.
2. AWS launches a desktop for agents
3. Anthropic Sonnet 5: It closes the gap with Opus 4.8, and is cheap until August
 - Developers are now validating code they didn't write — and may not understand

INSIGHTS FROM OUR SPONSORS



A 15-Second Health Check for Your Heroku Connect Data Pipeline

24 June 2026

Heroku March 2026 Update

19 March 2026

Modernizing the Command Line: Heroku CLI v11

19 March 2026



Kepler, re-architected: Improved power accuracy and a community call to action!

30 June 2026

Dragonfly v2.5.0 is released

30 June 2026

OTel and mesh-derived metrics: A 2026 reference

29 June 2026



Securing AI agents: When AI tools move from reading to acting

30 June 2026

Small and medium businesses aren't waiting for an AI invitation — they're already leading

^^ June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in
your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions

 roadmap.sh

Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



© The New Stack 2026

Disclosures
Terms of Use
Advertising Terms & Conditions
Privacy Policy
Cookie Policy

