

ИИ / АГЕНТЫ ИСКУССТВЕННОГО ИНТЕЛЛЕКТА

Anthropic предоставляет @Claude постоянное место в ваших каналах Slack

Claude Tag позволяет корпоративным командам использовать в Slack многопользовательский ИИ, который работает под собственным именем.

23 июня 2026, 14:50 автор: [Пол Созэрс](#) и [Фредерик Лардинуа](#)



Мохаммеду Нохасси за незапятнанность+

Anthropic **анонсировала Claude Tag**, новый продукт, который встраивает Claude непосредственно в Slack в качестве постоянного члена совместной команды — такого, который со временем накапливает институциональные знания, работает асинхронно и может действовать без подсказок.


Продукт, доступный в бета-версии для клиентов с тарифными планами Enterprise и Team, заменяет существующее приложение Claude от Anthropic в Slack, которое **дебютировало в октябре 2025 года**. Эта первоначальная интеграция позволяла пользователям отправлять личные сообщения Claude, вызывать его через панель ИИ-помощника или пометать его в каналах для получения помощи по запросу — это было полезно, но несколько

сессию программирования в веб-версии, а по завершении опубликовать пул-реквест в трее.

В этом смысле Claude Tag делает для Slack примерно то же, что Cowork сделал для работников умственного труда: он берет агентский цикл, зародившийся в Claude Code, и распространяет его на новую среду и более широкий спектр задач, а не ограничивает его инженерными рабочими процессами.

Claude Tag во многом является следующим шагом: теперь Claude не присоединяется к разговору по запросу, а остается в канале постоянно. Задачи можно запланировать на несколько часов или дней, и Claude будет выполнять их самостоятельно — накапливая контекст, изучая терминологию работы и подбирая незавершенные нити, которые в противном случае остались бы незамеченными.

Этот контекст не обязательно ограничивается одним каналом: если администраторы предоставят доступ, Claude сможет использовать информацию из других каналов для работы с ними.



Claude Tag в действии

Чувство идентичности

каждого подключенного инструмента, он публикует сообщения в Slack от имени приложения Claude, открывает пул-реквесты от имени собственного приложения на GitHub или запрашивает данные из хранилища под учетной записью службы, созданной администратором. Если включено «фоновое» поведение, Claude вообще не нужно помечать тегами — он отслеживает обсуждения, отмечает то, что, по его мнению, должно быть известно команде, и следит за неактивными обсуждениями, как и в режиме «авто», уже знакомом пользователям Claude Code.


Beyond scoping tool and data access, administrators can also set a ceiling on how many tokens each channel's Claude is allowed to spend; it's a guardrail against an autonomous agent running up costs while working unattended.

There's one Claude per channel, shared by everyone in it, so any team member can see what's being worked on and pick up where a colleague left off. That multiplayer dynamic is arguably one of the most significant departures from how AI assistants have hitherto worked.

TRENDING STORIES

1. [Greptile, Cursor, and Devin agree that agents should run their code. What they run it against matters.](#)
2. [The Anthropic leader who built Claude Code says he ditched prompting — now he just writes loops.](#)
3. ["Bring it to our shop": Workday's pitch for keeping AI agents close to your most valuable data](#)
4. [Claude Code vs. Cursor vs. Codex vs. Antigravity — six months in](#)
5. [Managing The Expectations And Reality of What AI Agents Can Do](#)

For instance, a task tagged to @Claude by one person can be steered, corrected, built upon by the rest of the team in real time.



Multiplayer steering in action

Underpinning all of this is what Anthropic calls “agent identity” — the access construct that makes multiplayer AI possible without creating a security mess. **Noah Zweben**, a member of technical staff on the Claude Code team, **notes in a companion blog post** that this model is a necessary foundation for the kind of autonomous, team-wide work Claude Tag is designed to handle.

“The shift from single player to multiplayer AI in products like Claude Tag makes long-running, team-based work possible.”

“The shift from single player to multiplayer AI in products like Claude Tag makes long-running, team-based work possible,” Zweben writes. “Agent identity ensures that Claude’s access to tools is broad enough to be useful, but scoped enough to be secure at enterprise scale.”

In terms of how this works, Zweben explains why the traditional model — where an AI assistant borrows the permissions of whoever invoked it — needed to be rethought. Agents now schedule tasks and respond to events long after the person who asked has logged off, and in a shared channel with multiple people steering, there’s no single user whose credentials make sense to inherit.

stead, administrators define what Claude can access at the workspace level, with each channel inheriting a baseline set of permissions that can be tightened or

“Agent identity replaces the question ‘what can this user do?’ with ‘what can this agent do in this compartment?’”

An engineering channel might get read and write access to GitHub and the data warehouse; a general channel might get read-only. Crucially, because Claude operates under its own accounts rather than anyone’s personal credentials, a shared channel can never serve as a backdoor into someone’s private documents.

“Agent identity replaces the question ‘what can this user do?’ with ‘what can this agent do in this compartment?’” Zweben writes.

Race for the Slack layer

Team messaging platforms are fast becoming the primary control layer for AI agents in the enterprise — the place where work gets assigned, monitored, and reported back on. GitHub has **brought Copilot into Microsoft Teams**, for instance, where OpenAI’s **Codex includes** a native Slack integration that lets teams delegate tasks directly from threads. And Cognition’s **Devin has been built around Slack** from the get-go.

Block, the Jack Dorsey-led company behind Square and Cash App, **has also spent some time building something akin to this from scratch**, using its open-source Goose framework to develop an agent orchestration system its engineers manage entirely from a single Slack thread.

Whichever AI ends up as the default presence in the channel where work is coordinated gains both a distribution advantage and a compounding data advantage, making it progressively harder to displace. For a company already running Claude Code in its engineering org, adding Claude Tag to its Slack channels is a natural next step — and one that deepens that dependency considerably.

TNS



Paul is an experienced technology journalist covering some of the biggest stories from Europe and beyond, most recently at TechCrunch where he covered startups, enterprise, Big Tech, infrastructure, open source, AI, regulation, and more. Based in London, these days Paul...



Before joining The New Stack as its senior editor for AI, Frederic was the enterprise editor at TechCrunch, where he covered everything from the rise of the cloud and the earliest days of Kubernetes to the advent of quantum computing....

[Read more from Frederic Lardinois →](#)

TNS owner Insight Partners is an investor in: OpenAI, Anthropic.

SHARE THIS STORY



TRENDING STORIES

1. Greptile, Cursor, and Devin agree that agents should run their code. What they run it against matters.
2. The Anthropic leader who built Claude Code says he ditched prompting — now he just writes loops.
3. "Bring it to our shop": Workday's pitch for keeping AI agents close to your most valuable data
4. Claude Code vs. Cursor vs. Codex vs. Antigravity — six months in
5. Managing The Expectations And Reality of What AI Agents Can Do

INSIGHTS FROM OUR SPONSORS



YugabyteDB: The Data Backbone for Thousands of Agents

22 June 2026

[PostgreSQL Database for Every Agent](#)

June 2026



AI chat stream resumption: when Redis is enough, and when you need durable sessions

30 June 2026

How durable sessions unify human-to-human and human-to-agent messages

26 June 2026

Is AI making your teams better, or just busier?

26 June 2026



Measuring OSPO Value

29 June 2026

Strengthening LFX Mentoring: A Fairer Stipend Model and a Shared Standard of Excellence

26 June 2026

Who will be the senior engineers of 2036?

24 June 2026

TNS DAILY NEWSLETTER

**Receive a free roundup of the
most recent TNS articles in
your inbox each day.**

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors



Community created roadmaps, articles,
resources and journeys for developers

Frontend Developer Roadmap

Backend Developer Roadmap

Devops Roadmap

FOLLOW TNS



© The New Stack 2026

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)