


ДАННЫЕ / БАЗЫ ДАННЫХ / С ОТКРЫТЫМ ИСХОДНЫМ КОДОМ

В течение многих лет Apache Cassandra перекладывала эту работу на вашу команду — в версии 6.0 она возвращается к нам

В альфа-версии Apache Cassandra 6.0 добавлены ACID-транзакции Accord, метаданные транзакционного кластера и функция восстановления, благодаря чему операционная работа переносится в базу данных.

8 июня 2026 г. в 10:00 [Анил Инамдар](#)



Milhad для Unsplash+

[NetApp](#) спонсировала этот пост.

Выпуски Apache Cassandra обычно оцениваются примерно по одним и тем же критериям. Что она может делать сейчас, чего не могла раньше? На версию

не столько связаны с новыми возможностями. Скорее, в версии 6.0 почтенная база данных NoSQL возвращает себе работу, которую всегда негласно перекладывала на кого-то другого.

Я тесно сотрудничаю с разработчиками Cassandra и **два года назад** писал о Cassandra 5.0 на этих (виртуальных) страницах. В этом релизе команды получили несколько потрясающих новых инструментов, таких как **собственная векторная индексация** и индексы с подключением к хранилищу. Это укрепило позиции Cassandra в сфере ИИ-нагрузок. Мы воочию убедились в прогрессе: возможности проекта в области ИИ используются в производственных кластерах, которыми мы управляем в NetApp Instacustr. Но я считаю, что в версии 6.0 приводятся другие аргументы в пользу использования Cassandra, и, на мой взгляд, с этим различием стоит разобраться до того, как команды начнут планировать пути обновления.

Проблема координации

Любой, кто создавал многосекционные рабочие процессы в Cassandra, знает, в чем заключается настоящая сложность. Это всегда код приложения, потому что именно там всегда требовалась координация между секциями. База данных не определяет, как должно выглядеть корректное состояние, что должно происходить, когда из-за частичного сбоя работа не доведена до конца, и как обеспечить согласованность всех сервисов, обращающихся к кластеру.

«С практической точки зрения часть этой работы по координации теперь переходит из кода приложения в саму базу данных».

Accord в Cassandra 6.0 — это протокол консенсуса без лидера, который обеспечивает семантику **транзакций ACID** со строгой сериализуемой изоляцией между несколькими разделами. С практической точки зрения часть работы по координации теперь переносится из кода приложения в саму базу данных. Текущая реализация требует тщательного планирования схемы и не заменяет существующую логику приложения. Однако командам, которые точно знают, сколько пользовательского кода для координации они используют, стоит присмотреться к этому решению.

this release. I would also bet it gets the least attention, because “improved cluster metadata coordination” does not make for an especially click-happy headline.



Spot by NetApp enables cloud operations teams to deliver scalability, performance and security for cloud infrastructure and applications — at the lowest possible cost — through continuous automation and optimization combined with deep visibility and governance. Learn more at spot.io.

[Learn More →](#)

Cassandra has always propagated metadata through Gossip Protocol and eventual consistency. Token ownership and schema changes (as well as data placement decisions) were all spread across nodes over time without strong ordering guarantees. While a functional model, certain classes of failure are genuinely hard to diagnose, and certain operations hard to sequence predictably. (Anyone who has untangled a messy Cassandra cluster that got into a bad state during a schema migration or topology change can now mentally relive the experience I am describing.)

TCM offers a Cluster Metadata Service to maintain a tidy log of state transitions, replacing Gossip as the primary mechanism for this work. Gossip stays in the picture but is no longer doing the heavy lifting. Your end result is a cluster whose operators can see what changed and when, and operations that have historically behaved unpredictably will now be more deterministic.

TRENDING STORIES

1. [Benchmarking PostgreSQL vs MongoDB for GenAI](#)
2. [After the Fork, Redis 8.0 Charges Forward with Perf Improvements.](#)
3. [The database storage problem is solved. Here's what comes next.](#)
4. [Why MotherDuck refuses to fork DuckDB](#)

Repair, which is no longer someone else's problem

Our team at NetApp Instaclustr has managed repair orchestration on behalf of customers for many a year. The approaches vary, from cron jobs to custom scripts to third-party tooling of different quality levels and beyond. The underlying reality, however, does not. Repair is how Cassandra maintains replica consistency over time; it has to happen, and *until now* Cassandra itself had nothing to say about how or when.

“Repair is how Cassandra maintains replica consistency over time. Until now Cassandra itself had nothing to say about how or when.”

I have watched teams under-run repair and pay for it gradually. I have also watched teams run it too aggressively and pay for it all at once. Neither outcome is exactly Cassandra's fault, because the tooling to do it right has always existed. The challenge was that it lived outside the database, which meant it lived outside what the database could operationally guarantee. With 6.0, full and incremental repair is now live inside Cassandra, with scheduling controls and treatment as an operational concern rather than something bolted on from the outside looking in. A parallel to **Unified Compaction Strategy in 5.0**, you have another background process that used to require external infrastructure now within the database's own responsibility. Teams running their own clusters without dedicated Cassandra expertise on staff are the ones most likely to feel this on day one.

Also pay attention to...

Three other changes are worth looking at before you finalize any upgrade plan, including the constraints framework. This update brings column-level validation directly into the schema, covering range checks, null constraints, regex patterns, and JSON validation, so that logic does not have to be replicated across every service writing to the cluster. Application-level validation still has its place; this just adds an important layer of defense closer to the data itself.

Zstandard dictionary compression trains dictionaries against repetitive data patterns in SSTables. Expect storage savings to vary slightly depending on how uniform the data is, and managing the dictionary lifecycle will inevitably add some

reasonably decide otherwise.

New cursor-based compaction processes SSTable data more incrementally, reducing heap pressure and GC overhead during compaction. Operators will not configure this directly. For workloads with heavy compaction activity, the effect should manifest as more stable node behavior under sustained load.

What I'm watching

I have been working with Cassandra clusters at NetApp Instaclustr since 2019. During that time, I have watched the project mature through two major releases and seen firsthand which **operational gaps** kept recurring in the clusters we manage. Version 6.0 addresses several of the ones that came up most often.

The features in this release share a theme that I find more interesting than any one of them individually. The project is taking ownership of operational work it has long delegated elsewhere. Coordination logic that lived in application code, metadata consistency managed through eventual propagation, and replica reconciliation handed off to external tooling are all challenges the Cassandra community has been working toward solving for years. Version 6.0 is where that work lands.

The maturity here is different than expanding what the database can store or index, as it's the database acknowledging that work it has been handing off to ops teams and application developers is properly its own work to do. Whether that plays out the way it should depends largely on how Accord holds up as real production workloads start hitting it. That's the part of 6.0 I'll be watching most closely. But of course, as always, test against your actual workloads before upgrading anything that matters. The best upgrades are the ones you understand before they hit production.

TNS



Anil Inamdar is the global head of Data Services at NetApp Instaclustr, which provides a managed platform around open source data technologies including Cassandra, Kafka, Postgres, ClickHouse and OpenSearch. Anil has more than 20 years of experience in data and...

[Read more from Anil Inamdar →](#)



TRENDING STORIES

1. Benchmarking PostgreSQL vs MongoDB for GenAI
2. After the Fork, Redis 8.0 Charges Forward with Perf Improvements.
3. The database storage problem is solved. Here's what comes next.
4. Why MotherDuck refuses to fork DuckDB
5. Microsoft's Azure HorizonDB Takes on Aurora, AlloyDB

INSIGHTS FROM OUR SPONSOR



Spot by NetApp enables cloud operations teams to deliver scalability, performance and security for cloud infrastructure and applications — at the lowest possible cost — through continuous automation and optimization combined with deep visibility and governance. Learn more at spot.io.

[Learn More →](#)

TNS DAILY NEWSLETTER

most recent TNS articles in your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter

THE NEW STACK

[About / Contact](#)
[Sponsors](#)
[Advertise With Us](#)
[Contributions](#)



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

[Frontend Developer Roadmap](#)
[Backend Developer Roadmap](#)
[Devops Roadmap](#)

FOLLOW TNS



© The New Stack 2026

[Disclosures](#)
[Terms of Use](#)
[Advertising Terms & Conditions](#)
[Privacy Policy](#)
[Cookie Policy](#)