




БОЛЬШИЕ ЯЗЫКОВЫЕ МОДЕЛИ

Насколько эффективен Claude Fable (Mythos) 5 для программирования?

Узнайте о преимуществах и недостатках Claude Fable 5

Эйвинд Кьосбаккен

18 июня 2026 9 мин на чтение



В этой статье я поделюсь своим мнением о Claude Fable 5, расскажу, что у него получается хорошо, а что не очень. Изображение предоставлено ChatGPT.

На прошлой неделе компания Anthropic запустила свою последнюю модель Claude Fable 5, которая представляет собой защищенную версию модели Claude Mythos.

Я тестировал эту модель днем и ночью с момента ее выпуска в течение 72 часов, пока она была доступна.


К сожалению, в настоящее время модель недоступна, так как ее использование было приостановлено по распоряжению правительства США. Однако, поскольку у меня была возможность тщательно протестировать эту модель, я хочу поделиться своими мыслями о ней, о том, что я смог с ней сделать, о ее ограничениях и так далее. Несмотря на то, что модель сейчас недоступна, я верю и надеюсь, что она скоро снова появится в продаже. Я также думаю, что через несколько месяцев мы увидим другие модели с такими же возможностями.

Конечно, было много статей, посвященных Claude Fable и его возможностям. Я считаю, что эта статья интересна тем, что я каждый день работаю с Claude Code и делаю это практически каждый день с начала 2026 года. Когда вышел Claude Fable, я сразу же протестировал модель на предмет ее возможностей. У меня был ряд задач, с которыми Opus не справлялся или которые не мог решить быстро, и я протестировал Claude Fable, чтобы по-настоящему оценить его возможности.

Вам стоит прочитать эту статью, чтобы узнать мое мнение о возможностях Claude Fable, особенно в сравнении с другими передовыми моделями, такими как Claude Opus 4.8 и GPT-5.5.

Claude Fable 5: My Review

Testing the most powerful
coding model yet



1 Why It Matters



The most anticipated AI model ever; available only 72 hours before being suspended.

2 Tasks End To End



One-shots complex multi-repo-features, understands intent, and runs autonomously for longer

3 Finds Hidden Bugs



The same prompt surfaces bugs, security flaws, and refactors that Opus missed

4 Costly To Run



Heavy token use, hits subscription limits-fast expensive on the API

5 Overly Eager

Sometimes makes-over-co-mplex changes and flags minor issues



Bottom Line

The most powerful coding model so far—hopefully it returns soon

This infographic highlights the main contents of this article. I'll discuss Claude Fable 5, how it compares to Anthropic's previous frontier model Opus 4.8, and I'll cover what Claude Fable does well and what it doesn't do as well. Image by ChatGPT.

Why care about Claude Fable 5

First of all, we should start with why you should care about Claude Fable 5. This is perhaps the most anticipated LLM ever, as it had a lot of hype going around it for months before it was launched.

Anthropic itself spent a lot of time hyping up the model, talking about its capabilities and how dangerous it could be in the wrong hands.

A lot of people were thus excited for the launch of the model, and it finally launched last week with full access to anyone with a Claude Pro or Max subscription.

• • •

The model was made widely available to anyone, to everyone. Personally, I did not experience any issues using the model until Saturday morning, Norwegian time. Even though the model was not available for too long (I believe around 72 hours), I feel like I got to test the model extensively and formed a good opinion on both its upsides, downsides, and overall capabilities.

What Claude Fable does well

First of all, I want to cover what Claude Fable did well. My overall impression is that it's significantly better than Claude Opus 4.8. I have read other people online mentioning how they felt that the capabilities didn't go that much further than Claude Opus 4.8. In my opinion, this was obviously not true. I believe that people reporting this have not tested Claude Fable on complex enough tasks.

Of course, if you try Claude Fable on an already easy task that Claude Opus is able to do, you'll not experience its full capabilities. Where Claude Fable really shines is when you apply it to super complex coding tasks.

Completing complex tasks end-to-end

I had several tasks I was working on where I had spent some time with Claude Opus 4.8 to implement it. Opus was definitely able to implement it, but it wasn't a one-shot implementation, and I had to manually guide Claude Opus through some of the implementations.

These tasks were, for example:

- A feature implementation where I had to change code in multiple repos
- Fixing an issue/bug that a customer experienced, where I had to improve the information extraction capabilities of an LLM

Unfortunately, I can't go into more detail about these issues because it's work that's in a closed codebase. What I did to compare Claude Opus versus Claude Fable is that I had previously applied Claude Opus to these problems and solved them with quite a bit of manual direction. I then applied it to Claude Fable, and it was able to one-shot the problems. A clear sign that Fable is a more powerful model than Opus.

• • •

As a more generic note on Fable's capabilities, I would say the following:

Claude Fable is more able to complete tasks end-to-end, with both better understanding of ambiguities and user intent, better at implementing the actual planned solution, and better at verifying the solution is actually correct through navigating on the computer or running integration tests.

I simply found that Claude Fable was willing to run for longer periods of time, completing more complex tasks without giving up or ending up in recurring issues. I simply felt that tasks were now being done more autonomously, and I didn't have to give that much direction to ensure the model was aligned with my intentions.

Discovering issues in codebases

Another incredible capability I noticed in Claude Fable was that it was much better than Claude Opus at discovering issues in codebases, either finding bugs or looking for refactoring opportunities or potential future issues.

I constantly run a prompt similar to the one below to discover issues in my codebase.

```
Scan thoroughly through the codebase to identify any potential bug issues, or refactoring opportunities, and come back to me with an HTML report with issues prioritized from most severe to least severe
```

With Claude Opus, I ran this same prompt and didn't get good results either. Claude Opus wasn't able to discover any more refactoring opportunities or bugs, or the issues it discovered weren't really relevant. (Of course, note that this was the case after I had already been doing a lot of refactoring and bug detection with Claude Opus in a specific repository.)

However, when I then applied Claude Fable with the exact same prompt, it started finding a lot of severe issues, both security-wise and actual bugs, and also finding a lot of good refactoring opportunities that Claude Opus was not able to see.

I immediately started going through all the repositories with Claude Fable running this prompt and fixing all the issues. I pushed a lot of code that vastly increased the quality of my codebases.

I think this is probably the single clearest sign that Fable is a more powerful model than Opus. You could run the exact same prompt in the same codebase where Fable is able to detect a bunch of issues that Claude Opus was not able to detect.

I'm just happy that I got to run this refactoring, bug detection, and fixing a lot before the model was unfortunately suspended.

What Claude Fable doesn't do well

In the previous section, I covered what Claude Fable does really well. I think it's important to also highlight some downsides of Claude Fable, considering it's not a perfect model.

Price

Claude Fable is definitely the most powerful coding model I've ever used. However, one of the major issues is how many tokens it spends to complete tasks.

Naturally, this is not an issue with the model itself. It's more of an issue with the rate limits that you have with Anthropic. Using Claude Fable with a subscription, I started hitting the subscription limit way faster.

This is definitely a limitation, as you cannot just run the model infinitely anymore. Furthermore, I would argue that the subscription pricing of Claude Fable is very prohibitive for almost all companies. Running a model that costs \$10 per million in and \$50 per million out is not feasible for basically anyone but the largest of companies.

Of course, you could argue that you can use Claude Fable only for planning and only for bug detection, and then use Claude Opus for actual implementations. I agree this could probably be done, and you would still reap most of the benefits from Claude Fable; however, spending a lot of time optimizing which model you use in which situations is very time-consuming and definitely something you want to avoid if you want to be as productive and effective as possible.

This is one of the major downsides, I'd argue, rate limits and how much you can use the model/the cost of the model if you're using API pricing.

Sometimes the model is overly eager

Another small downside I would also like to cover with Claude Fable is that the model is sometimes overly eager to find issues or perform implementations. Sometimes I found that the model was implementing things in an overly complex manner. For example, changing way more lines of code than it really needed to, or finding more issues in a codebase where a lot of the issues weren't really that severe.

I find this slightly annoying sometimes, but I also believe that this is a trade-off that Anthropic has accepted. You, of course, want a model to always look for issues and constantly try to fix

them, and you want the fix to work immediately, of course. It's hard to balance this eagerness and, at the same time, avoid the model becoming overly eager to find issues and fix them.

Overall, however, this is quite a small downside. It's just one small thing I noticed when using Claude Fable. By far the biggest downside is the prohibitive pricing of the model.

Conclusion

In this article, I covered my thoughts on Claude Fable. I've compared it against the previous frontier model from Anthropic, which was Claude Opus 4.8. The Claude Fable 5 model is incredible, but also has some downsides:

1. One being that it's overly eager, which I believe is quite a small downside.
2. The pricing. The pricing is, of course, an important downside. However, it's not an issue with the model itself, and the best models are always going to cost a lot of money. People can make their own consideration if the model is worth what you have to pay for it.

All in all, however, it's a very powerful model. I hope that it becomes available again and that other models, both from other frontier labs and open source models, will reach these capabilities in a few months, so that we have even more powerful coding agents to perform software engineering tasks.

Also, check out [How to Effectively Run Many Claude Code Agents in Parallel](#).

[Get in Touch](#)

👉 My free eBook and Webinar:

🚀 [10x Your Engineering with LLMs \(Free 3-Day Email Course\)](#)

📖 [Get my free Vision Language Models ebook](#)

💻 [My webinar on Vision Language Models](#)

👉 Find me on socials:

✉️ [Substack](#)

🔗 [LinkedIn](#)

🐦 [X / Twitter](#)

• • •

WRITTEN BY

Eivind Kjosbakken

[See all from Eivind Kjosbakken](#)

Agentic Ai

coding agents

LLM Agents

Machine Learning

Programming

Share This Article



Towards Data Science is a community publication. Submit your insights to reach our global audience and earn through the TDS Author Payment Program.

Write for TDS

Related Articles


LARGE LANGUAGE MODELS

Deep Dive into LLaMA 3 by Hand 🖋️

Explore the nuances of the transformer architecture behind Llama 3 and its prospects for the...

Srijanie Dey, PhD

May 3, 2024 12 min read




LARGE LANGUAGE MODELS

Deep Dive into Transformers by Hand 🖋️

Explore the details behind the power of transformers

Srijanie Dey, PhD

April 12, 2024 6 min read



LARGE LANGUAGE MODELS




ARTIFICIAL INTELLIGENCE

UniFliXsg: AI-Powered Undergraduate Program Recommendations for Singapore Universities

How could AI suggest your majors?

Phanuphat (Oad) Srisukhawasu

August 14, 2024 8 min read



ARTIFICIAL INTELLIGENCE

Semantic Search Engine for Emojis in 50+ Languages Using AI



If you are on social media like Twitter or LinkedIn, you have probably noticed that...

Badr Alabsi, PhD

July 17, 2024 14 min read



LARGE LANGUAGE MODELS

Building a Conformal Chatbot in Julia

Large Language Models (LLM) are all the buzz right now. They are used for a...

nk Altmeyer

May 2, 2023 8 min read

Beware of Unreliable Data in Model Evaluation: A LLM Prompt Selection case study with Flan-T5

You may choose suboptimal prompts for your LLM (or make other suboptimal choices via model...

Chris Mauck

June 16, 2023 12 min read

LARGE LANGUAGE MODELS

Deep Dive into Anthropic's Sparse Autoencoders by Hand 🖋️

Explore the concepts behind the interpretability quest for LLMs

Srijanie Dey, PhD

May 31, 2024 12 min read



Your home for data science and AI. The world's leading publication for data science, data analytics, data engineering, machine learning, and artificial intelligence professionals.

© Insight Media Group, LLC 2026

[Subscribe to Our Newsletter](#)

[WRITE FOR TDS](#) · [ABOUT](#) · [ADVERTISE](#) · [PRIVACY POLICY](#) · [TERMS OF USE](#)