

Продукт

Представляем Claude Sonnet 5

30 июня 2026 г.



Claude Sonnet 5 — самая «агентная» модель Sonnet. Она может строить планы, использовать такие инструменты, как браузеры и терминалы, и работать автономно на уровне, который еще несколько месяцев назад требовал более крупных и дорогих моделей.

Для многих разработчиков эра агентного ИИ началась с моделей класса Sonnet: Claude Sonnet 3.5, 3.6 и 3.7 были первыми моделями, продемонстрировавшими впечатляющие навыки в программировании и использовании инструментов. Однако в последнее время наибольший прогресс в агентных возможностях наблюдается у наших моделей класса Opus.

Sonnet 5 сокращает отставание: по производительности она близка к Opus 4.8, но стоит дешевле. По сравнению со своим предшественником, Sonnet

4.6, эта модель значительно улучшена в таких важных аспектах, как способность к логическому мышлению, использование инструментов, кодирование и работа со знаниями:

	Sonnet 5	Sonnet 4.6	Opus 4.8 For reference
Agentic coding SWE-bench Pro	63.2%	58.1%	69.2%
Agentic coding Terminal-Bench 2.1	80.4%	67.0%	82.7%
Multidisciplinary reasoning Humanity's Last Exam	43.2% no tools	34.6% no tools	49.8% no tools
	57.4% with tools	46.8% with tools	57.9% with tools
Computer use OSWorld-Verified	81.2%	78.5%	83.4%
Knowledge work GDPval-AA v2	1618	1395	1615

Баллы Sonnet 5 по различным тестам в сравнении с Sonnet 4.6 и Opus 4.8 (более универсальной моделью, для сравнения). В системной карте Claude Sonnet 5 подробно описан более широкий спектр тестов.


По результатам наших оценок безопасности, Sonnet 5 демонстрирует в целом более низкий уровень нежелательного поведения, чем Sonnet 4.6, и в целом более безопасен для использования в агентных контекстах. Оценки также показывают, что он гораздо хуже справляется с задачами в области кибербезопасности, чем наши текущие модели Opus.

С сегодняшнего дня Claude Sonnet 5 доступен во всех тарифных планах: он является моделью по умолчанию для тарифных планов Free и Pro, а также доступен пользователям тарифных планов Max, Team и Enterprise. Он также доступен в Claude Code и на платформе Claude, где до 31 августа 2026 года действует вводная цена: 2 доллара за миллион входных токенов и 10 долларов за миллион выходных токенов. После этой даты цена составит 3 доллара за миллион входных токенов и 15 долларов за миллион выходных токенов. Разработчики могут использовать `claude-sonnet-5` через Claude API.

Работа с Claude Sonnet 5

На графиках ниже сравнивается производительность Sonnet 5 с Sonnet 4.6 и Opus 4.8 при различных усилиях в оценке агентного поиска BrowseComp и оценке использования компьютера OSWorld-Verified. Sonnet 5 (оранжевая линия) значительно превосходит Sonnet 4.6 (серая линия) и охватывает гораздо более широкий диапазон соотношения стоимости и производительности, чем Opus 4.8 (желтая линия). Он обеспечивает значительно более высокую экономическую эффективность при средних затратах, а при больших затратах может сравниться по производительности с Opus 4.8 при выполнении некоторых задач. В отличие от Sonnet 5 и Opus 4.8, пользователи могут регулировать уровень затрат, чтобы найти оптимальный баланс между стоимостью и производительностью.

- Агентный поиск
- Агентное использование компьютера



Cost-performance curves at different effort levels. The previous best Sonnet model (Sonnet 4.6) fell well short of Opus 4.8. Sonnet 5 offers a wider range of cost-performance options than Sonnet 4.6, and in some cases matches Opus 4.8's capability levels. The charts show Sonnet 5 priced at \$3 per million input tokens and \$15 per million output tokens. With the introductory launch pricing through August 31 (\$2/MTok input and \$10/MTok output), the effective cost of Sonnet 5 is even lower than shown here. Opus 4.8 is priced at \$5/MTok input and \$25/MTok output. xhigh = extra high effort level.

Feedback from our early access partners has been consistent: Sonnet 5 is much more agentic than its predecessors. Testers described how it finishes complex tasks where previous Sonnet models would stop short, how it checks its own output without explicitly being asked, and how it does all this agentic work at an attractive price point:

FACTORY

“

Claude Sonnet 5 gives our agents a strong execution layer for multi-step software engineering work. It handles sustained coding, tool use, and debugging well across messy technical contexts, and has been especially useful for workflows where follow-through and technical grounding matter.

Zimu Li
Member of Technical Staff

zapier

“

We handed Claude Sonnet 5 a two-part job—update Salesforce account tiers, send a launch announcement to enterprise contacts—and it finished end to end. That used to stall halfway. For day-to-day automation, it’s a no-brainer.

Daniel Shepard
Senior Engineer

Safety evaluations

Our pre-deployment safety evaluations found that Sonnet 5 was overall an improvement on Sonnet 4.6. On agentic safety, the model is better at refusing malicious requests and resisting hijack attempts in prompt injection attacks. The


model shows lower rates of hallucination and sycophancy than Sonnet 4.6. On our automated behavioral audit, which tests a wide range of misaligned behaviors such as cooperation with misuse and deception, Sonnet 5 scored lower (that is, safer) overall. However, it did show somewhat higher rates of misaligned behavior on this assessment compared to the more capable Opus 4.8 and Claude Mythos Preview.



Rates of misaligned behavior on our automated behavioral audit, which tests for a very wide range of undesirable behaviors across many situations and contexts (see Section 6.4 of the [Sonnet 5 System Card](#) for a complete list and results for each specific behavior). Sonnet 5 shows an overall lower rate of misaligned behavior than Sonnet 4.6, though a higher rate than Mythos Preview and Opus 4.8.

We did not deliberately train Sonnet 5 on cybersecurity tasks. It can perform some routine, non-harmful cyber tasks, but on evaluations testing potentially dangerous cyber skills, such as developing software exploits, it shows substantially poorer performance than models such as Opus 4.8 and Mythos 5. Scores from one evaluation, which tested models’ ability to develop exploits for vulnerabilities in the Firefox browser, are shown in the chart below. Sonnet 5 was never able to develop a full working exploit, but it does show a slightly higher rate of *partial* success than Sonnet 4.6. This latter change is likely due to improvements in general intelligence rather than specific training.

Firefox 147 exploit development



Scores measuring models' success at developing exploits for software vulnerabilities in Firefox 147 (this evaluation was developed in collaboration with Mozilla; all vulnerabilities have been patched in Firefox 148). For each model, the left-hand bar shows how often the model (without safeguards) developed a working exploit; the right-hand bar shows how often the model had partial success. Neither of the Sonnet models could successfully develop a working exploit (both scored 0.0%); Sonnet 5 showed a slightly higher partial success rate than Sonnet 4.6. Both Sonnet models have substantially poorer cyber capabilities than Opus 4.8 and Mythos 5. For full details, see Section 3.2.4 of the Sonnet 5 System Card.

Since Sonnet 5 is somewhat stronger than its predecessor on these tasks, we've launched it with cyber safeguards enabled by default. These safeguards—which detect and block dangerous cyber usage in real time—are the same as those present in Claude Opus 4.7 and 4.8 (because we judged that the overall level of cybersecurity risk from Sonnet 5 was low, the safeguards are less strict than those launched with Fable 5, which block a much wider range of cybersecurity tasks).¹

Our full assessment of Sonnet 5 across many safety and capability evaluations is reported in the Claude Sonnet 5 System Card.

Availability and pricing

Claude Sonnet 5 is available everywhere today at an introductory price of \$2 per million input tokens and \$10 per million output tokens through August 31, 2026. It then moves to standard pricing at \$3 per million input tokens and \$15 per million

output tokens.² We've increased rate limits across Chat, Cowork, Claude Code, and the Claude Platform³ to accommodate the higher token usage of higher effort levels; users can select whichever level makes sense for their particular project.

Changelog

Edit June 30, 2026: In the original version of this post, we included a cost-performance chart for the BrowseComp evaluation that was based on data from a simpler methodology that did not reflect the standard methodology we use for agentic search evaluations. This had the result of underestimating Sonnet 5's performance on the evaluation.

We have now updated the chart so that it matches the methodology that we used and discussed in the Sonnet 5 system card (which used a 10M token budget with compaction and programmatic tool calling). We have also updated the surrounding text.

Footnotes

¹ Sonnet 5 is part of our Cyber Verification Program, which is available today on the native Claude Platform, the Claude Platform on AWS, and Claude in Microsoft Foundry (hosted on Azure and Anthropic), and coming soon on Claude in Google Vertex. Organizations that are already enrolled in the Cyber Verification Program automatically have the same access on Sonnet 5, with no need to reapply. Overall, we recommend Claude Opus 4.8 for cybersecurity work that requires reduced guardrails.

² Sonnet 5 is an upgrade to Sonnet 4.6, but it uses an updated tokenizer that changes how the model processes text to improve performance (this is similar to the tokenizer change we introduced with Claude Opus 4.7). The tradeoff is that the same input can map to more tokens: roughly 1.0–1.35× depending on the content type. The introductory pricing is set so that the transition to Sonnet 5 is roughly cost-neutral.

³ On April 26, 2026, we raised Sonnet and Haiku rate limits at every usage tier and simplified to three tiers (Start, Build, and Scale) on the native Claude Platform. You can view your tier and current limits in the Claude Console or read the documentation to learn more.

- **Humanity's Last Exam:** We updated the grader model for Humanity's Last Exam and have updated the Sonnet 4.6 score to 34.6% (no tools) and 46.8% (with tools). This is the reason the score differs from that reported in the Sonnet 4.6 launch blog.

- **OSWorld-Verified:** We made changes to how we run the OSWorld-Verified evaluation to more accurately reflect the model's performance in the real world, and have updated the Sonnet 4.6 score to 78.5%. This is the reason the score differs from that reported in the [Sonnet 4.6 launch blog](#).
-



Related content

Redeploying Fable 5

Fable 5 returns globally July 1. We're also proposing an industry-wide framework for scoring jailbreak severity, together with Amazon, Microsoft, Google, and other Glasswing partners.

[Read more](#) →

Claude Science, an AI workbench for scientists, is now available

Claude Science is a customizable app that integrates the tools and packages researchers most often use, produces auditable artifacts, and provides flexible access to computing resources.

[Read more](#) →

Introducing Claude Tag

Claude Tag is a new way for teams to work with Claude.

[Read more](#) →



Products

[Claude](#)

[Claude Code](#)

[Claude Code Enterprise](#)

[Claude Cowork](#)

[@Claude](#)

Claude Design

Claude Science

Claude Security

Claude for Chrome

Claude for Microsoft 365

Skills

Download app

Pricing

Log in to Claude

Models

Mythos

Fable

Opus

Sonnet

Haiku

Solutions

AI agents

Code modernization

Coding

Customer support

Education

Enterprise

Financial services

Government

Healthcare

Legal

Life sciences

Nonprofits

Security

Small business

Startups

Claude Platform

[Overview](#)

[Developer docs](#)

[Pricing](#)

[Ecosystem](#)

[Marketplace](#)

[Regional compliance](#)

[Claude on AWS](#)

[Google Cloud](#)

[Microsoft Foundry](#)

[Console login](#)

Resources

[Blog](#)

[Claude partner network](#)

[Community](#)

[Connectors](#)

[Courses](#)

[Customer stories](#)

[Engineering at Anthropic](#)

[Events](#)

[Inside Claude Code](#)

[Inside Claude Cowork](#)

[Inside Claude Enterprise](#)

[Plugins](#)

[Powered by Claude](#)

[Service partners](#)

[Tutorials](#)

[Use cases](#)

Help and security

[Availability](#)

[Status](#)

[Support center](#)

Company

[Anthropic](#)

[Careers](#)

[Policy](#)

[Economic Futures](#)

[Research](#)

[News](#)

[Claude's Constitution](#)

[Claude Corps](#)

[Policy on the AI Exponential](#)

[Responsible Scaling Policy](#)

[Security and compliance](#)

[Transparency](#)

Terms and policies

[Privacy choices](#)

[Privacy policy](#)

[Consumer health data privacy policy](#)

[Responsible disclosure policy](#)

[Terms of service: Commercial](#)

[Terms of service: Consumer](#)

[Usage policy](#)

© 2026 Anthropic PBC

