


ИНФРАСТРУКТУРА ИИ / СТРАТЕГИЯ ИИ / АППАРАТНОЕ ОБЕСПЕЧЕНИЕ

OpenAI хочет расширить свой стек ИИ с помощью Jalapeño — своего первого специализированного чипа

Ускоритель логического вывода, разработанный совместно с Broadcom, — это первый шаг в многолетней программе развития, который открывает неопределенное будущее для гибкости разработчиков.

24 июня 2026 года, 16:55 [Мерedit Шубл](#)



Источник: Summerizze, Unsplash+

В среду компания OpenAI анонсировала Jalapeño — свой первый специализированный ускоритель логического вывода, разработанный совместно с Broadcom при поддержке канадского производителя

Компания, занимающаяся разработкой искусственного интеллекта, **утверждает**, что Jalaréño был создан для работы со всеми большими языковыми моделями (БЯМ) и поможет сделать ИИ быстрее, качественнее и дешевле. За этой радужной перспективой скрывается желание OpenAI владеть полным стекком ИИ, к чему уже стремятся другие гиганты в этой области.

«Те, кто серьезно относится к платформам, должны серьезно относиться и к кремнию».

Как написал Бен Баджарин, генеральный директор и главный аналитик исследовательской компании Creative Strategies, специализирующейся на потребительских технологиях, в **X**: «Те, кто серьезно относится к платформам, должны серьезно относиться и к кремнию».

Но поскольку технических подробностей мало, разработчикам остается только гадать, будет ли расширение присутствия OpenAI способствовать развитию или, наоборот, сдерживать его.

Заходите, крупные технологические компании. Мы все сейчас разрабатываем собственные чипы.

OpenAI — не единственная крупная технологическая компания, которая создает собственные чипы для искусственного интеллекта.

в 2016 году Google разработала и создала собственное специализированное оборудование для TensorFlow — программного

специализированный чип для искусственного интеллекта и машинного обучения. **Trainium** появился в 2022 году, а в 2023-м — **Microsoft Azure Maia AI Accelerator**. Пока ничего не ясно, но в апреле **Reuters** сообщило, что Anthropic рассматривает возможность разработки собственных чипов, хотя компания, занимающаяся искусственным интеллектом, пока не дает никаких комментариев, по крайней мере публично.

Why is everyone jumping on the custom-chip bandwagon?

TRENDING STORIES

1. [GitLab just surveyed 1,500 developers. Here's why it matters for your codebase.](#)
2. [AWS, Microsoft, and Google agree the session is the new unit of compute. They disagree on how to isolate it.](#)
3. [Neoclouds, sovereign AI and Postgres: The new operating model for regulated enterprises](#)
4. [OpenAI wants to claim more of the AI stack with Jalapeño, its first custom chip](#)
5. ["A data lake of nuance for AI agents to swim in": AWS Context gets shipshape on reasoning](#)

Blame the compute gold rush, as AI companies increasingly clamor for compute power — "a compute-powered economy," as **Greg Brockman**, president, chairman, and co-founder, OpenAI, puts it. And the numbers surely back it; **Stanford's 2025 AI Index Report** says, "training compute doubles every five months."

While building custom AI chips in-house doesn't completely alleviate compute pressures, it is one way for OpenAI and its Big Tech brethren to expand compute capacity while potentially lowering costs and reducing reliance on third-party suppliers.

Exciting claims, no proof

In its announcement blog post, OpenAI describes its new chip as "designed to be best inference platform for LLMs."

“We optimized the architecture around the kernels, memory movement, networking, and serving patterns that matter most for frontier AI models. Based on early testing, Jalapeño will efficiently execute our most important workloads close to the hardware’s theoretical limits.”

But the AI company remains tight-lipped on any real technical details.

While it claims current tests put Jalapeño’s performance “substantially better than current state-of-the-art,” it doesn’t provide benchmarks to back that up. Instead, it tells developers to expect a detailed technical report “in the coming months.”

What OpenAI does divulge is that engineering samples of the chip are currently running on ML workloads in its lab, including GPT-5.3-Codex-Spark.

Will Jalapeño serve developers, or is OpenAI’s desire to own the AI stack?

OpenAI makes no qualms about its quest for full-stack control. In doing so, the AI company claims it will make its models “faster, more reliable, and more affordable for users.”

Its logic goes a little something like this: Better infrastructure means more efficient compute, which means better training, which means better models, which means better products, which means more revenue. Then, it explains, it can reinvest that revenue in its infrastructure to make intelligence better for everyone.

But given how little the AI company has revealed about the chip’s specs, it seems developers will have to sit back and watch where the chips fall.

Jalapeño, then, is simply the next move in OpenAI’s quest to control the whole AI chessboard, moving beyond models and products to the underlying infrastructure itself.

For developers, OpenAI seems adamant on insisting its full-stack strategy will lead to better performance and pricing for everyone and ultimately empower “anyone

Several times in its announcement, OpenAI reiterates that it designed Jalapeño for current and future LLMs — all of them. But given how little the AI company has revealed about the chip's specs, it seems developers will have to sit back and watch where the chips fall.

Built fast with a long roadmap ahead

The few behind-the-scenes details OpenAI does choose to share boast about its development speed, stating it brought Jalapeño from design to manufacturing tape-out in nine months — “what we believe to be the fastest ASIC development cycle ever achieved in high-performance advanced semiconductors.”

The AI company chalks up that fast timeline, in part, to its own models accelerating parts of the design and optimization processes.

Looking ahead, Jalapeño is slated for deployment at a gigawatt scale in Microsoft's and other partners' data centers by the end of the year.

That's just the beginning. OpenAI hints at an upcoming multi-generation roadmap, posing the question: What will it seek to control next?

TNS



Meredith Shubel is a technical writer covering cloud infrastructure and enterprise software. She has contributed to The New Stack since 2022, profiling startups and exploring how organizations adopt emerging technologies. Beyond The New Stack, she ghostwrites white papers, executive bylines,...

[Read more from Meredith Shubel →](#)

TNS owner Insight Partners is an investor in: OpenAI, Unit, Anthropic.

SHARE THIS STORY



2. AWS, Microsoft, and Google agree the session is the new unit of compute. They disagree on how to isolate it.
3. Neoclouds, sovereign AI and Postgres: The new operating model for regulated enterprises
4. OpenAI wants to claim more of the AI stack with Jalapeño, its first custom chip
5. "A data lake of nuance for AI agents to swim in": AWS Context gets shipshape on reasoning

INSIGHTS FROM OUR SPONSORS



Beyond the DBaaS Trap: Claiming Operational Independence for PostgreSQL

26 February 2026

Why the Future of Postgres Extensions in Kubernetes is Immutable and Flexible

18 March 2025

Cloud Neutral Postgres Databases Made Simple | EDB

19 December 2024



Moving from Probabilistic Reasoning to Deterministic Execution

24 June 2026

Kong Mesh 2.14: Finer Zone Proxy Control and Tighter Security

23 June 2026

Why We Need to Stop Prompt Hacking



AI Test Automation: Fix the 3 New Testing Bottlenecks

17 June 2026

Test Observability: Fix Flaky Kubernetes Tests (2026)

11 June 2026

Testkube AI Launch: New Use Cases, OAuth, and Open Source Analysis

9 June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the
most recent TNS articles in
your inbox each day.

svo1975pvn1982@gmail.com

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

Databases

Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

API Management

Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors
Advertise With Us
Contributions



Community created roadmaps, articles, resources and journeys for developers to help you choose your path and grow in your career.

Frontend Developer Roadmap
Backend Developer Roadmap
Devops Roadmap

FOLLOW TNS



[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)