


ИИ

OpenClaw и Hermes сходятся во мнении о том, что такое агент. Но расходятся в вопросе о том, что им управляет.

Agent превращает модели искусственного интеллекта в автономные системы. OpenClaw делает ставку на широкие возможности шлюзов, а Hermes — на память, и поставщики платформ хотят получить и то, и другое.

23 июня 2026 года, 13:56 [Джанакiram MSV](#)



В этом месяце на конференции Microsoft Build генеральный директор компании Сатья Наделла рассказал о переходе от операционных систем и приложений к агентному искусственному интеллекту, который не ждет, пока его откроет пользователь. Затем он показал уровень, который делает этот переход возможным: **OpenClaw**, набор инструментов с открытым исходным кодом и независимый проект, выпущенный всего несколько месяцев назад и изначально работающий на Windows в новых контейнерах выполнения Microsoft. На его основе был создан **Scout, постоянно работающий поративный агент Microsoft**. Шлейка, которой едва исполнился год, нацене в качестве управляемой инфраструктуры.

персонального ИИ». Аналогия работает: фреймворк делает для языковой модели то же, что операционная система делает для процессора. Модель может самостоятельно отвечать на вопросы, а фреймворк позволяет ей работать непрерывно, запоминать то, что она изучает, и вызывать инструменты для выполнения действий.

В этом году два открытых проекта создали этот уровень, отталкиваясь от разных отправных точек. OpenClaw построен вокруг шлюза — той части, которая подключает агента к уже используемым каналам. Hermes Agent от Nous Research построен вокруг памяти — той части, которая позволяет агенту изучать работу разработчика и совершенствоваться в ней. Состязание между ними происходит на уровне управления, а не в том, какую модель они используют.

Агентная платформа на простом английском

По своей сути агентная платформа — это программное обеспечение, которое превращает модель в систему, способную работать автономно. Она объединяет в себе среду выполнения, которая поддерживает работу агента между задачами, шлюз, через который проходят входящие и исходящие сообщения, и память, сохраняющуюся между сеансами. Кроме того, в нее входят инструменты, к которым агент обращается для выполнения действий, идентификатор, под которым он работает, навыки, которые он может расширять, а также средства управления политиками и наблюдаемостью, которые определяют, к чему он может прикасаться и что должен фиксировать. Помощник по написанию кода, такой как Claude Code или Codex, решает эту задачу лишь частично. Он работает в рамках интерактивного сеанса и теряет большую часть своего рабочего контекста после завершения сеанса. Платформа обеспечивает работу, память и управление, чтобы агент мог работать без участия пользователя.

A harness keeps the runtime, memory, and governance in place so the agent can run unattended.

Nous Research and OpenClaw agree on that anatomy. They differ on which part they treat as the primary control point. OpenClaw starts with the gateway, so one agent can answer on WhatsApp, Discord, Slack, and other channels from a single

OpenClaw's gateway-first design

OpenClaw began as an independent open-source project from Peter Steinberger, a developer known for his earlier work in PDF tooling, who released an early version late in 2025 and renamed it twice before settling on OpenClaw in January. It is built for breadth, centered on a central gateway that connects the agent to dozens of messaging channels. **ClawHub**, its public skills marketplace, holds thousands of community skills that extend what the agent can do. The open-source momentum was real, with the repository near 380,000 GitHub stars by late June, though stars measure visibility rather than production use.

The more important development is who adopted that gateway. Steinberger **joined** OpenAI in February, and the project moved to an independent foundation with OpenAI as a sponsor rather than an owner.

TRENDING STORIES

1. **OpenClaw and Hermes agree on what an agent is. They disagree on what controls it.**
2. **MCP gets its missing enterprise authorization layer**
3. **Developers are now validating code they didn't write — and may not understand**
4. **Can DNS become the basis for AI agent identity?**
5. **AI agents need infrastructure: Why Europe's regional cloud strategy matters**

At GTC in March, **Nvidia wrapped OpenClaw in NemoClaw**, an OpenShell runtime that sandboxes each agent and enforces policy from outside the agent's reach.

At Build in June, Microsoft made OpenClaw native to Windows execution containers and shipped Scout, an agent on the OpenClaw gateway with its own Entra identity and connections into Teams, Outlook, and SharePoint. In each case, the platform vendor kept OpenClaw's breadth and added the governance and identity the original project lacked.

which stay hidden, rather than granting the broad access that made early OpenClaw deployments risky.

For enterprises, that changes the calculus. A security team can now scope which folders an agent reads and which stay hidden, rather than granting the broad access that made early OpenClaw deployments risky. Platform teams can offer a single governed agent that staff can access from the tools they already use. Breadth gave OpenClaw distribution, and the platform vendors supplied the controls that distribution needed to enter production.

Hermes' memory-first design

Hermes Agent takes the other path. Nous Research, the lab behind the Hermes, Nomos, and Psyche model families, released it on February 25 under an MIT license, written in Python and built to run persistently on infrastructure the team owns: a VPS, a home server, or a laptop.

The defining capability of Hermes is persistent memory across sessions. It keeps a layered memory, develops new skills after a hard task, and refines those skills as it uses them. It also builds a profile of the developer it works for, so each session starts with more context than the last. The skills follow the **agentskills.io** standard, which keeps them portable across agents rather than locked to one.

That depth has translated into measurable usage. Hermes surpassed 100,000 GitHub stars by mid-May and reached roughly 160,000 by late in the month. It overtook OpenClaw in OpenRouter's daily token rankings on May 10, with the day's total reported at 224 billion tokens, up from 186 billion, and placed Hermes first by total tokens.

By late June, OpenRouter's app rankings also placed Hermes first by total tokens, past 22 trillion. GitHub stars, token volume, and platform endorsements measure different kinds of adoption, and they rarely move together. Nous has also made portability part of the pitch, shipping a `hermes claw migrate` command that imports an OpenClaw user's settings, memories, skills, and keys in one step.

Developers can keep an agent that holds a codebase, its conventions, and prior decisions across weeks, instead of rebuilding that context each morning. Teams

Since the team that runs Hermes also has to secure and maintain the infrastructure it lives on.

Breadth, depth, and where each fits

The choice resembles the familiar tradeoff between a managed cloud service and self-managed infrastructure. A managed service is convenient and vendor-governed, while self-managed infrastructure offers full control and operational responsibility.

Many enterprises will run both, depending on the workload. Neither project is limited to a single capability. OpenClaw includes memory and skills, and Hermes speaks across twenty-odd channels, so the distinction is one of emphasis rather than exclusivity. The table below maps the common cases, with the caveat that both are emerging platforms rather than finished products.

Scenario	Stronger fit	Why, with the tradeoff
Regulated enterprise that needs audit and policy control	OpenClaw under NemoClaw, or Microsoft Scout	Governance and identity are wrapped around the agent by Nvidia or Microsoft, though both are early and tie the buyer to their stack
Developer who wants an agent that learns their work and stays portable	Hermes	Persistent memory and self-improving skills are the design center, at the cost of running your own infrastructure
Team reaching users across many chat platforms	OpenClaw	The gateway and large skills marketplace cover breadth no rival matches, though skill quality varies and supply-chain risk is real
Organization standardizing on one provider's cloud	Scout inside Microsoft 365	Deepest integration within that ecosystem, and the least portability outside it

Real deployments will not standardize on one approach. Nvidia's NemoClaw blueprints already run Hermes agents under OpenShell as readily as they run OpenClaw. The governance layer is being built to sit beneath multiple agent projects rather than to pick one.

Why the harness layer matters

rewrite its own memory and skills between sessions, as Hermes does, the team needs to know who can explain a change in behavior and where that change is logged. The second question concerns ownership. When governance and identity come from a platform vendor, as they do with NemoClaw and Scout, the policy engine and the identity belong to that vendor, not to the team running the agent.

Nvidia and Microsoft are competing to place governance, identity, and observability around whichever agent a customer chooses.

For the platform vendors, the prize is the runtime layer, which will outlast any single foundation model. Nvidia and Microsoft are competing to place governance, identity, and observability around whichever agent a customer chooses, which is why NemoClaw supports Hermes alongside OpenClaw.

Security is part of the same logic. Audits of OpenClaw's **skill marketplace** flagged 341 malicious entries among the skills they scanned, and security firms reported tens of thousands of exposed instances earlier this year – the gap that governed runtimes are meant to close.

What's next

The agent market is moving past model selection into the runtime, governance, and memory layers. OpenClaw showed that a broad gateway and a large skills ecosystem can attract developers and draw in OpenAI, Nvidia, and Microsoft. Hermes showed that persistent memory and self-improving skills can drive heavy daily usage without the same platform backing. Whether breadth and depth will remain in separate projects is already uncertain, since NemoClaw runs under a single set of controls and Hermes can import an OpenClaw setup.

The next phase will turn on ownership.

Enterprises will need to know who controls the memory an agent accumulates, who governs the tools it can call, and who owns the runtime that keeps it alive. An agent that has learned a year's worth of a developer's habits creates a higher switching cost than one that merely connects to many applications. Memory, more than channel reach, is becoming the durable form of lock-in, which is why the runtime, governance, and memory layers are where the platform vendors will compete as they settle in beneath both projects.



Janakiram MSV (Jani) is a practicing architect, research analyst, and advisor to Silicon Valley startups. He focuses on the convergence of modern infrastructure powered by cloud-native technology and machine intelligence driven by generative AI. Before becoming an entrepreneur, he spent...

[Read more from Janakiram MSV →](#)

TNS owner Insight Partners is an investor in: OpenAI, Real.

SHARE THIS STORY



TRENDING STORIES

1. OpenClaw and Hermes agree on what an agent is. They disagree on what controls it.
2. MCP gets its missing enterprise authorization layer
3. Developers are now validating code they didn't write — and may not understand
4. Can DNS become the basis for AI agent identity?
5. AI agents need infrastructure: Why Europe's regional cloud strategy matters

INSIGHTS FROM OUR SPONSORS



10 Years of MongoDB Atlas: Built for What's Next

25 June 2026

Build Trust in Agentic AI: From POC to Production

23 June 2026



Comparing Kubernetes Development Tools in 2026: Telepresence, mirrord, Okteto, and Signadot

18 June 2026

Watch Claude Code Validate a Microservices Change in a Closed Loop

4 June 2026

Signadot Plans

20 May 2026



Announcing Advanced Security for SonarQube Cloud Team plan

24 June 2026

Now available: SonarQube plugin for Codex

24 June 2026

Now available: SonarQube plugin for Cursor

24 June 2026

TNS DAILY NEWSLETTER

Receive a free roundup of the most recent TNS articles in your inbox each day.

SUBSCRIBE

The New Stack does not sell your information or share it with unaffiliated third parties. By continuing, you agree to our [Terms of Use](#) and [Privacy Policy](#).

ARCHITECTURE

Cloud Native Ecosystem
Containers
Databases
Edge Computing
Infrastructure as Code
Linux
Microservices
Open Source
Networking
Storage

ENGINEERING

AI
AI Engineering
API Management
Backend development
Data
Frontend Development
Large Language Models
Security
Software Development
WebAssembly

OPERATIONS

AI Operations
CI/CD
Cloud Services
DevOps
Kubernetes
Observability
Operations
Platform Engineering

CHANNELS

Podcasts
Ebooks
Events
Webinars
Newsletter
TNS RSS Feeds

THE NEW STACK

About / Contact
Sponsors



Community created roadmaps, articles,
resources and journeys for developers

Frontend Developer Roadmap

Backend Developer Roadmap

Devops Roadmap

FOLLOW TNS



© The New Stack 2026

[Disclosures](#)

[Terms of Use](#)

[Advertising Terms & Conditions](#)

[Privacy Policy](#)

[Cookie Policy](#)